

THE EXPERT'S VOICE® IN OPEN SOURCE

Pro Docker

Learn how to use Containers as a Service
for development and deployment

Deepak Vohra

Apress®

Pro Docker



Deepak Vohra

Apress®

Pro Docker

Copyright © 2016 by Deepak Vohra

This work is subject to copyright. All rights are reserved by the Publisher, whether the whole or part of the material is concerned, specifically the rights of translation, reprinting, reuse of illustrations, recitation, broadcasting, reproduction on microfilms or in any other physical way, and transmission or information storage and retrieval, electronic adaptation, computer software, or by similar or dissimilar methodology now known or hereafter developed. Exempted from this legal reservation are brief excerpts in connection with reviews or scholarly analysis or material supplied specifically for the purpose of being entered and executed on a computer system, for exclusive use by the purchaser of the work. Duplication of this publication or parts thereof is permitted only under the provisions of the Copyright Law of the Publisher's location, in its current version, and permission for use must always be obtained from Springer. Permissions for use may be obtained through RightsLink at the Copyright Clearance Center. Violations are liable to prosecution under the respective Copyright Law.

ISBN-13 (pbk): 978-1-4842-1829-7

ISBN-13 (electronic): 978-1-4842-1830-3

Trademarked names, logos, and images may appear in this book. Rather than use a trademark symbol with every occurrence of a trademarked name, logo, or image we use the names, logos, and images only in an editorial fashion and to the benefit of the trademark owner, with no intention of infringement of the trademark.

The use in this publication of trade names, trademarks, service marks, and similar terms, even if they are not identified as such, is not to be taken as an expression of opinion as to whether or not they are subject to proprietary rights.

While the advice and information in this book are believed to be true and accurate at the date of publication, neither the authors nor the editors nor the publisher can accept any legal responsibility for any errors or omissions that may be made. The publisher makes no warranty, express or implied, with respect to the material contained herein.

Managing Director: Welmoed Spahr

Lead Editor: Michelle Lowman

Technical Reviewer: Massimo Nardone

Editorial Board: Steve Anglin, Pramila Balan, Louise Corrigan, Jonathan Gennick, Robert Hutchinson, Celstine Suresh John, Michelle Lowman, James Markham, Susan McDermott, Matthew Moodie, Jeffrey Pepper, Douglas Pundick, Ben Renow-Clarke, Gwenan Spearling

Coordinating Editor: Mark Powers

Composer: SPi Global

Indexer: SPi Global

Artist: SPi Global

Distributed to the book trade worldwide by Springer Science+Business Media New York,
233 Spring Street, 6th Floor, New York, NY 10013. Phone 1-800-SPRINGER, fax (201) 348-4505, e-mail
orders-ny@springer-sbm.com, or visit www.springeronline.com. Apress Media, LLC is a California LLC and
the sole member (owner) is Springer Science + Business Media Finance Inc (SSBM Finance Inc).
SSBM Finance Inc is a Delaware corporation.

For information on translations, please e-mail rights@apress.com, or visit www.apress.com.

Apress and friends of ED books may be purchased in bulk for academic, corporate, or promotional use.
eBook versions and licenses are also available for most titles. For more information, reference our Special
Bulk Sales-eBook Licensing web page at www.apress.com/bulk-sales.

Any source code or other supplementary material referenced by the author in this text is available to
readers at www.apress.com/9781484218297. For additional information about how to locate and download
your book's source code, go to www.apress.com/source-code/. Readers can also access source code at
SpringerLink in the Supplementary Material section for each chapter.

Contents at a Glance

About the Author	xiii
About the Technical Reviewer	xv
■ Chapter 1: Hello Docker	1
■ Chapter 2: Installing Linux	19
■ Chapter 3: Using Oracle Database	31
■ Chapter 4: Using MySQL Database	41
■ Chapter 5: Using MongoDB	57
■ Chapter 6: Using Apache Cassandra	81
■ Chapter 7: Using Couchbase Server	95
■ Chapter 8: Using Apache Hadoop	117
■ Chapter 9: Using Apache Hive	131
■ Chapter 10: Using Apache HBase	141
■ Chapter 11: Using Apache Sqoop	151
■ Chapter 12: Using Apache Kafka	185
■ Chapter 13: Using Apache Solr	195
■ Chapter 14: Using Apache Spark	219
■ Appendix A: Using the Amazon EC2.....	229
Index.....	253

Contents

About the Author	xiii
About the Technical Reviewer	xv
■Chapter 1: Hello Docker.....	1
Setting the Environment.....	2
Installing Docker on Red Hat 7	5
Uninstalling Docker	7
Installing a Specific Docker Version	7
Installing Docker on Ubuntu	7
Starting the Docker Service	8
Finding the Docker Service Status	9
Running a Docker Hello World Application	9
Downloading a Docker Image	11
Running an Application in a Docker Container	12
Listing Running Docker Containers	13
Accessing the Application Output on Command Line.....	14
Accessing the Application Output in a Browser	15
Stopping a Docker Container.....	16
Removing a Docker Container.....	17
Removing a Docker Image	17
Stopping the Docker Service.....	18
Summary.....	18

■ Chapter 2: Installing Linux	19
Setting the Environment.....	19
Downloading the Docker Image	21
Listing Docker Images.....	21
Running a Container in Detached Mode.....	22
Running a Container in Foreground	23
Listing Docker Containers	23
Finding Oracle Linux Container Information.....	24
Listing the Container Processes.....	25
Starting an Interactive Shell.....	26
Creating a Container.....	28
Stopping a Container.....	29
Removing a Container	29
Summary.....	30
■ Chapter 3: Using Oracle Database	31
Setting the Environment.....	31
Starting Oracle Database	33
Listing Container Logs.....	34
Starting SQL* Plus	36
Creating a User.....	37
Creating a Database Table.....	38
Removing Oracle Database	39
Summary.....	40
■ Chapter 4: Using MySQL Database	41
Setting the Environment.....	42
Starting MySQL Server	44
Starting MySQL CLI Shell.....	46
Setting the Database to Use.....	46

Creating a Database Table.....	47
Adding Table Data.....	47
Querying a Table.....	48
Listing Databases and Tables.....	48
Exiting TTY Terminal	49
Stopping a Docker Container.....	49
Starting Another MySQL Server Instance	50
Listing Docker Container Logs	54
Summary.....	55
Chapter 5: Using MongoDB.....	57
Setting the Environment.....	58
Starting MongoDB	59
Starting an Interactive Terminal	60
Starting a Mongo Shell.....	60
Creating a Database	63
Creating a Collection	64
Creating a Document.....	65
Finding Documents	66
Adding Another Document.....	66
Querying a Single Document.....	68
Dropping a Collection	69
Adding a Batch of Documents	69
Updating a Document.....	71
Outputting Documents as JSON	72
Making a Backup of the Data	73
Removing Documents	75
Stopping and Restarting the MongoDB Database	78
Exiting the Mongo Shell.....	80
Summary.....	80

■ Chapter 6: Using Apache Cassandra	81
Setting the Environment.....	82
Starting Apache Cassandra	83
Starting the TTY.....	84
Connecting to CQL Shell.....	85
Creating a Keyspace.....	85
Altering A Keyspace	86
Using A Keyspace.....	86
Creating a Table.....	87
Adding Table Data.....	87
Querying a Table.....	88
Deleting from a Table	89
Truncating a Table	90
Dropping A Table	90
Dropping a Keyspace	91
Exiting CQL Shell	91
Stopping Apache Cassandra.....	92
Starting Multiple Instances of Apache Cassandra.....	92
Summary.....	93
■ Chapter 7: Using Couchbase Server	95
Setting the Environment.....	95
Starting Couchbase	98
Accessing Couchbase Web Console	99
Configuring Couchbase Server Cluster	101
Adding Documents	109
Starting Interactive Terminal	114
Running Couchbase CLI Tools.....	114
Stopping Couchbase Server and Container.....	115
Summary.....	115

■ Chapter 8: Using Apache Hadoop	117
Setting the Environment.....	117
Starting Hadoop	119
Starting the Interactive Shell.....	120
Creating Input Files for a MapReduce Word Count Application	121
Running a MapReduce Word Count Application	124
Stopping the Hadoop Docker Container	128
Using a CDH Docker Image	128
Summary.....	130
■ Chapter 9: Using Apache Hive	131
Setting the Environment.....	131
Starting Apache Hive	132
Connecting to Beeline CLI Shell	132
Connecting to HiveServer2.....	133
Creating a Hive Table.....	135
Loading Data into the Hive Table.....	136
Querying Hive Table.....	138
Stopping Apache Hive	138
Summary.....	139
■ Chapter 10: Using Apache HBase	141
Setting the Environment.....	141
Starting CDH.....	143
Starting Interactive Shell.....	143
Starting HBase Shell	144
Creating a HBase Table	144
Listing HBase Tables	146
Getting A Single Table Row.....	147
Getting A Single Row Column.....	147

Scanning a Table	148
Stopping CDH	149
Summary	150
Chapter 11: Using Apache Sqoop	151
Setting the Environment.....	152
Starting Docker Containers	153
Starting Interactive Terminals	155
Creating a MySQL Tables.....	155
Adding MySQL JDBC Jar to Sqoop Classpath.....	160
Setting the JAVA_HOME Environment Variable	160
Configuring Apache Hadoop	163
Importing MySQL Table Data into HDFS with Sqoop	167
Listing Data Imported into HDFS	174
Exporting from HDFS to MySQL with Sqoop.....	175
Querying Exported Data	181
Stopping and Removing Docker Containers.....	182
Summary	183
Chapter 12: Using Apache Kafka	185
Setting the Environment.....	186
Starting Docker Containers for Apache Kafka	188
Finding IP Addresses	189
Listing the Kafka Logs.....	190
Creating a Kafka Topic.....	190
Starting the Kafka Producer	191
Starting the Kafka Consumer	191
Producing and Consuming Messages	192
Stopping and Removing the Docker Containers.....	193
Summary	194

■ Chapter 13: Using Apache Solr.....	195
Setting the Environment.....	195
Starting Docker Container for Apache Solr Server	197
Starting the Interactive Shell.....	199
Logging in to the Solr Admin Console	200
Creating a Core Index.....	201
Loading Sample Data	204
Querying Apache Solr in Solr Admin Console	206
Querying Apache Solr using REST API Client.....	210
Deleting Data.....	214
Listing Logs	216
Stopping Apache Solr Server.....	217
Summary.....	218
■ Chapter 14: Using Apache Spark.....	219
Setting the Environment.....	219
Running the Docker Container for CDH	220
Running Apache Spark Job in yarn-cluster Mode	221
Running Apache Spark Job in yarn-client Mode	224
Running the Apache Spark Shell	226
Summary.....	228
■ Appendix A: Using the Amazon EC2.....	229
Creating an Amazon EC2 Instance.....	229
Creating a Key Pair.....	235
Starting an Amazon EC2 Instance	237
Connecting to an Amazon EC2 Instance	238
Finding the Public IP Address.....	240
Finding the Public DNS.....	240

■ CONTENTS

Adding the default Security Group	244
Stopping an Amazon EC2 Instance.....	249
Changing the Instance Type	250
Summary.....	252
Index.....	253

About the Author



Deepak Vohra is a consultant and a principal member of the NuBean.com software company. Deepak is a Sun-certified Java programmer and Web component developer. He has worked in the fields of XML, Java programming, and Java EE for over seven years. Deepak is the coauthor of *Pro XML Development with Java Technology* (Apress, 2006). Deepak is also the author of the *JDBC 4.0* and *Oracle JDeveloper for J2EE Development*, *Processing XML Documents with Oracle JDeveloper 11g*, *EJB 3.0 Database Persistence with Oracle Fusion Middleware 11g*, and *Java EE Development in Eclipse IDE* (Packt Publishing). He also served as the technical reviewer on *WebLogic: The Definitive Guide* (O'Reilly Media, 2004) and *Ruby Programming for the Absolute Beginner* (Cengage Learning PTR, 2007).

About the Technical Reviewer



Massimo Nardone holds a Master of Science degree in Computing Science from the University of Salerno, Italy. He worked as a PCI QSA and Senior Lead IT Security/Cloud/SCADA Architect for many years and currently works as Security, Cloud and SCADA Lead IT Architect for Hewlett Packard Enterprise. He has more than 20 years of work experience in IT including Security, SCADA, Cloud Computing, IT Infrastructure, Mobile, Security and WWW technology areas for both national and international projects. Massimo has worked as a Project Manager, Cloud/SCADA Lead IT Architect, Software Engineer, Research Engineer, Chief Security Architect, and Software Specialist. He worked as visiting lecturer and supervisor for exercises at the Networking Laboratory of the Helsinki University of Technology (Aalto University). He has been programming and teaching how to program with Perl, PHP, Java, VB, Python, C/C++ and MySQL for more than 20 years. He holds four international patents (PKI, SIP, SAML and Proxy areas).

He is the author of *Pro Android Games* (Apress, 2015).

Massimo dedicates his work on this book to Roberto Salvato, Roberto Franzese and Michele Romano, who are like brothers to him and are always there when he needs them.

CHAPTER 1



Hello Docker

Docker is an open standards platform for developing, packaging and running portable distributed applications. Using Docker, developers and sysadmins may build, ship and run applications on any platform such as a PC, the cloud, data center or a virtual machine. Getting all the required dependencies for a software application including the code, the runtime libraries, and the system tools and libraries is often a challenge when developing and running an application. Docker simplifies the application development and execution by packaging all the required software for an application including the dependencies into a single software unit called a Docker image that may be run on any platform and environment.

What makes Docker images unique and different from virtual appliances, which are also software images (virtual machine images), is that while each virtual machine image runs on a separate guest OS, the Docker images run within the same OS kernel. Docker software runs in an isolated environment called a Docker container that includes its own filesystem and environment variables. Docker containers are isolated from each other and from the underlying OS.

A Docker container for a software application includes all that is required to run the software, and files may be copied from the host OS to a Docker container if required. As an application could require other software to develop a linked application, Docker containers may be linked, which makes the environment variables and software from another Docker container available to a Docker container.

Docker makes use of a Dockerfile to build an image. A Dockerfile consists of all the instructions such as what software to download, which commands to run, which network ports to expose, which files and directories to add to the filesystem, and which environment variables to set. A Docker image may be made an executable by providing an entrypoint. A Docker image may be built by providing a Dockerfile, or pre-built Docker images may be downloaded from the Docker Hub (<https://hub.docker.com/>). The complete instruction set supported by Dockerfile can be found at <http://docs.docker.com/engine/reference/builder/>.

In this chapter, we shall install the Docker engine on Linux, download a Hello World Docker image, and run a Docker container for a Hello World application. We have used Linux because some of the other software we have used, such as Apache Hadoop, is supported (both in development and production) only on Linux. We have used two commonly used distributions of Linux, Red Hat 7 and Ubuntu 14, but any of the supported installations (<https://docs.docker.com/v1.8/installation/>) could be used.

Setting the Environment

Installing Docker on Red Hat 7

Uninstalling Docker

Installing a Specific Docker Version

Installing Docker on Ubuntu

Starting the Docker Service

Finding the Docker Service Status

- Running the Docker Hello World Application
- Downloading a Docker Image
- Running an Application in a Docker Container
- Listing Running Docker Containers
- Accessing the Application Output on Command Line
- Accessing the Application Output in a Browser
- Stopping a Docker Container
- Removing a Docker Container
- Removing a Docker Image
- Stopping the Docker Service

Setting the Environment

We shall use Amazon EC2 instances based on Linux for deploying Docker and Docker images. Linux is required to support 64 bit software. We have made use of two different 64 bit (required) Amazon Machine Images (AMIs):

1. Ubuntu Server 14.04 LTS (HVM), SSD Volume Type - ami-d05e75b8 64 bit
2. Red Hat Enterprise Linux version 7.1 (HVM), EBS General Purpose (SSD) Volume Type (ami-12663b7a) 64 bit

An Amazon EC2 instance based on the Ubuntu AMI is shown in Figure 1-1.




Figure 1-1. Amazon EC2 Instance Based on Ubuntu AMI

To connect to an Amazon EC2 instance, the public IP address is used. The public IP address may be obtained from the EC2 Console as shown in Figure 1-2.

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS	Public IP
	i-036601bc	m3.medium	us-east-1b	stopped				
Docker	i-e25afa40	t2.micro	us-east-1b	stopped				
Docker	i-816ee03e	t2.micro	us-east-1b	stopped				
Docker	i-836fe93c	m3.medium	us-east-1b	stopped				
Docker	i-f85265d	m3.medium	us-east-1b	stopped				
Docker	i-2ed3718c	m3.medium	us-east-1b	stopped				
Docker	i-024c364d	t2.micro	us-east-1b	running	2/2 checks...	None	ec2-54-86-12-113.compute-1.amazonaws.com	54.86.12.113
Docker	i-6a1354d5	m3.medium	us-east-1b	stopped				
Ubuntu test	i-2247029d	t2.micro	us-east-1b	stopped				

Instance: **i-024c364d (Docker)** Public DNS: **ec2-54-86-12-113.compute-1.amazonaws.com**

Description Status Checks Monitoring Tags

Instance ID	i-024c364d	Public DNS	ec2-54-86-12-113.compute-1.amazonaws.com
Instance state	running	Public IP	54.86.12.113
Instance type	t2.micro	Elastic IP	-

Figure 1-2. Obtaining the Public IP Address

Connect to an Amazon EC2 Ubuntu instance using SSH and the public IP address with the following command in which docker.pem is the private key format (.pem) generated by Amazon EC2.

```
ssh -i "docker.pem" ubuntu@54.86.12.113
```

The Ubuntu instance gets connected to as shown in Figure 1-3.

```

ubuntu@ip-172-30-1-190: ~
File Edit View Search Terminal Help
[root@localhost ~]# ssh -i "docker.pem" ubuntu@54.86.12.113
Welcome to Ubuntu 14.04.3 LTS (GNU/Linux 3.13.0-66-generic x86_64)

 * Documentation:  https://help.ubuntu.com/

 System information as of Fri Oct 23 16:09:08 UTC 2015

 System load:  0.0          Processes:           143
 Usage of /:   76.2% of 7.74GB  Users logged in:    0
 Memory usage: 17%          IP address for eth0:  172.30.1.190
 Swap usage:   0%          IP address for docker0: 172.17.42.1

 Graph this data and manage this system at:
  https://landscape.canonical.com/

 Get cloud support with Ubuntu Advantage Cloud Guest:
  http://www.ubuntu.com/business/services/cloud

12 packages can be updated.
6 updates are security updates.

Last login: Fri Oct 23 16:09:08 2015 from d75-157-54-139.bchsia.telus.net
ubuntu@ip-172-30-1-190:~$ 
```

Figure 1-3. Connecting to Ubuntu Instance on Amazon EC2 from Local Host

If a Red Hat AMI is used the command to connect to the Amazon EC2 instance is slightly different. Instead of the user “ubuntu” use the “ec2-user” user. For example, connect to the Linux instance using the following command in which docker.pem is the private key format (.pem) generated by Amazon EC2.

```
ssh -i "docker.pem" ec2-user@54.175.182.96
```

The RHEL 7.1 instance gets connected to as shown in Figure 1-4.

```

[root@localhost ~]# ssh -i "docker.pem" ec2-user@54.175.182.96
[ec2-user@ip-172-30-1-61 ~]$ 
```

Figure 1-4. Connecting to RHEL Instance

Run the following command to find if the Linux architecture supports 64 bit software.

```
uname -r
```

The x86_64 in the output as shown in Figure 1-5 indicates that 64 bit is supported.

```
[root@localhost ~]# ssh -i "docker.pem" ec2-user@54.175.182.96
[ec2-user@ip-172-30-1-61 ~]$ uname -r
3.10.0-229.el7.x86_64
[ec2-user@ip-172-30-1-61 ~]$ █
```

Figure 1-5. Finding Architecture Support

Installing Docker on Red Hat 7

Two different methods for installing Docker on Red Hat are available: install with yum or install with script. Installing with yum requires a user to add the yum repo, which could be more involved than the script option. We have used the Docker installation script to install Docker.

As a user with sudo or root privileges, update the local repository packages with the following command.

```
sudo yum update
```

Run the Docker installation script to install Docker Engine.

```
curl -sSL https://get.docker.com/ | sh
```

Docker Engine gets installed as shown in Figure 1-6.

```
[ec2-user@ip-172-30-1-80 ~]$ curl -sSL https://get.docker.com/ | sh
+ sudo -E sh -c 'sleep 3; yum -y -q install docker-engine'
warning: /var/cache/yum/x86_64/7Server/docker-main-repo/packages/docker-engine-1
.8.2-1.el7.centos.x86_64.rpm: Header V4 RSA/SHA1 Signature, key ID 2c52609d: NOK
EY
Public key for docker-engine-1.8.2-1.el7.centos.x86_64.rpm is not installed
Importing GPG key 0x2C52609D:
  Userid : "Docker Release Tool (releasedocker) <docker@docker.com>"
  Fingerprint: 5811 8e89 f3a9 1289 7c07 0adb f762 2157 2c52 609d
  From    : https://yum.dockerproject.org/gpg
```

If you would like to use Docker as a non-root user, you should now consider adding your user to the "docker" group with something like:

```
sudo usermod -aG docker ec2-user
```

Remember that you will have to log out and back in for this to take effect!

```
[ec2-user@ip-172-30-1-80 ~]$ █
```

Figure 1-6. Installing Docker Engine

Before starting the Docker service, you should modify the `docker.service` file to disable the Docker start timeout. The `docker.service` file is in the `/usr/lib/systemd/system` directory, which has permissions set. Either run a sudo command or copy the file to a directory which does not have permissions set. For example, copy the `docker.service` to the root directory with the following command.

```
cp /usr/lib/systemd/system/docker.service .
```

Open the `docker.service` file in vi editor.

```
vi docker.service
```

Alternatively open the `docker.service` file as sudo.

```
sudo vi /usr/lib/systemd/system/docker.service
```

Add the following line to `docker.service` in the [Service] header.

```
TimeoutStartSec=0
```

The updated `docker.service` is shown in Figure 1-7.



A screenshot of a terminal window titled "ec2-user@ip-172-30-1-80:~". The window displays the contents of the `docker.service` file. The configuration includes the `[Unit]`, `[Service]`, and `[Install]` sections. In the `[Service]` section, the `TimeoutStartSec=0` line has been added. The file also contains comments for `ExecStart` and other parameters like `MountFlags`, `LimitNOFILE`, `LimitNPROC`, and `LimitCORE`.

```
[Unit]
Description=Docker Application Container Engine
Documentation=https://docs.docker.com
After=network.target docker.socket
Requires=docker.socket

[Service]
Type=notify
TimeoutStartSec=0
ExecStart=/usr/bin/docker daemon -H fd://
MountFlags=slave
LimitNOFILE=1048576
LimitNPROC=1048576
LimitCORE=infinity

[Install]
WantedBy=multi-user.target
```

Figure 1-7. Updated `docker.service`

If the `docker.service` was copied to another directory copy the file back to the `/usr/lib/systemd/system` directory with the following command.

```
sudo cp docker.service /usr/lib/systemd/system/docker.service
```

Flush changes to load the new configuration.

```
sudo systemctl daemon-reload
```

All the options for installing Docker on Red Hat are discussed at <http://docs.docker.com/engine/installation/rhel/>.

Uninstalling Docker

This section may be skipped if Docker is to be made use of in this chapter and later chapters. To uninstall Docker, run the following command to list the Docker engines installed.

```
yum list installed | grep docker
```

Remove the Docker engine and Docker directory with the following commands.

```
sudo yum -y remove docker-engine.x86_64
rm -rf /var/lib/docker
```

Installing a Specific Docker Version

To install a specific version of Docker download and install the rpm for the version. For example, install Docker 1.7.0 as follows.

```
curl -O -sSL https://get.docker.com/rpm/1.7.0/centos-6/RPMS/x86\_64/
docker-engine-1.7.0-1.el6.x86_64.rpm
sudo yum localinstall --nogpgcheck docker-engine-1.7.0-1.el6.x86_64.rpm
```

Installing Docker on Ubuntu

Docker is supported on the following versions of Ubuntu: Ubuntu Wily 15.10, Ubuntu Vivid 15.04, Ubuntu Trusty 14.04 (LTS) and Ubuntu Precise 12.04 (LTS). Regardless of version, Docker requires a 64 bit OS with a minimum Linux kernel version of 3.10. To find the kernel version, run the following command in Ubuntu terminal.

```
uname -r
```

The kernel version output is 3.13, as shown in Figure 1-8, which is fine to install Docker.



```
ubuntu@ip-172-30-1-190:~$ uname -r
3.13.0-48-generic
ubuntu@ip-172-30-1-190:~$ █
```

Figure 1-8. Outputting Kernel Version

Before installing the Docker engine on Ubuntu, update the apt sources starting with the following commands.

```
sudo apt-key adv --keyserver hkp://pgp.mit.edu:80 --recv-keys
58118E89F3A912897C070ADBF76221572C52609D
```

In the “Update your apt sources” (<http://docs.docker.com/engine/installation/ubuntu/>) Section 6. requires you to update the /etc/apt/sources.list.d/docker.list based on the Ubuntu version. The Ubuntu distribution may be found with the following command.

```
lsb_release -a
```

For Ubuntu Trusty, the following line was added to the /etc/apt/sources.list.d/docker.list file.

```
deb https://apt.dockerproject.org/repo ubuntu-trusty main
```

Run the following commands after updating the /etc/apt/sources.list.d/docker.list file.

```
sudo apt-get update
sudo apt-get purge lxc-docker*
sudo apt-cache policy docker-engine
```

Install the pre-requisites for Ubuntu with the following commands.

```
sudo apt-get update
sudo apt-get install linux-image-generic-lts-trusty
```

Reboot the system.

```
sudo reboot
```

After the host system reboots, install Docker with the following commands.

```
sudo apt-get update
sudo apt-get install docker-engine
```

Starting the Docker Service

Regardless of the Linux distribution, start the Docker service with the following command.

```
sudo service docker start
```

Docker gets started via systemctl as indicated by the OK message in Figure 1-9.

```
[ec2-user@ip-172-30-1-80 ~]$ sudo systemctl daemon-reload
[ec2-user@ip-172-30-1-80 ~]$ sudo service docker start
Starting docker (via systemctl):
[ OK ] [ec2-user@ip-172-30-1-80 ~]$ █
```

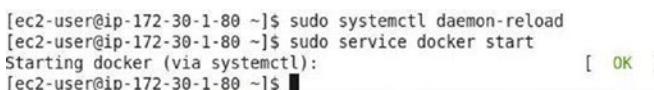


Figure 1-9. Starting Docker Service

Finding the Docker Service Status

To verify the status of the Docker service run the following command.

```
sudo service docker status
```

If the Docker service is running, the message Active: **active (running)** should be output as shown in Figure 1-10.



```
[ec2-user@ip-172-30-1-61 ~]$ sudo service docker status
docker.service - Docker Application Container Engine
  Loaded: loaded (/usr/lib/systemd/system/docker.service; enabled)
  Active: active (running) since Mon 2015-10-12 14:43:46 EDT; 9min ago
    Docs: https://docs.docker.com
   Main PID: 724 (docker)
     CGrou... /system.slice/docker.service
             └─724 /usr/bin/docker daemon -H fd://

Oct 12 14:43:46 ip-172-30-1-61.ec2.internal docker[724]: time="2015-10-12T14:...
Oct 12 14:43:46 ip-172-30-1-61.ec2.internal systemd[1]: Started Docker Applic...
Hint: Some lines were ellipsized, use -l to show in full.
[ec2-user@ip-172-30-1-61 ~]$
```

Figure 1-10. Finding Docker Service Status

Running a Docker Hello World Application

To test Docker, run the Hello World application with the following `docker run` command.

```
sudo docker run hello-world
```

The `docker run` command is introduced in a later section. If the `hello-world` application runs fine, the output in Figure 1-11, which was generated on Red Hat 7, should be generated.

```
[ec2-user@ip-172-30-1-80 ~]$ sudo service docker start
Starting docker (via systemctl): [ OK ]
[ec2-user@ip-172-30-1-80 ~]$ sudo docker run hello-world
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
535020c3e8ad: Pull complete
af340544ed62: Pull complete
Digest: sha256:a68868bfe696c00866942e8f5ca39e3e31b79c1e50feaee4ce5e28df2f051d5c
Status: Downloaded newer image for hello-world:latest
```

Hello from Docker.

This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:

```
$ docker run -it ubuntu bash
```

Share images, automate workflows, and more with a free Docker Hub account:

```
https://hub.docker.com
```

For more examples and ideas, visit:

```
https://docs.docker.com/userguide/
```

```
[ec2-user@ip-172-30-1-80 ~]$ █
```

Figure 1-11. Running hello-world Application

On Ubuntu, run the same command for hello-world.

```
sudo docker run hello-world
```

The “Hello from Docker” message gets output as shown in Figure 1-12.

The screenshot shows a terminal window titled "ubuntu@ip-172-30-1-190: ~". The terminal displays the following text:

```
ubuntu@ip-172-30-1-190:~$ sudo docker run hello-world
sudo: unable to resolve host ip-172-30-1-190
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
b901d36b6f2f: Pull complete
0a6ba66e537a: Pull complete
Digest: sha256:517f03be3f8169d84711c9ffb2b3235a4d27c1eb4ad147f6248c8040adb93113
Status: Downloaded newer image for hello-world:latest

Hello from Docker.
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:
 1. The Docker client contacted the Docker daemon.
 2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
 3. The Docker daemon created a new container from that image which runs the
 executable that produces the output you are currently reading.
 4. The Docker daemon streamed that output to the Docker client, which sent it
 to your terminal.

To try something more ambitious, you can run an Ubuntu container with:
 $ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker Hub account:
 https://hub.docker.com

For more examples and ideas, visit:
 https://docs.docker.com/userguide/
```

At the bottom of the terminal window, the prompt "ubuntu@ip-172-30-1-190:~\$" is visible.

Figure 1-12. Running `hello-world` on Ubuntu

Downloading a Docker Image

When we ran the `hello-world` application using the `docker run` command, the Docker image `hello-world` got downloaded and a Docker container for the `HelloWorld` application started. A Docker image may be downloaded automatically when a Docker container for the Docker image is started, or the Docker image may be downloaded separately. The `docker pull` command is used to download a Docker image. For example, run the following command to download the Docker image `tutum/hello-world`, which is a different `HelloWorld` application packaged as a Docker image.

```
sudo docker pull tutum/hello-world
```

The Docker image is pre-built and is not required to be built. Docker image `tutum/hello-world:latest` gets downloaded as shown in Figure 1-13. The suffix `:latest` is a label for the Docker image specifying the image version, and by default the latest version gets downloaded.

```
ubuntu@ip-172-30-1-190:~$ sudo service docker status
docker start/running, process 896
ubuntu@ip-172-30-1-190:~$ sudo docker pull tutum/hello-world
Using default tag: latest
latest: Pulling from tutum/hello-world
e8ab10ad658e: Pull complete
4fbclb79196a: Pull complete
b4aed9498114: Pull complete
6afabb57d3b2: Pull complete
00615b0849a1: Pull complete
fde01e38127f: Pull complete
5f98efe94a4e: Pull complete
5f75a51d4b66: Pull complete
Digest: sha256:bae77beaf4c2938d80e6745c788327dcacbcdad15e048980538f61ef8beca394
Status: Downloaded newer image for tutum/hello-world:latest
ubuntu@ip-172-30-1-190:~$ █
```

Figure 1-13. Downloading tutum:hello-world:latest

List the downloaded Docker images using the following command.

```
sudo docker images
```

The tutum/hello-world Docker image gets listed as shown in Figure 1-14 in addition to other images that might have been installed previously.

REPOSITORY	TAG	IMAGE ID	CREATED
couchbase	latest	ff61ecf3bacb	6 days ago
tutum/hello-world	latest	5f75a51d4b66	5 months ago

Figure 1-14. Listing Docker Images

Running an Application in a Docker Container

The docker run command is used to run a process, which is another term for an application, in a separate container. The syntax for the docker run command is as follows.

```
docker run [OPTIONS] IMAGE[:TAG|@DIGEST] [COMMAND] [ARG...]
```

The only required command parameter is a Docker image. A Docker container may be started in a detached mode (or background) or foreground mode. In detached mode the process's stdin, stdout and stderr streams are detached from the command line from which the docker run command is run. To start a container in detached mode, set -d=true or just -d. The default mode is the foreground mode in which the container starts in the foreground, and the stdin, stdout and stderr streams are attached to the host command line console. The -name option may be used to specify a name for the Docker container. The

-p option is used to specify a port for the process running in the container. As an example, start a Docker container for the tutum/hello-world image in detached mode using the **-d** parameter, with container name as helloapp and port on which the application runs as 80 using the **-p** parameter.

```
sudo docker run -d -p 80 --name helloapp tutum/hello-world
```

The Docker container gets started as shown in Figure 1-15.

```
ubuntu@ip-172-30-1-190:~$ sudo docker run -d -p 80 --name helloapp tutum/hello-world
82171f7ade462aa940f7176ea92d0a095409d42285a0b62bfa83be82ffbc57b0
```

Figure 1-15. Running an Application in a Docker Container

An interactive shell or terminal (tty) may be started to run commands applied to the process running in a container. An interactive terminal is started with the **-i** and **-t** command parameters used together or combined as **-it**. For a complete syntax of the `docker run` command, refer to <http://docs.docker.com/engine/reference/run/>.

Listing Running Docker Containers

To list running Docker container run the following command.

```
sudo docker ps
```

The helloapp container gets listed as shown in Figure 1-16. A container id is also assigned to the container. In all docker commands such as `docker stop`, `docker start` either the container name or the container id may be used.

```
ubuntu@ip-172-30-1-190:~$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND                  CREATED             STATUS              PORTS
STATUS            NAMES
82171f7ade46      tutum/hello-world   "/bin/sh -c 'php-fpm'"   34 seconds ago
Up 33 seconds          0.0.0.0:32768->80/tcp    helloapp
ubuntu@ip-172-30-1-190:~$
```

Figure 1-16. Listing only the Docker Containers that are Running

In the PORTS column, the external port allocated to the process running on port 80 in the container is listed as 32768. When accessing the helloapp application from outside the container, the 32768 port has to be used (not port 80). The external port may also be listed using the `docker port` command.

```
sudo docker port 82171f7ade46
```

The port 32768 gets listed as shown in Figure 1-17. The 0.0.0.0 host IP Address implies all IP Addresses on the local machine.

```
ubuntu@ip-172-30-1-190:~$ sudo docker port 82171f7ade46 80
0.0.0.0:32768
ubuntu@ip-172-30-1-190:~$ █
```

Figure 1-17. Listing Port

To list all Docker containers, running or exited, run the following command.

```
sudo docker ps -a
```

Accessing the Application Output on Command Line

The curl tool may be used to connect to the host and port on which the helloapp is running. Run the following command to access the application on external port 32768.

```
curl http://localhost:32768
```

The HTML generated by the helloapp gets output in the host as shown in Figure 1-18.

```
ubuntu@ip-172-30-1-190:~$ curl http://localhost:32768/
<html>
<head>
    <title>Hello world!</title>
    <link href='http://fonts.googleapis.com/css?family=Open+Sans:400,700' rel='stylesheet' type='text/css'>
        <style>
            body {
                background-color: white;
                text-align: center;
                padding: 50px;
                font-family: "Open Sans", "Helvetica Neue", Helvetica, Arial, sans-serif;
            }
            #logo {
                margin-bottom: 40px;
            }
        </style>
</head>
<body>
    
    <h1>Hello world!</h1>
    <h3>My hostname is 82171f7ade46</h3>
</body>
</html>
ubuntu@ip-172-30-1-190:~$ █
```

Figure 1-18. Output from helloapp Application

Accessing the Application Output in a Browser

However, accessing an application that generates HTML output using a curl tool is not always the best method. In this section we shall access the helloapp in a browser. If the browser is on the same machine as the host running the Docker container, the url <http://localhost:32768> may be used to display the application output. But if the browser is on a different host as in the example used in this chapter, the public DNS of the Amazon EC2 instance must be used to access the application. The public DNS may be obtained from the Amazon EC2 Console as shown in Figure 1-19.

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS	Public IP
	i-036b01bc	m3.medium	us-east-1b	stopped		None		
Docker	i-e25afa40	t2.micro	us-east-1b	stopped		None		
Docker	i-816ee83e	t2.micro	us-east-1b	stopped		None		
Docker	i-836fe93c	m3.medium	us-east-1b	stopped		None		
Docker	i-ff95265d	m3.medium	us-east-1b	stopped		None		
Docker	i-2ed3718c	m3.medium	us-east-1b	stopped		None		
Docker	i-024c364d	t2.micro	us-east-1b	running	2/2 checks...	None	ec2-54-86-12-113.comp...	54.86.12.113
Docker	i-6a1354d5	m3.medium	us-east-1b	stopped		None		
Ubuntu test	i-2247029d	t2.micro	us-east-1b	stopped		None		

Instance: **i-024c364d (Docker)** Public DNS: ec2-54-86-12-113.compute-1.amazonaws.com

Description Status Checks Monitoring Tags

Instance ID	i-024c364d	Public DNS	ec2-54-86-12-113.compute-1.amazonaws.com
Instance state	running	Public IP	54.86.12.113
Instance type	t2.micro	Elastic IP	-

Figure 1-19. Finding Public DNS

Using the public DNS, access the helloapp in a remote browser, which could be running on Windows OS, with the URL <http://ec2-54-86-12-113.compute-1.amazonaws.com:32768>. The output generated by the application running in the Docker container helloapp gets displayed in the browser as shown in Figure 1-20.




Figure 1-20. Displaying Output from helloapp in a Browser

Stopping a Docker Container

A Docker container may be stopped with the `docker stop` command. For example, stop the `helloapp` container with the following command.

```
sudo docker stop helloapp
```

The Docker container gets stopped. Subsequently run the `docker ps` command to list the running containers. The `helloapp` container does not get listed as shown in Figure 1-21.

```
ubuntu@ip-172-30-1-190:~$ sudo docker stop helloapp
helloapp
ubuntu@ip-172-30-1-190:~$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND       CREATED
STATUS              PORTS
ubuntu@ip-172-30-1-190:~$
```

Figure 1-21. Stopping a Container

Removing a Docker Container

A Docker container may be removed with the docker `rm` command. For example, remove the `helloapp` container with the following command.

```
sudo docker rm helloapp
```

A Docker container must be stopped before removing the container.

Removing a Docker Image

To remove a Docker image, run the docker `rmi` command. For example, run the following command to remove the Docker image `tutum/hello-world`.

```
sudo docker rmi tutum/hello-world
```

All containers accessing a Docker image must be stopped and removed before a Docker image can be removed. Sometimes some incompletely downloaded Docker images could get listed with the docker `images` command. Such Docker images do not have a name assigned to them and instead are listed as `<>`. All such dangling images may be removed with the following command.

```
sudo docker rmi $(sudo docker images -f "dangling=true" -q)
```

As indicated in the output in Figure 1-22, multiple Docker images get removed.

```
ubuntu@ip-172-30-1-190:~$ sudo docker rmi $(sudo docker images -f "dangling=true" -q)
Deleted: d02b48686af0eeb6ac5b22c78df4493b3a50e99c98f4cf4a4452ad1ab968d85f
Deleted: b9e0df4394839f2f474af56294e163aa8de7e41c321eb697d3487e231ef3a015
Deleted: 61c9f33b4ccee13666b55a00b27803c0a32ad764ec961a717c4c90249ed87fb4
Deleted: 6cc920ef2ab61d96b976680de2f977c318bed511b81d63ecdb44201b22083810
Deleted: 151412b3996cc4f5865f51708da0e9bf88977930b9b9663f02b96eba04aed811
Deleted: ecf0db68805f1c65163be5f791394f0ebf35a04c1033c15477af49f24fc93ee
Deleted: 94669ea371331cfe58e90e7d275d7c4afb6037d70f9ceb279188536d91a0a710
Deleted: b73d97ec6a2a56a9523d8689ccaf02f54876ddfe2ad9b9f7654adef006c2b74
Deleted: 1d073211c498fd5022699b46a936b4e4bdacb04f637ad64d3475f558783f5c3e
Deleted: 5a4526e952f0aa24f3fcc1b6971f7744eb5465d572a48d47c492cb6bbf9cbcda
Deleted: 99fcрафe76ef1aa4077b90a413af57fd17d19dce4e50d7964a273aae67055235
Deleted: c63fb41c2213f511f12f294dd729b9903a64d88f098c20d2350905ac1fdbcbba
```

Figure 1-22. Removing Dangling Docker Images

Stopping the Docker Service

To stop a Docker service, run the following command.

```
sudo service docker stop
```

The Docker service may be started again with the following command.

```
sudo service docker start
```

Alternatively, a running Docker service may be restarted with the following command.

```
sudo service docker restart
```

Summary

In this chapter we introduced the Docker engine. We installed Docker on two Linux distributions: Red Hat 7 and Ubuntu, but Docker may also be installed on other Linux distributions. For supported Docker installation operating systems, refer to <http://docs.docker.com/v1.8/installation/>. We discussed downloading a Docker image, running a Docker container using a Docker image, accessing the Docker container application from a remote browser, and stopping and removing a Docker container and a Docker image. In the next chapter, we shall run Linux in a Docker container.

CHAPTER 2



Installing Linux

Installing Linux is a task most developers and all Linux administrators are familiar with. Several Linux distributions are available including Red Hat Linux, Ubuntu, openSuse and Oracle Linux. Some of the options for installing Linux include using the Amazon Linux AMIs, ISO images and virtual machine images. Linux could also be installed using a Docker image. Several Docker images for Linux distributions are available from the Docker public repository (<https://hub.docker.com/>). In this chapter we will install Oracle Linux using a Docker image.

Setting the Environment

Downloading the Docker Image

Listing Docker Images

Running a Container in Detached Mode

Running a Container in Foreground

Listing Docker Containers

Finding Oracle Linux Container Information

Listing the Container Processes

Starting an Interactive Shell

Creating a Container

Stopping a Container

Removing a Container

Setting the Environment

The following software is required for this chapter:

-Docker (version 1.8.x used)

-Docker Image for Oracle Linux

-Host Linux OS (Amazon EC2 AMI used)

For Host OS we have used the Red Hat Enterprise Linux 7.1 (HVM), SSD Volume Type - ami-12663b7a on Amazon EC2. Login to the Amazon EC2 instance using the following command; the IP address (54.165.251.73) will be different for different users and may be obtained as explained in Appendix A.

```
ssh -i "docker.pem" ec2-user@54.165.251.73
```


Install Docker as explained in Chapter 1. Start Docker with the following command.

```
sudo service docker start
```

An OK message indicates that Docker has started. To confirm that Docker has started run the following command.

```
sudo service docker status
```

If the Active: label has the **active (running)** value as shown in Figure 2-1, Docker has started and is ready to deploy applications in Docker containers.



```
ec2-user@ip-172-30-1-192:~
```

```
File Edit View Search Terminal Help
```

```
[root@localhost ~]# ssh -i "docker.pem" ec2-user@52.23.242.61
Last login: Fri Oct 16 16:20:04 2015 from d75-157-54-139.bchsys.telus.net
[ec2-user@ip-172-30-1-192 ~]$ sudo service docker status
docker.service - Docker Application Container Engine
  Loaded: loaded (/usr/lib/systemd/system/docker.service; disabled)
  Active: active (running) since Fri 2015-10-16 14:55:59 EDT; 2h 6min ago
    Docs: https://docs.docker.com
 Main PID: 2014 (docker)
   CGroup: /system.slice/docker.service
           └─2014 /usr/bin/docker daemon -H fd://

Oct 16 16:38:48 ip-172-30-1-192.ec2.internal docker[2014]: time="2015-10-16T1...
Oct 16 16:38:48 ip-172-30-1-192.ec2.internal docker[2014]: time="2015-10-16T1...
Oct 16 16:38:48 ip-172-30-1-192.ec2.internal docker[2014]: time="2015-10-16T1...
Oct 16 16:39:40 ip-172-30-1-192.ec2.internal docker[2014]: time="2015-10-16T1...
Oct 16 16:39:47 ip-172-30-1-192.ec2.internal docker[2014]: time="2015-10-16T1...
Oct 16 16:40:04 ip-172-30-1-192.ec2.internal docker[2014]: time="2015-10-16T1...
Oct 16 16:40:48 ip-172-30-1-192.ec2.internal docker[2014]: time="2015-10-16T1...
Oct 16 16:40:48 ip-172-30-1-192.ec2.internal docker[2014]: time="2015-10-16T1...
Oct 16 16:40:48 ip-172-30-1-192.ec2.internal docker[2014]: time="2015-10-16T1...
Oct 16 16:41:59 ip-172-30-1-192.ec2.internal docker[2014]: time="2015-10-16T1...
Hint: Some lines were ellipsized, use -l to show in full.
[ec2-user@ip-172-30-1-192 ~]$
```

Figure 2-1. Finding Docker Status

Downloading the Docker Image

We have used the Docker image `oraclelinux` available from the Docker Hub Repository (https://hub.docker.com/_/oraclelinux/). Download the latest version of the `oraclelinux` Docker image with the following command.

```
sudo docker pull oraclelinux
```

Docker images are tagged to the image name to differentiate the variants (or versions) of the image. For example, to download the `oraclelinux 6.6` version, run the following command.

```
sudo docker pull oraclelinux:6.6
```

To download the `oraclelinux 7` version run the following command.

```
sudo docker pull oraclelinux:7
```

The Docker images for `oraclelinux 6.6` and `7` versions get downloaded as indicated by the output in Figure 2-2.



```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker pull oraclelinux:6.6
6.6: Pulling from library/oraclelinux
f359075ce4d8: Already exists
be320e81e49b: Already exists
e87e65fe4000: Already exists
Digest: sha256:63085b5ebc89c4d537254adf6be3b54b6029f450dcfb605974b0f2f73d98d88e
Status: Image is up to date for oraclelinux:6.6
[ec2-user@ip-172-30-1-192 ~]$ sudo docker pull oraclelinux:7.0
7.0: Pulling from library/oraclelinux
2b2532654289: Pull complete
707f44423637: Pull complete
f359075ce4d8: Already exists
Digest: sha256:a104c349bdb5153373ae979247df91b12a0ec6742a1d83f575ed237a5ff61dff
Status: Downloaded newer image for oraclelinux:7.0
```

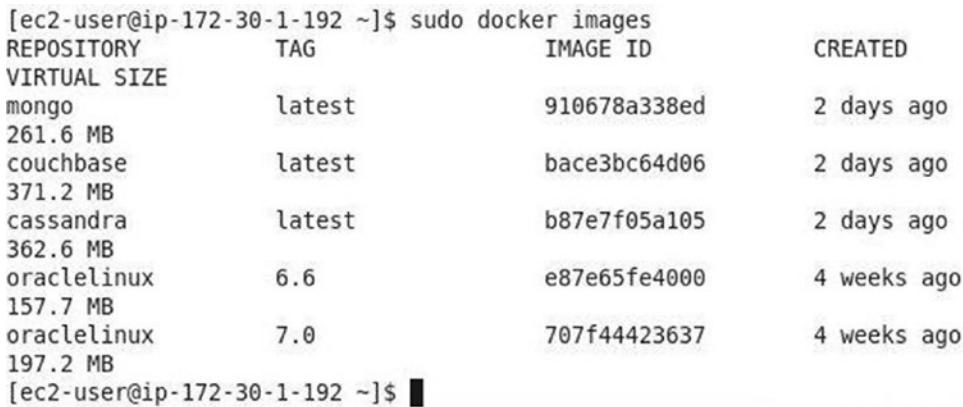
Figure 2-2. Downloading Docker Images

Listing Docker Images

The Docker images downloaded and available to run applications may be listed with the following command.

```
sudo docker images
```

The two `oraclelinux` images; versions `6.6` and `7` are listed as shown in Figure 2-3. The TAG column lists the version (or variant) of the image.



REPOSITORY	TAG	IMAGE ID	CREATED
VIRTUAL SIZE			
mongo	latest	910678a338ed	2 days ago
261.6 MB			
couchbase	latest	bace3bc64d06	2 days ago
371.2 MB			
cassandra	latest	b87e7f05a105	2 days ago
362.6 MB			
oraclelinux	6.6	e87e65fe4000	4 weeks ago
157.7 MB			
oraclelinux	7.0	707f44423637	4 weeks ago
197.2 MB			

Figure 2-3. Listing Docker Images

Running a Container in Detached Mode

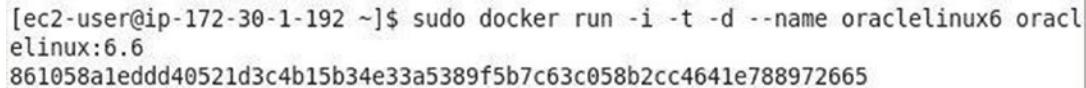
The `docker run` command is used to run a process in a container. The `docker run` command may be run in *detached mode* or *attached mode*. In detached mode the container is detached from the command line and the I/O is done through networking and shared volumes. The following command syntax would run a Docker container in a detached mode as indicated by the `-d` option. The `--name` option sets the name of the container.

```
sudo docker run -d --name <container-name> <image-name>
```

The `-i -t` options if specified with the `-d` option do not start an interactive terminal or shell. For example run the following command to start a container in detached mode with name `oraclelinux` using the `oraclelinux` Docker image with tag 6.6.

```
sudo docker run -i -t -d --name oraclelinux6 oraclelinux:6.6
```

Even though the `-i` and `-t` options are specified, the container runs in detached mode as shown in Figure 2-4.



```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker run -i -t -d --name oraclelinux6 oraclelinux:6.6
861058aleddd40521d3c4b15b34e33a5389f5b7c63c058b2cc4641e788972665
```

Figure 2-4. Starting Docker Container in Detached Mode

In detached mode, the Docker container is detached from the STDIN, STDOUT and STDERR streams. The `-rm` option cannot be used in the detached mode. For `docker run` command syntax detail, refer to <https://docs.docker.com/engine/reference/run/>.

Running a Container in Foreground

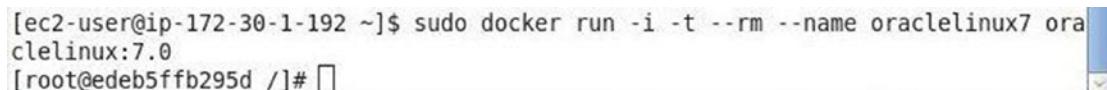
To run a Docker container in attached mode, omit the `-d` option.

```
sudo docker run <image-name>
```

In attached mode, a container process is started and attached to all the standard streams (STDIN, STDOUT and STDERR). The `-name` option may also be used in attached mode to specify a container name. To start an interactive terminal, use the `-i` and `-t` options, which allocates a tty to the container process. The `--rm` option if specified cleans up the container resources including the filesystem allocated to the container after the container has exited. Run the following command to run a container process using the `oraclelinux:7.0` Docker image; the `-name` option specifies a name to the container, the `-i -t` options start an interactive terminal (tty) and the `--rm` option cleans up the container after the container has exited.

```
sudo docker run -i -t --rm --name oraclelinux7 oraclelinux:7.0
```


The Docker container process using the `oraclelinux` image starts and attaches to an interactive shell or tty as shown in Figure 2-5.



```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker run -i -t --rm --name oraclelinux7 oraclelinux:7.0
[root@edeb5ffb295d /]# 
```

Figure 2-5. Starting Docker Container in Attached Mode

A container name must be unique. If a container with the same name as a running container is started, an error is generated as indicated in Figure 2-6.



```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker run -i -t --rm --name oraclelinux7 oraclelinux:7.0
Error response from daemon: Conflict. The name "oraclelinux7" is already in use by container edeb5ffb295d. You have to delete (or rename) that container to be able to reuse that name.
```

Figure 2-6. Container Name must be Unique

Listing Docker Containers

Docker containers can be running or not running. Run the following command to list Docker containers that are running.

```
sudo docker ps
```

The only running containers, `oraclelinux:6.6` and `oraclelinux:7.0`, get listed as shown in Figure 2-7. The STATUS column indicates whether the container is “Up” and running or “Exited”. The CONTAINER ID column lists the container ID.

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
edeb5ffb295d	oraclelinux:7.0	"/bin/bash" oraclelinux7	About a minute ago
Up About a minute			
861058aleddd	oraclelinux:6.6	"/bin/bash" oraclelinux6	5 minutes ago
Up 5 minutes			

Figure 2-7. Listing Running Docker Containers

To list all containers running or exited, run the following command.

```
sudo docker ps -a
```

The containers that have exited also get listed as shown in Figure 2-8.

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
4a453d3ebe8d	oraclelinux:6.6	"/bin/bash" orcl6	2 minutes ago
Exited (137) 30 seconds ago			
edeb5ffb295d	oraclelinux:7.0	"/bin/bash" oraclelinux7	8 minutes ago
Up 8 minutes			
861058aleddd	oraclelinux:6.6	"/bin/bash" oraclelinux6	12 minutes ago
Up 12 minutes			
a037a210d3f2	couchbase	"/entrypoint.sh couch" couchbasedb	2 hours ago
Exited (0) 37 minutes ago			
d965cbf2ad18	cassandra	"/docker-entrypoint.s" cassandradb3	22 hours ago
Exited (143) 22 hours ago			
3629909b411b	cassandra	"/docker-entrypoint.s" cassandradb2	22 hours ago
Exited (143) 22 hours ago			
dfade563f871	cassandra	"/docker-entrypoint.s" cassandradb	23 hours ago
Exited (143) 22 hours ago			
68fe88ca79fe	mongo	"/entrypoint.sh mongo" mongodb	26 hours ago
Exited (0) 26 hours ago			

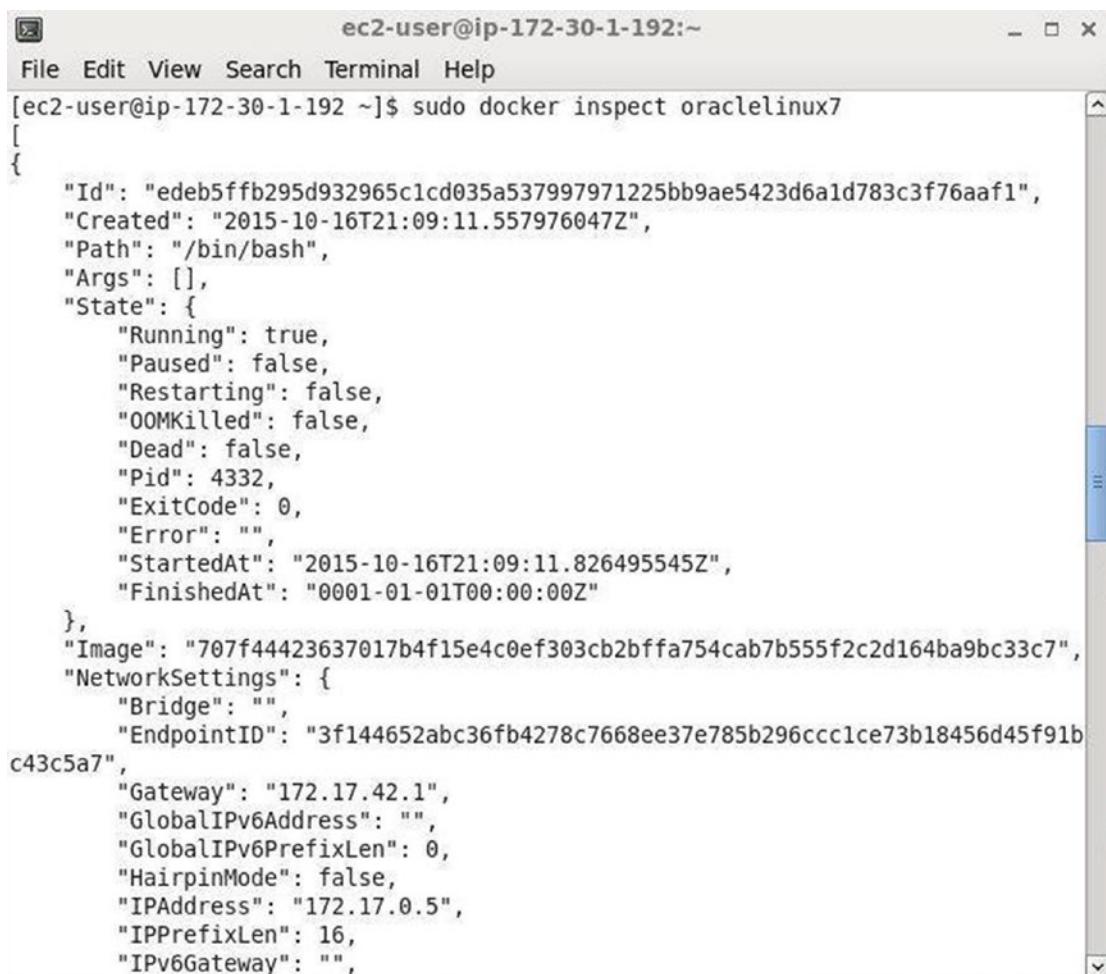
Figure 2-8. Listing All Docker Containers

Finding Oracle Linux Container Information

Information about a container can be listed with the docker inspect command. Run the following command to list information about container oraclelinux7.

```
sudo docker inspect oraclelinux7
```

The container detail gets listed in JSON format as shown in Figure 2-9.



```
ec2-user@ip-172-30-1-192:~$ sudo docker inspect oraclelinux7
[{"Id": "edeb5ffb295d932965c1cd035a537997971225bb9ae5423d6a1d783c3f76aaf1",
 "Created": "2015-10-16T21:09:11.557976047Z",
 "Path": "/bin/bash",
 "Args": [],
 "State": {
     "Running": true,
     "Paused": false,
     "Restarting": false,
     "OOMKilled": false,
     "Dead": false,
     "Pid": 4332,
     "ExitCode": 0,
     "Error": "",
     "StartedAt": "2015-10-16T21:09:11.826495545Z",
     "FinishedAt": "0001-01-01T00:00:00Z"
 },
 "Image": "707f44423637017b4f15e4c0ef303cb2bffa754cab7b555f2c2d164ba9bc33c7",
 "NetworkSettings": {
     "Bridge": "",
     "EndpointID": "3f144652abc36fb4278c7668ee37e785b296ccc1ce73b18456d45f91bc43c5a7",
     "Gateway": "172.17.42.1",
     "GlobalIPv6Address": "",
     "GlobalIPv6PrefixLen": 0,
     "HairpinMode": false,
     "IPAddress": "172.17.0.5",
     "IPPrefixLen": 16,
     "IPv6Gateway": ""
}
```

Figure 2-9. Output from `docker inspect`

Listing the Container Processes

List the processes that a container is running with the `docker top` command. The following command lists the processes run by the `oraclelinux6` container.

```
sudo docker top oraclelinux6
```

The UID and PID are among the columns listed for the processes as shown in Figure 2-10.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker top oraclelinux6
UID          PID    PPID   C   STIME   TTY      TIME     CMD
root        4135    1      0   2014   pts/2    00:00:00 /bin/bash
root        4441    4      0   2014   pts/4    00:00:00 bash
[ec2-user@ip-172-30-1-192 ~]$
```

Figure 2-10. Listing Container Processes

Starting an Interactive Shell

The interactive shell or tty may be started when the container process is started with the `docker run` command using the attached mode and the `-i -t` options to indicate an interactive terminal.

```
sudo docker run -i -t --rm <image-name>
```

Run the following command to run a container for the `oraclelinux:7.0` image and start a tty terminal.

```
sudo docker run -i -t --rm --name oraclelinux7 oraclelinux:7.0
```

An interactive shell gets started and the container process gets attached to the terminal as shown in Figure 2-11.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker run -i -t --rm --name oraclelinux7 oraclelinux:7.0
[root@edeb5ffb295d /]#
```

Figure 2-11. The interactive shell gets started when a Docker container is started in Attached Mode

If a container process has already been started in detached mode using the `-d` option, the interactive terminal may be started with the following command syntax.

```
docker exec -i -t <container> bash
```

The `-i` and `-t` options could be combined into `-it`. Run the following command to start a tty for the `oraclelinux6` container.

```
sudo docker exec -it oraclelinux6 bash
```

An interactive tty gets started as shown in Figure 2-12.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker exec -it oraclelinux6 bash  
[root@861058a1eddd /]#
```

Figure 2-12. Starting an Interactive Terminal for a Docker Docker Container running in Detached Mode

Whether the tty is started when a container process is started using the `-rm`, `-it` options or subsequently using the preceding command, container commands may be run in the interactive shell. Commands run in an interactive shell are directed at the software or application that is running in the container. For example, if the Docker container is running Oracle Linux, the tty commands are for the Oracle Linux platform. For example, output the Oracle release using the following command.

```
cat /etc/oracle-release
```

The Oracle Linux Server release 7.0 gets listed as shown in Figure 2-13.

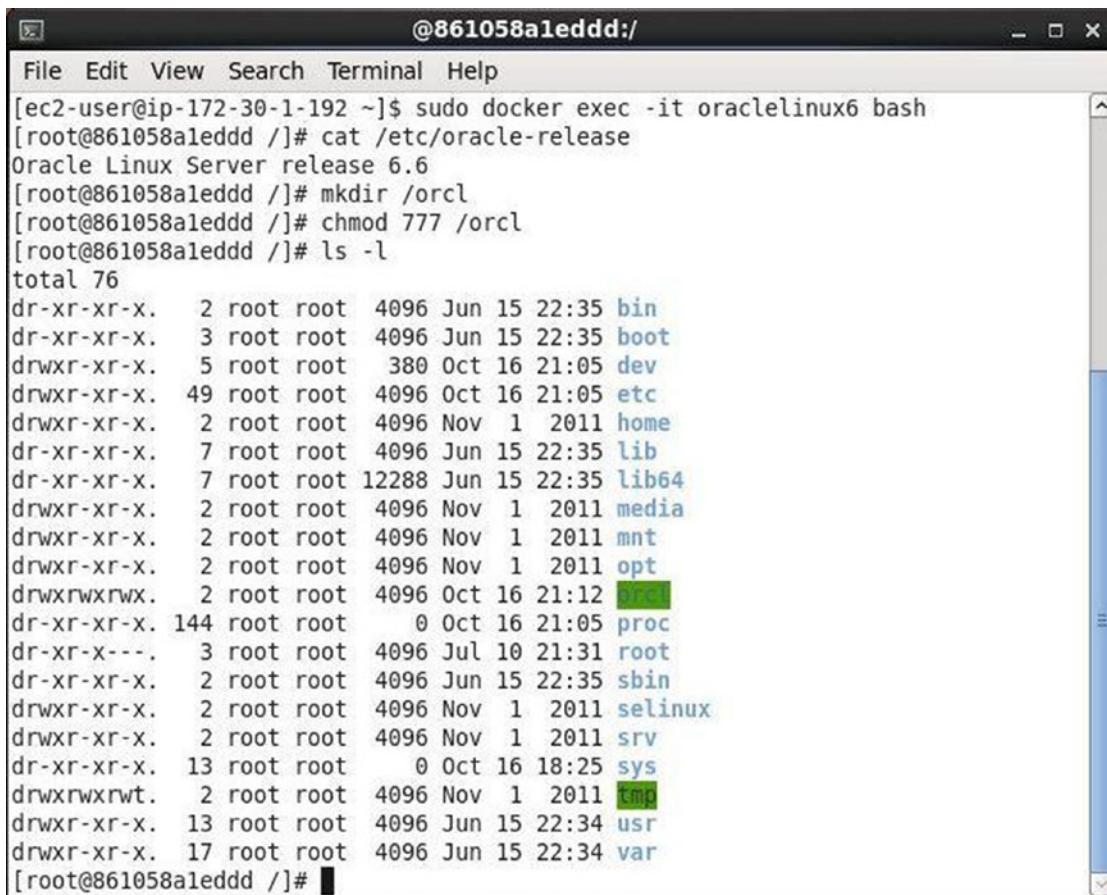
```
[root@1a402fa196a0 /]# cat /etc/oracle-release  
Oracle Linux Server release 7.0  
[root@1a402fa196a0 /]#
```

Figure 2-13. Outputting Oracle Release

Run some other Linux commands to create a directory, set the permissions on the directory, and list the files and directories.

```
mkdir /orcl  
chmod 777 /orcl  
ls -l
```


The `/orcl` directory gets created and gets listed as shown in Figure 2-14.



```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker exec -it oraclelinux6 bash
[root@861058a1eddd /]# cat /etc/oracle-release
Oracle Linux Server release 6.6
[root@861058a1eddd /]# mkdir /orcl
[root@861058a1eddd /]# chmod 777 /orcl
[root@861058a1eddd /]# ls -l
total 76
dr-xr-xr-x.  2 root root  4096 Jun 15 22:35 bin
dr-xr-xr-x.  3 root root  4096 Jun 15 22:35 boot
drwxr-xr-x.  5 root root   380 Oct 16 21:05 dev
drwxr-xr-x.  49 root root  4096 Oct 16 21:05 etc
drwxr-xr-x.  2 root root  4096 Nov  1 2011 home
dr-xr-xr-x.  7 root root  4096 Jun 15 22:35 lib
dr-xr-xr-x.  7 root root 12288 Jun 15 22:35 lib64
drwxr-xr-x.  2 root root  4096 Nov  1 2011 media
drwxr-xr-x.  2 root root  4096 Nov  1 2011 mnt
drwxr-xr-x.  2 root root  4096 Nov  1 2011 opt
drwxrwxrwx.  2 root root  4096 Oct 16 21:12 proc
dr-xr-xr-x. 144 root root     0 Oct 16 21:05 root
dr-xr-x---.  3 root root  4096 Jul 10 21:31 sbin
dr-xr-xr-x.  2 root root  4096 Jun 15 22:35 selinux
drwxr-xr-x.  2 root root  4096 Nov  1 2011 srv
dr-xr-xr-x.  13 root root     0 Oct 16 18:25 sys
drwxrwxrwt.  2 root root  4096 Nov  1 2011 tmp
drwxr-xr-x.  13 root root  4096 Jun 15 22:34 usr
drwxr-xr-x.  17 root root  4096 Jun 15 22:34 var
[root@861058a1eddd /]#
```

Figure 2-14. Listing Files and Directories

Run the `exit` command to exit the interactive shell as shown in Figure 2-15.



```
[root@la402fa196a0 /]# exit
exit
[ec2-user@ip-172-30-1-192 ~]$ sudo docker ps
```

Figure 2-15. Running the `exit` Command

Creating a Container

The `docker create` command is used to create a container. Run the following command to create a container called `orcl6` for the `oraclelinux:6.6` image. Even though the `-i -t` options are specified, an interactive shell does not get started.

```
docker create -i -t --name orcl6 oraclelinux:6.6 /bin/bash
```

To start the Docker container `orcl6` and an interactive shell for the `orcl6` container, run the `docker start` command. The `-a` and `-i` options attach the current shell's standard input, standard output and standard error streams to those of the container. All signals are forwarded to the container.

```
sudo docker start -a -i orcl6
```

The Docker container `orcl6` and an interactive shell get started as shown in Figure 2-16.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker create -i -t --name orcl6 oraclelinux:6.6 /bin/bash
4a453d3ebe8df5d16948a89c1736d0eb1fa5297d531661496588bd0895576bf1
[ec2-user@ip-172-30-1-192 ~]$ sudo docker start -a -i orcl6
[root@4a453d3ebe8d ~]#
```

Figure 2-16. Starting an Interactive Shell with `docker start`

Stopping a Container

To stop a running container, run the `docker stop` command. Run the following command to stop the `orcl6` container.

```
sudo docker stop orcl6
```

The `orcl6` container gets stopped as shown in Figure 2-17.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker stop orcl6
orcl6
```

Figure 2-17. Stopping a Docker Container

Subsequently, the `docker ps -a` command should list the `orcl6` container as “Exited” as shown in Figure 2-18.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker ps -a
CONTAINER ID        IMAGE               COMMAND                  CREATED             NAMES
STATUS              PORTS              NAMES
4a453d3ebe8d        oraclelinux:6.6   "/bin/bash"           2 minutes ago      orcl6
Exited (137) 30 seconds ago
```

Figure 2-18. Listing an Exited Container

Removing a Container

To remove a container, run the `docker rm` command. The container first must be stopped before removing, or the `docker rm` command will not remove the container. Run the following command to remove the `orcl6` container.

```
sudo docker rm orcl6
```

The orcl6 container gets removed as shown in Figure 2-19.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker rm orcl6  
orcl6  
[ec2-user@ip-172-30-1-192 ~]$ █
```

Figure 2-19. Removing A Docker Container

Summary

In this chapter we installed Oracle Linux in a Docker container. We discussed how to download the Docker image and run a container process. We also discussed using the different image tags, starting an interactive shell, the different modes of running a container, and starting, stopping and removing a container. In the next chapter we shall discuss running Oracle database in a Docker container.

CHAPTER 3



Using Oracle Database

Oracle Database is the most commonly used relational database. Relational databases are based on a fixed schema with the basic unit of storage being a table. Docker Hub has several Docker images for Oracle Database in the Public repository. In this chapter we shall use a Docker image for Oracle Database to install and use the database on Linux. This chapter has the following sections.

- Setting the Environment
- Starting Oracle Database
- Listing Container Logs
- Starting SQL* Plus
- Creating a User
- Creating a Database Table
- Removing Oracle Database

Setting the Environment

The following software is required for this chapter.

- Docker Engine (version 1.8 used)
- Docker Image for Oracle Database

We have used an Amazon EC2 instance with Red Hat Linux 7 as the OS. First, SSH login to the Amazon EC2 instance. The IP Address would be different for different users.

```
ssh -i "docker.pem" ec2-user@54.175.172.33
```

Find the status of the Docker engine.

```
sudo service docker status
```

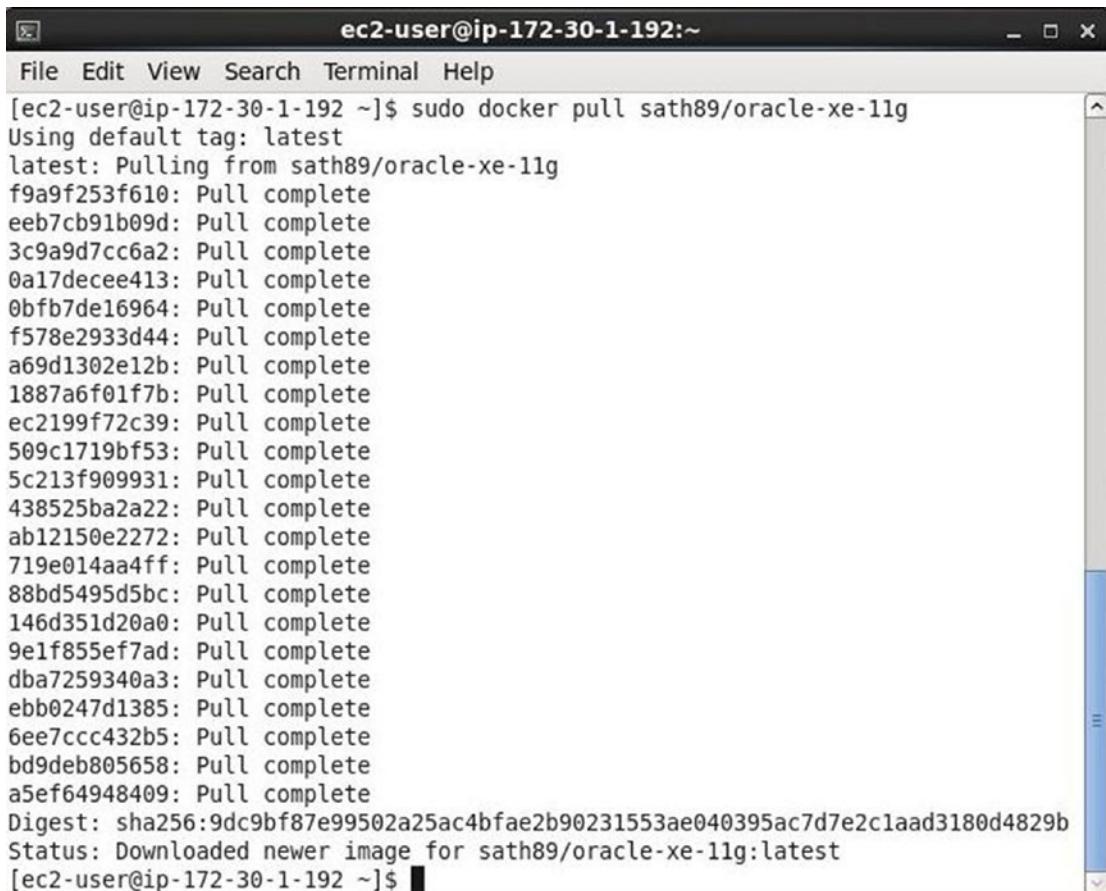
If the Docker engine is not running, start the Docker service.

```
sudo service docker start
```

Download the sath89/oracle-xe-11g Docker image.

```
sudo docker pull sath89/oracle-xe-11g
```

The latest image of sath89/oracle-xe-11g gets downloaded as shown in Figure 3-1.



A screenshot of a terminal window titled "ec2-user@ip-172-30-1-192:~". The window shows the command "sudo docker pull sath89/oracle-xe-11g" being run. The output of the command is displayed, showing the progress of pulling multiple layers of the image. The layers listed include: f9a9f253f610: Pull complete, eeb7cb91b09d: Pull complete, 3c9a9d7cc6a2: Pull complete, 0a17deceee413: Pull complete, 0bfb7de16964: Pull complete, f578e2933d44: Pull complete, a69d1302e12b: Pull complete, 1887a6f01f7b: Pull complete, ec2199f72c39: Pull complete, 509c1719bf53: Pull complete, 5c213f909931: Pull complete, 438525ba2a22: Pull complete, ab12150e2272: Pull complete, 719e014aa4ff: Pull complete, 88bd5495d5bc: Pull complete, 146d351d20a0: Pull complete, 9ef855ef7ad: Pull complete, dba7259340a3: Pull complete, ebb0247d1385: Pull complete, 6ee7ccc432b5: Pull complete, bd9deb805658: Pull complete, a5ef64948409: Pull complete. The command concludes with a Digest and Status message, both of which point to the image sath89/oracle-xe-11g:latest.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker pull sath89/oracle-xe-11g
Using default tag: latest
latest: Pulling from sath89/oracle-xe-11g
f9a9f253f610: Pull complete
eeb7cb91b09d: Pull complete
3c9a9d7cc6a2: Pull complete
0a17deceee413: Pull complete
0bfb7de16964: Pull complete
f578e2933d44: Pull complete
a69d1302e12b: Pull complete
1887a6f01f7b: Pull complete
ec2199f72c39: Pull complete
509c1719bf53: Pull complete
5c213f909931: Pull complete
438525ba2a22: Pull complete
ab12150e2272: Pull complete
719e014aa4ff: Pull complete
88bd5495d5bc: Pull complete
146d351d20a0: Pull complete
9ef855ef7ad: Pull complete
dba7259340a3: Pull complete
ebb0247d1385: Pull complete
6ee7ccc432b5: Pull complete
bd9deb805658: Pull complete
a5ef64948409: Pull complete
Digest: sha256:9dc9bf87e99502a25ac4bfae2b90231553ae040395ac7d7e2c1aad3180d4829b
Status: Downloaded newer image for sath89/oracle-xe-11g:latest
[ec2-user@ip-172-30-1-192 ~]$
```

Figure 3-1. Downloading Docker Image for Oracle Database

List the Docker images.

```
sudo docker images
```

The sath89/oracle-xe-11g image gets listed as shown in Figure 3-2.

REPOSITORY	VIRTUAL SIZE	TAG	IMAGE ID	CREATED
mongo	261.6 MB	latest	910678a338ed	2 days ago
sath89/oracle-xe-11g	792.3 MB	latest	a5ef64948409	2 days ago
couchbase	371.2 MB	latest	bace3bc64d06	2 days ago
cassandra	362.6 MB	latest	b87e7f05a105	2 days ago
oraclelinux	6.6		e87e65fe4000	4 weeks ago
o	157.7 MB			
oraclelinux	7.0		707f44423637	4 weeks ago
o	197.2 MB			
wscherphof/oracle-xe-11g-r2	latest		d4f75d4d9566	13 months ago
[ec2-user@ip-172-30-1-192 ~]\$				

Figure 3-2. Listing Docker Images

Starting Oracle Database

Next, start an Oracle Database instance in a Docker container with the `docker run` command. Specify the 8080 port for the Oracle Application Express admin console and the 1521 port for the Oracle Database listener. Specify the container name with the `-name` option.

```
docker run --name orcldb -d -p 8080:8080 -p 1521:1521 sath89/oracle-xe-11g
```

Oracle Database gets started in a Docker container as shown in Figure 3-3.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker run -d -p 8080:8080 -p 1521:1521 sath89/oracle-xe-11g
d8fb6c478d14af5bbd40a5ea7ab953ff456484511bd67c20be34a1815cd383c8
[ec2-user@ip-172-30-1-192 ~]$
```

Figure 3-3. Starting Oracle Database in a Docker Container

List the Docker containers with the following command.

```
sudo docker ps
```

The `orcldb` container gets listed as shown in Figure 3-4.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            NAMES
STATUS              PORTS              COMMAND
958aa1d91772        sath89/oracle-xe-11g   "/entrypoint.sh"   13 seconds ago   orcldb
Up 12 seconds      0.0.0.0:1521->1521/tcp, 0.0.0.0:8080->8080/tcp
361058a1eddd       oraclelinux:6.6      "/bin/bash"        About an hour ago
Up About an hour
nux6
[ec2-user@ip-172-30-1-192 ~]$
```

Figure 3-4. Listing Docker Containers that are Running

The Oracle Database hostname, port, SID, user name and password are as follows.

```
hostname: localhost
port: 1521
sid: xe
username: system
password: oracle
```

Listing Container Logs

To list the container logs, run the `docker logs` command.

```
sudo docker logs -f c0fa107a43d2
```

The container logs get listed as shown in Figure 3-5. The Oracle Database logs include the database initialization and configuration.

```

ec2-user@ip-172-30-1-192:~$ sudo docker logs -f c0fa107a43d2
ls: cannot access /u01/app/oracle/oradata: No such file or directory
Database not initialized. Initializing database.
Setting up:
processes=500
sessions=555
transactions=610
If you want to use different parameters set processes, sessions, transactions env variables and consider this formula:
processes=x
sessions=x*1.1+5
transactions=sessions*1.1

Oracle Database 11g Express Edition Configuration
-----
This will configure on-boot properties of Oracle Database 11g Express Edition. The following questions will determine whether the database should be starting upon system boot, the ports it will use, and the passwords that will be used for database accounts. Press <Enter> to accept the defaults. Ctrl-C will abort.

Specify the HTTP port that will be used for Oracle Application Express [8080]:
Specify a port that will be used for the database listener [1521]:
Specify a password to be used for database accounts. Note that the same password will be used for SYS and SYSTEM. Oracle recommends the use of different passwords for each database account. This can be done after initial configuration:
Confirm the password:

Do you want Oracle Database 11g Express Edition to be started on boot (y/n) [y]: Starting Oracle Net Listener...Done

```

Figure 3-5. Listing Docker Container Log

A more detailed Docker container log is as follows.

```

[ec2-user@ip-172-30-1-192 ~]$ sudo docker logs -f cofa107a43d2
Database not initialized. Initializing database.
Setting up:
processes=500
sessions=555
transactions=610
If you want to use different parameters set processes, sessions, transactions env variables and consider this formula:
processes=x
sessions=x*1.1+5
transactions=sessions*1.1

```

Oracle Database 11g Express Edition Configuration

This will configure on-boot properties of Oracle Database 11g Express Edition. The following questions will determine whether the database should be starting upon system boot, the ports it will use, and the passwords that will be used for database accounts. Press <Enter> to accept the defaults. Ctrl-C will abort.

Specify the HTTP port that will be used for Oracle Application Express [8080]:

Specify a port that will be used for the database listener [1521]:

Specify a password to be used for database accounts. Note that the same password will be used for SYS and SYSTEM. Oracle recommends the use of different passwords for each database account. This can be done after initial configuration:

Confirm the password:

Do you want Oracle Database 11g Express Edition to be started on boot (y/n) [y]:

Starting Oracle Net Listener...Done

Configuring database...Done

Starting Oracle Database 11g Express Edition instance...Done

Installation completed successfully.

Database initialized. Please visit <http://#containereer:8080/apex> to proceed with configuration

Oracle Database 11g Express Edition instance is already started

Database ready to use. Enjoy! ;)

[ec2-user@ip-172-30-1-192 ~]\$

Starting SQL* Plus

Start an interactive shell using the following command. The container ID would most likely be different.

```
sudo docker exec -it c0fa107a43d2 bash
```

For more detail on bash refer to <http://www.gnu.org/software/bash/manual/bash.html#Bash-Startup-Files>. Run the following command in the tty. The terms “tty”, “interactive shell” and “interactive terminal” have been used interchangeably.

```
sqlplus
```

When prompted for a user-name as shown in Figure 3-6, specify “system”

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker exec -it d8fb6c478d14 bash
root@d8fb6c478d14:/# sqlplus
SQL*Plus: Release 11.2.0.2.0 Production on Fri Oct 16 21:55:50 2015
Copyright (c) 1982, 2011, Oracle. All rights reserved.

Enter user-name: [REDACTED]
```

Figure 3-6. Starting SQL*Plus

When prompted for a password, specify “oracle”. A connection gets established with Oracle Database 11g Express. SQL*Plus gets started and the SQL> prompt gets displayed as shown in Figure 3-7.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker exec -it d8fb6c478d14 bash
root@d8fb6c478d14:/# sqlplus
SQL*Plus: Release 11.2.0.2.0 Production on Fri Oct 16 21:55:50 2015
Copyright (c) 1982, 2011, Oracle. All rights reserved.

Enter user-name: system
Enter password:

Connected to:
Oracle Database 11g Express Edition Release 11.2.0.2.0 - 64bit Production
SQL> [REDACTED]
```

Figure 3-7. SQL*Plus Shell Prompt

We used the container id to start the interactive tty terminal. Alternatively, the container name may be used as follows.

```
sudo docker exec -it orcldb bash
```

Creating a User

To create a user called OE with unlimited quota on SYSTEM tablespace and password as “OE”, run the following command.

```
SQL> CREATE USER OE QUOTA UNLIMITED ON SYSTEM IDENTIFIED BY OE;
Grant the CONNECT and RESOURCE roles to the OE user.
GRANT CONNECT, RESOURCE TO OE;
```

User “OE” gets created and the roles get granted as shown in Figure 3-8.

```
SQL> CREATE USER OE QUOTA UNLIMITED ON SYSTEM IDENTIFIED BY OE;
GRANT CONNECT,RESOURCE TO OE;
```

User created.

```
SQL>
Grant succeeded.
```

Figure 3-8. Creating User OE

Creating a Database Table

Create a database called “Catalog” in the “OE” schema with the following SQL statement.

```
SQL> CREATE TABLE OE.Catalog(CatalogId INTEGER PRIMARY KEY,Journal VARCHAR2(25),Publisher
VARCHAR2(25),Edition VARCHAR2(25),Title VARCHAR2(45),Author VARCHAR2(25));
```

Table “Catalog” gets created as shown in Figure 3-9.

```
SQL> CREATE TABLE OE.Catalog(CatalogId INTEGER PRIMARY KEY,Journal VARCHAR2(25),
Publisher VARCHAR2(25),Edition VARCHAR2(25),Title VARCHAR2(45),Author VARCHAR2(2
5));
```

Table created.

```
SQL> □
```

Figure 3-9. Creating Oracle Database Table OE.Catalog

Add data to the Catalog table with the following INSERT SQL statement.

```
SQL> INSERT INTO OE.Catalog VALUES('1','Oracle Magazine','Oracle Publishing','November
December 2013','Engineering as a Service','David A. Kelly');
```

One row of data gets added as shown in Figure 3-10.

```
SQL> INSERT INTO OE.Catalog VALUES('1','Oracle Magazine','Oracle Publishing','No
vember December 2013','Engineering as a Service','David A. Kelly');
```

1 row created.

Figure 3-10. Adding Data to OE.Catalog Table

Run a SQL query with the following SELECT statement.

```
SQL> SELECT * FROM OE.CATALOG;
```

The one row of data added gets listed as shown in Figure 3-11.

```
SQL> SELECT * FROM OE.CATALOG;
```

CATALOGID	JOURNAL	PUBLISHER
EDITION	TITLE	
AUTHOR		
1	Oracle Magazine November December 2013	Oracle Publishing Engineering as a Service
	David A. Kelly	

Figure 3-11. Running a SQL Query

To exit from SQL*Plus, specify the exit command as shown in Figure 3-12.

```
SQL> exit
Disconnected from Oracle Database 11g Express Edition Release 11.2.0.2.0 - 64bit
Production
root@c0fa107a43d2:/#
```

Figure 3-12. Exiting SQL*Plus

Removing Oracle Database

To remove the container running the Oracle Database instance, run the following docker rm command.

```
sudo docker rm c0fa107a43d2
```

To remove the Docker image sath89/oracle-xe-11g, run the following command.

```
sudo docker rmi sath89/oracle-xe-11g
```

The Docker container and image get removed as shown in Figure 3-13.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker rm c0fa107a43d2
c0fa107a43d2
[ec2-user@ip-172-30-1-192 ~]$ sudo docker rmi sath89/oracle-xe-11g;
Untagged: sath89/oracle-xe-11g:latest
Deleted: 64ba943f4946a3708406cd85bfdedcfe92db6c88390cc28610624e88dadffec8
Deleted: 129333db1a2bccf888380860cdc657c676284895d2183c5f0a79e1d42b165d71
Deleted: 3deda6a7eb42960b7644d6f40ba6cd4f372d66e79d5ed8efc3c342aa7466204c
Deleted: 3f115617bdbbf1262cdc9d6d5a251b97500b30dbb6e0f57f77907261d7bd8f2d
Deleted: ea069a6319bc52509d9fd396649804925a5a1b9eacec42504b98f8d9c814affd
Deleted: c7dd68b919effac9889cde331b727e3d424fc969586e6576db01f816390cf6
Deleted: 1ca1d4664e455015c91d2b5d84380b38f264d326ce8b905e6c1e056560a9aca2
Deleted: 1c0948b818253d845e5bf8d610eeaaca70165801f88c4ffd902e1d75d20527dc
Deleted: 9a32c4a673dcf0d7c322b27d9ad641e3856f0fc1411ab37aab0efe3d8347e3c5
Deleted: 5cb89504dc26dc126623a92585ae81ab32b6230980ccfb9a5b0743c682758ec3
Deleted: 451173061496e3dbc当地9c1cbc当地836b2e61a2868ca10aba3542af0e753d951
Deleted: b2840dc7b7e9095fb6d1527a2067d37d7708bebcc697715df1d91d8788581ba1
Deleted: 61c9e7b435e700c08b5e3803c37bcfdffea4218723e09e2a1f47a9b58440e404
Deleted: c55cb1f33d2b44539acfe2630f55c965b56d9c9146e97215d2af0721b8822eb4
Deleted: 21b6bd3f7d2c2ebabbcc449b3e7a0e24f57963be53907138b0b58193706db137e
Deleted: e108ae7ceb2b0c5d46ec1d206eb5b6e629bf9120572e77974c68e94077ab216c
Deleted: 032e7ae8eda85b20ab5e7778839819c727fbce2ec5fc9b39fc010b73ae42163d
Deleted: 9969fa71ca0c96b58365c45d3fa8704679dc当地7a234d276c49e44881d6e7956aa
Deleted: 0a17decee4139b0de68478f149cc16346f5e711c5ae3bb969895f22dd6723751
Deleted: 3c9a9d7cc6a235eb2de58ca9ef3551c67ae42a991933ba4958d207b29142902b
Deleted: eeb7cb91b09d5de9edb2798301aeedf50848eacc2123e98538f9d014f80f243c
Deleted: f9a9f253f6105141e0f8e091a6bcd当地19e3f27af949842db93acba9048ed2410b
[ec2-user@ip-172-30-1-192 ~]$ █
```

Figure 3-13. Removing Docker Image

Summary

In this chapter we used a Docker image to install Oracle Database 11g XE on an Amazon EC2 instance. We logged into SQL*Plus and created a database table to demonstrate the use of Oracle Database running in a Docker container. In the next chapter, we shall run the MySQL Database in a Docker container.

CHAPTER 4



Using MySQL Database

MySQL is the most commonly used open source relational database. MySQL is similar to Oracle Database in some regards such as users are kept in grant tables by the database. But MySQL is different from Oracle Database in some regards too:

1. MySQL does not have roles and privileges have to be granted individually to users.
2. Database and table names are case-insensitive in Oracle but are case sensitive if the underlying OS is case-sensitive.
3. MySQL provides a default value for columns that do not allow a NULL value and a value is not provided explicitly in the INSERT statement, if the strict mode is not enabled. Oracle database does not generate a default value for columns with the NOT NULL constraint.
4. MySQL database supports AUTO_INCREMENT for a column while a Sequence is used in Oracle Database.
5. Some of the data types in MySQL are different. For example, MySQL does not support the VARCHAR2 data type.

In this chapter we shall run MySQL database in a Docker container. This chapter has the following sections.

Setting the Environment

Starting MySQL CLI Shell

Setting the Database to Use

Creating a Database Table

Adding Table Data

Querying a Table

Listing Databases and Tables

Exiting TTY Terminal

Starting Another MySQL Server Instance

Listing Docker Container Log

Setting the Environment

The following software is required for this chapter.

- Docker Engine (version 1.8 used)
- Docker image for MySQL Database

Login to an Amazon EC2 instance using the public IP address of the instance.

```
ssh -i "docker.pem" ec2-user@52.91.169.69
```

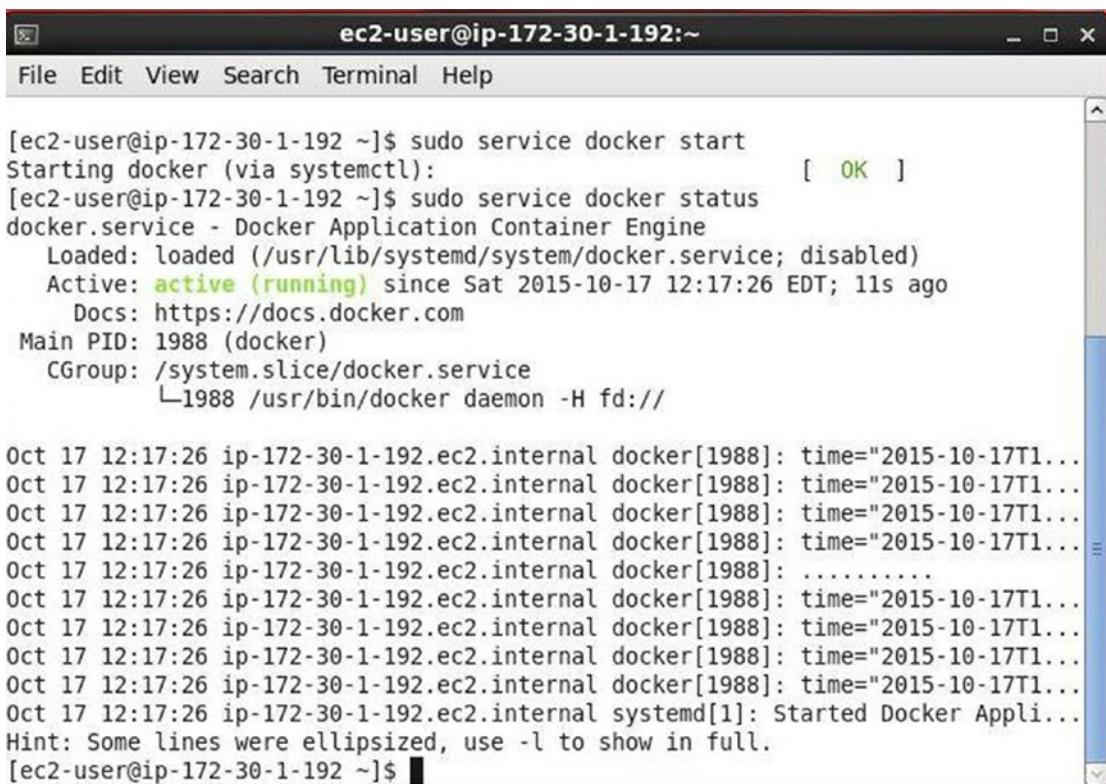
Start the Docker service.

```
sudo service docker start
```

Verify that the Docker service is running.

```
sudo service docker status
```

The output from the `docker start` command should be OK and the output from the `docker status` command should be **active (running)** for the Active field as shown in Figure 4-1.



```
[ec2-user@ip-172-30-1-192 ~]$ sudo service docker start
Starting docker (via systemctl): [OK]
[ec2-user@ip-172-30-1-192 ~]$ sudo service docker status
docker.service - Docker Application Container Engine
  Loaded: loaded (/usr/lib/systemd/system/docker.service; disabled)
  Active: active (running) since Sat 2015-10-17 12:17:26 EDT; 11s ago
    Docs: https://docs.docker.com
   Main PID: 1988 (docker)
     CGroup: /system.slice/docker.service
             └─1988 /usr/bin/docker daemon -H fd://

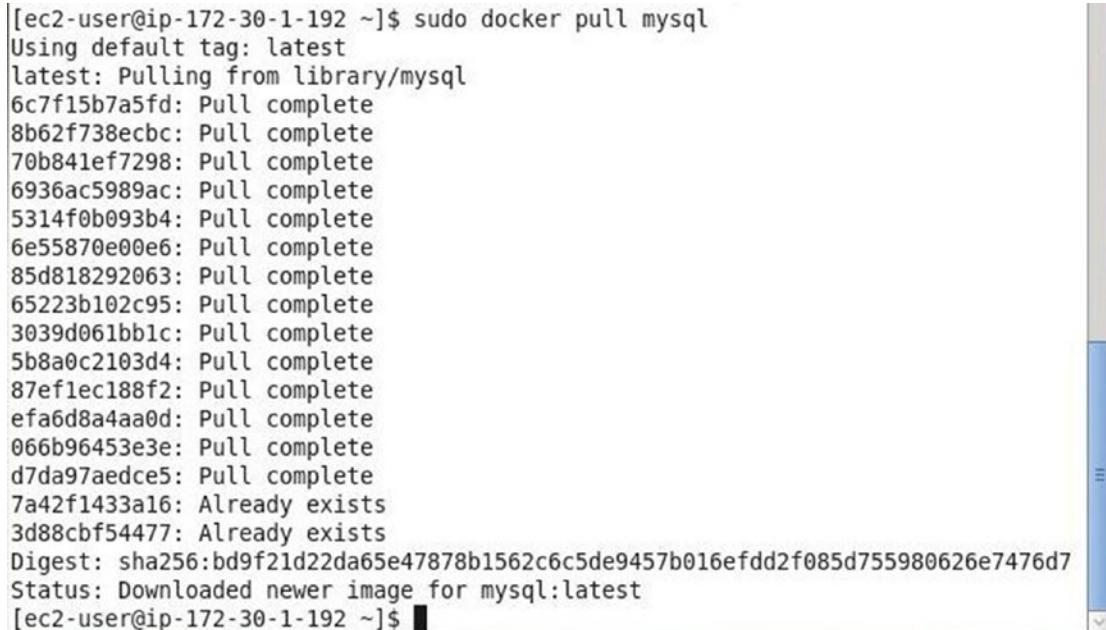
Oct 17 12:17:26 ip-172-30-1-192.ec2.internal docker[1988]: time="2015-10-17T1...
Oct 17 12:17:26 ip-172-30-1-192.ec2.internal docker[1988]: .....
Oct 17 12:17:26 ip-172-30-1-192.ec2.internal docker[1988]: time="2015-10-17T1...
Oct 17 12:17:26 ip-172-30-1-192.ec2.internal systemd[1]: Started Docker Appli...
Hint: Some lines were ellipsized, use -l to show in full.
[ec2-user@ip-172-30-1-192 ~]$
```

Figure 4-1. Starting Docker Service and verifying Status

Docker Hub provides an official Docker image. Download the Docker image with the following command.

```
sudo docker pull mysql
```

The latest Docker image `mysql:latest` gets downloaded as shown in Figure 4-2.



```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker pull mysql
Using default tag: latest
latest: Pulling from library/mysql
6c7f15b7a5fd: Pull complete
8b62f738ecbc: Pull complete
70b841ef7298: Pull complete
6936ac5989ac: Pull complete
5314f0b093b4: Pull complete
6e55870e00e6: Pull complete
85d818292063: Pull complete
65223b102c95: Pull complete
3039d061bb1c: Pull complete
5b8a0c2103d4: Pull complete
87ef1ec188f2: Pull complete
efa6d8a4aa0d: Pull complete
066b96453e3e: Pull complete
d7da97aedce5: Pull complete
7a42f1433a16: Already exists
3d88cbf54477: Already exists
Digest: sha256:bd9f21d22da65e47878b1562c6c5de9457b016efdd2f085d755980626e7476d7
Status: Downloaded newer image for mysql:latest
[ec2-user@ip-172-30-1-192 ~]$
```

Figure 4-2. Downloading Docker Image for MySQL Database

List the Docker images with the following command.

```
sudo docker images
```

The `mysql` image gets listed as shown in Figure 4-3.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker images
REPOSITORY          TAG      IMAGE ID      CREATED
VIRTUAL SIZE
mongo              latest   910678a338ed  2 days ago
261.6 MB
couchbase          latest   bace3bc64d06  3 days ago
371.2 MB
cassandra          latest   b87e7f05a105  3 days ago
362.6 MB
mysql              latest   d7da97aedce5  3 days ago
324.3 MB
oraclelinux         6.6     e87e65fe4000  4 weeks ago
157.7 MB
oraclelinux         7.0     707f44423637  4 weeks ago
197.2 MB
[ec2-user@ip-172-30-1-192 ~]$
```

Figure 4-3. Listing Docker Image for MySQL Database

Starting MySQL Server

In this section we shall run MySQL database in a Docker container. MySQL database uses the `/var/lib/mysql` directory by default for storing data, but another directory may also be used. We shall use the `/mysql/data` directory for storing MySQL data. Create the `/mysql/data` directory and set its permissions to global (777).

```
sudo mkdir -p /mysql/data
sudo chmod -R 777 /mysql/data
```

The `/mysql/data` directory gets created as shown in Figure 4-4.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo mkdir -p /mysql/data
[ec2-user@ip-172-30-1-192 ~]$ sudo chmod -R 777 /mysql/data
```

Figure 4-4. Creating the Data Directory

When the `docker run` command is run to start MySQL in a Docker container, certain environment variables may be specified as discussed in the following table.

Env Variable	Description	Required
<code>MYSQL_ROOT_PASSWORD</code>	Password for the “root” user.	Yes
<code>MYSQL_DATABASE</code>	Creates a database	No
<code>MYSQL_USER</code> , <code>MYSQL_PASSWORD</code>	Specify the username and password to create a new user. The user is granted superuser privileges on the database specified in the <code>MYSQL_DATABASE</code> variable. Both the user name and password must be set if either is set.	No
<code>MYSQL_ALLOW_EMPTY_PASSWORD</code>	Specifies whether the “root” user is permitted to have an empty password.	No

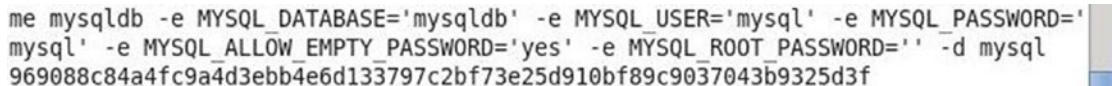
Other than the `MYSQL_ROOT_PASSWORD` environment variable, all the other variables are optional, but we shall run a MySQL instance container using all the environment variables. We shall run the `docker run` command using the following command parameters.

Command Parameter	Value
<code>MYSQL_ROOT_PASSWORD</code>	''
<code>MYSQL_DATABASE</code>	<code>mysqldb</code>
<code>MYSQL_USER, MYSQL_PASSWORD</code>	<code>mysql, mysql</code>
<code>MYSQL_ALLOW_EMPTY_PASSWORD</code>	<code>yes</code>
<code>-v</code>	<code>/mysql/data:/var/lib/mysql</code>
<code>--name</code>	<code>mysqldb</code>
<code>-d</code>	

The environment variables are specified with `-e`. Run the following `docker run` command to start a MySQL instance in a Docker container.

```
sudo docker run -v /mysql/data:/var/lib/mysql --name mysqldb -e MYSQL_DATABASE='mysqldb' -e MYSQL_USER='mysql' -e MYSQL_PASSWORD='mysql' -e MYSQL_ALLOW_EMPTY_PASSWORD='yes' -e MYSQL_ROOT_PASSWORD='' -d mysql
```

The output from the `docker run` command is shown in Figure 4-5.



```
me mysqldb -e MYSQL_DATABASE='mysqldb' -e MYSQL_USER='mysql' -e MYSQL_PASSWORD='mysql' -e MYSQL_ALLOW_EMPTY_PASSWORD='yes' -e MYSQL_ROOT_PASSWORD='' -d mysql
969088c84a4fc9a4d3ebb4e6d133797c2bf73e25d910bf89c9037043b9325d3f
```

Figure 4-5. Running MySQL Database in a Docker Container

Run the following command to list the Docker containers that are running.

```
sudo docker ps
```

The Docker container `mysqldb` that is running the MySQL database instance gets listed as shown in Figure 4-6.



CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
969088c84a4f	mysql	"/entrypoint.sh mysql"	26 seconds ago
Up 25 seconds	3306/tcp	mysqldb	

Figure 4-6. Listing Docker Containers

Starting MySQL CLI Shell

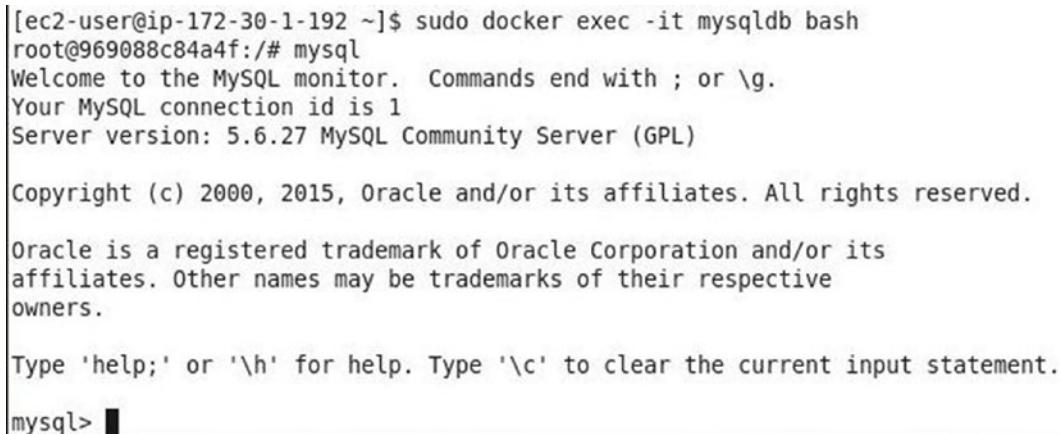
Next, we shall log into the MySQL CLI shell. But first we need to start an interactive terminal to run the `mysql` command to start the MySQL CLI. Start the interactive terminal or shell with the following command.

```
sudo docker exec -it mysqlDb bash
```

In the interactive terminal run the following command.

```
mysql
```

The MySQL CLI gets started as shown in Figure 4-7.



A screenshot of a terminal window showing the MySQL CLI starting up. The terminal output is as follows:

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker exec -it mysqlDb bash
root@969088c84a4f:/# mysql
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 1
Server version: 5.6.27 MySQL Community Server (GPL)

Copyright (c) 2000, 2015, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> ■
```

Figure 4-7. Starting MySQL CLI

The interactive terminal may also be started using the container id instead of the container name.

```
sudo docker exec -it 969088c84a4f bash
```

Setting the Database to Use

Set the database with the “use” command. The “test” database is not provided by the MySQL database started in a Docker container by default. If the “use test” command is run, the following error message is output.

```
mysql> use test
ERROR 1049 (42000): Unknown database 'test'
```

We created a database called “mysqlDb” when we started the Docker container for MySQL database with the `docker run` command. Set the database to “mysqlDb” with the following command.

```
mysql> use mysqlDb
```

The output from the preceding commands is as follows. The database gets set to “mysqldb” as shown in Figure 4-8.

```
mysql> use test
ERROR 1049 (42000): Unknown database 'test'
mysql> use mysqldb
Database changed
mysql>
```

Figure 4-8. Setting Database to mysqldb

Creating a Database Table

Next, create a database table called “Catalog” with columns CatalogId, Journal, Publisher, Edition, Title and Author. Run the following SQL statement.

```
mysql> CREATE TABLE Catalog(CatalogId INTEGER PRIMARY KEY,Journal VARCHAR(25),Publisher
VARCHAR(25),Edition VARCHAR(25),Title VARCHAR(45),Author VARCHAR(25));
```

The Catalog table gets created as shown in Figure 4-9.

```
mysql> use mysqldb
Database changed
mysql> CREATE TABLE Catalog(CatalogId INTEGER PRIMARY KEY,Journal VARCHAR(25),Pu
blisher VARCHAR(25),Edition VARCHAR(25),Title VARCHAR(45),Author VARCHAR(25));
Query OK, 0 rows affected (0.02 sec)
```

Figure 4-9. Creating a MySQL Database Table

Adding Table Data

Add data to the Catalog table with the following INSERT statement.

```
mysql> INSERT INTO Catalog VALUES('1','Oracle Magazine','Oracle Publishing','November
December 2013','Engineering as a Service','David A. Kelly');
```

A row of data gets added to the Catalog table as shown in Figure 4-10.

```
mysql> INSERT INTO Catalog VALUES('1','Oracle Magazine','Oracle Publishing','Nov
ember December 2013','Engineering as a Service','David A. Kelly');
Query OK, 1 row affected (0.01 sec)
```

Figure 4-10. Adding a Row of Data to MySQL Table

Querying a Table

Next, query the Catalog table with a SQL query. The following SELECT statement selects all the data in the Catalog table.

```
mysql> SELECT * FROM Catalog;
```

The one row of data added gets listed as shown in Figure 4-11.

```
mysql> SELECT * FROM Catalog;
+-----+-----+-----+-----+-----+
| CatalogId | Journal           | Publisher        | Edition       | Title
le          | Author            |                 |
+-----+-----+-----+-----+-----+
|      1 | Oracle Magazine | Oracle Publishing | November December 2013 | Eng
ineering as a Service | David A. Kelly |
+-----+-----+-----+-----+
1 row in set (0.00 sec)

mysql> █
```

Figure 4-11. Running a SQL Query

MySQL table name is case sensitive on the OS (RHEL 7.1 OS) used in this chapter. If a variation of the table name Catalog is used, an error is generated. For example, use table name CATALOG in the SQL query and the following error gets generated as shown in Figure 4-12.

```
mysql> SELECT * FROM CATALOG;
ERROR 1146 (42S02): Table 'mysqldb.CATALOG' doesn't exist
mysql> █
```

Figure 4-12. The table name is Case-sensitive in MySQL

Listing Databases and Tables

The databases in a MySQL server instance may be listed with the following command in MySQL CLI.

```
mysql> show databases;
```

The databases get listed, including the newly created database “mysqldb” as shown in Figure 4-13.

```
mysql> show databases;
+-----+
| Database      |
+-----+
| information_schema |
| mysql          |
| mysql ldb      |
| performance_schema |
+-----+
4 rows in set (0.00 sec)
```

```
mysql> █
```

Figure 4-13. Listing MySQL Databases

Exiting TTY Terminal

Exit the MySQL CLI with the “exit” command.

```
mysql> exit
Bye
```

Exit the interactive shell or tty with the “exit” command.

```
root@969088c84a4f:/# exit
exit
```

The output from the preceding commands is shown in Figure 4-14.

```
mysql> exit
Bye
root@969088c84a4f:/# exit
exit
[ec2-user@ip-172-30-1-192 ~]$ █
```

Figure 4-14. Exiting MySQL CLI

Stopping a Docker Container

Stop the Docker container with the docker stop command.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker stop 969088c84a4f
969088c84a4f
```

Subsequently, list the running Docker containers with the `docker ps` command. The `mysqldb` container does not get listed.

```
sudo docker ps
[ec2-user@ip-172-30-1-192 ~]$ sudo docker ps
CONTAINER ID        IMAGE       COMMAND       CREATED          STATUS          PORTS          NAMES
```

In the next section we shall create another MySQL Server instance just as we created the MySQL server instance earlier in this chapter. But we cannot use the same container name as an existing container. Another Docker container running a MySQL database, or any other software, may be started if the Docker container name is different. If we created a Docker container to run another MySQL server with the same name “`mysqldb`”, an error gets generated. For example, run the following `docker run` command to create another container called “`mysqldb`”.

```
sudo docker run --name mysqldb -e MYSQL_ROOT_PASSWORD=mysql -d mysql
```

The following error gets output.

```
Error response from daemon: Conflict. The name "mysqldb" is already in use by container 969088c84a4f. You have to delete (or rename) that container to be able to reuse that name.
```

To create a new Docker container called “`mysqldb`” first remove the “`mysqldb`” container already created with the `docker rm` command. Either the container id or the container name may be used in `docker` commands for a container such as `stop`, `start`, and `rm`.

```
sudo docker rm 969088c84a4f
```

Starting Another MySQL Server Instance

Having removed the “`mysqldb`” container, create the container again with the `docker run` command. We shall create the new “`mysqldb`” container differently. Specify different environment variables for the second run of the `docker run` command. Specify only the required environment variable `MYSQL_ROOT_PASSWORD` and set its value to “`mysql`”.

```
sudo docker run --name mysqldb -e MYSQL_ROOT_PASSWORD=mysql -d mysql
```

Subsequently, start the interactive shell with the following command.

```
sudo docker exec -it 113458c31ce5 bash
```

Login to the MySQL CLI with the following command in the interactive shell.

```
mysql -u root -p mysql
```

Specify the password for the “root” user, which is `mysql`. MySQL CLI gets started as shown in Figure 4-15.

The screenshot shows a terminal window titled "ec2-user@ip-172-30-1-192:~". The window contains the following text:

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED          STATUS              PORTS
113458c31ce5        mysql              "/entrypoint.sh mysql"   30 seconds ago
Up 28 seconds        mysql              "mysqld"           mysqldb
[ec2-user@ip-172-30-1-192 ~]$ sudo docker exec -it 113458c31ce5 bash
root@113458c31ce5:/# mysql -u root -p mysql
Enter password:
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 1
Server version: 5.6.27 MySQL Community Server (GPL)

Copyright (c) 2000, 2015, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> ]
```

Figure 4-15. Using a Password to Start MySQL CLI

The `mysql` command may also be issued as follows.

```
mysql -u root -p
```

Specify the password for the “mysql” user. MySQL CLI gets started as shown in Figure 4-16.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker exec -it 113458c31ce5 bash
root@113458c31ce5:/# mysql -u root -p
Enter password:
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 3
Server version: 5.6.27 MySQL Community Server (GPL)

Copyright (c) 2000, 2015, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> █
```

Figure 4-16. Alternative mysql Login command

The following mysql command does not start a MySQL CLI.

```
root@113458c31ce5:/# mysql -u root
```

The following error is generated.

```
ERROR 1045 (28000): Access denied for user 'root'@'localhost' (using password: NO)
```

List the databases with the show databases command. The default databases include the “mysql” database as shown in Figure 4-17. Previously, the “mysqldb” database also got listed with the show databases command because the “mysqldb” database was created when the docker run command was run.

```
ec2-user@ip-172-30-1-192:~  
File Edit View Search Terminal Help  
Enter password:  
Welcome to the MySQL monitor. Commands end with ; or \g.  
Your MySQL connection id is 3  
Server version: 5.6.27 MySQL Community Server (GPL)  
  
Copyright (c) 2000, 2015, Oracle and/or its affiliates. All rights reserved.  
  
Oracle is a registered trademark of Oracle Corporation and/or its  
affiliates. Other names may be trademarks of their respective  
owners.  
  
Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.  
  
mysql> show databases;  
+-----+  
| Database      |  
+-----+  
| information_schema |  
| mysql          |  
| performance_schema |  
+-----+  
3 rows in set (0.00 sec)  
  
mysql> █
```

Figure 4-17. Listing the Default Databases

Set the database as the “mysql” database with the “use mysql” command as shown in Figure 4-18.

```
mysql> use mysql;  
Reading table information for completion of table and column names  
You can turn off this feature to get a quicker startup with -A  
  
Database changed  
mysql> █
```

Figure 4-18. Using the mysql Database

List the database tables in the mysql database with the `show tables` command as shown in Figure 4-19.

```
mysql> show tables;
+-----+
| Tables_in_mysql |
+-----+
| columns_priv
| db
| event
| func
| general_log
| help_category
| help_keyword
| help_relation
| help_topic
| innodb_index_stats
| innodb_table_stats
| ndb_binlog_index
| plugin
| proc
| procs_priv
| proxies_priv
| servers
|
```

Figure 4-19. Listing Tables

Listing Docker Container Logs

Next, list the logs for the `mysqldb` container with the `docker logs` command.

```
sudo docker logs -f mysqldb
```

The logs for the `mysqldb` container get listed as shown in Figure 4-20.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker logs -f mysqlDb
Running mysql_install_db
2015-10-17 16:30:35 0 [Note] /usr/sbin/mysqld (mysqld 5.6.27) starting as process 15 ...
2015-10-17 16:30:35 15 [Note] InnoDB: Using atomics to ref count buffer pool pages
2015-10-17 16:30:35 15 [Note] InnoDB: The InnoDB memory heap is disabled
2015-10-17 16:30:35 15 [Note] InnoDB: Mutexes and rw_locks use GCC atomic builtins
2015-10-17 16:30:35 15 [Note] InnoDB: Memory barrier is not used
2015-10-17 16:30:35 15 [Note] InnoDB: Compressed tables use zlib 1.2.8
2015-10-17 16:30:35 15 [Note] InnoDB: Using Linux native AIO
2015-10-17 16:30:35 15 [Note] InnoDB: Using CPU crc32 instructions
2015-10-17 16:30:35 15 [Note] InnoDB: Initializing buffer pool, size = 128.0M
2015-10-17 16:30:35 15 [Note] InnoDB: Completed initialization of buffer pool
2015-10-17 16:30:35 15 [Note] InnoDB: The first specified data file ./ibdata1 did not exist: a new database to be created!
2015-10-17 16:30:35 15 [Note] InnoDB: Setting file ./ibdata1 size to 12 MB
2015-10-17 16:30:35 15 [Note] InnoDB: Database physically writes the file full: wait...
2015-10-17 16:30:35 15 [Note] InnoDB: Setting log file ./ib_logfile101 size to 48 MB
```

Figure 4-20. Listing Docker Container Log

Summary

In this chapter we used a Docker image to run MySQL Server in a Docker container. We ran two different variations of the `docker run` command; one included all the environment variables that may be set for the “mysql” image and the other included only the required environment variable/s. In the next chapter we shall discuss running MongoDB on Docker.

CHAPTER 5



Using MongoDB

MongoDB is the most commonly used NoSQL database. MongoDB is based on the Document store data model and stores data as BSON (Binary JSON) documents. MongoDB provides a flexible schema-less storage format in which different records could have different fields, implying that no fixed data structure is applied. Field values have no data types associated with them and different fields could be of different data types. With the JSON format, hierarchies of data structures become feasible, and a field could store multiple values using an array. In this chapter we shall use a Docker image to run MongoDB in a Docker container. This chapter has the following sections.

- Setting the Environment
- Starting MongoDB
- Starting an Interactive Terminal
- Starting a Mongo Shell
- Creating a Database
- Creating a Collection
- Creating a Document
- Finding Documents
- Adding Another Document
- Dropping a Collection
- Adding a Batch of Documents
- Updating a Document
- Querying a Single Document
- Querying All the Documents
- Making a Backup of the Data
- Stopping and Restarting the MongoDB Database
- Removing Documents
- Exiting Mongo Shell

Setting the Environment

The following software is required for this chapter:

- Docker Engine (version 1.8)

- Docker image for MongoDB

We have used an Amazon EC2 instance (Amazon Machine Image Red Hat Enterprise Linux 7.1 (HVM), SSD Volume Type - ami-12663b7a) to install the Docker image and run MongoDB in a Docker container. SSH login to the Amazon EC2 instance.

```
ssh -i "docker.pem" ec2-user@54.174.254.96
```

Start the Docker service.

```
sudo service docker start
```

Verify the Docker service status.

```
sudo service docker status
```

Docker service should be **active (running)** as shown in Figure 5-1.

```
[ec2-user@ip-172-30-1-61 ~]$ sudo service docker start
Starting docker (via systemctl): [ OK ]
[ec2-user@ip-172-30-1-61 ~]$ sudo service docker status
docker.service - Docker Application Container Engine
  Loaded: loaded (/usr/lib/systemd/system/docker.service; disabled)
  Active: active (running) since Thu 2015-10-15 13:14:55 EDT; 50s ago
    Docs: https://docs.docker.com
 Main PID: 2149 (docker)
   CGroup: /system.slice/docker.service
           └─2149 /usr/bin/docker daemon -H fd://

Oct 15 13:14:55 ip-172-30-1-61.ec2.internal docker[2149]: time="2015-10-15T13...
Oct 15 13:14:55 ip-172-30-1-61.ec2.internal systemd[1]: Started Docker Appli...
Oct 15 13:14:55 ip-172-30-1-61.ec2.internal docker[2149]: time="2015-10-15T13...
Hint: Some lines were ellipsized, use -l to show in full.
[ec2-user@ip-172-30-1-61 ~]$ █
```

Figure 5-1. Starting Docker Service and verifying Status

Download the official Docker image for MongoDB database.

```
sudo docker pull mongo:latest
```

List the Docker images.

```
sudo docker images
```

The Docker image called “mongo” gets listed as shown in Figure 5-2.

The screenshot shows a terminal window titled "ec2-user@ip-172-30-1-61:~". The user runs the command "sudo docker pull mongo:latest", which triggers a series of "Pull complete" messages for various MongoDB image versions. After the pull is finished, the user runs "docker images" to list the available images. The output shows one image entry for "mongo" with tag "latest", having an image ID of "910678a338ed" and being created 20 hours ago. The terminal window has a standard Linux-style interface with a menu bar and a scroll bar on the right.

```
[ec2-user@ip-172-30-1-61 ~]$ sudo cd /mongodb
[ec2-user@ip-172-30-1-61 ~]$ docker pull mongo:latest
latest: Pulling from library/mongo
1e20f9afa674: Pull complete
3ce45fcdfa8: Pull complete
44ae943f170a: Pull complete
eb4ffde66e1e: Pull complete
db6b9bfcc218: Pull complete
40e1ba7c3bb5: Pull complete
68d8ab161b83: Pull complete
4f69e92307b3: Pull complete
dbef49e93b9d: Pull complete
b1ae3e437bf0: Pull complete
711a5d59ecf5: Pull complete
664e23ecb316: Pull complete
2de189aa7810: Pull complete
c43884a3c774: Pull complete
a447e9d02e75: Pull complete
26abe62ffb50: Pull complete
910678a338ed: Pull complete
Digest: sha256:703623552c5477d9eb7427ca4e4db65609618f8fc17841b814c547df6b113502
Status: Downloaded newer image for mongo:latest
[ec2-user@ip-172-30-1-61 ~]$ docker images
REPOSITORY      TAG          IMAGE ID       CREATED
VIRTUAL SIZE
mongo           latest        910678a338ed   20 hours ago
261.6 MB
[ec2-user@ip-172-30-1-61 ~]$
```

Figure 5-2. Downloading Docker Image mongo

Starting MongoDB


Next, start MongoDB in a Docker container. MongoDB stores data in the /data/db directory in the Docker container by default. A directory could be mounted from the underlying host system to the container running the MongoDB database. For example, create a directory /data on the host.

```
sudo mkdir -p /data
```

Start the Docker container using the `docker run` command on the `mongo` image with the `/data` directory in the container mounted as `/data` directory on the host. Specify container name as “`mongodb`”.

```
sudo docker run -t -i -v /data:/data --name mongodb -d mongo
```

The Docker container, and the MongoDB server in the container, gets started as shown in Figure 5-3.




```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker run -t -i -v /data:/data --name mongodb -d mongo
68fe88ca79fe5c8606a8576e60aadbec5218ea314dcec49888316d11ef81c4ca
[ec2-user@ip-172-30-1-192 ~]$ █
```

Figure 5-3. Starting Docker Container for MongoDB

List the running Docker containers.

```
sudo docker ps
```

The `mongodb` container gets listed as running on port 27017 as shown in Figure 5-4.



CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
68fe88ca79fe	mongo	"/entrypoint.sh mongo"	37 seconds ago
Up 36 seconds	27017/tcp	mongodb	

Figure 5-4. Listing Docker Container for MongoDB

The MongoDB port could also be specified explicitly using the `-p` option.

```
docker run -t -i -v /data:/data -p 27017:27017 --name mongodb -d mongo
```

The container logs may be listed using the `docker logs` command.

```
sudo docker logs mongodb
```

Starting an Interactive Terminal

Start an interactive terminal (`tty`) using the following command.

```
sudo docker exec -it mongodb bash
```

Starting a Mongo Shell

To start the MongoDB shell, run the following command.

```
mongo
```

The MongoDB shell gets started and the > prompt gets displayed as shown in Figure 5-5.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker exec -it mongodb bash
root@68fe88ca79fe:/# mongo
MongoDB shell version: 3.0.7
connecting to: test
Welcome to the MongoDB shell.
For interactive help, type "help".
For more comprehensive documentation, see
    http://docs.mongodb.org/
Questions? Try the support group
    http://groups.google.com/group/mongodb-user
Server has startup warnings:
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten]
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten] ** WARNING: /sys/kernel/mm/transparent_hugepage/enabled is 'always'.
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten] ** We suggest setting it to 'never'
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten]
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten] ** WARNING: /sys/kernel/mm/transparent_hugepage/defrag is 'always'.
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten] ** We suggest setting it to 'never'
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten]
> █
```

Figure 5-5. Starting MongoDB Shell from TTY

The MongoDB shell may also be started on a specific host and port as follows.

```
mongo -host localhost -port 27017
```

MongoDB shell gets started on host localhost, port 27017 as shown in Figure 5-6. The “test” database instance gets connected to.

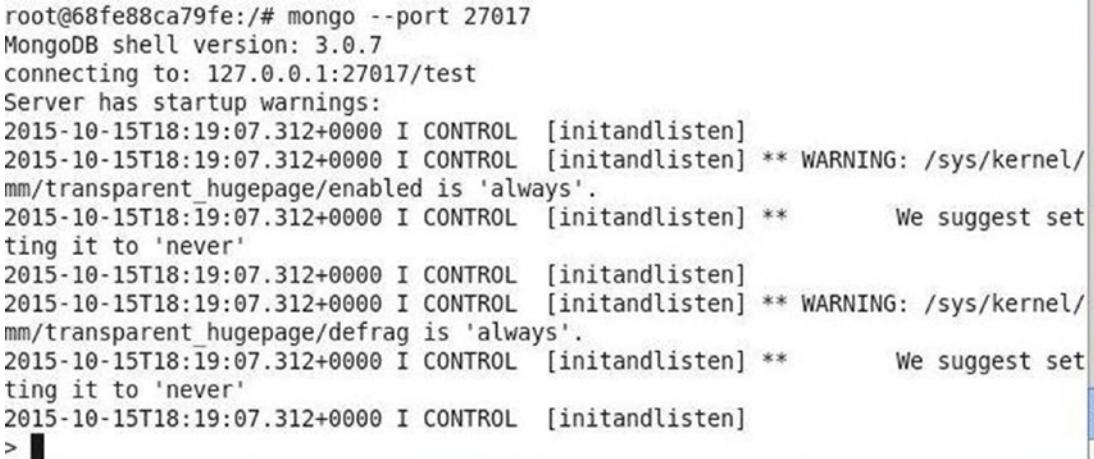
```
root@68fe88ca79fe:/# mongo --host localhost --port 27017
MongoDB shell version: 3.0.7
connecting to: localhost:27017/test
Server has startup warnings:
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten]
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten] ** WARNING: /sys/kernel/mm/transparent_hugepage/enabled is 'always'.
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten] ** We suggest setting it to 'never'
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten]
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten] ** WARNING: /sys/kernel/mm/transparent_hugepage/defrag is 'always'.
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten] ** We suggest setting it to 'never'
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten]
> █
```

Figure 5-6. Starting MongoDB Shell using Host and Port

Alternatively, only one of the host or the port may be specified to start the MongoDB shell.

```
mongo -port 27017
```

MongoDB shell gets started and gets connected to MongoDB server on 127.0.0.1:27071/test as shown in Figure 5-7.



```
root@68fe88ca79fe:/# mongo --port 27017
MongoDB shell version: 3.0.7
connecting to: 127.0.0.1:27017/test
Server has startup warnings:
2015-10-15T18:19:07.312+0000 I CONTROL [initandlisten]
2015-10-15T18:19:07.312+0000 I CONTROL [initandlisten] ** WARNING: /sys/kernel/mm/transparent_hugepage/enabled is 'always'.
2015-10-15T18:19:07.312+0000 I CONTROL [initandlisten] ** We suggest setting it to 'never'
2015-10-15T18:19:07.312+0000 I CONTROL [initandlisten]
2015-10-15T18:19:07.312+0000 I CONTROL [initandlisten] ** WARNING: /sys/kernel/mm/transparent_hugepage/defrag is 'always'.
2015-10-15T18:19:07.312+0000 I CONTROL [initandlisten] ** We suggest setting it to 'never'
2015-10-15T18:19:07.312+0000 I CONTROL [initandlisten]
> █
```

Figure 5-7. Starting MongoDB Shell using only the Port

Another form of specifying the host and port is host:port. For example, start the MongoDB shell and connect to localhost:27017 with the following command.

```
mongo localhost:27017
```

MongoDB Shell gets connected to localhost:27017/test database as shown in Figure 5-8.



```
root@68fe88ca79fe:/# mongo localhost:27017
MongoDB shell version: 3.0.7
connecting to: localhost:27017/test
Server has startup warnings:
2015-10-15T18:19:07.312+0000 I CONTROL [initandlisten]
2015-10-15T18:19:07.312+0000 I CONTROL [initandlisten] ** WARNING: /sys/kernel/mm/transparent_hugepage/enabled is 'always'.
2015-10-15T18:19:07.312+0000 I CONTROL [initandlisten] ** We suggest setting it to 'never'
2015-10-15T18:19:07.312+0000 I CONTROL [initandlisten]
2015-10-15T18:19:07.312+0000 I CONTROL [initandlisten] ** WARNING: /sys/kernel/mm/transparent_hugepage/defrag is 'always'.
2015-10-15T18:19:07.312+0000 I CONTROL [initandlisten] ** We suggest setting it to 'never'
2015-10-15T18:19:07.312+0000 I CONTROL [initandlisten]
> █
```

Figure 5-8. Starting MongoDB Shell using host:port Format

Creating a Database

List the databases from the MongoDB shell with the following command help method (also called command helper).

```
show dbs
```

A new database is created implicitly when the database name is set to the database to be created. For example, set the database to “mongodb” with the following command.

```
use mongodb
```

The `show dbs` command help method does not list the `mongodb` database till the database is used. Use the `db.createCollection()` method to create a collection called “catalog”. Subsequently, run the `show dbs` command again.

```
show dbs
db.createCollection("catalog")
show dbs
```

The `show dbs` command does not list the “`mongodb`” database before the “catalog” collection is created, but lists the “`mongodb`” database after the collection has been created as shown in Figure 5-9.



```
> show dbs
local 0.078GB
> use mongodb
switched to db mongodb
> show dbs
local 0.078GB
> db.createCollection("catalog")
{ "ok" : 1 }
> show dbs
local 0.078GB
mongodb 0.078GB
> █
```

Figure 5-9. Creating a Database

List the collections in the `mongodb` database with the following command.

```
show collections
```

The “catalog” collection gets listed in addition to the system collection `system.indexes` as shown in Figure 5-10.

```
> show dbs
local 0.078GB
mongodb 0.078GB
> show collections
catalog
system.indexes
> █
```



Figure 5-10. Listing Collections

Creating a Collection

In the previous section we created a collection called “catalog” using the `db.createCollection` command. Next, create a capped collection “catalog_capped” by setting the `capped` option field to true. A capped collection is a fixed size collection that keeps track of the insertion order while adding and getting a document, and as a result provides high throughput.

```
db.createCollection("catalog_capped", {capped: true, autoIndexId: true, size: 64 * 1024,
max: 1000} )
```

A capped collection called “catalog_capped” gets created as shown in Figure 5-11.

```
> show collections
catalog
system.indexes
> db.createCollection("catalog_capped", {capped: true, autoIndexId: true, size:
64 * 1024, max: 1000} )
{ "ok" : 1 }
> █
```



Figure 5-11. Creating a Capped Collection

A collection may also be created using the `db.runCommand` command. Create another capped collection called “catalog_capped_2” using the `db.runCommand` command.

```
db.runCommand( { create: "catalog_capped_2", capped: true, size: 64 * 1024, max: 1000 } )
```

Capped collection catalog_capped_2 gets created as shown in Figure 5-12.

```
> db.createCollection("catalog_capped", {capped: true, autoIndexId: true, size: 64 * 1024, max: 1000} )
{ "ok" : 1 }
> db.runCommand( { create: "catalog_capped_2", capped: true, size: 64 * 1024, max: 1000 } )
{ "ok" : 1 }
> show collections
catalog
catalog_capped
catalog_capped_2
system.indexes
> █
```

Figure 5-12. Creating a Capped Collection using `db.runCommand()`

Creating a Document

Next, we shall add documents to a MongoDB collection. Initially the `catalog` collection is empty. Run the mongo shell method `db.<collection>.count()` to count the documents in the `catalog` collection. Substitute `<collection>` with the collection name “`catalog`”.

```
db.catalog.count()
```

The number of documents in the `catalog` collection gets listed as 0 as shown in Figure 5-13.

```
> db.catalog.count()
0
```

Figure 5-13. Finding Document Count

Next, we shall add a document to the `catalog` collection. Create a JSON document structure with fields `catalogId`, `journal`, `publisher`, `edition`, `title` and `author`.

```
doc1 = {"catalogId": "catalog1", "journal": 'Oracle Magazine', "publisher": 'Oracle Publishing', "edition": 'November December 2013', "title": 'Engineering as a Service', "author": 'David A. Kelly'}
```

Add the document to the `catalog` collection using the `db.<collection>.insert()` method.

```
db.catalog.insert(doc1)
```

Subsequently output the document count again.

```
db.catalog.count()
```

The output from the `db.catalog.insert()` method, shown in Figure 5-14, is an object of type `WriteResult` with `nInserted` as 1, which implies that one document got added. The document count is listed as 1.

```
> doc1 = {"catalogId" : "catalog1", "journal" : 'Oracle Magazine', "publisher" : "Oracle Publishing", "edition" : 'November December 2013', "title" : 'Engineering as a Service', "author" : 'David A. Kelly'}
{
    "catalogId" : "catalog1",
    "journal" : "Oracle Magazine",
    "publisher" : "Oracle Publishing",
    "edition" : "November December 2013",
    "title" : "Engineering as a Service",
    "author" : "David A. Kelly"
}
>
> db.catalog.insert(doc1)
WriteResult({ "nInserted" : 1 })
> db.catalog.count()
1
> ■
```

Figure 5-14. Adding a Document

Finding Documents

The `db.collection.find(query, projection)` method is used to find document/s. The `query` parameter of type document specifies selection criteria using query operators. The `projection` parameter also of type document specifies the fields to return. Both the parameters are optional. To select all documents do not specify any args or specify an empty document `{}`. For example, find all documents in the `catalog` collection.

```
db.catalog.find()
```

The one document added previously gets listed as a JSON document as shown in Figure 5-15. The `_id` field is added to the documented automatically if not specified explicitly.

```
> db.catalog.find()
{ "_id" : ObjectId("561fefea7380a18f6587b0aa4"), "catalogId" : "catalog1", "journal" : "Oracle Magazine", "publisher" : "Oracle Publishing", "edition" : "November December 2013", "title" : "Engineering as a Service", "author" : "David A. Kelly" }
> ■
```

Figure 5-15. Running a Query using `find()` Method

Adding Another Document

Similarly, create the JSON structure for another document. The same document may be added again if the `_id` is unique. In the JSON include the `_id` field as an explicit field/attribute. The `_id` field value must be an object of type `ObjectId` and not a string literal.

```
doc2 = {"_id": ObjectId("507f191e810c19729de860ea"), "catalogId" : "catalog1", "journal" : 'Oracle Magazine', "publisher" : 'Oracle Publishing', "edition" : 'November December 2013', "title" : 'Engineering as a Service', "author" : 'David A. Kelly'};
```

Add the document using the db.<collection>.insert() method.

```
db.catalog.insert(doc2)
```

Another document gets added to the catalog collection as indicated by the nInserted value of 1 shown in Figure 5-16.

```
> doc2 = {"_id": ObjectId("507f191e810c19729de860ea"), "catalogId" : "catalog1",
  "journal" : 'Oracle Magazine', "publisher" : 'Oracle Publishing', "edition" : 'November December 2013', "title" : 'Engineering as a Service', "author" : 'David A. Kelly'};
{
  "_id" : ObjectId("507f191e810c19729de860ea"),
  "catalogId" : "catalog1",
  "journal" : "Oracle Magazine",
  "publisher" : "Oracle Publishing",
  "edition" : "November December 2013",
  "title" : "Engineering as a Service",
  "author" : "David A. Kelly"
}
>
> db.catalog.insert(doc2)
WriteResult({ "nInserted" : 1 })
> █
```

Figure 5-16. Adding Another Document

Subsequently query the catalog collection using db.<collection>.find() method.

```
db.catalog.find()
```

The two documents added to the catalog collection get listed as shown in Figure 5-17. The two documents have all the same name/value pairs in the JSON except the _id field, which has a unique value.

```
> db.catalog.find()
{ "_id" : ObjectId("561fefef7380a18f6587b0aa4"), "catalogId" : "catalog1", "journal" : "Oracle Magazine", "publisher" : "Oracle Publishing", "edition" : "November December 2013", "title" : "Engineering as a Service", "author" : "David A. Kelly" }
{ "_id" : ObjectId("507f191e810c19729de860ea"), "catalogId" : "catalog1", "journal" : "Oracle Magazine", "publisher" : "Oracle Publishing", "edition" : "November December 2013", "title" : "Engineering as a Service", "author" : "David A. Kelly" }
> █
```

Figure 5-17. Running the find() Method

Querying a Single Document

The `db.<collection>.findOne()` method is used to find a single document. Find a single document from the catalog collection.

```
db.catalog.findOne()
```

One of the documents gets output by the query as shown in Figure 5-18.

```
> db.catalog.findOne()
{
    "_id" : ObjectId("561ff033380a18f6587b0aa5"),
    "catalogId" : "catalog1",
    "journal" : "Oracle Magazine",
    "publisher" : "Oracle Publishing",
    "edition" : "November December 2013",
    "title" : "Engineering as a Service",
    "author" : "David A. Kelly"
}
```



Figure 5-18. Using the `findOne()` Method

The `db.collection.findOne(query, projection)` method also takes two args both of type document and both optional. The `query` parameter specifies the query selection criteria and the `projection` parameter specifies the fields to select. For example, select the `edition`, `title` and `author` fields and specify the query document as {}.

```
db.catalog.findOne(
    { },
    { edition: 1, title: 1, author: 1 }
)
```

The `edition`, `title` and `author` fields get listed. The `_id` field is always output by a query as shown in Figure 5-19.

```
> db.catalog.findOne(
...     { },
...     { edition: 1, title: 1, author: 1 }
... )
{
    "_id" : ObjectId("561ff033380a18f6587b0aa5"),
    "edition" : "November December 2013",
    "title" : "Engineering as a Service",
    "author" : "David A. Kelly"
}
```



Figure 5-19. Using a Query Projection

Dropping a Collection

The `db.collection.drop()` method drops or removes a collection. For example, remove the `catalog` collection.

```
db.catalog.drop()
```

Subsequently, the `show collections` method does not list the `catalog` collection as shown in Figure 5-20.

```
> db.catalog.drop()
true
> show collections
catalog_capped
catalog_capped_2
system.indexes
> █
```

Figure 5-20. Dropping a Collection

Adding a Batch of Documents

Previously, we added a single document at a time. Next, we shall add a batch of documents. Drop the `catalog` collection if not already dropped in the previous section.

```
db.catalog.drop()
```

Add an array of documents using the `db.catalog.insert()` method invocation with the `doc1` and `doc2` being the same as earlier. The `writeConcern` option specifies the guarantee MongoDB provides and a value of “majority” implies that the `insert()` method does not return till the write has been propagated to the majority of the nodes. Setting the `ordered` option to `true` adds the documents in the order specified.

```
db.catalog.insert([doc1, doc2], { writeConcern: { w: "majority", wtimeout: 5000 },
ordered:true })
```

The full syntax of the `insert` method is made use of in the preceding method invocation and is as follows.

```
db.collection.insert(
  <document or array of documents>,
  {
    writeConcern: <document>,
    ordered: <boolean>
  }
)
```

The first parameter is a single document or an array of documents. The second parameter is a document with fields `writeConcern` and `ordered`. The `writeConcern` specifies the write concern or the guarantee that MongoDB provides on the success of an insert. The `ordered` parameter is set to `true`, which implies that the documents are added in the order specified and if an error occurs with one of the documents none of the documents are added. The `nInserted` in the output is 2 for the two documents added as shown in Figure 5-21.

```
> db.catalog.insert([doc1, doc2], { writeConcern: { w: "majority", wtimeout: 5000 }, ordered:true })
BulkWriteResult({
    "writeErrors" : [ ],
    "writeConcernErrors" : [ ],
    "nInserted" : 2,
    "nUpserted" : 0,
    "nMatched" : 0,
    "nModified" : 0,
    "nRemoved" : 0,
    "upserted" : [ ]
})
> █
```

Figure 5-21. Adding a Batch of Documents

Run the `db.catalog.find()` method to query the documents in the `catalog` collection as shown in Figure 5-22.

```
> db.catalog.insert([doc1, doc2], { writeConcern: { w: "majority", wtimeout: 5000 }, ordered:true })
BulkWriteResult({
    "writeErrors" : [ ],
    "writeConcernErrors" : [ ],
    "nInserted" : 2,
    "nUpserted" : 0,
    "nMatched" : 0,
    "nModified" : 0,
    "nRemoved" : 0,
    "upserted" : [ ]
})
> db.catalog.find()
{ "_id" : ObjectId("561ff033380a18f6587b0aa5"), "catalogId" : "catalog1", "journal" : "Oracle Magazine", "publisher" : "Oracle Publishing", "edition" : "November December 2013", "title" : "Engineering as a Service", "author" : "David A. Kelly" }
{ "_id" : ObjectId("507f191e810c19729de860ea"), "catalogId" : "catalog1", "journal" : "Oracle Magazine", "publisher" : "Oracle Publishing", "edition" : "November December 2013", "title" : "Engineering as a Service", "author" : "David A. Kelly" }
> █
```

Figure 5-22. Running the `find()` Method to list Documents added in a Batch

Updating a Document

The `db.collection.save()` method has the following syntax and updates a document if the document already exists, and adds a new document if the document does not exist.

```
db.collection.save(
  <document>,
  {
    writeConcern: <document>
  }
)
```

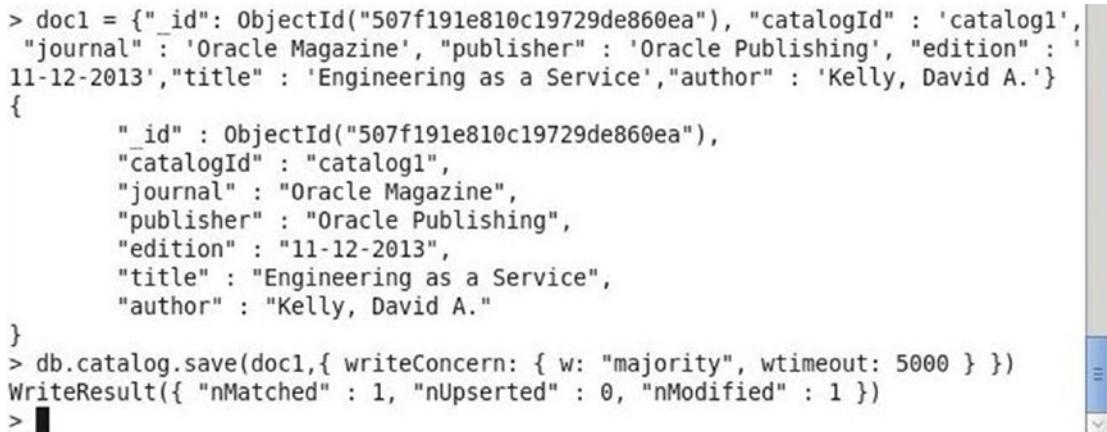
A document is identified by the unique `_id` of type `ObjectId`. Next, we shall update document with `_id` as `ObjectId("507f191e810c19729de860ea")`. Create an updated JSON document with some of the field values modified.

```
doc1 = {"_id": ObjectId("507f191e810c19729de860ea"), "catalogId" : 'catalog1', "journal" : 'Oracle Magazine', "publisher" : 'Oracle Publishing', "edition" : '11-12-2013', "title" : 'Engineering as a Service', "author" : 'Kelly, David A.'}
```

Save the document using the `db.collection.save()` method in the `catalog` collection.

```
db.catalog.save(doc1, { writeConcern: { w: "majority", wtimeout: 5000 } })
```

The document gets saved by updating an existing document. The `nMatched` is 1 and `nUpserted` is 0, and `nModified` is 1 in the `WriteResult` object returned as shown in Figure 5-23. The `nUpserted` field refers to the number of new documents added in contrast to modifying an existing document.



```
> doc1 = {"_id": ObjectId("507f191e810c19729de860ea"), "catalogId" : 'catalog1', "journal" : 'Oracle Magazine', "publisher" : 'Oracle Publishing', "edition" : '11-12-2013', "title" : 'Engineering as a Service', "author" : 'Kelly, David A.'}
{
  "_id" : ObjectId("507f191e810c19729de860ea"),
  "catalogId" : "catalog1",
  "journal" : "Oracle Magazine",
  "publisher" : "Oracle Publishing",
  "edition" : "11-12-2013",
  "title" : "Engineering as a Service",
  "author" : "Kelly, David A."
}
> db.catalog.save(doc1, { writeConcern: { w: "majority", wtimeout: 5000 } })
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
> █
```

Figure 5-23. Using the `save()` Method to Update a Document

Query the `catalog` collection using the `find()` method.

```
db.catalog.find()
```

The updated document gets listed as one of the documents as shown in Figure 5-24.

```
> db.catalog.save(doc1, { writeConcern: { w: "majority", wtimeout: 5000 } })
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
> db.catalog.find()
{ "_id" : ObjectId("561ff033380a18f6587b0aa5"), "catalogId" : "catalog1", "journal" : "Oracle Magazine", "publisher" : "Oracle Publishing", "edition" : "November December 2013", "title" : "Engineering as a Service", "author" : "David A. Kelly" }
{ "_id" : ObjectId("507f191e810c19729de860ea"), "catalogId" : "catalog1", "journal" : "Oracle Magazine", "publisher" : "Oracle Publishing", "edition" : "11-12-2013", "title" : "Engineering as a Service", "author" : "Kelly, David A." }
> █
```

Figure 5-24. Querying Updated Document

Outputting Documents as JSON

The `db.collection.find(query, projection)` method returns a cursor over the documents that are selected by the query. Invoke the `forEach(printjson)` method on the cursor to output the documents as formatted JSON.

```
db.catalog.find().forEach(printjson)
```

The documents get output as JSON as shown in Figure 5-25.

```
> db.catalog.find().forEach(printjson)
{
  "_id" : ObjectId("561ff033380a18f6587b0aa5"),
  "catalogId" : "catalog1",
  "journal" : "Oracle Magazine",
  "publisher" : "Oracle Publishing",
  "edition" : "November December 2013",
  "title" : "Engineering as a Service",
  "author" : "David A. Kelly"
}
{
  "_id" : ObjectId("507f191e810c19729de860ea"),
  "catalogId" : "catalog1",
  "journal" : "Oracle Magazine",
  "publisher" : "Oracle Publishing",
  "edition" : "11-12-2013",
  "title" : "Engineering as a Service",
  "author" : "Kelly, David A."
}
> █
```

Figure 5-25. Outputting JSON

Making a Backup of the Data

The mongodump utility is used for creating a binary export of the data in a database. The mongorestore utility is used in conjunction with mongodump to restore a database from backup. The mongorestore utility either creates a new database instance or adds to an existing database.

Run the following mongodump command to export the test database to the /data/backup directory.

```
mongodump --db test --out /data/backup
```

The test database gets exported to the /data/backup directory as shown in Figure 5-26.

```
root@68fe88ca79fe:/# mongodump --db test --out /data/backup
root@68fe88ca79fe:/#
```

Figure 5-26. Exporting the test Database

List the directories in the /data/backup directory. The test database directory gets listed as shown in Figure 5-27.

```
[ec2-user@ip-172-30-1-192 ~]$ ls -la /data/backup
total 0
drwxr-xr-x. 3 root root 17 Oct 15 14:41 .
drwxr-xr-x. 4 root root 28 Oct 15 14:41 ..
drwxr-xr-x. 2 root root 6 Oct 15 14:41 test
[ec2-user@ip-172-30-1-192 ~]$
```

Figure 5-27. Listing the test Database

Run the following mongorestore command to restore the exported data from /data/backup/test to the testrestore database.

```
mongorestore --db testrestore /data/backup/test
```

The /data/backup/test directory data gets restored in the testrestore database as shown in Figure 5-28.

```
root@68fe88ca79fe:/# mongodump --db test --out /data/backup
root@68fe88ca79fe:/# mongorestore --db testrestore /data/backup/test
2015-10-15T18:43:18.829+0000      building a list of collections to restore from /
/data/backup/test dir
2015-10-15T18:43:18.829+0000      done
root@68fe88ca79fe:/#
```

Figure 5-28. Restoring a Database

Connect to the MongoDB shell with the following command.

```
mongo localhost:27017/testrestore
```

The MongoDB shell gets started as shown in Figure 5-29.

```
root@68fe88ca79fe:/# mongo localhost:27017/testrestore
MongoDB shell version: 3.0.7
connecting to: localhost:27017/testrestore
Server has startup warnings:
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten]
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten] ** WARNING: /sys/kernel/mm/transparent_hugepage/enabled is 'always'.
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten] **           We suggest setting it to 'never'
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten]
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten] ** WARNING: /sys/kernel/mm/transparent_hugepage/defrag is 'always'.
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten] **           We suggest setting it to 'never'
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten]
> █
```

Figure 5-29. Connecting to the Restored Database

List the databases with the following command.

```
show dbs
```

As we restored the backup to the testrestore database the `mongodb` database, which was previously exported, gets listed as shown in Figure 5-30.

```
root@68fe88ca79fe:/# mongo localhost:27017/testrestore
MongoDB shell version: 3.0.7
connecting to: localhost:27017/testrestore
Server has startup warnings:
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten]
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten] ** WARNING: /sys/kernel/mm/transparent_hugepage/enabled is 'always'.
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten] **           We suggest setting it to 'never'
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten]
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten] ** WARNING: /sys/kernel/mm/transparent_hugepage/defrag is 'always'.
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten] **           We suggest setting it to 'never'
2015-10-15T18:19:07.312+0000 I CONTROL  [initandlisten]
> show dbs
local      0.078GB
mongodb    0.078GB
> █
```

Figure 5-30. Listing the Restored Database

Set the database name as `mongodb`.

```
use mongodb
```

List the collections.

```
show collections
```

Query the documents in the `catalog` collection.

```
db.catalog.find()
```

Output from the preceding commands is shown in Figure 5-31.



```
> show dbs
local      0.078GB
mongodb    0.078GB
> use mongodb
switched to db mongodb
> show collections
catalog
catalog_capped
catalog_capped_2
system.indexes
> db.catalog.find()
{ "_id" : ObjectId("561ff033380a18f6587b0aa5"), "catalogId" : "catalog1", "journal" : "Oracle Magazine", "publisher" : "Oracle Publishing", "edition" : "November December 2013", "title" : "Engineering as a Service", "author" : "David A. Kelly" }
{ "_id" : ObjectId("507f191e810c19729de860ea"), "catalogId" : "catalog1", "journal" : "Oracle Magazine", "publisher" : "Oracle Publishing", "edition" : "11-12-2013", "title" : "Engineering as a Service", "author" : "Kelly, David A." }
> █
```

Figure 5-31. Listing and Querying the Restored Collection

Removing Documents

The `db.collection.remove` method is used to remove document/s and has the following syntax.

```
db.collection.remove(
  <query>,
  <justOne>
)
```

For example, remove the document with `ObjectId("561ff033380a18f6587b0aa5")`.

```
db.catalog.remove({ _id: ObjectId("561ff033380a18f6587b0aa5") })
```

The nRemoved in the WriteResult is 1 indicating that one document got removed. Run the db.catalog.find() method before and after the db.catalog.remove() method invocation. Before the db.catalog.remove() method is invoked, two documents get listed, and afterward only one document gets listed as shown in Figure 5-32.

```
> db.catalog.find()
{ "_id" : ObjectId("561ff033380a18f6587b0aa5"), "catalogId" : "catalog1", "journal" : "Oracle Magazine", "publisher" : "Oracle Publishing", "edition" : "November December 2013", "title" : "Engineering as a Service", "author" : "David A. Kelly" }
{ "_id" : ObjectId("507f191e810c19729de860ea"), "catalogId" : "catalog1", "journal" : "Oracle Magazine", "publisher" : "Oracle Publishing", "edition" : "11-12-2013", "title" : "Engineering as a Service", "author" : "Kelly, David A." }
> db.catalog.remove({ _id: ObjectId("561ff033380a18f6587b0aa5") })
WriteResult({ "nRemoved" : 1 })
> db.catalog.find()
{ "_id" : ObjectId("507f191e810c19729de860ea"), "catalogId" : "catalog1", "journal" : "Oracle Magazine", "publisher" : "Oracle Publishing", "edition" : "11-12-2013", "title" : "Engineering as a Service", "author" : "Kelly, David A." }
> █
```

Figure 5-32. Removing a Single Document

To remove all documents, provide an empty document {} to the db.catalog.remove() method invocation.

```
db.catalog.remove({})
```

Multiple documents get removed as indicated by nRemoved value of 2 as shown in Figure 5-33.

```

File Edit View Search Terminal Help
"edition" : "November December 2013",
"title" : "Engineering as a Service",
"author" : "David A. Kelly"
}
>
> db.catalog.insert(doc1)
WriteResult({ "nInserted" : 1 })
> doc2 = {"_id": ObjectId("507f191e810c19729de860ea"), "catalogId" : "catalog1",
"journal" : 'Oracle Magazine', "publisher" : 'Oracle Publishing', "edition" : 'November December 2013', "title" : 'Engineering as a Service', "author" : 'David A. Kelly'};
{
    "_id" : ObjectId("507f191e810c19729de860ea"),
    "catalogId" : "catalog1",
    "journal" : "Oracle Magazine",
    "publisher" : "Oracle Publishing",
    "edition" : "November December 2013",
    "title" : "Engineering as a Service",
    "author" : "David A. Kelly"
}
>
> db.catalog.insert(doc2)
WriteResult({ "nInserted" : 1 })
> db.catalog.remove({})
WriteResult({ "nRemoved" : 2 })
> █

```

Figure 5-33. Removing All Documents

An empty query document must be supplied to the `db.catalog.remove()` method invocation. If an empty document `{}` is not supplied, an error is generated indicating that a query is needed as shown in Figure 5-34.

```

> db.catalog.remove()
2015-10-15T19:06:49.774+0000 E QUERY      Error: remove needs a query
at Error (<anonymous>)
at DBCollection._parseRemove (src/mongo/shell/collection.js:305:32)
at DBCollection.remove (src/mongo/shell/collection.js:328:23)
at (shell):1:12 at src/mongo/shell/collection.js:305
> db.catalog.remove({})
WriteResult({ "nRemoved" : 1 })
> █

```

Figure 5-34. An empty document must be provided to the `remove()` method to remove all documents

Stopping and Restarting the MongoDB Database

The Docker container running the MongoDB instance may be stopped with the `docker stop` command.

```
sudo docker stop mongo
```

List the running Docker containers with the following command.

```
sudo docker ps
```

Start the Docker container again with the `docker start` command.

```
sudo docker start mongo
```

Run the following command again to list the running containers.

```
sudo docker ps
```

The output from the preceding commands is shown in Figure 5-35. The Docker container `mongodb` is again listed as running.



```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker stop mongodb
mongodb
[ec2-user@ip-172-30-1-192 ~]$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED
STATUS              PORTS              NAMES
[ec2-user@ip-172-30-1-192 ~]$ sudo docker start mongodb
mongodb
[ec2-user@ip-172-30-1-192 ~]$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED
STATUS              PORTS              NAMES
68fe88ca79fe        mongo              "/entrypoint.sh mongo"   51 minutes ago
Up 5 seconds         27017/tcp          mongodb
[ec2-user@ip-172-30-1-192 ~]$ █
```

Figure 5-35. Listing a Docker Container after Restarting the Container

Start the interactive terminal with the following command in which the container ID is used instead of the container name.

```
sudo docker exec -it 68fe88ca79fe bash
```

Start the MongoDB shell with the `mongo` command in the interactive shell as shown in Figure 5-36.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND                  CREATED             STATUS              NAMES
68fe88ca79fe        mongo              "/entrypoint.sh mongo"   51 minutes ago   Up 5 seconds      mongodb
[ec2-user@ip-172-30-1-192 ~]$ sudo docker exec -it 68fe88ca79fe bash
root@68fe88ca79fe:/# mongo
MongoDB shell version: 3.0.7
connecting to: test
Server has startup warnings:
2015-10-15T19:10:31.640+0000 I CONTROL  [initandlisten]
2015-10-15T19:10:31.640+0000 I CONTROL  [initandlisten] ** WARNING: /sys/kernel/mm/transparent_hugepage/enabled is 'always'.
2015-10-15T19:10:31.640+0000 I CONTROL  [initandlisten] **           We suggest setting it to 'never'
2015-10-15T19:10:31.640+0000 I CONTROL  [initandlisten]
2015-10-15T19:10:31.640+0000 I CONTROL  [initandlisten] ** WARNING: /sys/kernel/mm/transparent_hugepage/defrag is 'always'.
2015-10-15T19:10:31.640+0000 I CONTROL  [initandlisten] **           We suggest setting it to 'never'
2015-10-15T19:10:31.640+0000 I CONTROL  [initandlisten]
> █
```

Figure 5-36. Starting the MongoDB Shell

Set the database to local and list the collections with the `show collections` command. Subsequently set the database to `mongodb` and list the collections. The `db.catalog.find()` method does not list any documents as shown in Figure 5-37.

```
> use local
switched to db local
> show collections
startup_log
system.indexes
> use mongodb
switched to db mongodb
> show collections
catalog
catalog_capped
catalog_capped_2
system.indexes
> db.catalog.find()
> █
```

Figure 5-37. Listing Documents in the catalog Collection in local Database

Exiting the Mongo Shell

To exit the interactive terminal use the “exit” command and exit the MongoDB shell with the “exit” command also as shown in Figure 5-38.

```
> db.catalog.insert(doc2)
WriteResult({ "nInserted" : 1 })
> db.catalog.remove({})
WriteResult({ "nRemoved" : 2 })
> ^C
bye
root@68fe88ca79fe:/# ^C
root@68fe88ca79fe:/# exit
exit
[ec2-user@ip-172-30-1-192 ~]$
```

Figure 5-38. Exiting MongoDB Shell and TTY

Summary

In this chapter we used a Docker image for MongoDB to run a MongoDB instance in a Docker container. We created a database, added collections to the database, and added documents to the collections. We also queried the documents in MongoDB. We demonstrated stopping and starting the Docker container. We also made a backup of a MongoDB database and subsequently restored the database from the backup. In the next chapter we shall discuss running another NoSQL database, Apache Cassandra, in a Docker container.

CHAPTER 6



Using Apache Cassandra

Apache Cassandra is a wide-column, open source NoSQL database and the most commonly used NoSQL database in its category. The container of data, equivalent to a database schema in a relational database, in Apache Cassandra is a *Keyspace*. The basic unit of storage is a *column family* (also called *table*), and each record in a table is stored in a row with the data being stored in columns. A column has a name, a value, and a timestamp associated with it. A column is not required to store a value and the column could be empty. Apache Cassandra is based on a flexible schema (or schema-free or dynamic schema) data model in which different rows could have different columns and the columns are not required to be pre-specified in a table definition. Apache Cassandra supports data types for column names (called *comparators*) and column values (called *validators*), but does not require the data types (validators and comparators) to be specified. The validators and comparators may be added or modified after a table (column family) has been defined. Apache Cassandra provides a Cassandra Query Language (CQL) for CRUD (add, get, update, delete) operations on a table. Apache Cassandra installation includes a `cqlsh` utility, which is an interactive shell, from which CQL commands may be run. An official Docker image for Apache Cassandra is available and in this chapter we shall run Apache Cassandra in a Docker container.

Setting the Environment

Starting Apache Cassandra

Starting the TTY

Connecting to CQL Shell

Creating a Keyspace

Altering A Keyspace

Using A Keyspace

Creating a Table

Adding Table Data

Querying a Table

Deleting from a Table

Truncating a Table

Dropping A Table

Dropping a Keyspace

- Exiting CQLSh
- Stopping Apache Cassandra
- Starting Multiple Instances of Apache Cassandra

Setting the Environment

The following software is required for this chapter.

- Docker (version 1.8)
- Docker image for Apache Cassandra

We have used an Amazon EC2 AMI as in other chapters to install Docker and the Docker image. First, SSH to the Amazon EC2 instance.

```
ssh -i "docker.pem" ec2-user@54.86.243.122
```

Installing Docker is discussed in Chapter 1. Start the Docker service. The following command should output an OK message.

```
sudo service docker start
```

Verify that the Docker service has been started. The following command should output active (running) in the Active field.

```
sudo service docker status
```

Output from the preceding commands is shown in Figure 6-1.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo service docker start
Starting docker (via systemctl): [ OK ]
[ec2-user@ip-172-30-1-192 ~]$ sudo service docker status
docker.service - Docker Application Container Engine
  Loaded: loaded (/usr/lib/systemd/system/docker.service; disabled)
  Active: active (running) since Thu 2015-10-15 18:02:38 EDT; 9s ago
    Docs: https://docs.docker.com
 Main PID: 2405 (docker)
   CGroup: /system.slice/docker.service
           └─2405 /usr/bin/docker daemon -H fd://

Oct 15 18:02:38 ip-172-30-1-192.ec2.internal docker[2405]: time="2015-10-15T1...
Oct 15 18:02:38 ip-172-30-1-192.ec2.internal docker[2405]: ..
Oct 15 18:02:38 ip-172-30-1-192.ec2.internal docker[2405]: time="2015-10-15T1...
Oct 15 18:02:38 ip-172-30-1-192.ec2.internal systemd[1]: Started Docker Appli...
Hint: Some lines were ellipsized, use -l to show in full.
[ec2-user@ip-172-30-1-192 ~]$
```

Figure 6-1. Starting Docker Service and verifying Status

Next, download the latest cassandra Docker image.

```
sudo docker pull cassandra:latest
```

List the Docker images downloaded.

```
sudo docker images
```

The cassandra image should get listed as shown in Figure 6-2.

```
ec2-user@ip-172-30-1-192:~$ sudo docker pull cassandra:latest
latest: Pulling from library/cassandra
7a42f1433a16: Pull complete
3d88cbf54477: Pull complete
16053478214c: Pull complete
3bfd7a1b6b2a: Pull complete
c54b76ffbde3: Pull complete
5dc57e6b79ff: Pull complete
b44c9df2ed0a: Pull complete
8e3ab86c884c: Pull complete
a53f45dc6723: Pull complete
7a42e23533cd: Pull complete
ed1e959c4c76: Pull complete
81228f51eac7: Pull complete
b87e7f05a105: Pull complete
Digest: sha256:25dcb1262c4eb2146718491516bbeaf51f78febc37c59d0cdfc217a61db42e8
Status: Downloaded newer image for cassandra:latest
[ec2-user@ip-172-30-1-192 ~]$ sudo docker images
REPOSITORY          TAG           IMAGE ID            CREATED             VIRTUAL SIZE
mongo              latest        910678a338ed   25 hours ago      261.6 MB
cassandra          latest        b87e7f05a105   41 hours ago      362.6 MB
[ec2-user@ip-172-30-1-192 ~]$
```

Figure 6-2. Listing Docker Image *cassandra*

Starting Apache Cassandra

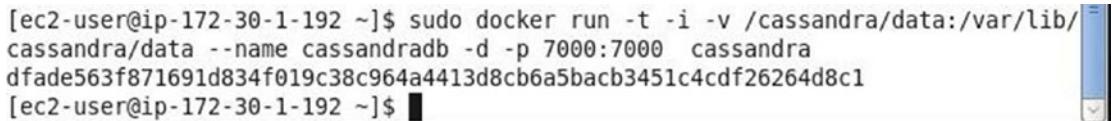
Start the Apache Cassandra server process in a Docker container with the following command in which the inter-node Apache Cassandra cluster communication port is specified as 7000 and the directory in which Apache Cassandra stores data is /cassandra/data. The container name is specified with the `-name` option as `cassandradb`. The syntax to start a Cassandra instance in detached mode is as follows.

```
docker run --name some-cassandra -d cassandra:tag
```

The `-d` parameter starts the container in a detached mode, implying that an interactive shell is not connected to with the docker run command even if the `-t -i` options are specified.

```
sudo docker run -t -i -v /cassandra/data:/var/lib/cassandra/data --name cassandradb -d -p 7000:7000 cassandra
```

A Docker container running an Apache Cassandra server process gets started as shown in Figure 6-3.



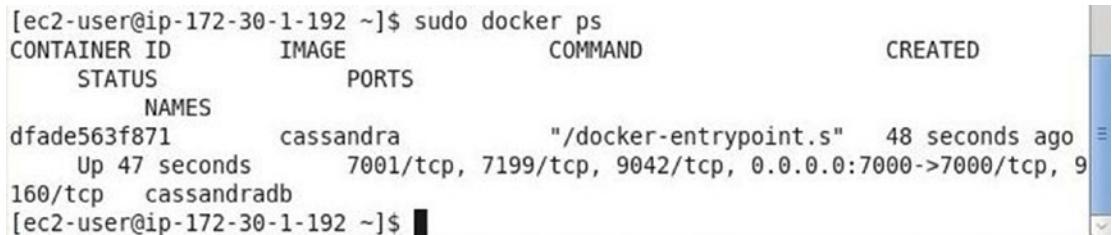
```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker run -t -i -v /cassandra/data:/var/lib/cassandra/data --name cassandradb -d -p 7000:7000 cassandra
dfade563f871691d834f019c38c964a4413d8cb6a5bacb3451c4cdf26264d8c1
[ec2-user@ip-172-30-1-192 ~]$ █
```

Figure 6-3. Starting Docker Container for Apache Cassandra

List the running Docker containers with the following command.

```
sudo docker ps
```

The `cassandradb` container, which is running an Apache Cassandra server instance, gets listed. The container id is also listed. By default, port 9042 is the client port on which Apache Cassandra listens for client connections. Port 9160 is Thrift API as shown in Figure 6-4.



```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND                  CREATED             STATUS              PORTS
NAMES
dfade563f871        cassandra          "/docker-entrypoint.s"   48 seconds ago      Up 47 seconds       7001/tcp, 7199/tcp, 9042/tcp, 0.0.0.0:7000->7000/tcp, 9160/tcp
   cassandradb
[ec2-user@ip-172-30-1-192 ~]$ █
```

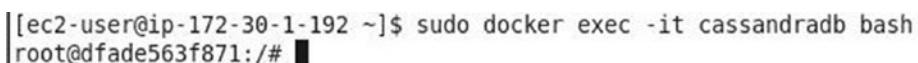
Figure 6-4. Listing Docker Containers that are Running

Starting the TTY

Start the interactive terminal (tty) with the following command.

```
sudo docker exec -it cassandradb bash
```

The tty gets connected to and the command prompt gets set to `user@containerid`. If the user is root and the container id is `dfade56f871`, the command prompt becomes `root@dfade56f871` as shown in Figure 6-5.



```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker exec -it cassandradb bash
root@dfade563f871:/# █
```

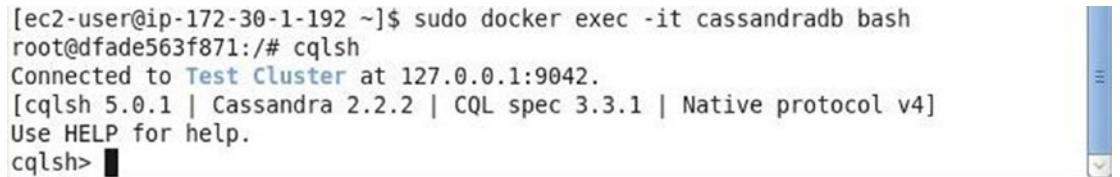
Figure 6-5. Starting the TTY

Connecting to CQL Shell

The cqlsh terminal is used to connect to an Apache Cassandra instance and run CQL commands. Start the cqlsh terminal with the following command.

```
cqlsh
```

A connection gets established to the Test Cluster at 127.0.0.1:9042. The Apache Cassandra version gets output as 2.2.2 and the CQL spec version as 3.3.1. The cqlsh> command prompt gets displayed as shown in Figure 6-6.



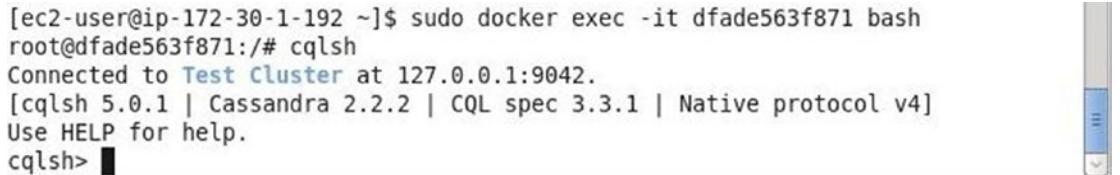
```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker exec -it cassandradb bash
root@dfade563f871:/# cqlsh
Connected to Test Cluster at 127.0.0.1:9042.
[cqlsh 5.0.1 | Cassandra 2.2.2 | CQL spec 3.3.1 | Native protocol v4]
Use HELP for help.
cqlsh> █
```

Figure 6-6. Connecting the CQL Shell

We started the interactive terminal using the container name, but the tty may also be started using the container id. The cqlsh shell is started with the cqlsh command regardless of how the tty is started.

```
sudo docker exec -it dfade563f871 bash
cqlsh
```

The cqlsh> command prompt gets displayed as before as shown in Figure 6-7.



```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker exec -it dfade563f871 bash
root@dfade563f871:/# cqlsh
Connected to Test Cluster at 127.0.0.1:9042.
[cqlsh 5.0.1 | Cassandra 2.2.2 | CQL spec 3.3.1 | Native protocol v4]
Use HELP for help.
cqlsh> █
```

Figure 6-7. Connecting to CQL Shell using the Container ID

Creating a Keyspace

A Keyspace is the container of application data and is used to group column families. Replication is set at a per-keyspace basis. The DDL command for creating a Keyspace is as follows.

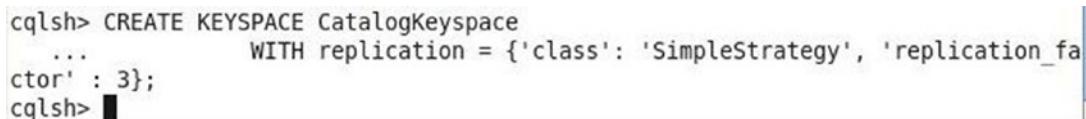
```
CREATE KEYSPACE (IF NOT EXISTS)? <identifier> WITH <properties>
```

By default, the keyspace name is case-insensitive and may consist exclusively of alpha-numeric characters with a maximum length of 32. To make a keyspace name case-sensitive add quotes. The supported properties by the CREATE KEYSPACE statement, which creates a top-level keyspace, are replication for

specifying the replication strategy and options and `durable_writes` for whether a commit log is to be used for updates on the keyspace, with the `replication` property being mandatory. As an example, create a keyspace called `CatalogKeyspace` with replication strategy class as `SimpleStrategy` and replication factor as 3.

```
CREATE KEYSPACE CatalogKeyspace
    WITH replication = {'class': 'SimpleStrategy', 'replication_factor' : 3};
```

The `CatalogKeyspace` keyspace gets created as shown in Figure 6-8.



```
cqlsh> CREATE KEYSPACE CatalogKeyspace
...           WITH replication = {'class': 'SimpleStrategy', 'replication_factor' : 3};
cqlsh> █
```

Figure 6-8. Creating a Keyspace

Altering A Keyspace

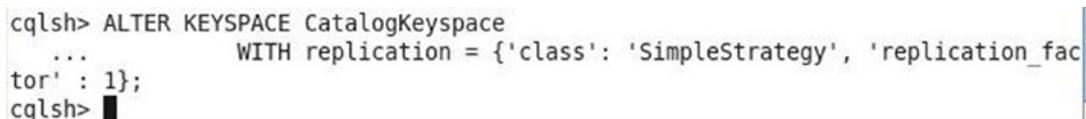
The `ALTER KEYSPACE` statement is used to alter a keyspace and has the following syntax with the supported properties being the same as for the `CREATE KEYSPACE` statement.

```
ALTER KEYSPACE <identifier> WITH <properties>
```

As an example, alter the `CatalogKeyspace` keyspace to make the replication factor 1.

```
ALTER KEYSPACE CatalogKeyspace
    WITH replication = {'class': 'SimpleStrategy', 'replication_factor' : 1};
```

The replication factor gets set to 1 as shown in Figure 6-9.



```
cqlsh> ALTER KEYSPACE CatalogKeyspace
...           WITH replication = {'class': 'SimpleStrategy', 'replication_factor' : 1};
cqlsh> █
```

Figure 6-9. Altering a Keyspace

Using A Keyspace

The `USE` statement is used to set the current keyspace and has the following syntax.

```
USE <identifier>
```

All subsequent commands are run in the context of the Keyspace set with the `USE` statement. As an example, set the current Keyspace as `CatalogKeyspace`.

```
use CatalogKeyspace;
```

The `cqlsh>` command prompt becomes `cqlsh:catalogkeyspace>` as shown in Figure 6-10.

```
cqlsh> use CatalogKeyspace;
cqlsh:catalogkeyspace>
```

Figure 6-10. Using a Keyspace

Creating a Table

A TABLE is also called a COLUMN FAMILY, and the CREATE TABLE or CREATE COLUMN FAMILY statement is used to create a table (column family).

```
CREATE ( TABLE | COLUMNFAMILY ) ( IF NOT EXISTS )? <tablename>
    (' <column-definition> ( ',' <column-definition> )* ')*
    ( WITH <option> ( AND <option>)* )?
```

For the complete syntax of the CREATE TABLE statement refer to <https://cassandra.apache.org/doc/cql3/CQL.html#createTableStmt>. As an example create a table called ‘catalog’ with columns catalog_id, journal, publisher, edition, title and author all of type text. Specify the primary key as catalog_id and set the compaction class as LeveledCompactionStrategy.

```
CREATE TABLE catalog(catalog_id text,journal text,publisher text,edition text,title
text,author text,PRIMARY KEY (catalog_id)) WITH compaction = { 'class' :
'LeveledCompactionStrategy' };
```

The catalog table gets created as shown in Figure 6-11.

```
cqlsh:catalogkeyspace> CREATE TABLE catalog(catalog_id text,journal text,publish
er text,edition text,title text,author text,PRIMARY KEY (catalog_id)) WITH com
paction = { 'class' : 'LeveledCompactionStrategy' };
```

Figure 6-11. Creating a Table

Adding Table Data

The INSERT DML statement is used to add data into a table and has the following syntax.

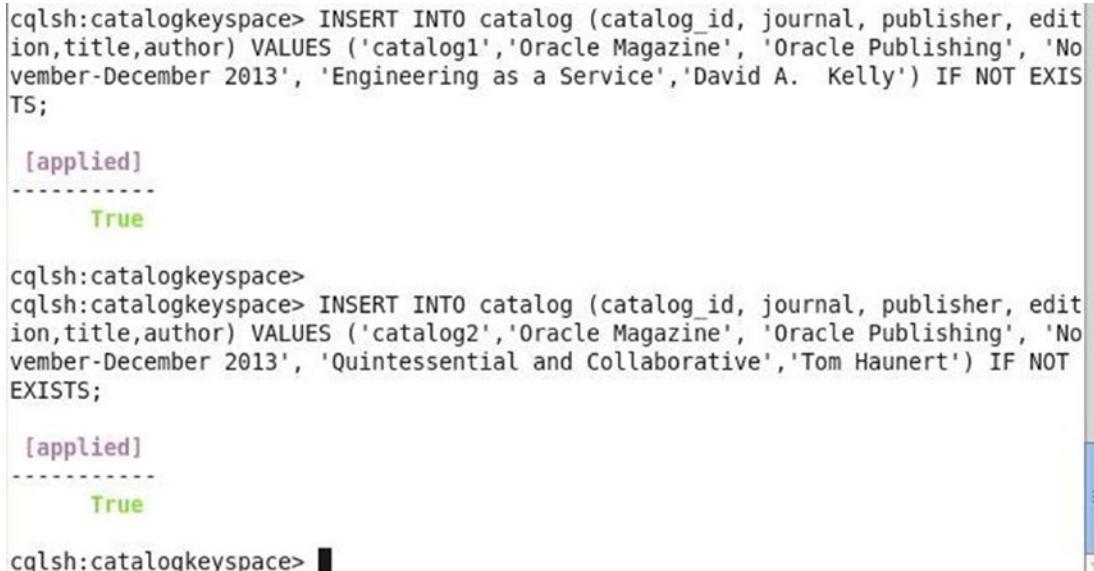
```
INSERT INTO <tablename>
    (' <identifier> ( ',' <identifier> )* ')*
    VALUES (' <term-or-literal> ( ',' <term-or-literal> )* ')*
    ( IF NOT EXISTS )?
    ( USING <option> ( AND <option> )* )?
```

For complete syntax for the INSERT statement refer <https://cassandra.apache.org/doc/cql3/CQL.html#insertStmt>. As an example add two rows of data to the catalog table and include the IF NOT EXISTS clause to add a row if a row identified by the primary key does not exist.

```
INSERT INTO catalog (catalog_id, journal, publisher, edition,title,author) VALUES
('catalog1','Oracle Magazine', 'Oracle Publishing', 'November-December 2013', 'Engineering
as a Service','David A. Kelly') IF NOT EXISTS;
```

```
INSERT INTO catalog (catalog_id, journal, publisher, edition,title,author) VALUES
('catalog2','Oracle Magazine', 'Oracle Publishing', 'November-December 2013',
'Quintessential and Collaborative','Tom Haunert') IF NOT EXISTS;
```

As indicated by the [applied] True output, two rows of data get added as shown in Figure 6-12.



```
cqlsh:catalogkeyspace> INSERT INTO catalog (catalog_id, journal, publisher, edit
ion,title,author) VALUES ('catalog1','Oracle Magazine', 'Oracle Publishing', 'No
vember-December 2013', 'Engineering as a Service','David A. Kelly') IF NOT EXIS
TS;
[applied]
-----
True

cqlsh:catalogkeyspace>
cqlsh:catalogkeyspace> INSERT INTO catalog (catalog_id, journal, publisher, edit
ion,title,author) VALUES ('catalog2','Oracle Magazine', 'Oracle Publishing', 'No
vember-December 2013', 'Quintessential and Collaborative','Tom Haunert') IF NOT
EXISTS;
[applied]
-----
True

cqlsh:catalogkeyspace> █
```

Figure 6-12. Adding Table Data

Querying a Table


The SELECT statement, which has the following syntax, is used to query a table.

```
SELECT <select-clause>
      FROM <tablename>
      ( WHERE <where-clause> )?
      ( ORDER BY <order-by> )?
      ( LIMIT <integer> )?
      ( ALLOW FILTERING )?
```

For the complete syntax for the SELECT statement refer to <https://cassandra.apache.org/doc/cql3/CQL.html#selectStmt>. As an example select all columns from the catalog table.

```
SELECT * FROM catalog;
```

The two rows of data added previously get listed as shown in Figure 6-13.



```
cqlsh:catalogkeyspace> SELECT * FROM catalog;
catalog_id | author           | edition          | journal          | publi
sher
-----+-----+-----+-----+-----+
catalog1 | David A. Kelly | November-December 2013 | Oracle Magazine | Oracle
e Publishing | Engineering as a Service
catalog2 | Tom Haunert | November-December 2013 | Oracle Magazine | Oracle
e Publishing | Quintessential and Collaborative
(2 rows)
cqlsh:catalogkeyspace>
```

Figure 6-13. Querying Table

Deleting from a Table

The DELETE statement is used to delete columns and rows and has the following syntax.

```
DELETE ( <selection> ( ',' <selection> )* )?
        FROM <tablename>
        ( USING TIMESTAMP <integer> )?
        WHERE <where-clause>
        ( IF ( EXISTS | ( <condition> ( AND <condition> )* ) ) )?
```

For complete syntax for the DELETE statement refer to <https://cassandra.apache.org/doc/cql3/CQL.html#deleteStmt>. As an example, delete all columns from the row with catalog_id as catalog1.

```
DELETE catalog_id, journal, publisher, edition, title, author from catalog WHERE catalog_id='catalog1';
```

Subsequently, query the catalog table with the SELECT statement.

```
SELECT * FROM catalog;
```

Column values from the row with catalog_id as catalog1 get deleted, but the row itself including the primary key column value do not get deleted even though the primary key catalog_id is listed as one of the columns to delete. Subsequent query lists the primary key column value but lists the column values for the other columns as null as shown in Figure 6-14.

```
cqlsh:catalogkeyspace> DELETE journal, publisher, edition, title, author from catalog WHERE catalog_id='catalog1';
cqlsh:catalogkeyspace> SELECT * FROM catalog;

catalog_id | author      | edition          | journal      | publisher
           | title
-----+-----+-----+-----+-----+
catalog1 | null       | null            | null         | null
null    |             | null            |             |
catalog2 | Tom Haunert | November-December 2013 | Oracle Magazine | Oracle Publishing | Quintessential and Collaborative
          |             |                |             |
(2 rows)
```

Figure 6-14. Deleting Table Data

Truncating a Table

The TRUNCATE statement removes all data from a table and has the following syntax.

```
TRUNCATE <tablename>
```

As an example, truncate the catalog table. Subsequently, run a query with the SELECT statement.

```
TRUNCATE catalog;
SELECT * from catalog;
```

As the output of the query indicates, no data is listed because the TRUNCATE statement has removed all data as shown in Figure 6-15.

```
cqlsh:catalogkeyspace> TRUNCATE catalog;
cqlsh:catalogkeyspace> SELECT * FROM catalog;

catalog_id | author | edition | journal | publisher | title
-----+-----+-----+-----+-----+
(0 rows)
```

Figure 6-15. Truncating a Table

Dropping A Table

The DROP TABLE or DROP COLUMN FAMILY statement is used to drop a table and has the following syntax.

```
DROP TABLE ( IF EXISTS )? <tablename>
```

As an example, drop the catalog table.

```
DROP TABLE IF EXISTS catalog;
```

If the IF EXISTS clause is not specified and the table does not exist, an error is generated. But with the IF EXISTS clause, an error is not generated as indicated by two consecutively run DROP TABLE statements with the IF EXISTS clause included in Figure 6-16.

```
cqlsh:catalogkeyspace> DROP TABLE IF EXISTS catalog;
cqlsh:catalogkeyspace> DROP TABLE IF EXISTS catalog;
```

Figure 6-16. Dropping a Table

Dropping a Keyspace

The DROP KEYSPACE statement, which has the following syntax, removes the specified key space including the column families in the key space and the data in the column families, and the keyspace does not have to be empty before being dropped.

```
DROP KEYSPACE ( IF EXISTS )? <identifier>
```

As an example, drop the CatalogKeyspace keyspace.

```
DROP KEYSPACE IF EXISTS CatalogKeyspace;
```

If the IF EXISTS clause is not specified and the keyspace does not exist, an error is generated. But with the IF EXISTS clause, an error is not generated as indicated by two consecutively run DROP KEYSPACE statements with the IF EXISTS clause included as shown in Figure 6-17.

```
cqlsh:catalogkeyspace> DROP KEYSPACE IF EXISTS CatalogKeyspace;
cqlsh:catalogkeyspace> DROP KEYSPACE IF EXISTS CatalogKeyspace;
```

Figure 6-17. Dropping a Keyspace

Exiting CQL Shell

To exit the cqlsh shell specify the exit command as shown in Figure 6-18. Subsequently exit the tty with the exit command also.

```
cqlsh:catalogkeyspace> exit
root@dfade563f871:/# exit
exit
[ec2-user@ip-172-30-1-192 ~]$
```

Figure 6-18. Exiting CQL Shell

Stopping Apache Cassandra

To stop Apache Cassandra, stop the Docker container running the Apache Cassandra server.

```
sudo docker stop cassandradb
```

Subsequently, run the following command to list the running containers.

```
sudo docker ps
```

The cassandradb container does not get listed as running as shown in Figure 6-19.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker stop cassandradb
cassandradb
[ec2-user@ip-172-30-1-192 ~]$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED
STATUS              PORTS              NAMES
[ec2-user@ip-172-30-1-192 ~]$
```

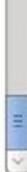


Figure 6-19. Stopping Cassandra DB Docker Container

Starting Multiple Instances of Apache Cassandra

Multiple Docker containers running Apache Cassandra instances may be started, but the container name has to be unique. As an example, start a new Docker container also called cassandradb to run another instance of Apache Cassandra database.

```
sudo docker run -t -i -v /cassandra/data:/var/lib/cassandra/data --name cassandradb -d -p 7000:7000 cassandra
```

Because a Docker container with the same name (cassandradb) was already created earlier, an error is generated even though the container has been stopped as shown in Figure 6-20. A container has to be removed with the docker rm command to be able to create a new container with the same name.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker run -t -i --name cassandradb -d cassandra
Error response from daemon: Conflict. The name "cassandradb" is already in use by container dfade563f871. You have to delete (or rename) that container to be able to reuse that name.
```



Figure 6-20. Duplicate Docker Container name error

Another container with a different name, cassandradb2 for example, may be started.

```
sudo docker run -t -i -v /cassandra/data:/var/lib/cassandra/data --name cassandradb2 -d -p 7000:7000 cassandra
```

Start a third container and specify the CASSANDRA_SEEDS environment variable for the IP address/es to be used to run multiple nodes in the cluster if required.

```
sudo docker run -t -i -v /cassandra/data:/var/lib/cassandra/data --name cassandradb3 -d -p 7000:7000 -e CASSANDRA_SEEDS=52.91.214.50,54.86.243.122,54.86.205.95 cassandra
```

Subsequently, run the following command to list the running containers.

```
sudo docker ps
```

The cassandradb2 and cassandradb3 containers get listed as running as shown in Figure 6-21.

```
[ec2-user@ip-172-30-1-192 ~]$ sudo docker run -t -i --name cassandradb2 -d cassandra
3629909b411b927ccde3d5a42fcfbe1754e6f8c7f3acf90cea2505827c29eac8
[ec2-user@ip-172-30-1-192 ~]$ sudo docker run -t -i -v /cassandra/data:/var/lib/cassandra/data --name cassandradb3 -d -p 7000:7000 -e CASSANDRA_SEEDS=52.91.214.50,54.86.243.122,54.86.205.95 cassandra
d965cbf2ad18931d84f4df32dc2832704c92523f9f5e5199f5cd5ae1326dd308
[ec2-user@ip-172-30-1-192 ~]$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS
NAMES
d965cbf2ad18        cassandra          "/docker-entrypoint.s"   27 seconds ago
Up 26 seconds       7001/tcp, 7199/tcp, 9042/tcp, 0.0.0.0:7000->7000/tcp, 9
160/tcp    cassandradb3
3629909b411b        cassandra          "/docker-entrypoint.s"   51 seconds ago
Up 50 seconds       7000-7001/tcp, 7199/tcp, 9042/tcp, 9160/tcp
cassandradb2
[ec2-user@ip-172-30-1-192 ~]$
```

Figure 6-21. Running Multiple Docker Containers for Instances of Apache Cassandra

Summary

In this chapter we use the Docker image for Apache Cassandra to run Apache Cassandra in a Docker container. We used the different CQL statements in a cqlsh shell to create a Keyspace, create a table in the Keyspace and add data to the table. We also ran CQL statements to query a table, delete data from the table, truncate a table, drop a table, and drop a keyspace. We also demonstrated creating multiple Docker containers to run multiple instances of Apache Cassandra. In the next chapter we shall run Couchbase Server in Docker.

CHAPTER 7



Using Couchbase Server

Couchbase Server is a distributed NoSQL database. Couchbase is a JSON (JavaScript Object Notation) based document store. Couchbase, like other NoSQL datastores, does not have a fixed schema for data storage. Couchbase differs from MongoDB in that MongoDB is based on the BSON (binary JSON) document data model. Couchbase provides a Web Console for accessing the Couchbase server from a graphical user interface (GUI). Couchbase also provides a command-line interface (CLI) including several tools to run in the CLI. In this chapter we shall run Couchbase server in a Docker container.

- Setting the Environment
- Starting Couchbase
- Accessing Couchbase Web Console
- Configuring Couchbase Server
- Adding Documents
- Starting Interactive Terminal
- Running Couchbase CLI Tools
- Stopping Couchbase Server

Setting the Environment

The following software is required for this chapter.

- Docker (version 1.8)
- Docker image for Couchbase (version latest)

We have used the Ubuntu Server AMI shown in Figure 7-1 for running software in this chapter. Installing and configuring an Amazon EC2 instance is discussed in Appendix A.

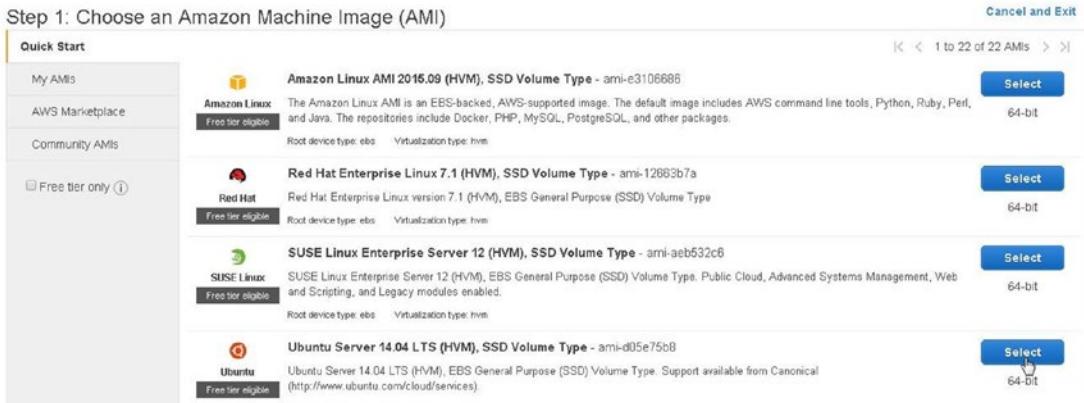


Figure 7-1. Ubuntu Server AMI

SSH Login to the Ubuntu Amazon EC2 instance using user as “ubuntu” and the public IP address of the Amazon EC2 instance. The public IP address would be different for different users (multiple public IP addresses are also used in this chapter based on multiple runs of the sample discussed).

```
ssh -i "docker.pem" ubuntu@54.152.90.139
```

We need to modify the IP address setting for localhost in the hosts IP addresses file /etc/hosts. Set the IP address to the public IP address of the Amazon EC2 instance. Obtaining the public IP address of an Amazon EC2 instance is discussed in Appendix A. Open the /etc/hosts file in a vi editor.

```
sudo vi /etc/hosts
```

Replace “127.0.0.1” with the public IP address; replace the following line:

```
127.0.0.1 localhost
```

with:

```
54.152.90.139 localhost
```

Install Docker on Ubuntu as discussed in Chapter 1. Run the hello-world Docker image to test the Docker installation.

```
sudo docker run hello-world
```

The output from the hello-world application is shown in Figure 7-2.

The screenshot shows a terminal window titled "ubuntu@ip-172-30-1-190: ~". The terminal displays the following command and its output:

```
ubuntu@ip-172-30-1-190:~$ sudo docker run hello-world
sudo: unable to resolve host ip-172-30-1-190
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
b901d36b6f2f: Pull complete
0a6ba66e537a: Pull complete
Digest: sha256:517f03be3f8169d84711c9ffb2b3235a4d27c1eb4ad147f6248c8040adb93113
Status: Downloaded newer image for hello-world:latest

Hello from Docker.
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:
1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:
$ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker Hub account:
https://hub.docker.com

For more examples and ideas, visit:
https://docs.docker.com/userguide/
```

The terminal prompt at the bottom is "ubuntu@ip-172-30-1-190:~\$".

Figure 7-2. Output from *hello-world*

Download the official Couchbase Docker image called “couchbase”

```
sudo docker pull couchbase
```

The latest Docker image gets downloaded as shown in Figure 7-3.

```
ubuntu@ip-172-30-1-190:~$ sudo docker pull couchbase
sudo: unable to resolve host ip-172-30-1-190
Using default tag: latest
latest: Pulling from library/couchbase
c0de77b824d9: Pull complete
7f183271ade4: Pull complete
02cc7e934fcc: Pull complete
fd97fd3cdea7: Pull complete
d28869c6aca6: Pull complete
63bf59e0c713: Pull complete
b33db0ff9d8b: Pull complete
6aa99012b457: Pull complete
8828e29d95f0: Pull complete
28087ef36a81: Pull complete
618bca76e9ce: Pull complete
1e040ef5e848: Pull complete
da7f5c40fc89: Pull complete
541fbf32b9aa: Pull complete
ff61ecf3bacb: Pull complete
Digest: sha256:261cab57c6e40f1eb414abec1a70ea4ecce01d4640414b40868a1bdaee79a88
Status: Downloaded newer image for couchbase:latest
ubuntu@ip-172-30-1-190:~$
```

Figure 7-3. Downloading Docker Image couchbase

Starting Couchbase

Next, run a Docker container for Docker image “couchbase”, which would start a Couchbase server process in the Docker container. Run the following docker command in which the port for the Couchbase Web Console to connect to Couchbase Server is specified as 8091. The container name is specified as “couchbasedb”.

```
sudo docker run --name couchbasedb -d -p 8091:8091 couchbase
```

Couchbase server could require non-default ulimit settings.

Ulimit Setting	Value	Description
ulimit -n	40960	nofile: max number of open files
ulimit -c	100000000	core: max core file size. The 100000000 setting is equivalent to “unlimited”, which is not directly supported.
ulimit -l	100000000	memlock: maximum locked-in-memory address space. The 100000000 setting is equivalent to “unlimited”, which is not directly supported.

A Docker container stores all persistent data in the /opt/couchbase/var directory, which could be mounted from the host using the -v command parameter. The -ulimit command parameter is used to set the docker run command. Run the following command to run a Docker container to run a Couchbase server as shown in Figure 7-4.

```
sudo docker run --name couchbasedb -v ~/couchbase/data:/opt/couchbase/var -d --ulimit nofile=40960:40960 --ulimit core=100000000:100000000 --ulimit memlock=100000000:100000000 -p 8091:8091 couchbase
```

Subsequently, list the running Docker containers.

```
sudo docker ps
```

The couchbasedb container gets listed as shown in Figure 7-4.

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
bff916e55a52	couchbase	"/entrypoint.sh couch"	6 seconds ago
Up 5 seconds	8092-8093/tcp, 11207/tcp, 11210-11211/tcp, 0.0.0.0:8091->8091/tcp, 18091-18092/tcp	couchbasedb	

Figure 7-4. Running Docker Container for Couchbase

Output the logs for the container with the docker logs command.

```
sudo docker logs couchbasedb
```

The message shown in Figure 7-5 gets displayed.

```
ubuntu@ip-172-30-1-251:~$ sudo docker logs couchbasedb
Starting Couchbase Server -- Web UI available at http://<ip>:8091
```

Figure 7-5. Listing Docker Container Log

Accessing Couchbase Web Console

Next, we shall access the Couchbase Web Console from the URL indicated in the logs: <http://<ip>:8091>. The <ip> address to use would vary from which host system the Web Console is accessed. If on the same host as on which the Docker container is running, use the public IP address of the host Amazon EC2 instance. If on a remote host system as we have accessed, use the public DNS for the Amazon EC2 instance. Obtaining the public IP address and the public DNS are discussed in Appendix A. If the public DNS is ec2-54-152-90-139.compute-1.amazonaws.com, the URL to access the Couchbase WebConsole becomes the following.

<http://ec2-54-152-90-139.compute-1.amazonaws.com:8091>

Open a browser at the preceding URL. The Couchbase Console gets displayed as shown in Figure 7-6. In the next section we shall setup a Couchbase server cluster.

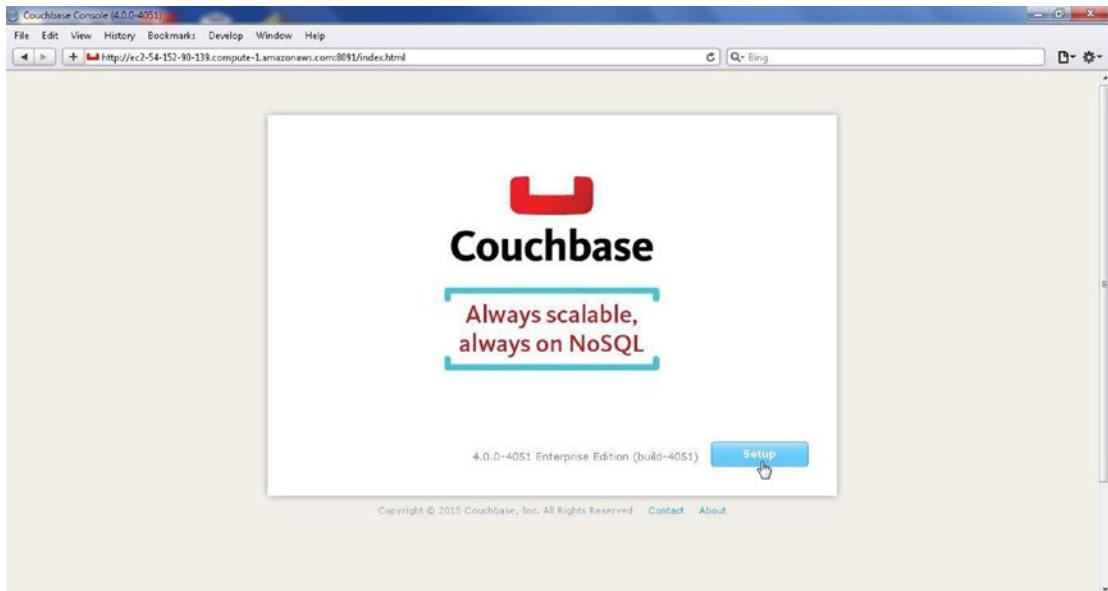


Figure 7-6. Accessing Couchbase Admin Console

If a Couchbase cluster has already been configured, the Couchbase Console URL would display the login page as shown in Figure 7-7.




Figure 7-7. Login Page

Specify the Username (Administrator) and the Password and click on Sign In as shown in Figure 7-8.




Figure 7-8. Specifying Username and Password

Configuring Couchbase Server Cluster

In this section we shall configure the Couchbase Server cluster. Access the Couchbase Web Console as discussed in previous section and shown in Figure 7-6, with URL <http://ec2-54-152-90-139.compute-1.amazonaws.com:8091>. Click on Setup in the Web Console; the “Setup” page is displayed only the first time the Web Console is accessed. Subsequently, after a cluster has been configured the Login page is displayed as discussed in the previous section.

Use the default settings for the Configure Disk Storage section. In Configure Server Hostname specify the Hostname as the Public IP Address of the Amazon EC2 instance, which would be different for different users, as shown in Figure 7-9. Short names are not acceptable for the Hostname field and at least one dot is required in the host name.

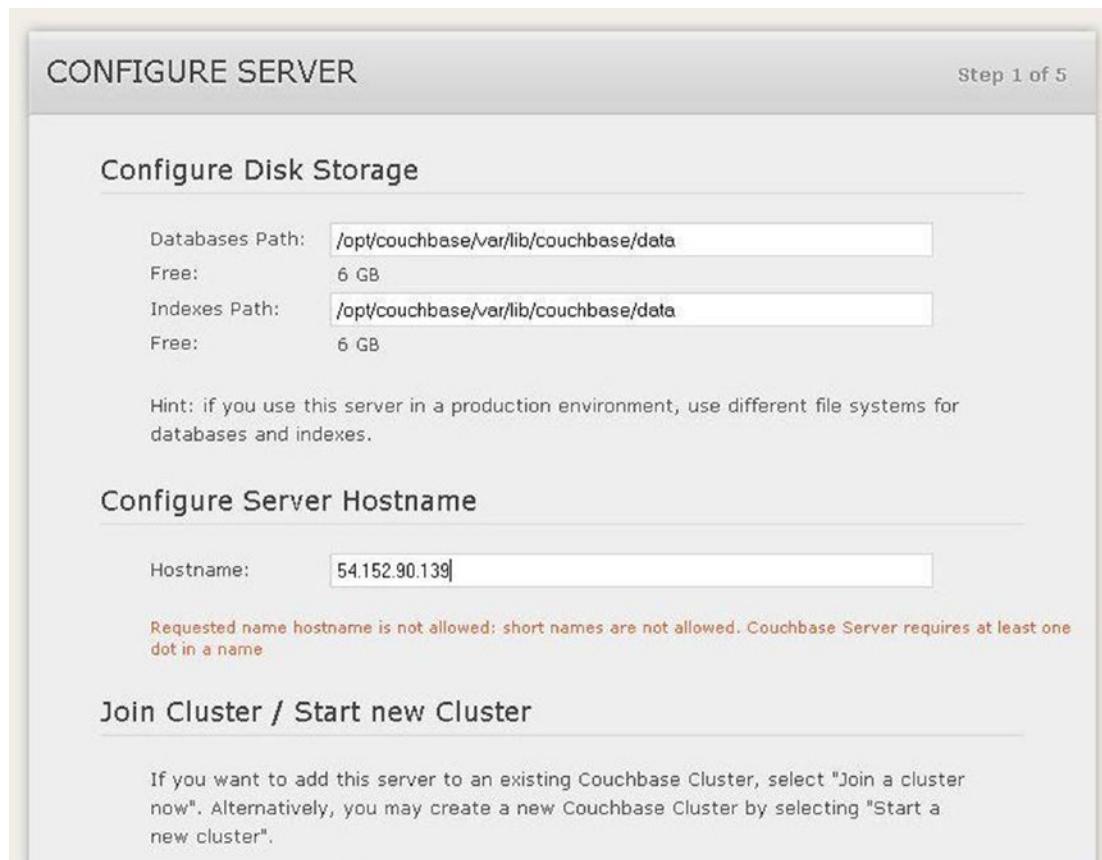


Figure 7-9. Configuring Server

Two options are provided in the Join Cluster/ Start new cluster section. As we are configuring a new cluster, select Start a new cluster as shown in Figure 7-10. Select the default settings or modify the settings keeping in consideration the total RAM configurable per server. Click on Next.

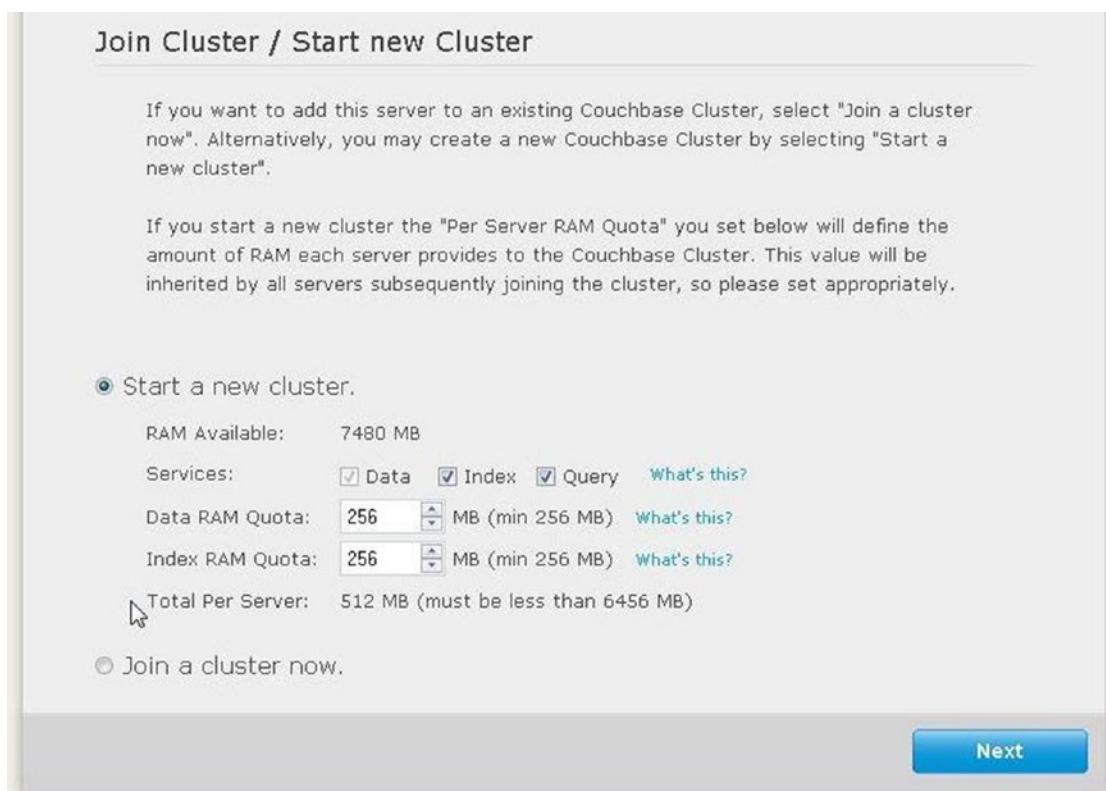


Figure 7-10. Starting a New Cluster

Couchbase server stores data in data buckets. In the Sample Buckets section the sample buckets are listed. A sample bucket is not required to be selected. Click on Next. In the Create Default Bucket screen the Bucket Name is pre-specified as “default”. Select Bucket Type as “Couchbase”. Select the default Memory Size & Replicas settings. Also select the default Disk I/O Optimization setting.

CREATE DEFAULT BUCKET Step 3 of 5

Bucket Settings

Bucket Name: **default**

Bucket Type: Couchbase Memcached

Memory Size

Per Node RAM Quota: MB Cluster quota (256 MB)

Other Buckets (0 B) This Bucket (256 MB) Free (0 B)

Total bucket size = 256 MB (256 MB x 1 node)

Cache Metadata: Value Ejection [What's this?](#) Full Ejection

Replicas

Enable 1 Number of replica (backup) copies
 View index replicas

Figure 7-11. Configuring the Default Cluster

In Flush select Enable and click on Next as shown in Figure 7-12. To be able to flush (delete) data from a bucket, 'Flush' must be enabled.

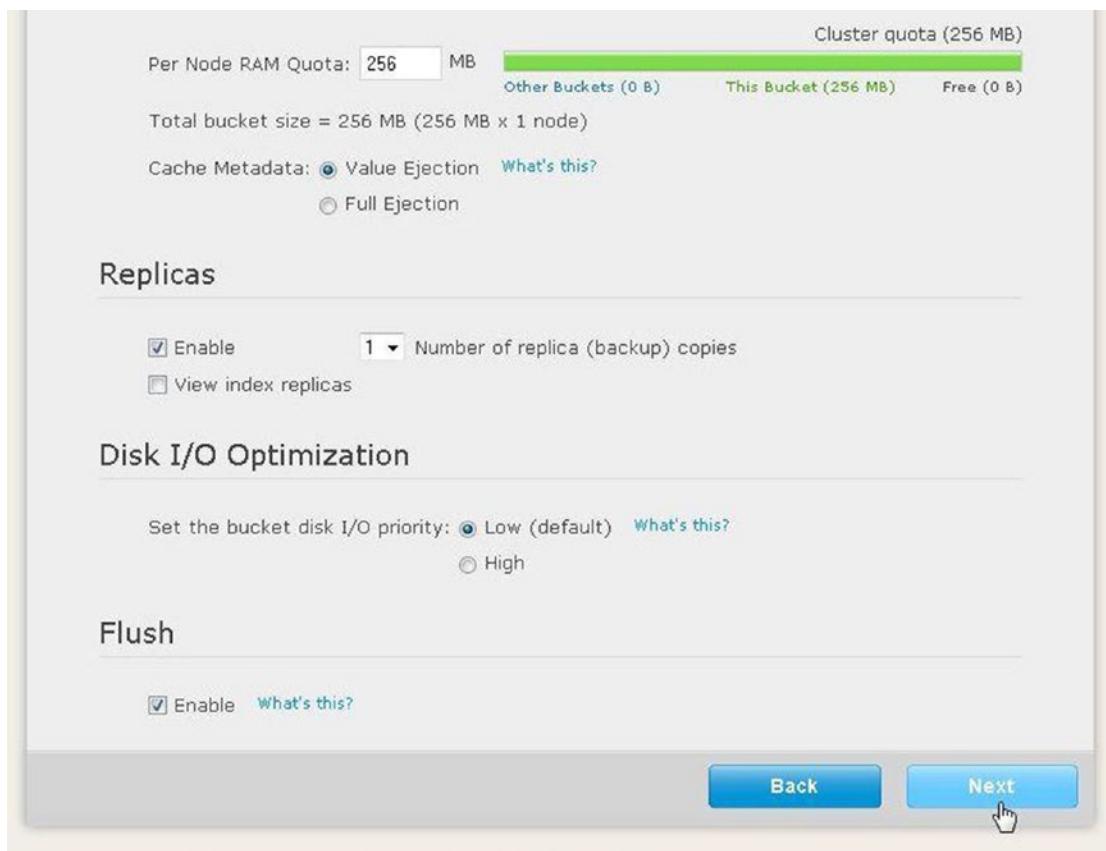


Figure 7-12. Enabling Flush

In Notifications, select the default settings and the “I agree...” checkbox and click on Next as shown in Figure 7-13.

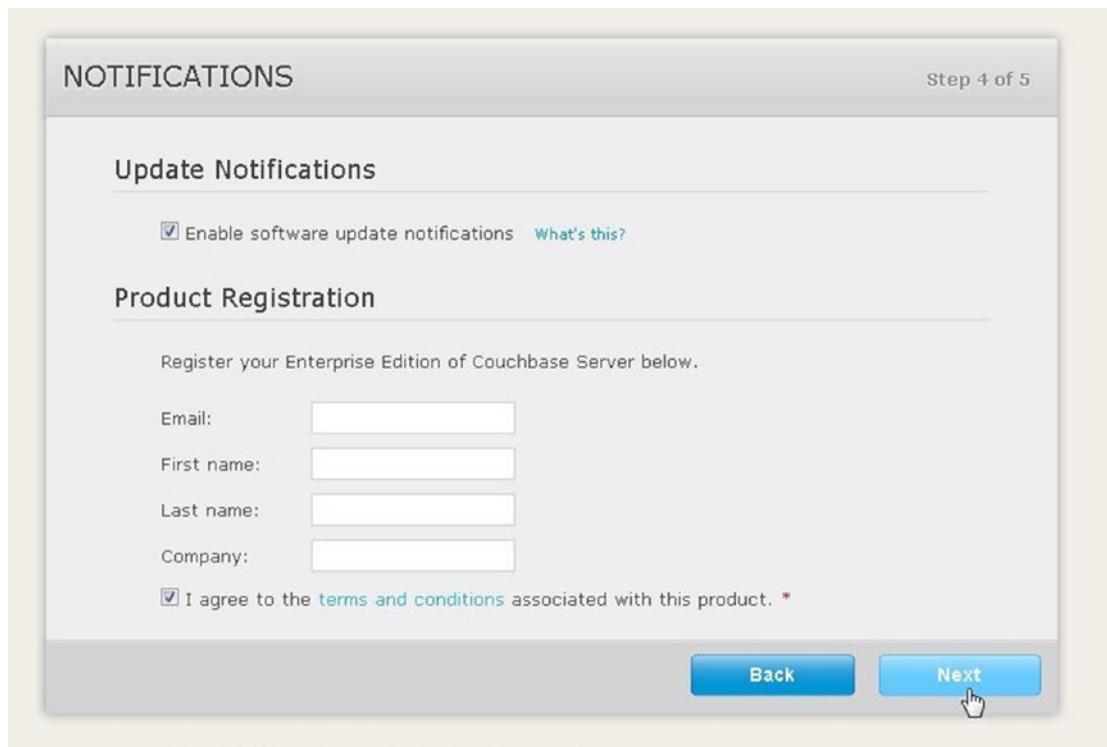


Figure 7-13. Configuring Notifications

In Secure this Server screen specify the Username as Administrator (default setting) as shown in Figure 7-14. Specify a password in the Password field and specify the same password in Verify Password field. Click on Next.



Figure 7-14. Specifying Username and Password

Click on the Cluster Overview tab to display the Cluster summary including the RAM allocated and in use, and the Disk storage allocated and in use as shown in Figure 7-15.

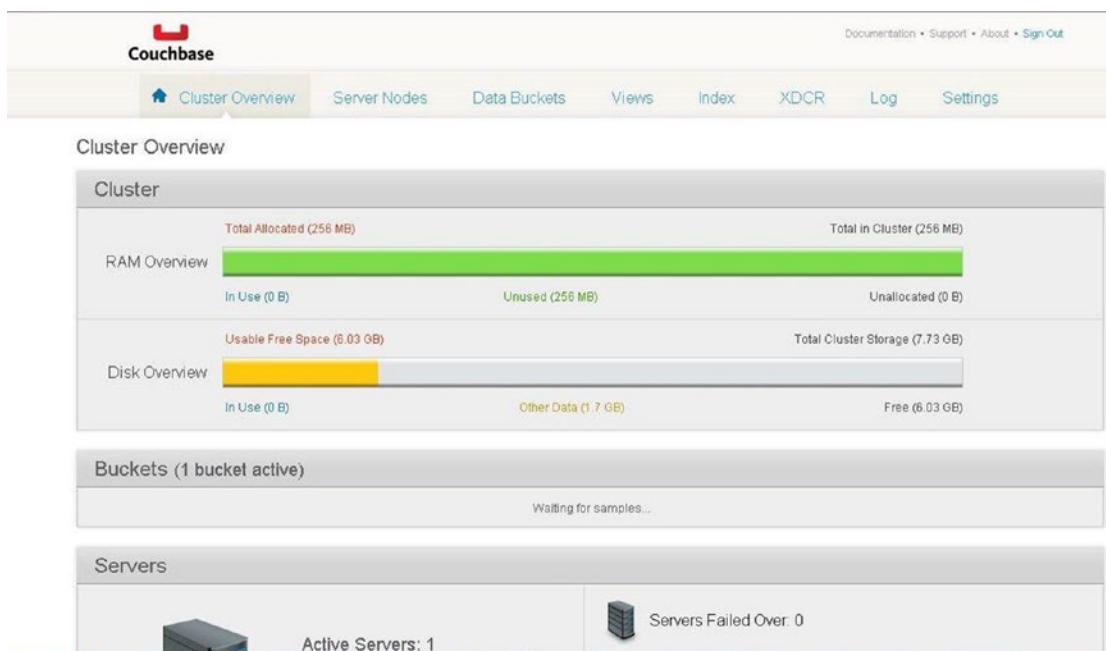


Figure 7-15. Displaying Cluster Summary

One bucket is shown as Active and one server is shown as Active in Figure 7-16.

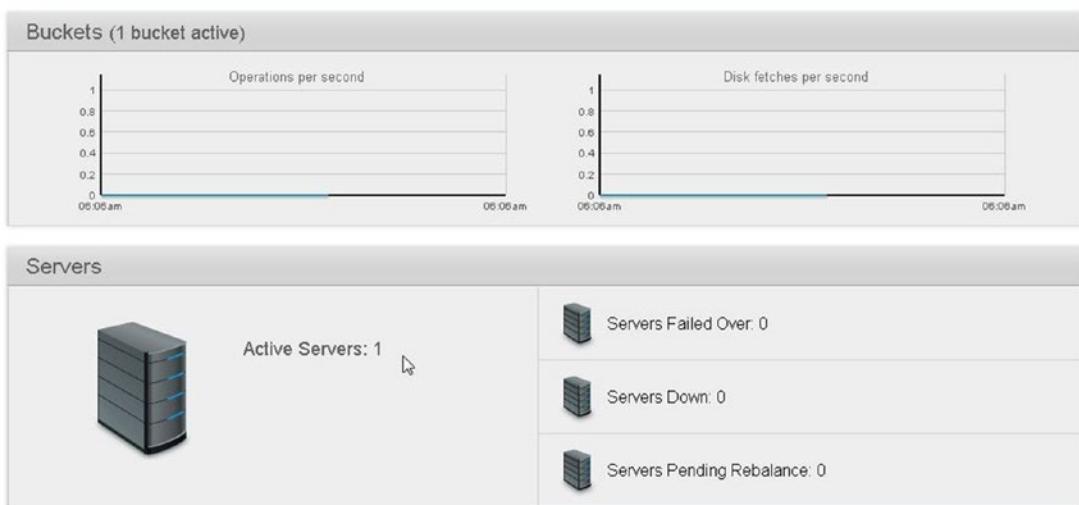


Figure 7-16. Displaying Servers Summary

Click on Server Nodes to list the server nodes. The server running at IP address 172.17.0.1 gets listed as shown in Figure 7-17.

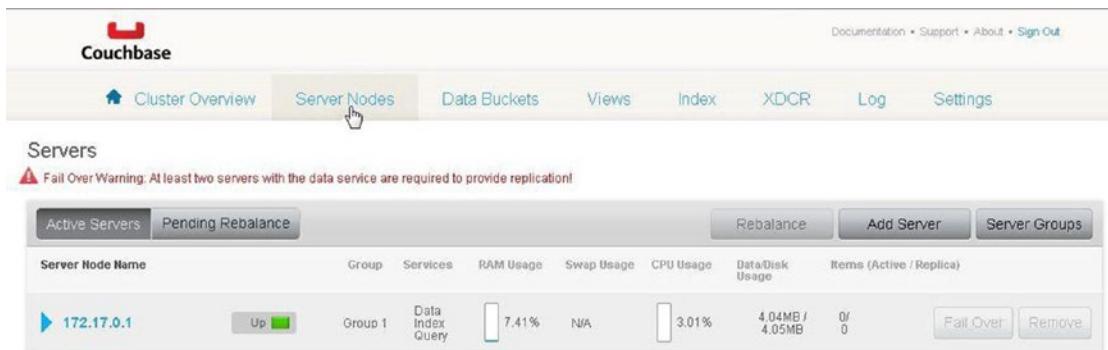


Figure 7-17. Listing Server IP Address

Click on the Data Buckets tab. The “default” bucket gets listed as shown in Figure 7-18.

The screenshot shows the Couchbase Data Buckets interface. At the top, there's a navigation bar with links for Cluster Overview, Server Nodes, Data Buckets (which is the active tab), Views, Index, XDCR, Log, and Settings. On the right side of the header, there are links for Documentation, Support, About, and Sign Out. Below the header, the title "Data Buckets" is displayed. A sub-header "Couchbase Buckets" is followed by a table. The table has columns: Bucket Name, Data Nodes, Item Count, Ops/sec, Disk Fetches/sec, RAM/Quota Usage, and Data/Disk Usage. One row is visible for the "default" bucket, which has 1 data node, 0 items, 0 ops/sec, 0 disk fetches/sec, 36.4MB / 256MB RAM usage, and 4.04MB / 4.05MB data/disk usage. To the right of the table are two buttons: "Documents" and "Views".

Figure 7-18. Listing the Default Buckets

Adding Documents

In this section we shall add documents to the Couchbase server from the Couchbase Console. Click on the Documents button for the default bucket as shown in Figure 7-19.

This screenshot is identical to Figure 7-18, showing the Couchbase Data Buckets page with the "default" bucket listed. However, a mouse cursor is hovering over the "Documents" button located to the right of the bucket details table. The browser's address bar shows the URL as <http://ec2-54-152-90-139.compute-1.amazonaws.com:8091/index.html#sec=buckets>.

Figure 7-19. Clicking on the Documents button

In the default ► Documents no document is listed to start with. Click on Create Document button as shown in Figure 7-20.



Figure 7-20. Clicking on ‘Create Document’

In the Create Document dialog specify a Document ID, catalog1 for example and click on Create as shown in Figure 7-21.

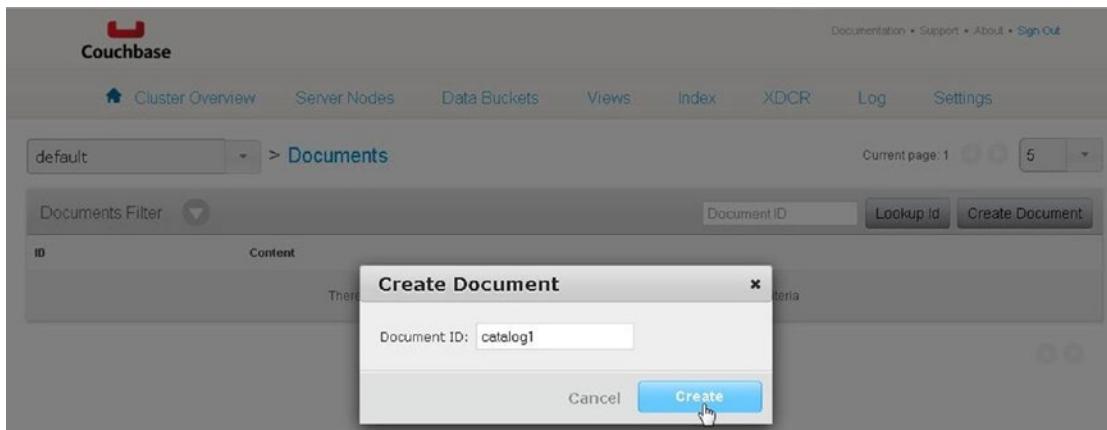


Figure 7-21. Creating a Document

A JSON document with Id catalog1 gets added to the default bucket as shown in Figure 7-22. The new document has some default fields, which would probably be required to be modified.

The screenshot shows the Couchbase Console interface. At the top, there's a navigation bar with links for Documentation, Support, About, and Sign Out. Below the navigation bar, there's a menu bar with Cluster Overview, Server Nodes, Data Buckets, Views, Index, XDCR, Log, and Settings. A dropdown menu is open under 'Cluster Overview' with 'default' selected. To the right of the dropdown, it says 'Documents'. In the main content area, there's a table with one row. The row has a 'catalog1' icon, the ID 'catalog1', and a preview pane containing a JSON object. The JSON object is:


```

1  {
2      "click": "to edit",
3      "new in 2.0": "there are no reserved field names"
4  }
    
```

 To the right of the preview pane are three buttons: Delete, Save As..., and Save. The 'Save' button is highlighted with a mouse cursor.

Figure 7-22. New Document with ID as catalog1

Replace the sample JSON document with the following JSON document.

```
{
  "journal": "Oracle Magazine",
  "publisher": "Oracle Publishing",
  "edition": "November-December 2013",
  "title": "Quintessential and Collaborative",
  "author": "Tom Haunert"
}
```

Click on Save to save the modified JSON document as shown in Figure 7-23.

This screenshot is similar to Figure 7-22, showing the Couchbase Console with the 'catalog1' document. The JSON content has been replaced with the new document. The 'Save' button is now highlighted with a mouse cursor, indicating the action of saving the document.

Figure 7-23. Saving a Couchbase Document

The catalog1 JSON document gets saved and also formatted in the Couchbase Console as shown in Figure 7-24.

The screenshot shows the Couchbase Web UI. The top navigation bar includes links for Documentation, Support, About, and Sign Out. Below the navigation is a menu bar with Cluster Overview, Server Nodes, Data Buckets, Views, Index, XDCR, Log, and Settings. A dropdown menu is open under Cluster Overview, showing 'default' selected. To the right of the dropdown is a breadcrumb trail: default > Documents. The main content area displays a JSON document titled 'catalog1'. The JSON content is:

```

catalog1
1 {
2   "journal": "Oracle Magazine",
3   "publisher": "Oracle Publishing",
4   "edition": "November-December 2013",
5   "title": "Quintessential and Collaborative",
6   "author": "Tom Haunert"
7 }

```

Buttons for Delete, Save As..., and Save are visible on the right.

Figure 7-24. Formatted JSON Document

In Couchbase Buckets, the Item Count for the “default” bucket gets listed as 1 as shown in Figure 7-25. Click on the Documents button to display the documents in the default bucket.

The screenshot shows the Couchbase Web UI. The top navigation bar includes links for Documentation, Support, About, and Sign Out. Below the navigation is a menu bar with Cluster Overview, Server Nodes, Data Buckets, Views, Index, XDCR, Log, and Settings. The 'Data Buckets' tab is selected. A sub-menu for 'Data Buckets' is open, showing 'default' selected. The main content area displays a table titled 'Couchbase Buckets' with one row for the 'default' bucket. The table columns are: Bucket Name, Data Nodes, Item Count, Ops/sec, Disk Fetches/sec, RAM/Quota Usage, and Data/Disk Usage. The 'default' row shows 1 data node, 1 item count, 0 ops/sec, 0 disk fetches/sec, 36.4MB / 256MB RAM usage, and 4.05MB / 4.08MB data/disk usage. Buttons for 'Documents' and 'Views' are shown on the right.

Bucket Name	Data Nodes	Item Count	Ops/sec	Disk Fetches/sec	RAM/Quota Usage	Data/Disk Usage
default	1	1	0	0	36.4MB / 256MB	4.05MB / 4.08MB

Figure 7-25. Item Count for default Bucket

The catalog1 document gets listed as shown in Figure 7-26. Click on the Edit Document button to display the document JSON if required.

The screenshot shows the Couchbase Web UI. The top navigation bar includes links for Documentation, Support, About, and Sign Out. Below the navigation is a menu bar with Cluster Overview, Server Nodes, Data Buckets, Views, Index, XDCR, Log, and Settings. A dropdown menu is open under Cluster Overview, showing 'default' selected. To the right of the dropdown is a breadcrumb trail: default > Documents. The main content area displays a table titled 'Documents' with one row for the 'catalog1' document. The table columns are: ID and Content. The 'catalog1' row shows the JSON content: { "journal": "Oracle Magazine", "publisher": "Oracle Publishing", "ed...". Buttons for 'Edit Document' and 'Delete' are visible on the right. Navigation controls for page 1 of 5 are at the bottom right.

ID	Content
catalog1	{ "journal": "Oracle Magazine", "publisher": "Oracle Publishing", "ed...

Figure 7-26. Listing Documents in the default Bucket

Similarly add another document with document id as catalog2. The JSON for catalog2 document is as follows.

```
{
  "journal": "Oracle Magazine",
  "publisher": "Oracle Publishing",
  "edition": "November December 2013",
  "title": "Engineering as a Service",
  "author": "David A. Kelly",
}
```

Add the JSON the sample document for catalog2 as we did for the catalog1 document and click on Save as shown in Figure 7-27.

The screenshot shows the Couchbase Web UI interface. At the top, there's a navigation bar with links for Documentation, Support, About, and Sign Out. Below the navigation is a main menu with Cluster Overview, Server Nodes, Data Buckets, Views, Index, XDCR, Log, and Settings. Under the Data Buckets menu, a dropdown shows 'default' selected. To its right, a breadcrumb trail reads 'Documents'. In the center, there's a text input field containing the JSON for 'catalog2'. Above the input field are three buttons: Delete, Save As..., and Save. The JSON content is:

```

1  {
2    "journal": "Oracle Magazine",
3    "publisher": "Oracle Publishing",
4    "edition": "November December 2013",
5    "title": "Engineering as a Service",
6    "author": "David A. Kelly"
7  }

```

Figure 7-27. Adding another JSON Document

The two documents catalog1 and catalog2 get listed as shown in Figure 7-28.

The screenshot shows the Couchbase Web UI interface after the documents have been saved. The layout is similar to Figure 7-27, with the same navigation bar and main menu. The 'default' bucket is selected in the dropdown. The breadcrumb trail shows 'Documents'. The central area displays a table of documents. The table has columns for ID and Content. Two rows are visible: one for 'catalog1' and one for 'catalog2'. Each row contains the JSON content and includes 'Edit Document' and 'Delete' buttons. The table also features a 'Documents Filter' dropdown and search fields for Document ID and Lookup Id. A pagination control at the bottom indicates 'Current page: 1' and a dropdown for '5' items per page.

ID	Content
catalog1	{ "journal": "Oracle Magazine", "publisher": "Oracle Publishing", "ed...
catalog2	{ "journal": "Oracle Magazine", "publisher": "Oracle Publishing", "ed...


Figure 7-28. Listing the Two Documents Added

Starting Interactive Terminal

To access the Couchbase server from a command line, start the interactive terminal (tty).

```
sudo docker exec -it couchbasedb bash
```

The interactive shell gets started as shown in Figure 7-29.



```
ubuntu@ip-172-30-1-251:~$ sudo docker exec -it couchbasedb bash
root@bff916e55a52:/#
```

Figure 7-29. Starting the Interactive Shell

The interactive terminal may also be started using the container id instead of the container name.

```
sudo docker exec -it bff916e55a52 bash
```

Running Couchbase CLI Tools

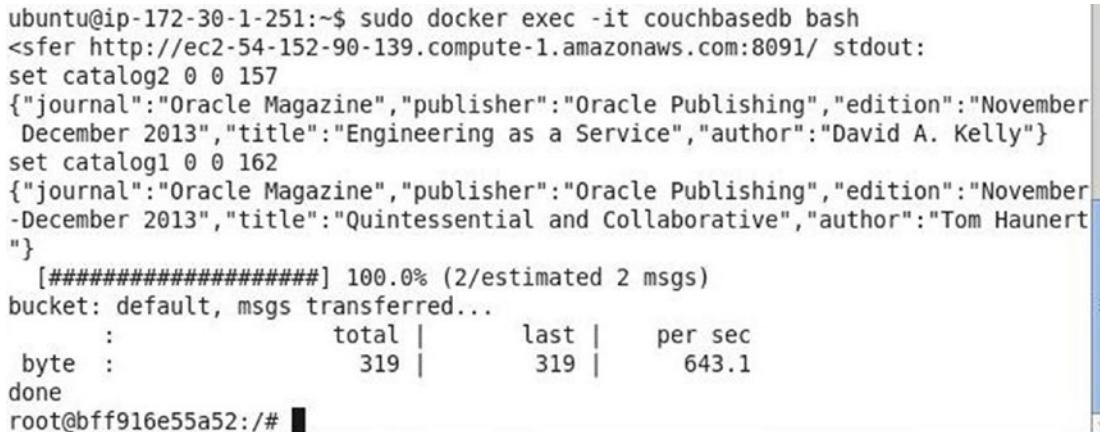
Couchbase Server provides several command-line interface tools (CLI) to monitor and manage Couchbase server buckets, nodes and cluster.

Some of these CLI tools are the `couchbase-cli` tool for operations on the entire cluster, the `cbbackup` tool to create a backup, the `cbdocloader` tool to load JSON documents, and the `cbtransfer` tool to transfer data between clusters and data files on the host.

As an example, run the `cbtransfer` tool to transfer data from the Couchbase server to the stdout with the following command run from the tty.

```
cbtransfer http://ec2-54-152-90-139.compute-1.amazonaws.com:8091/ stdout:
```

The two JSON documents previously added to the Couchbase cluster from the Couchbase Console get output the stdout as shown in Figure 7-30.



```
ubuntu@ip-172-30-1-251:~$ sudo docker exec -it couchbasedb bash
<sfer http://ec2-54-152-90-139.compute-1.amazonaws.com:8091/ stdout:
set catalog2 0 0 157
{"journal":"Oracle Magazine","publisher":"Oracle Publishing","edition":"November December 2013","title":"Engineering as a Service","author":"David A. Kelly"}
set catalog1 0 0 162
{"journal":"Oracle Magazine","publisher":"Oracle Publishing","edition":"November -December 2013","title":"Quintessential and Collaborative","author":"Tom Haunert "}
[########################################] 100.0% (2/estimated 2 msgs)
bucket: default, msgs transferred...
      :          total |      last |    per sec
  byte :            319 |       319 |     643.1
done
root@bff916e55a52:/#
```

Figure 7-30. Running cbtransfer

Stopping Couchbase Server and Container

To stop the Couchbase Server and container, exit the interactive terminal with `exit` command as shown in Figure 7-31.

```
root@bff916e55a52:/# exit
exit
ubuntu@ip-172-30-1-251:~$
```

Figure 7-31. Stopping Couchbase Server

In the host system, run the `docker stop` command to stop the Docker container.

```
sudo docker stop couchbasedb
```

Subsequently, list the running Docker containers.

```
sudo docker ps
```

The `couchbasedb` container does not get listed as shown in Figure 7-32.

```
ubuntu@ip-172-30-1-251:~$ sudo docker stop couchbasedb
couchbasedb
ubuntu@ip-172-30-1-251:~$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND       CREATED
STATUS              PORTS              NAMES
ubuntu@ip-172-30-1-251:~$
```

Figure 7-32. The Docker Container for `couchbasedb` does not get listed

Summary

In this chapter we used the official Docker image for Couchbase Server to run a Couchbase Server instance in a Docker container. We accessed the Couchbase Sever from the Couchbase Console and added some JSON documents. Subsequently, we used the `cbtransfer` CLI tool to output the documents stored to the `stdout`. In the next chapter we shall discuss using Apache Hadoop.

CHAPTER 8



Using Apache Hadoop

Apache Hadoop is the de facto framework for processing large data sets. Apache Hadoop is a distributed software application that runs across several (up to hundreds and thousands) of nodes across a cluster. Apache Hadoop comprises of two main components: Hadoop Distributed File System (HDFS) and MapReduce. The HDFS is used for storing large data sets and MapReduce is used for processing the large data sets. Hadoop is linearly scalable without degradation in performance and makes use of commodity hardware rather than any specialized hardware. Hadoop is designed to be fault tolerant and makes use of data locality by moving the computation to the data rather than data to the computation. MapReduce framework has two versions MapReduce1 (MR1) and MapReduce2 (MR2) (also called YARN). MR1 is the default MapReduce framework in earlier versions of Hadoop (Hadoop 1.x) and YARN is the default in latter versions of Hadoop (Hadoop 2.x).

Setting the Environment

Starting Hadoop

Starting the Interactive Shell

Creating Input Files for a MapReduce Word Count Application

Running a MapReduce Word Count Application

Stopping the Hadoop Docker Container

Using a CDH Docker Image

Setting the Environment

The following software is used in this chapter.

-Docker (version 1.8)

-Apache Hadoop Docker Image

-Cloudera Hadoop (CDH) Docker Image

As in other chapters we have used an Amazon EC2 instance based on Red Hat Enterprise Linux 7.1 (HVM), SSD Volume Type - ami-12663b7a for installing the software. SSH login to the Amazon EC2 instance.

```
ssh -i "docker.pem" ec2-user@52.23.207.240
```

Install Docker as discussed in Chapter 1. Start the Docker service.

```
sudo service docker start
```

An OK message indicates that the Docker service has been started as shown in Figure 8-1.

```
[ec2-user@ip-172-30-1-211 ~]$ sudo service docker start
Starting docker (via systemctl): [ OK ]
```

Figure 8-1. Starting the Docker Service

Add a group called “hadoop” and a user called “hadoop”.

```
groupadd hadoop
useradd -g hadoop hadoop
```

Several Docker images are available for Apache Hadoop. We have used the sequenceiq/hadoop-docker Docker image available from the Docker Hub. Download the Docker image with label 2.7.0 or the latest tag image if different.

```
sudo docker pull sequenceiq/hadoop-docker:2.7.0
The docker pull command is shown in Figure 8-2.
```

```
[ec2-user@ip-172-30-1-211 ~]$ sudo docker pull sequenceiq/hadoop-docker:2.7.0
2.7.0: Pulling from sequenceiq/hadoop-docker
bf8622e5890e: Pulling fs layer
f378deb69e60: Pulling fs layer
e72ac57a5ca4: Pulling fs layer
5d32476d095f: Pulling fs layer
e8ef4721ce78: Pulling fs layer
4ba1516b7324: Pulling fs layer
9e15cf33fd16: Pulling fs layer
3c19d23c1a9a: Pulling fs layer
75161165b85c: Pulling fs layer
590c0e5d7008: Pulling fs layer
dbdbb1b52f3da: Pulling fs layer
0c031babbb191: Pulling fs layer
54ed83ba6569: Pulling fs layer
e1c1f4a38c78: Pulling fs layer
01c749264ddc: Pulling fs layer
368c32f0e407: Pulling fs layer
936e391b7ce3: Pulling fs layer
590e41abc605: Pulling fs layer
647ab508657a: Pulling fs layer
22a3743806b7: Pulling fs layer
d3b3ece18e7: Pulling fs layer
bd6b151ceb45: Pulling fs layer
124bac6063a3: Pulling fs layer
6bec6cb9c97d: Pulling fs layer
a228dc98c941: Pulling fs layer
cf392acc6304: Pulling fs layer
8f59e3ab71c7: Pulling fs layer
a36807fd30e6: Pulling fs layer
```

Figure 8-2. Running the docker pull Command

The Docker image `sequenceiq/hadoop-docker` gets downloaded as shown in Figure 8-3.

```
ec2-user@ip-172-30-1-86:~$ 
File Edit View Search Terminal Help
a67dbdaa3483: Pull complete
7738d90ca912: Pull complete
999867c8eb92: Pull complete
b25731f9c681: Pull complete
d3758b55e6eb: Pull complete
13718e471b2d: Pull complete
b08aadcc5362: Pull complete
316acd54632c: Pull complete
7acef3eb3f48: Pull complete
8be9869b36a0: Pull complete
8ac20a7a614b: Pull complete
446af44e8851: Pull complete
de40cb5733bd: Pull complete
d70a4c5c7d80: Pull complete
d87ea7766237: Pull complete
7711f4bfb4d: Pull complete
2d5ae4dfb209: Pull complete
1c48867840a2: Pull complete
7e7e85248613: Pull complete
a6f24ed3591a: Pull complete
679730efaea5: Pull complete
Digest: sha256:a40761746eca036fee6aafdf9fdbd6878ac3dd9a7cd83c0f3f5d8a0e6350c76a
Status: Downloaded newer image for sequenceiq/hadoop-docker:2.7.0
[ec2-user@ip-172-30-1-86 ~]$ 
```

Figure 8-3. Downloading Docker Image `sequenceiq/hadoop-docker`

Starting Hadoop

Next, start the Hadoop components HDFS and MapReduce. The Docker image `sequenceiq/hadoop-docker` is configured by default to start the YARN or MR2 framework. Run the following `docker run` command, which starts a Docker container in detached mode, to start the HDFS (NameNode and DataNode) and YARN (ResourceManager and NodeManager).

```
sudo docker run -d --name hadoop sequenceiq/hadoop-docker:2.7.0
```

Subsequently, list the running Docker containers.

```
sudo docker ps
```

The output from the preceding two commands is shown in Figure 8-4 including the running Docker container for Apache Hadoop based on the `sequenceiq/hadoop-docker` image. The Docker container name is “`hadoop`” and container id is “`27436aa7c645`”.

```
[ec2-user@ip-172-30-1-211 ~]$ sudo docker run -d --name hadoop sequenceiq/hadoop-docker:2.7.0
27436aa7c645c0053dd5729ae73bb09bb988dc0341c496f55f30b3e708bfb92a
[ec2-user@ip-172-30-1-211 ~]$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED             STATUS              PORTS
NAMES
27436aa7c645      sequenceiq/hadoop-docker:2.7.0   "/etc/bootstrap.sh -d"   9
seconds ago         Up 8 seconds          2122/tcp, 8030-8033/tcp, 8040/tcp, 8042/tcp, 8088/tcp, 19888/tcp, 49707/tcp, 50010/tcp, 50020/tcp, 50070/tcp, 50075/tcp, 50090/tcp   hadoop
[ec2-user@ip-172-30-1-211 ~]$
```

Figure 8-4. Running Docker Container for Apache Hadoop

Starting the Interactive Shell

Start the interactive shell or terminal (tty) with the following command.

```
sudo docker exec -it hadoop bash
```

The interactive terminal prompt gets displayed as shown in Figure 8-5.

```
[ec2-user@ip-172-30-1-86 ~]$ sudo docker exec -it hadoop bash
bash-4.1#
```

Figure 8-5. Starting Interactive Terminal

The interactive shell may also be started using the container id instead of the container name.

```
sudo docker exec -it 27436aa7c645 bash
```

If the `-d` command parameter is omitted from the `docker run` command and the `-it` parameters (which is `-i` and `-t` supplied together) are supplied using the following command, the Docker container starts in foreground mode.

```
sudo docker run -it --name hadoop sequenceiq/hadoop-docker:2.7.0 /etc/bootstrap.sh -bash
```

The Hadoop components start and attach a console to the Hadoop `stdin`, `stdout` and `stderr` streams as shown in Figure 8-6. A message gets output to the console for each Hadoop component started. The `-it` parameter starts an interactive terminal (tty).

```
[ec2-user@ip-172-30-1-211 ~]$ sudo docker run -it --name hadoop sequenceiq/hadoop-p-docker:2.7.0 /etc/bootstrap.sh -bash
/
Starting sshd: [ OK ]
Starting namenodes on [ebb125a12e13]
ebb125a12e13: starting namenode, logging to /usr/local/hadoop/logs/hadoop-root-namenode-ebb125a12e13.out
localhost: starting datanode, logging to /usr/local/hadoop/logs/hadoop-root-data-node-ebb125a12e13.out
Starting secondary namenodes [0.0.0.0]
0.0.0.0: starting secondarynamenode, logging to /usr/local/hadoop/logs/hadoop-root-secondarynamenode-ebb125a12e13.out
starting yarn daemons
starting resourcemanager, logging to /usr/local/hadoop/logs/yarn--resourcemanager-ebb125a12e13.out
localhost: starting nodemanager, logging to /usr/local/hadoop/logs/yarn-root-nodemanager-ebb125a12e13.out
bash-4.1#
```

Figure 8-6. Starting Docker Container in Foreground

Creating Input Files for a MapReduce Word Count Application

In this section we shall create input files for a MapReduce Word Count application, which is included in the examples packaged with the Hadoop distribution. To create the input files, change the directory (cd) to the \$HADOOP_PREFIX directory.

```
bash-4.1# cd $HADOOP_PREFIX
```

The preceding command is to be run from the interactive terminal (tty) as shown in Figure 8-7.

```
[ec2-user@ip-172-30-1-86 ~]$ sudo docker exec -it hadoop bash
bash-4.1# cd $HADOOP_PREFIX
```

Figure 8-7. Setting Current Directory to \$HADOOP_PREFIX Directory

Create a directory called /input in the HDFS for the input files. Subsequently, set the directory permissions to global (777).

```
bash-4.1# bin/hdfs dfs -mkdir /input
bash-4.1# bin/hdfs dfs -chmod -R 777 /input
```

The preceding commands are also run from the interactive terminal as shown in Figure 8-8.

```
bash-4.1# bin/hdfs dfs -mkdir /input  
bash-4.1# bin/hdfs dfs -chmod -R 777 /input  
bash-4.1#
```

Figure 8-8. Creating Input Directory

Add two text files (`input1.txt` and `input2.txt`) with some sample text to the `/input` directory. To create a text file `input1.txt` run the following vi editor command in the tty.

```
vi input1.txt
```

Add the following two lines of text in the `input1.txt`.

```
Hello World Application for Apache Hadoop  
Hello World and Hello Apache Hadoop
```

Save the `input1.txt` file with the `:wq` command as shown in Figure 8-9.



The screenshot shows a terminal window titled "ec2-user@ip-172-30-1-86:~". The window contains the following text:

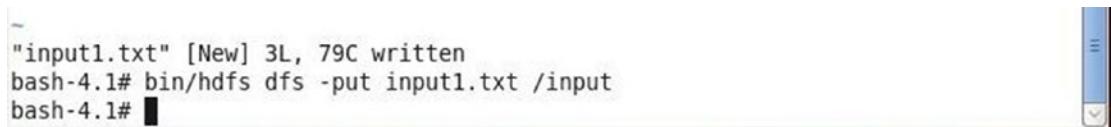
```
File Edit View Search Terminal Help  
Hello World Application for Apache Hadoop  
Hello World and Hello Apache Hadoop
```

The bottom of the terminal window shows the command `:wq` being typed.

Figure 8-9. The `input1.txt` File

Put the input1.txt file in the HDFS directory /input with the following command, also shown in Figure 8-10.

```
bin/hdfs dfs -put input1.txt /input
```



```
"input1.txt" [New] 3L, 79C written
bash-4.1# bin/hdfs dfs -put input1.txt /input
bash-4.1#
```

Figure 8-10. Putting the input1.txt in the HDFS

The input1.txt file gets added to the /input directory in the HDFS.

Similarly, open another new text file input2.txt with the following vi command.

```
vi input2.txt
```

Add the following two lines of text in the input2.txt file.

```
Hello World
Hello Apache Hadoop
```

Save the input2.txt file with the :wq command as shown in Figure 8-11.



```
ec2-user@ip-172-30-1-86:~$ vi input2.txt
File Edit View Search Terminal Help
Hello World
Hello Apache Hadoop

:wq
```

Figure 8-11. The input2.txt File

Put the `input2.txt` file in the HDFS directory `/input`.

```
bin/hdfs dfs -put input2.txt /input
```

Subsequently, run the following command to run the files in the `/input` directory.

```
bin/hdfs -ls /input
```

The two files added to the HDFS get listed as shown in Figure 8-12.

```
"input2.txt" [New] 3L, 33C written
bash-4.1# bin/hdfs dfs -put input2.txt /input
bash-4.1# bin/hdfs dfs -ls /input
Found 2 items
-rw-r--r-- 1 root supergroup      79 2015-10-18 15:43 /input/input1.txt
-rw-r--r-- 1 root supergroup      33 2015-10-18 15:45 /input/input2.txt
bash-4.1#
```

Figure 8-12. Listing the Input Files in the HDFS

Running a MapReduce Word Count Application

In this section we shall run a MapReduce application for word count; the application is packaged in the `hadoop-mapreduce-examples-2.7.0.jar` file and may be invoked with the arg “`wordcount`”. The `wordcount` application requires the input and output directories to be supplied. The input directory is the `/input` directory in the HDFS we created earlier and the output directory is `/output`, which must not exists before running the `hadoop` command. Run the following `hadoop` command from the interactive shell.

```
bin/hadoop jar $HADOOP_PREFIX/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.7.0.jar
wordcount /input /output
```

A MapReduce job gets started using the YARN framework as shown in Figure 8-13.

```
<mapreduce/hadoop-mapreduce-examples-2.7.0.jar wordcount /input /output
15/10/18 15:46:17 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0
:8032
15/10/18 15:46:19 INFO input.FileInputFormat: Total input paths to process : 2
15/10/18 15:46:19 INFO mapreduce.JobSubmitter: number of splits:2
15/10/18 15:46:20 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_14
45197241840_0001
15/10/18 15:46:21 INFO impl.YarnClientImpl: Submitted application application_14
45197241840_0001
15/10/18 15:46:21 INFO mapreduce.Job: The url to track the job: http://fb25c4cab
c55:8088/proxy/application_1445197241840_0001/
15/10/18 15:46:21 INFO mapreduce.Job: Running job: job_1445197241840_0001
```

Figure 8-13. Starting MapReduce Application with YARN Framework

The YARN job completes as shown in Figure 8-14, and the word count application gets output to the /output directory in the HDFS.

```

ec2-user@ip-172-30-1-86:~
File Edit View Search Terminal Help
Map input records=6
Map output records=17
Map output bytes=178
Map output materialized bytes=150
Input split bytes=212
Combine input records=17
Combine output records=11
Reduce input groups=7
Reduce shuffle bytes=150
Reduce input records=11
Reduce output records=7
Spilled Records=22
Shuffled Maps =2
Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=834
CPU time spent (ms)=2760
Physical memory (bytes) snapshot=540696576
Virtual memory (bytes) snapshot=2084392960
Total committed heap usage (bytes)=372310016
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=112
File Output Format Counters
Bytes Written=60
bash-4.1#

```

Figure 8-14. Output from the MapReduce Application

The complete output from the hadoop command is as follows.

```

<mapreduce/hadoop-mapreduce-examples-2.7.0.jar wordcount /input /output
15/10/18 15:46:17 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
15/10/18 15:46:19 INFO input.FileInputFormat: Total input paths to process : 2
15/10/18 15:46:19 INFO mapreduce.JobSubmitter: number of splits:2
15/10/18 15:46:20 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1445197241840_0001
15/10/18 15:46:21 INFO impl.YarnClientImpl: Submitted application
application_1445197241840_0001

```

```

15/10/18 15:46:21 INFO mapreduce.Job: The url to track the job: http://fb25c4cabcc5:8088/proxy/application\_1445197241840\_0001/
15/10/18 15:46:21 INFO mapreduce.Job: Running job: job_1445197241840_0001
15/10/18 15:46:40 INFO mapreduce.Job: Job job_1445197241840_0001 running in uber mode : false
15/10/18 15:46:40 INFO mapreduce.Job: map 0% reduce 0%
15/10/18 15:47:03 INFO mapreduce.Job: map 100% reduce 0%
15/10/18 15:47:17 INFO mapreduce.Job: map 100% reduce 100%
15/10/18 15:47:18 INFO mapreduce.Job: Job job_1445197241840_0001 completed successfully
15/10/18 15:47:18 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=144
    FILE: Number of bytes written=345668
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=324
    HDFS: Number of bytes written=60
    HDFS: Number of read operations=9
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=2
    Launched reduce tasks=1
    Data-local map tasks=2
    Total time spent by all maps in occupied slots (ms)=41338
    Total time spent by all reduces in occupied slots (ms)=11578
    Total time spent by all map tasks (ms)=41338
    Total time spent by all reduce tasks (ms)=11578
    Total vcore-seconds taken by all map tasks=41338
    Total vcore-seconds taken by all reduce tasks=11578
    Total megabyte-seconds taken by all map tasks=42330112
    Total megabyte-seconds taken by all reduce tasks=11855872
  Map-Reduce Framework
    Map input records=6
    Map output records=17
    Map output bytes=178
    Map output materialized bytes=150
    Input split bytes=212
    Combine input records=17
    Combine output records=11
    Reduce input groups=7
    Reduce shuffle bytes=150
    Reduce input records=11
    Reduce output records=7
    Spilled Records=22
    Shuffled Maps =2
    Failed Shuffles=0
    Merged Map outputs=2
    GC time elapsed (ms)=834
    CPU time spent (ms)=2760
    Physical memory (bytes) snapshot=540696576

```

```

Virtual memory (bytes) snapshot=2084392960
Total committed heap usage (bytes)=372310016
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=112
File Output Format Counters
  Bytes Written=60
bash-4.1#

```

List the output files in the /output directory in HDFS with the following command.

```
bin/hdfs dfs -ls /output
```

Two files get listed: _SUCCESS, which indicates that the YARN job completed successfully, and part-r-00000, which is the output from the wordcount application as shown in Figure 8-15.

```

bash-4.1# bin/hdfs dfs -ls /output
Found 2 items
-rw-r--r-- 1 root supergroup      0 2015-10-18 15:47 /output/_SUCCESS
-rw-r--r-- 1 root supergroup    60 2015-10-18 15:47 /output/part-r-00000
bash-4.1# █

```

Figure 8-15. Files Output by the YARN Application

List the output from the wordcount application using the following command.

```
hdfs dfs -cat /output/part-r-00000
```

The word count for each distinct word in the input files input1.txt and input2.txt gets output as shown in Figure 8-16.

```

bash-4.1# bin/hdfs dfs -ls /output
Found 2 items
-rw-r--r-- 1 root supergroup      0 2015-10-18 15:47 /output/_SUCCESS
-rw-r--r-- 1 root supergroup    60 2015-10-18 15:47 /output/part-r-00000
bash-4.1# bin/hdfs dfs -cat /output/part-r-00000
Apache 3
Application 1
Hadoop 3
Hello 5
World 3
and 1
for 1
bash-4.1# █

```

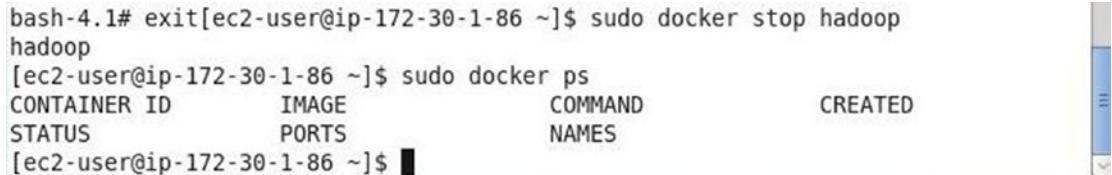
Figure 8-16. Listing the Word Count

Stopping the Hadoop Docker Container

The Docker container running the Hadoop processes may be stopped with the `docker stop` command.

```
sudo docker stop hadoop
```

Subsequently run the `docker ps` command and no container gets listed as running as shown in Figure 8-17.



```
bash-4.1# exit[ec2-user@ip-172-30-1-86 ~]$ sudo docker stop hadoop
hadoop
[ec2-user@ip-172-30-1-86 ~]$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED
STATUS              PORTS              NAMES
[ec2-user@ip-172-30-1-86 ~]$ █
```

Figure 8-17. Listing Running Docker Containers after stopping Apache Hadoop Container

Using a CDH Docker Image

As mentioned before several Docker images are available for Apache Hadoop. Another Docker image, which we shall also use in subsequent chapters based on the Apache Hadoop Ecosystem as packaged by the Cloudera Hadoop distribution called CDH, is the `svds/cdh` Docker image. The `svds/cdh` image includes not just Apache Hadoop but several frameworks in the Apache Hadoop ecosystem, some of which are discussed in later chapters. Download the `svds/cdh` image with the following command.

```
sudo docker pull svds/cdh
```

Start a Docker container running the CDH frameworks.

```
sudo docker run -d --name cdh svds/cdh
```

Start an interactive terminal to run commands for the CDH frameworks.

```
sudo docker exec -it cdh bash
```

In the tty, the Hadoop framework applications may be run without further configuration. For example, run the HDFS commands with “`hdfs`” on the command line. The `hdfs` commands usage is as listed as follows.

```
hdfs
```

The HDFS commands usage gets output as shown in Figure 8-18.



The screenshot shows a terminal window titled "root@86f0cf0a5c8d: /". The window contains the usage information for the hdfs command, which is used to run filesystem commands on HDFS. The usage information lists various sub-commands and their descriptions.

```

root@86f0cf0a5c8d:/# hdfs
Usage: hdfs [-config confdir] COMMAND
where COMMAND is one of:
  dfs          run a filesystem command on the file systems supported in
  Hadoop.
  namenode -format      format the DFS filesystem
  secondarynamenode    run the DFS secondary namenode
  namenode           run the DFS namenode
  journalnode        run the DFS journalnode
  zkfc              run the ZK Failover Controller daemon
  datanode          run a DFS datanode
  dfsadmin          run a DFS admin client
  haadmin           run a DFS HA admin client
  fsck              run a DFS filesystem checking utility
  balancer          run a cluster balancing utility
  jmxget            get JMX exported values from NameNode or DataNode.
  mover             run a utility to move block replicas across
                    storage types
  oiv               apply the offline fsimage viewer to an fsimage
  oiv_legacy        apply the offline fsimage viewer to an legacy fsimage
  oev               apply the offline edits viewer to an edits file
  fetchdt          fetch a delegation token from the NameNode
  getconf           get config values from configuration
  groups            get the groups which users belong to
  snapshotDiff     diff two snapshots of a directory or diff the
                    current directory contents with a snapshot
  lsSnapshottableDir list all snapshottable dirs owned by the current user
                     Use -help to see options
  portmap           run a portmap service
  nfs3              run an NFS version 3 gateway
  cacheadmin        configure the HDFS cache

```

Figure 8-18. *hdfs Command Usage*

The configuration files are available in the /etc/hadoop/conf symlink as shown in Figure 8-19.



The screenshot shows a terminal window listing the contents of the /etc/hadoop/conf directory. The output shows a symbolic link named "conf" pointing to /etc/alternatives/hadoop-conf.

```

root@86f0cf0a5c8d:/etc# cd hadoop
root@86f0cf0a5c8d:/etc/hadoop# ls -l
total 12
lrwxrwxrwx. 1 root root 29 Jul 5 13:34 conf -> /etc/alternatives/hadoop-conf
drwxr-xr-x. 2 root hadoop 4096 Jul 5 13:34 conf.empty
drwxr-xr-x. 2 root hadoop 4096 Jul 5 13:34 conf.impala
drwxr-xr-x. 2 root hadoop 4096 Jul 5 13:34 conf.pseudo
root@86f0cf0a5c8d:/etc/hadoop#

```

Figure 8-19. *Listing the Symlink for the Configuration Directory*

The configuration files in the `/etc/alternatives/hadoop-conf` directory to which the `conf` symlink points are listed as follows as shown in Figure 8-20.

```
root@86f0cf0a5c8d:/etc/hadoop# cd /etc/alternatives/hadoop-conf
root@86f0cf0a5c8d:/etc/alternatives/hadoop-conf# ls -l
total 40
-rw-r--r--. 1 root hadoop 1104 Jun 25 02:18 README
-rw-r--r--. 1 root hadoop 2133 Jun 25 02:18 core-site.xml
-rw-r--r--. 1 root hadoop 1366 Jun 25 02:18 hadoop-env.sh
-rw-r--r--. 1 root hadoop 2890 Jun 25 02:18 hadoop-metrics.properties
-rw-r--r--. 1 root hadoop 2324 Jun 25 02:56 hdfs-site.xml
-rw-r--r--. 1 root hadoop 11291 Jun 25 02:56 log4j.properties
-rw-r--r--. 1 root hadoop 1549 Jun 25 02:18 mapred-site.xml
-rw-r--r--. 1 root hadoop 2375 Jun 25 02:18 yarn-site.xml
root@86f0cf0a5c8d:/etc/alternatives/hadoop-conf#
```

Figure 8-20. Listing the Configuration Files

The `cdh` container may be stopped with the `docker stop` command.

```
sudo docker stop cdh
```

Summary

In this chapter we ran Apache Hadoop components in a Docker container. We created some files and put the files in the HDFS. Subsequently, we ran a MapReduce wordcount application packaged with the examples in the Hadoop distribution. We also introduced a Cloudera Hadoop distribution (CDH) based Docker image, which we shall also use in some of the subsequent chapters based on frameworks in the Apache Hadoop ecosystem.

CHAPTER 9



Using Apache Hive

Apache Hive is data warehouse framework for storing, managing and querying large data sets. The Hive query language HiveQL is a SQL-like language. Hive stores data in HDFS by default, and a Hive table may be used to define structure on the data. Hive supports two kinds of tables: managed tables and external tables. A managed table is managed by the Hive framework while an external table is not. When a managed table is deleted, the metadata and the table data are deleted. When a Hive external table is deleted, only the metadata is deleted, and the table data is not since the table data is not managed by the Hive framework. Hive makes use of a metastore to store metadata about Hive tables. A Hive metastore database is used for the metastore and is the Derby database by default. The metastore database may be run in *embedded mode* or *remote mode*; the default being embedded mode. In this chapter we shall use a Docker image to run Apache Hive in a Docker container.

- Setting the Environment
- Starting Apache Hive
- Connecting to Beeline CLI Shell
- Connecting to HiveServer2
- Creating a Hive Table
- Loading Data into Hive Table
- Querying Hive Table
- Stopping Apache Hive

Setting the Environment

The following software is required for this chapter.

- Docker (version 1.8 used)
- Docker image for Apache Hive

We have used an Amazon EC2 instance to install the software. Install Docker as discussed in Chapter 1. SSH connect to the Amazon EC2 instance.

```
ssh -i "docker.pem" ec2-user@52.23.241.186
```

Start the Docker service and verify status of the Docker service.

```
sudo service docker start
sudo service docker status
```

Download the svds/cdh Docker image, which is the same as used in some the other Apache Hadoop Ecosystem chapters on Apache HBase, Apache Sqoop and Apache Spark.

```
sudo docker pull svds/cdh
```

Starting Apache Hive

To start Apache Hive, start a Docker container running the cdh processes or components. Run the following docker run command, which starts a Docker container in detached mode and assigns the name “cdh” to the container.

```
sudo docker run -d --name cdh svds/cdh
```

List the running Docker containers; the “cdh” container should be listed.

```
sudo docker ps
```

Start an interactive terminal to run Apache Hive shell commands.

```
sudo docker exec -it cdh bash
```

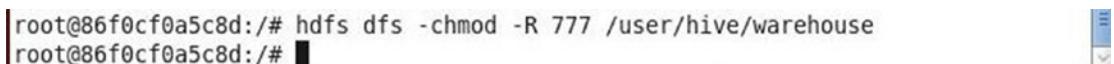
Connecting to Beeline CLI Shell

Apache Hive provides the Hive CLI to access HiveServer1 from a command line interface. In latter versions of Hive, HiveServer1 has been deprecated and replaced with HiveServer2, and Hive CLI has been deprecated and replaced with Beeline CLI. While Hive CLI is an Apache Thrift based client, Beeline is a JDBC client based on the SQLLine CLI. With Beeline, the Thrift API is still used but not directly from the client; the Thrift API is used by the JDBC driver to communicate with HiveServer2.

Before using the Hive CLI or the Beeline CLI, we need to modify the permissions for the directory in HDFS in which Hive stores its data, the /user/hive/warehouse directory. Set global permissions (777) on the /user/hive/warehouse directory.

```
hdfs dfs -chmod -R 777 /user/hive/warehouse
```

The preceding command is run in the interactive terminal as shown in Figure 9-1.




```
root@86f0cf0a5c8d:/# hdfs dfs -chmod -R 777 /user/hive/warehouse
root@86f0cf0a5c8d:/#
```

Figure 9-1. Setting Permissions on the Hive Warehouse Directory

If the Hive CLI is to be used, run the following command in the interactive terminal.

```
hive
```

The Hive CLI is started. A WARNING message is also output indicating that Hive CLI is deprecated and migration to Beeline is recommended as shown in Figure 9-2.




```
root@86f0cf0a5c8d:/# hive
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
hive> █
```

Figure 9-2. Message about Migration to Beeline

We shall use the Beeline CLI in this chapter. Exit from the Hive CLI with the exit or quit command. Start the Beeline CLI with the following command.

```
beeline
```

Beeline version 1.1.0 CDH 5.4.3 gets started as shown in Figure 9-3.



```
hive> root@86f0cf0a5c8d:/# beeline
Beeline version 1.1.0-cdh5.4.3 by Apache Hive
beeline> █
```

Figure 9-3. Starting Beeline

Connecting to HiveServer2

We started the Beeline CLI in the previous section, but we are not connected to the HiveServer2 yet. To demonstrate run the following commands.

```
use default;
show tables;
```

A “No current connection” message gets output as shown in Figure 9-4.



```
hive> root@86f0cf0a5c8d:/# beeline
Beeline version 1.1.0-cdh5.4.3 by Apache Hive
beeline> use default;
No current connection
beeline> show tables;
No current connection
beeline> █
```

Figure 9-4. Message “No Current Connection”

To connect to the HiveServer2, we need to run the !connect command. The !connect command usage may be output with the following command.

```
!connect
```

The !connect command usage gets output as shown in Figure 9-5.

```
beeline> !connect
Usage: connect <url> <username> <password> [driver]
```

Figure 9-5. Command Usage for !connect

HiveServer2 may be connected to in one of two modes: embedded or remote. The embedded mode may be used if the Beeline CLI is run on the same machine on which Hive is installed. The remote mode has to be used if the Beeline CLI is on a remote machine from the Hive. We shall use the embedded mode. The syntax for the connection url is the following in which the dbName is the Hive database and <host> and <port> are the hostname and port number for the HiveServer2.

```
jdbc:hive2://<host>:<port>/dbName
```

Run the following Beeline command !connect in which the connection url to HiveServer2 is specified first, followed by the username, password and the Hive JDBC driver. For the default username, password, and Hive JDBC driver specify an empty string "". The default Hive JDBC driver is org.apache.hive.jdbc.HiveDriver.

```
!connect jdbc:hive2://localhost:10000/default "" "" ""
```

A connection to Apache Hive 1.1.0 gets established as shown in Figure 9-6. Apache Hive 1.1.0 version is the renamed Hive 0.15.0 version.

```
root@86f0cf0a5c8d:/# beeline
Beeline version 1.1.0-cdh5.4.3 by Apache Hive
beeline> !connect jdbc:hive2://localhost:10000/default "" "" ""
scan complete in 16ms
Connecting to jdbc:hive2://localhost:10000/default
Connected to: Apache Hive (version 1.1.0-cdh5.4.3)
Driver: Hive JDBC (version 1.1.0-cdh5.4.3)
Transaction isolation: TRANSACTION_REPEATABLE_READ
0: jdbc:hive2://localhost:10000/default> █
```

Figure 9-6. Connecting with Hive2 Server

The Beeline commands that did not run previously get run after connecting to the HiveServer2. Run the following commands again to set the database as "default" and list the Hive tables.

```
use default
show tables
```

The database gets set to default and the Hive tables get listed. The database is already the “default” database as specified in the connection url and the `use default` command is run to demonstrate that the command gets run. No tables get listed as none have been created yet as shown in Figure 9-7. We shall create a table in the next section.

```
0: jdbc:hive2://localhost:10000/default> use default;
No rows affected (1.658 seconds)
0: jdbc:hive2://localhost:10000/default> show tables;
+-----+---+
| tab_name | 
+-----+---+
+-----+---+
No rows selected (0.57 seconds)
0: jdbc:hive2://localhost:10000/default>
```

Figure 9-7. Listing Tables

Creating a Hive Table

In this section we shall create a Hive table called “wlslog” with columns `time_stamp`, `category`, `type`, `servername`, `code` and `msg`, all of type `string`. Hive makes use of serializers/deserializers also called a *Serde*. A custom Serde may be used or the native Serde may be used. If a `ROW FORMAT` is not specified, the native Serde is used. If the `ROW FORMAT DELIMITED` is specified for delimited data files, the native Serde is used too. To separate fields with a `,` specify `FIELDS TERMINATED BY ','` and to terminate a line of data with a newline, specify `LINES TERMINATED BY '\n'`.

Run the following `CREATE TABLE` command to create a Hive managed table; the command for a Hive external table is `CREATE EXTERNAL TABLE`.

```
CREATE TABLE wlslog(time_stamp STRING,category STRING,type STRING,servername STRING,code STRING,msg STRING) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' LINES TERMINATED BY '\n';
```

A Hive table called `wlslog` gets created as shown in Figure 9-8. We have not used a `PRIMARY KEY` field in the `wlslog` table.

```
0: jdbc:hive2://localhost:10000/default> CREATE TABLE wlslog(time_stamp STRING ,category STRING, type STRING,servername STRING,code STRING,msg STRING) ROW FORM AT DELIMITED FIELDS TERMINATED BY ',' LINES TERMINATED BY '\n';
No rows affected (0.831 seconds)
0: jdbc:hive2://localhost:10000/default>
```

Figure 9-8. Creating Hive Table

Run the following command to describe the `wlslog` table.

```
desc wlslog;
```

The table structure consisting of the column names and data types gets listed as shown in Figure 9-9.

```
0: jdbc:hive2://localhost:10000/default> desc wlslog;
+-----+-----+-----+
| col_name | data_type | comment |
+-----+-----+-----+
| time_stamp | string |          |
| category | string |          |
| type | string |          |
| servername | string |          |
| code | string |          |
| msg | string |          |
+-----+-----+-----+
6 rows selected (0.389 seconds)
0: jdbc:hive2://localhost:10000/default>
```

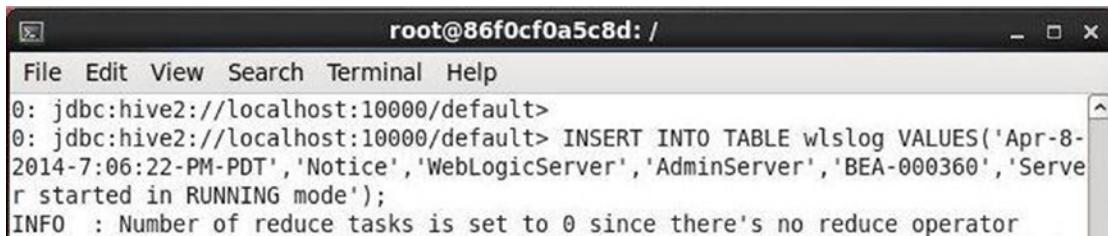
Figure 9-9. Describing Table Structure

Loading Data into the Hive Table

Next, we shall load data into the Hive table. Run the following INSERT HiveQL statement to add a row of data to the wlslog table.

```
INSERT INTO TABLE wlslog VALUES ('Apr-8-2014-7:06:16-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000365','Server state changed to STANDBY');
```

A MapReduce job gets started to load data into the Hive table as shown in Figure 9-10.



The screenshot shows a terminal window titled "root@86f0cf0a5c8d: /". The window contains the following text:

```
File Edit View Search Terminal Help
0: jdbc:hive2://localhost:10000/default>
0: jdbc:hive2://localhost:10000/default> INSERT INTO TABLE wlslog VALUES('Apr-8-2014-7:06:22-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000360','Server started in RUNNING mode');
INFO : Number of reduce tasks is set to 0 since there's no reduce operator
```

Figure 9-10. Running the INSERT Command

The MapReduce job consists of 1 mapper and 0 reducers. Data gets loaded into the default.wlslog table as shown in Figure 9-11.

```

INFO : Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
INFO : 2015-10-20 20:17:07,076 Stage-1 map = 0%, reduce = 0%
INFO : 2015-10-20 20:17:20,929 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.35 sec
INFO : MapReduce Total cumulative CPU time: 2 seconds 350 msec
INFO : Ended Job = job_1445367575035_0008
INFO : Stage-4 is selected by condition resolver.
INFO : Stage-3 is filtered out by condition resolver.
INFO : Stage-5 is filtered out by condition resolver.
INFO : Moving data to: hdfs://localhost:8020/user/hive/warehouse/wlslog/.hive-staging_hive_2015-10-20_20-16-49_585_3493713017925614717-2/-ext-10000 from hdfs://localhost:8020/user/hive/warehouse/wlslog/.hive-staging_hive_2015-10-20_20-16-49_585_3493713017925614717-2/-ext-10002
INFO : Loading data to table default.wlslog from hdfs://localhost:8020/user/hive/warehouse/wlslog/.hive-staging_hive_2015-10-20_20-16-49_585_3493713017925614717-2/-ext-10000
INFO : Table default.wlslog stats: [numFiles=8, numRows=8, totalSize=820, rawDataSize=812]
No rows affected (34.055 seconds)
0: jdbc:hive2://localhost:10000/default>

```

Figure 9-11. Loading Data into Hive Table

The data in a Hive table is not constrained to have unique column values if a PRIMARY KEY is not specified, which we did not. A row with the same data may be added without a PRIMARY KEY in the table definition. Run the following INSERT statements to add 7 more rows of data including a row of data with duplicate column data.

```

INSERT INTO TABLE wlslog VALUES ('Apr-8-2014-7:06:16-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000365','Server state changed to STANDBY');

INSERT INTO TABLE wlslog VALUES ('Apr-8-2014-7:06:17-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000365','Server state changed to STARTING');
INSERT INTO TABLE wlslog VALUES ('Apr-8-2014-7:06:18-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000365','Server state changed to ADMIN');
INSERT INTO TABLE wlslog VALUES ('Apr-8-2014-7:06:19-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000365','Server state changed to RESUMING');
INSERT INTO TABLE wlslog VALUES ('Apr-8-2014-7:06:20-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000331','Started WebLogic AdminServer');
INSERT INTO TABLE wlslog VALUES ('Apr-8-2014-7:06:21-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000365','Server state changed to RUNNING');
INSERT INTO TABLE wlslog VALUES ('Apr-8-2014-7:06:22-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000360','Server started in RUNNING mode');


```

Querying Hive Table

Having created a Hive table and loaded data into the table, we shall query the table using a SELECT HiveQL statement. Run the following query in the Beeline CLI.

```
select * from wlslog;
```

The 8 rows of data get listed as shown in Figure 9-12.



```
0: jdbc:hive2://localhost:10000/default> select * from wlslog;
+-----+-----+-----+-----+
| wlslog.time_stamp | wlslog.category | wlslog.type | wlslog.serve
rname | wlslog.code | wlslog.msg | |
+-----+-----+-----+-----+
| Apr-8-2014-7:06:16-PM-PDT | Notice | WebLogicServer | AdminServer
| BEA-000365 | Server state changed to STANDBY |
| Apr-8-2014-7:06:16-PM-PDT | Notice | WebLogicServer | AdminServer
| BEA-000365 | Server state changed to STANDBY |
| Apr-8-2014-7:06:17-PM-PDT | Notice | WebLogicServer | AdminServer
| BEA-000365 | Server state changed to STARTING |
| Apr-8-2014-7:06:18-PM-PDT | Notice | WebLogicServer | AdminServer
| BEA-000365 | Server state changed to ADMIN |
| Apr-8-2014-7:06:19-PM-PDT | Notice | WebLogicServer | AdminServer
| BEA-000365 | Server state changed to RESUMING |
| Apr-8-2014-7:06:20-PM-PDT | Notice | WebLogicServer | AdminServer
| BEA-000331 | Started WebLogic AdminServer |
| Apr-8-2014-7:06:21-PM-PDT | Notice | WebLogicServer | AdminServer
| BEA-000365 | Server state changed to RUNNING |
| Apr-8-2014-7:06:22-PM-PDT | Notice | WebLogicServer | AdminServer
| BEA-000360 | Server started in RUNNING mode |
+-----+-----+-----+-----+
8 rows selected (0.477 seconds)
0: jdbc:hive2://localhost:10000/default> █
```

Figure 9-12. Running a SELECT HiveQL Statement

Stopping Apache Hive

To stop the Apache Hive processes, run the `docker stop` command to stop the Docker container running the cdh frameworks.

```
sudo docker stop cdh
```

Summary

In this chapter we used a Docker image to run CDH frameworks including the Apache Hive framework in a Docker container. We started a Beeline CLI, which has replaced the Hive CLI and connected to the HiveServer2 from the Beeline CLI. We created a Hive managed table and loaded data into the Hive table. Subsequently, we queried the Hive table from the Beeline CLI. In the next chapter we shall use the Apache HBase Database in a Docker container.

CHAPTER 10



Using Apache HBase

Apache HBase is the Apache Hadoop database. Apache HBase is based on the wide column data store model with a table as the unit of storage. A table consists of one or more column families. Apache HBase is a schema-free NoSQL database. HBase stores data in the HDFS by default. In this chapter we shall use a Docker image to run Apache HBase in a Docker container. We shall use the svds/cdh Docker image, which we introduced in the Chapter 8.

Setting the Environment

Starting CDH

Starting Interactive Shell

Starting HBase Shell

Creating an HBase Table

Listing HBase Tables

Getting a Single Table Row

Getting a Single Row Column

Scanning a Table

Stopping CDH

Setting the Environment

The following software is required for this chapter.

-Docker (version 1.8 used)

-Docker Image for CDH

As in other chapters we have installed the software on an Amazon EC2 instance. SSH Login to the Amazon EC2 instance.

```
ssh -i "docker.pem" ec2-user@54.209.254.175
```

Start the Docker service.

```
sudo service docker start
```

Verify that Docker has started.

```
sudo service docker status
```

Download the svds/cdh Docker image if not already downloaded for the previous chapter.

```
sudo docker pull svds/cdh
```

The svds/cdh:latest Docker image gets downloaded as shown in Figure 10-1.

```
[ec2-user@ip-172-30-1-16 ~]$ sudo docker pull svds/cdh
Using default tag: latest
latest: Pulling from svds/cdh
de9c48daf08c: Pull complete
10de806794b2: Pull complete
031fd5268e85: Pull complete
0dc9ec408dd9: Pull complete
6adeac327f06: Pull complete
a4172715758c: Pull complete
6ec135e16988: Pull complete
d0aed77feb43: Pull complete
1e7f66f6f311: Pull complete
757082eae889: Pull complete
95119924abaa: Pull complete
4715b4ec211a: Pull complete
190d67765749: Pull complete
6bc9eaf71a3c: Pull complete
516bcbabdd00: Pull complete
c2ca5d25e887: Pull complete
81a9d08d9b11: Pull complete
8dc67cfb9d47: Pull complete
3f30741ab04b: Pull complete
b59a8b948d0b: Pull complete
6994ba23fe89: Pull complete
Digest: sha256:3e2a043497bdbb9b9b76f3193f830d255e1042ccb588b9052252818d7deaf83a
Status: Downloaded newer image for svds/cdh:latest
[ec2-user@ip-172-30-1-16 ~]$ █
```

Figure 10-1. Downloading the svds/cdh Docker Image

List the Docker images to verify that the svds/cdh image has been downloaded.

```
sudo docker images
```

Starting CDH

Start a Docker container to run the Apache Hadoop ecosystem frameworks, which include Apache HBase. Run the `docker run` command with the `-d` option, which starts the container in detached mode. The Docker container name is “cdh” as specified with the `-name` option.

```
sudo docker run -d --name cdh svds/cdh
```

Docker container gets started as shown in Figure 10-2.

```
[ec2-user@ip-172-30-1-16 ~]$ sudo docker run -d --name cdh svds/cdh
86f0cf0a5c8dd2d62878013bc47a1d9782538d8d93e4bceb4ea6e3ed3594add6
```

Figure 10-2. Starting Docker Container

List the running Docker containers.

```
sudo docker ps
```

The “cdh” container is listed as running as shown in Figure 10-3. The container id is also listed.

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
86f0cf0a5c8d	svds/cdh	"cdh_startup_script.s"	25 seconds ago
Up 24 seconds	8020/tcp, 8088/tcp, 8888/tcp, 9090/tcp, 11000/tcp, 1144	cdh	
3/tcp, 19888/tcp			

Figure 10-3. Listing the Running Docker Containers

Starting Interactive Shell

Next, start an interactive terminal (`tty`) to run the HBase shell in.

```
sudo docker exec -it cdh bash
```

An interactive terminal gets started and the command prompt becomes `root@86f0cf0a5c8d` as shown in Figure 10-4.

```
[ec2-user@ip-172-30-1-16 ~]$ sudo docker exec -it cdh bash
root@86f0cf0a5c8d:/#
```

Figure 10-4. Starting the Interactive Shell

The interactive shell may also be started using the container id instead of the container name.

```
sudo docker exec -it 86f0cf0a5c8d bash
```

Starting HBase Shell

Next, start the HBase shell with the following command run in the interactive terminal.

```
bin/hbase shell
```

HBase shell gets started as shown in Figure 10-5.

```
root@86f0cf0a5c8d:/# hbase shell
2015-10-20 19:07:03,381 INFO  [main] Configuration.deprecation: hadoop.native.lib is deprecated. Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.0.0-cdh5.4.3, rUnknown, Wed Jun 24 19:33:28 PDT 2015

hbase(main):001:0> █
```

Figure 10-5. Starting the HBase Shell

Creating a HBase Table

Create an HBase table using the “create” command. In addition to the table name, provide the column family or column families and a dictionary of specifications for each column family. Optionally, provide a dictionary of table configuration. As an example, create a table called ‘wlslog’ with a column family called ‘log’.

```
create 'wlslog' , 'log'
```

HBase table ‘wlslog’ gets created as shown in Figure 10-6.

```
hbase(main):001:0> create 'wlslog' , 'log'
0 row(s) in 1.5870 seconds

=> Hbase::Table - wlslog
hbase(main):002:0> █
```

Figure 10-6. Creating an HBase Table

Add cell values at table/row/column coordinates using the put command. Add 7 rows of data with the following put commands. Apache HBase and the other Apache Hadoop ecosystem software are designed for large quantities of data, which could be millions of rows of data, but only a sample of data is being added to demonstrate the use of Apache HBase.

```

put 'wlslog', 'log1', 'log:time_stamp', 'Apr-8-2014-7:06:16-PM-PDT'
put 'wlslog', 'log1', 'log:category', 'Notice'
put 'wlslog', 'log1', 'log:type', 'WeblogicServer'
put 'wlslog', 'log1', 'log:servername', 'AdminServer'
put 'wlslog', 'log1', 'log:code', 'BEA-000365'
put 'wlslog', 'log1', 'log:msg', 'Server state changed to STANDBY'

put 'wlslog', 'log2', 'log:time_stamp', 'Apr-8-2014-7:06:17-PM-PDT'
put 'wlslog', 'log2', 'log:category', 'Notice'
put 'wlslog', 'log2', 'log:type', 'WeblogicServer'
put 'wlslog', 'log2', 'log:servername', 'AdminServer'
put 'wlslog', 'log2', 'log:code', 'BEA-000365'
put 'wlslog', 'log2', 'log:msg', 'Server state changed to STARTING'

put 'wlslog', 'log3', 'log:time_stamp', 'Apr-8-2014-7:06:18-PM-PDT'
put 'wlslog', 'log3', 'log:category', 'Notice'
put 'wlslog', 'log3', 'log:type', 'WeblogicServer'
put 'wlslog', 'log3', 'log:servername', 'AdminServer'
put 'wlslog', 'log3', 'log:code', 'BEA-000365'
put 'wlslog', 'log3', 'log:msg', 'Server state changed to ADMIN'

put 'wlslog', 'log4', 'log:time_stamp', 'Apr-8-2014-7:06:19-PM-PDT'
put 'wlslog', 'log4', 'log:category', 'Notice'
put 'wlslog', 'log4', 'log:type', 'WeblogicServer'
put 'wlslog', 'log4', 'log:servername', 'AdminServer'
put 'wlslog', 'log4', 'log:code', 'BEA-000365'
put 'wlslog', 'log4', 'log:msg', 'Server state changed to RESUMING'

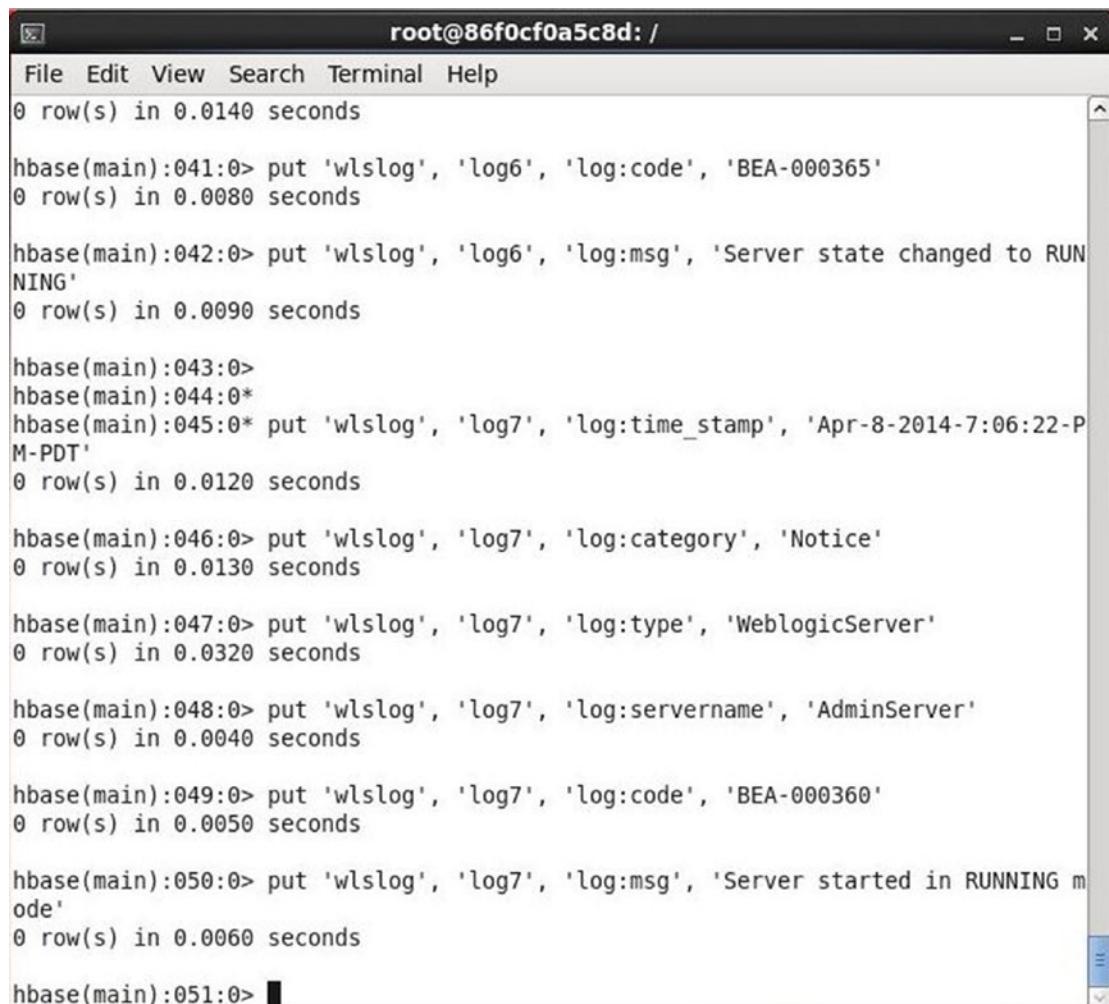
put 'wlslog', 'log5', 'log:time_stamp', 'Apr-8-2014-7:06:20-PM-PDT'
put 'wlslog', 'log5', 'log:category', 'Notice'
put 'wlslog', 'log5', 'log:type', 'WeblogicServer'
put 'wlslog', 'log5', 'log:servername', 'AdminServer'
put 'wlslog', 'log5', 'log:code', 'BEA-000331'
put 'wlslog', 'log5', 'log:msg', 'Started Weblogic AdminServer'

put 'wlslog', 'log6', 'log:time_stamp', 'Apr-8-2014-7:06:21-PM-PDT'
put 'wlslog', 'log6', 'log:category', 'Notice'
put 'wlslog', 'log6', 'log:type', 'WeblogicServer'
put 'wlslog', 'log6', 'log:servername', 'AdminServer'
put 'wlslog', 'log6', 'log:code', 'BEA-000365'
put 'wlslog', 'log6', 'log:msg', 'Server state changed to RUNNING'

put 'wlslog', 'log7', 'log:time_stamp', 'Apr-8-2014-7:06:22-PM-PDT'
put 'wlslog', 'log7', 'log:category', 'Notice'
put 'wlslog', 'log7', 'log:type', 'WeblogicServer'
put 'wlslog', 'log7', 'log:servername', 'AdminServer'
put 'wlslog', 'log7', 'log:code', 'BEA-000360'
put 'wlslog', 'log7', 'log:msg', 'Server started in RUNNING mode'

```

Data gets added to the 'wlslog' table as shown in Figure 10-7.



The screenshot shows a terminal window titled "root@86f0cf0a5c8d: /". The window contains the following HBase shell session:

```
File Edit View Search Terminal Help
0 row(s) in 0.0140 seconds
hbase(main):041:0> put 'wlslog', 'log6', 'log:code', 'BEA-000365'
0 row(s) in 0.0080 seconds
hbase(main):042:0> put 'wlslog', 'log6', 'log:msg', 'Server state changed to RUNNING'
0 row(s) in 0.0090 seconds
hbase(main):043:0>
hbase(main):044:0*
hbase(main):045:0* put 'wlslog', 'log7', 'log:time_stamp', 'Apr-8-2014-7:06:22-PM-PDT'
0 row(s) in 0.0120 seconds
hbase(main):046:0> put 'wlslog', 'log7', 'log:category', 'Notice'
0 row(s) in 0.0130 seconds
hbase(main):047:0> put 'wlslog', 'log7', 'log:type', 'WeblogicServer'
0 row(s) in 0.0320 seconds
hbase(main):048:0> put 'wlslog', 'log7', 'log:servername', 'AdminServer'
0 row(s) in 0.0040 seconds
hbase(main):049:0> put 'wlslog', 'log7', 'log:code', 'BEA-000360'
0 row(s) in 0.0050 seconds
hbase(main):050:0> put 'wlslog', 'log7', 'log:msg', 'Server started in RUNNING mode'
0 row(s) in 0.0060 seconds
hbase(main):051:0> █
```

Figure 10-7. Adding Data to HBase Table

Listing HBase Tables

List the tables with the following command run in HBase shell.

```
list
```

One table, the 'wlslog' table, gets listed as shown in Figure 10-8.

```
hbase(main):051:0> list
TABLE
wlslog
1 row(s) in 0.0430 seconds

=> [ "wlslog"]
hbase(main):052:0> ■
```

Figure 10-8. Listing HBase Tables

Getting A Single Table Row

The get command is used to get the data in a row or a column cell. Run the following get command to get the data in row 'log7' in table 'wlslog'.

```
get 'wlslog', 'log7'
```

A single row of data gets listed as shown in Figure 10-9.

```
hbase(main):052:0> get 'wlslog', 'log7'
COLUMN          CELL
log:category   timestamp=1445368199624, value=Notice
log:code        timestamp=1445368199778, value=BEA-000360
log:msg         timestamp=1445368199828, value=Server started in RUNNING m
                ode
log:servername timestamp=1445368199739, value=AdminServer
log:time_stamp  timestamp=1445368199589, value=Apr-8-2014-7:06:22-PM-PDT
log:type        timestamp=1445368199687, value=WeblogicServer
6 row(s) in 0.0370 seconds

hbase(main):053:0> ■
```

Figure 10-9. Getting a Single Table Row

Getting A Single Row Column

Optionally, a dictionary of columns may be supplied to the get command. For example, get the column data from the wlslog table from the 'log5' row in the log.msg column.

```
get 'wlslog', 'log5', {COLUMNS=>[ 'log:msg' ]}
```

The log.msg column data from row 'log5' from table 'wlslog' gets output as shown in Figure 10-10.

```
hbase(main):053:0> get 'wlslog', 'log5', {COLUMNS=>['log:msg']}
COLUMN          CELL
log:msg         timestamp=1445368199142, value=Started Weblogic AdminServer
                r
1 row(s) in 0.0090 seconds

hbase(main):054:0>
```

Figure 10-10. Getting a Single Row Column Value

Scanning a Table

The scan command is used to scan a table to get all the data in the table. Optionally a dictionary of scanner specifications may be provided, which are omitted from the following command.

```
scan 'wlslog'
```

Row ➤ column data for each row gets output as shown in Figure 10-11.

```
root@86f0cf0a5c8d: /
```

```
File Edit View Search Terminal Help
hbase(main):054:0> scan 'wlslog'
ROW          COLUMN+CELL
log1        column=log:category, timestamp=1445368158177, value=Notice
log1        column=log:code, timestamp=1445368158796, value=BEA-000365
log1        column=log:msg, timestamp=1445368158901, value=Server stat
e changed to STANDBY
log1        column=log:servername, timestamp=1445368158598, value=Admi
nServer
log1        column=log:time_stamp, timestamp=1445368157863, value=Apr-
8-2014-7:06:16-PM-PDT
log1        column=log:type, timestamp=1445368158336, value=WeblogicSe
rver
log2        column=log:category, timestamp=1445368159215, value=Notice
log2        column=log:code, timestamp=1445368159768, value=BEA-000365
log2        column=log:msg, timestamp=1445368159953, value=Server stat
e changed to STARTING
log2        column=log:servername, timestamp=1445368159600, value=Admi
nServer
log2        column=log:time_stamp, timestamp=1445368159031, value=Apr-
8-2014-7:06:17-PM-PDT
log2        column=log:type, timestamp=1445368159442, value=WeblogicSe
rver
log3        column=log:category, timestamp=1445368175925, value=Notice
log3        column=log:code, timestamp=1445368176305, value=BEA-000365
log3        column=log:msg, timestamp=1445368176360, value=Server stat
e changed to ADMIN
log3        column=log:servername, timestamp=1445368176160, value=Admi
nServer
log3        column=log:time_stamp, timestamp=1445368175884, value=Apr-
8-2014-7:06:18-PM-PDT
log3        column=log:type, timestamp=1445368175981, value=WeblogicSe
rver
```

Figure 10-11. Scanning a HBase Table

The 7 rows of data get output as shown in Figure 10-12.



```
root@86f0cf0a5c8d: / 
File Edit View Search Terminal Help
log5          column=log:msg, timestamp=1445368199142, value=Started Weblogic AdminServer
log5          column=log:servername, timestamp=1445368199067, value=AdminServer
log5          column=log:time_stamp, timestamp=1445368198909, value=Apr-8-2014-7:06:20-PM-PDT
log5          column=log:type, timestamp=1445368199027, value=WeblogicServer
log6          column=log:category, timestamp=1445368199277, value=Notice
log6          column=log:code, timestamp=1445368199497, value=BEA-000365
log6          column=log:msg, timestamp=1445368199530, value=Server state changed to RUNNING
log6          column=log:servername, timestamp=1445368199459, value=AdminServer
log6          column=log:time_stamp, timestamp=1445368199208, value=Apr-8-2014-7:06:21-PM-PDT
log6          column=log:type, timestamp=1445368199375, value=WeblogicServer
log7          column=log:category, timestamp=1445368199624, value=Notice
log7          column=log:code, timestamp=1445368199778, value=BEA-000360
log7          column=log:msg, timestamp=1445368199828, value=Server started in RUNNING mode
log7          column=log:servername, timestamp=1445368199739, value=AdminServer
log7          column=log:time_stamp, timestamp=1445368199589, value=Apr-8-2014-7:06:22-PM-PDT
log7          column=log:type, timestamp=1445368199687, value=WeblogicServer
7 row(s) in 0.4270 seconds
hbase(main):055:0> exit
root@86f0cf0a5c8d:/#
```

Figure 10-12. Output from the scan Command

Stopping CDH

To stop the Docker container, run the docker stop command for the “cdh” container.

```
sudo docker stop cdh
```

Alternatively, the container id may be specified.

```
sudo docker stop 86f0cf0a5c8d
```

Summary

In this chapter we used a Docker image to run CDH frameworks in a Docker container. We started an interactive terminal and started an HBase shell in the tty. In the HBase shell, we used the `create` command to create a table. We used the `put` command to put data in the table. Subsequently, we used the `get` command to get the data added. We also ran the `scan` command to scan the complete table and list all the data in the table. In the next chapter we shall run Apache Sqoop in a Docker container.

CHAPTER 11



Using Apache Sqoop

Apache Sqoop is a Hadoop ecosystem framework for transferring bulk data from a relational database (RDBMS) to Hadoop Distributed File System (HDFS), Apache HBase, and Apache Hive. Sqoop also supports bulk data transfer from HDFS to a RDBMS. The direct data transfer paths supported by Sqoop are shown in Figure 11-1. Sqoop supports HSQLDB (version 1.8.0+), MySQL (5.0+), Oracle (10.2.0) and PostgreSQL (8.3+) and may also be usable with other relational databases such as IBM DB2 database and versions. Sqoop makes use of JDBC for data transfer and requires Java to be installed and the JDBC driver jar to be in the runtime classpath.

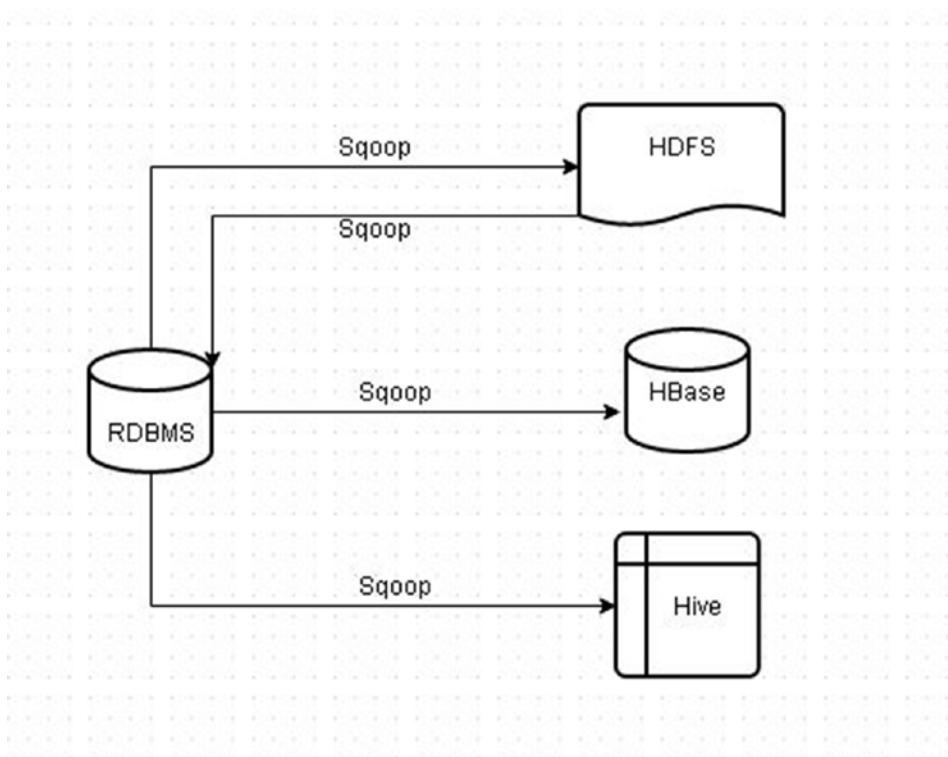


Figure 11-1. Direct Transfer Paths supported by Sqoop

In this chapter we shall use Apache Sqoop to import data into HDFS from MySQL database. We shall also export the data from HDFS back to a MySQL database table.

- Setting the Environment
- Starting Docker Containers
- Starting Interactive Terminals
- Creating a MySQL Tables
- Adding MySQL JDBC Jar to Sqoop Classpath
- Configuring Apache Hadoop
- Importing MySQL Table Data into HDFS with Sqoop
- Listing Data Imported into HDFS
- Exporting from HDFS to MySQL with Sqoop
- Querying Exported Data
- Stopping and Removing Docker Containers

Setting the Environment

The following software is required for this chapter.

- Docker Engine (version 1.8)
- Docker image for MySQL Database
- Docker image for CDH

SSH connect to an Amazon EC2 instance.

```
ssh -i "docker.pem" ec2-user@54.175.13.99
```

Install Docker if not already installed as discussed in Chapter 1. Start the Docker service and verify that Docker has been started.

```
sudo service docker start
sudo service docker status
```

Download jdk-8u65-linux-x64.gz from <http://www.oracle.com/technetwork/java/javase/downloads/jdk8-downloads-2133151.html>. As JDK download requires a BSD license to be accepted, downloading with wget or similar software for downloading files makes the download command a non-standard command. Download jdk-8u65-linux-x64.gz using a browser and copy to the EC2 instance using a scp command such as the following.

```
scp -i "docker.pem" /jdk-8u65-linux-x64.gz ec2-user@54.175.13.99:/
```

We need to download two Docker images for this chapter because the Docker image for CDH, which includes Apache Sqoop, does not include MySQL Server. Download the mysql Docker image with the docker pull command.

```
sudo docker pull mysql
```

Download the svds/cdh Docker image.

```
sudo docker pull svds/cdh
```

List the Docker images with the `docker images` command.

```
sudo docker images
```

Both the mysql and svds/cdh Docker images should get listed as shown in Figure 11-2.

REPOSITORY	TAG	IMAGE ID	CREATED
VIRTUAL SIZE			
couchbase	latest	ff61ecf3bacb	4 days ago
371.3 MB			
mysql	latest	d7da97aedce5	7 days ago
324.3 MB			
svds/cdh	latest	6994ba23fe89	3 months ago
2.647 GB			

Figure 11-2. Listing Docker Images Required for Apache Sqoop with MySQL Database

Starting Docker Containers

Both the mysql and svds/cdh Docker images have been discussed in earlier chapters separately and used to start Docker containers. But, using the two Docker images is slightly different and requires the two Docker containers to be linked. In this section we shall start two separate Docker containers: cdh for the cdh Docker image, and mysqldb for the mysql Docker image. For the mysqldb container, create a directory for the data stored by MySQL and set its permissions to global (777).

```
sudo mkdir -p /mysql/data
sudo chmod -R 777 /mysql/data
```

The preceding commands are to be run when connected to the Amazon EC2 instance as shown in Figure 11-3.

```
[ec2-user@ip-172-30-1-16 ~]$ sudo mkdir -p /mysql/data
[ec2-user@ip-172-30-1-16 ~]$ sudo chmod -R 777 /mysql/data
```

Figure 11-3. Creating Directory for MySQL Data

The environment variables used in the `docker run` command are discussed in the following table, Table 11-1.

Table 11-1. Environment Variables for a Docker container based on mysql Docker Image

Environment Variable	Description	Value
MYSQL_DATABASE	MySQL database instance to be created.	mysqldb
MYSQL_USER	Username for the database created.	mysql
MYSQL_PASSWORD	Password for the database created.	mysql
MYSQL_ALLOW_EMPTY_PASSWORD	Is empty password to be allowed.	no
MYSQL_ROOT_PASSWORD	Password for “root” user.	mysql

Run the following `docker run` command to start a Docker container for MySQL Database. The environment variables are only set in the `docker run` command and not in the bash shell.

```
sudo docker run -v /mysql/data:/var/lib/mysql --name mysqldb -e MYSQL_DATABASE='mysqldb'
-e MYSQL_USER='mysql' -e MYSQL_PASSWORD='mysql' -e MYSQL_ALLOW_EMPTY_PASSWORD='no'
-e MYSQL_ROOT_PASSWORD='mysql' -d mysql
```

Run the following `docker run` command to start a Docker container for svds/cdh image software, which includes Apache Sqoop, and link the container with the `mysqldb` container running the MySQL database using the `--link` command parameter.

```
sudo docker run -d --name cdh --link mysqldb svds/cdh
```

List the running Docker containers.

```
sudo docker ps
```

The output from the preceding commands is shown in Figure 11-4. Both the `cdh` and `mysqldb` containers are listed as started.

```
[ec2-user@ip-172-30-1-16 ~]$ sudo docker run -v /mysql/data:/var/lib/mysql --name mysql ldb -e MYSQL_DATABASE='mysql' -e MYSQL_USER='mysql' -e MYSQL_PASSWORD='mysql' -e MYSQL_ALLOW_EMPTY_PASSWORD='no' -e MYSQL_ROOT_PASSWORD='mysql' -d mysql e414f8c41d0b931f4ea7eec13d2c84ed8d50ea7c37e9f0078bcff56243d8d91
[ec2-user@ip-172-30-1-16 ~]$ sudo docker run -d --name cdh -e JAVA_HOME='./jdk1.8.0_65' --link mysql:svldb svds/cdh
49d774f8f1feae0c960999ced2025b48c115eb4d6b09cee21b931a3f989b431a
[ec2-user@ip-172-30-1-16 ~]$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS
NAMES
49d774f8f1fe        svds/cdh          "cdh_startup_script.s"   5 seconds ago
Up 4 seconds        8020/tcp, 8088/tcp, 8888/tcp, 9090/tcp, 11000/tcp, 1144
3/tcp, 19888/tcp    cdh
e414f8c41d0b        mysql              "/entrypoint.sh mysql"   18 seconds ago
Up 18 seconds       3306/tcp
mysqldb
[ec2-user@ip-172-30-1-16 ~]$ █
```

Figure 11-4. Starting Docker Containers for CDH and MySQL

Starting Interactive Terminals

Having started the Docker containers, start the interactive terminals (tty) for each of the Docker containers. Start the interactive shell for the mysqldb container with the following command.

```
sudo docker exec -it mysqldb bash
```

Start the interactive shell for the cdh container with the following command.

```
sudo docker exec -it cdh bash
```

Creating a MySQL Tables

In this section we shall login to the MySQL CLI and create a database table, which shall be imported into HDFS with Apache Sqoop. Run the following command to login into MySQL CLI.

```
mysql -u mysql -p
```

The `mysql>` prompt gets displayed as shown in Figure 11-5.

```
root@e414f8c41d0b:/# mysql -u mysql -p
Enter password:
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 4
Server version: 5.6.27 MySQL Community Server (GPL)

Copyright (c) 2000, 2015, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> ■
```

Figure 11-5. Starting the MySQL CLI Shell

Set the database to use as “mysqldb”.

```
use mysqldb
```

Grant all privileges on the mysqldb database to the mysql user with the GRANT option.

```
GRANT ALL PRIVILEGES ON mysqldb.* TO 'mysql'@'%' IDENTIFIED BY 'mysql' WITH GRANT OPTION;
```

Privileges get set on the mysqldb database as shown in Figure 11-6.

```
mysql> GRANT ALL PRIVILEGES ON mysqldb.* TO 'mysql'@'%' IDENTIFIED BY 'mysql' WI
TH GRANT OPTION;
Query OK, 0 rows affected (0.00 sec)

mysql> ■
```

Figure 11-6. Setting Privileges on mysqldb Database

Next, create a database table called wlslog with columns time_stamp, category, type, servername, code and msg. The PRIMARY KEY column is required to be included for sqoop import tool to import data into HDFS. Run the following SQL command in MySQL CLI.

```
CREATE TABLE wlslog(time_stamp VARCHAR(255) PRIMARY KEY,category VARCHAR(255),type
VARCHAR(255),servername VARCHAR(255), code VARCHAR(255),msg VARCHAR(255));
```

A database table called wlslog gets created as shown in Figure 11-7.



```
mysql> CREATE TABLE wlslog(time_stamp VARCHAR(255) PRIMARY KEY,category VARCHAR(255),type VARCHAR(255),servername VARCHAR(255), code VARCHAR(255),msg VARCHAR(255));
Query OK, 0 rows affected (0.02 sec)

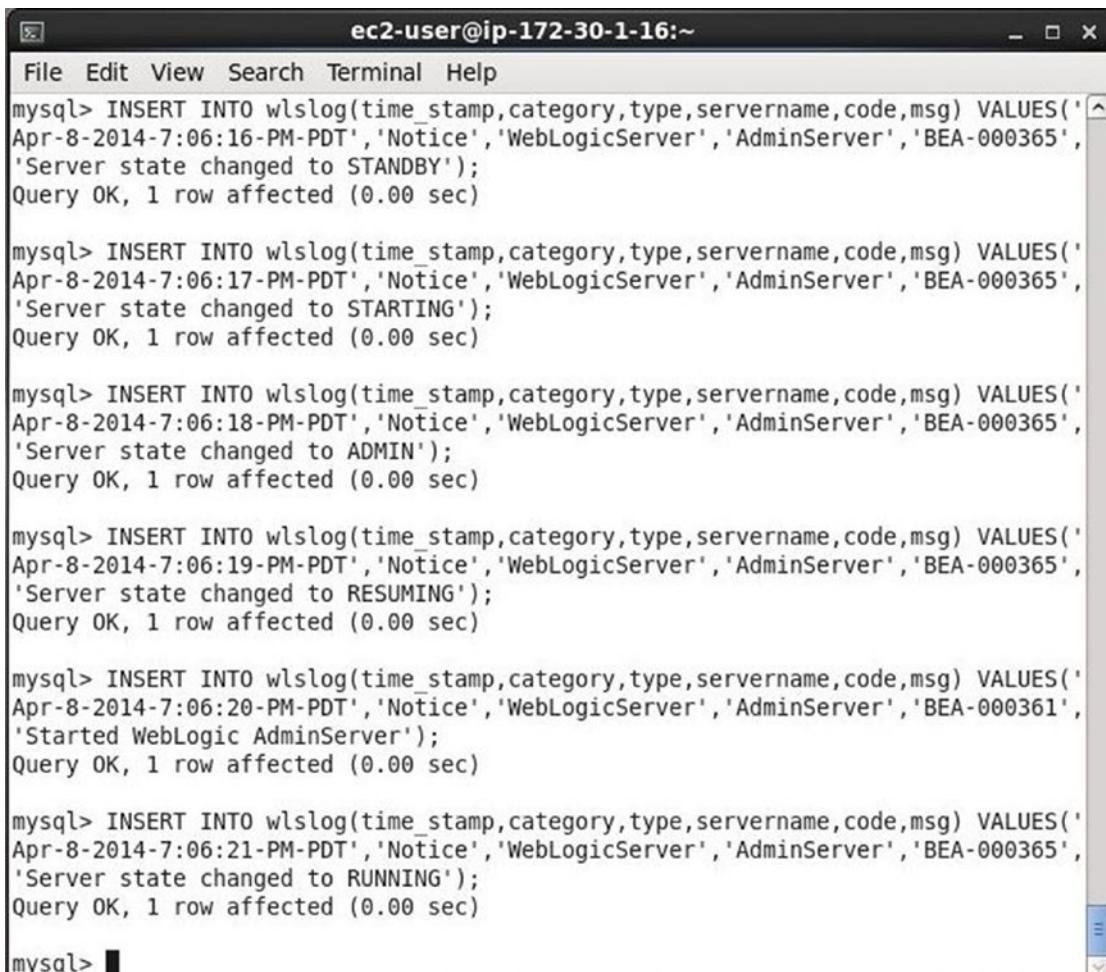
mysql>
```

Figure 11-7. Creating a MySQL Database Table

Add data to the wlslog table. Run the following INSERT SQL statements to add data to the wlslog table.

```
INSERT INTO wlslog(time_stamp,category,type,servername,code,msg) VALUES('Apr-8-2014-7:06:16-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000365','Server state changed to STANDBY');
INSERT INTO wlslog(time_stamp,category,type,servername,code,msg) VALUES('Apr-8-2014-7:06:17-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000365','Server state changed to STARTING');
INSERT INTO wlslog(time_stamp,category,type,servername,code,msg) VALUES('Apr-8-2014-7:06:18-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000365','Server state changed to ADMIN');
INSERT INTO wlslog(time_stamp,category,type,servername,code,msg) VALUES('Apr-8-2014-7:06:19-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000365','Server state changed to RESUMING');
INSERT INTO wlslog(time_stamp,category,type,servername,code,msg) VALUES('Apr-8-2014-7:06:20-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000361','Started WebLogic AdminServer');
INSERT INTO wlslog(time_stamp,category,type,servername,code,msg) VALUES('Apr-8-2014-7:06:21-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000365','Server state changed to RUNNING');
```

Output from the preceding SQL statements is shown in Figure 11-8.



The screenshot shows a terminal window titled "ec2-user@ip-172-30-1-16:~". The user is running several MySQL INSERT statements into a table named "wlslog". The statements describe the state changes of a WebLogic AdminServer over time. The log entries include transitions from STANDBY to STARTING, then to ADMIN, then to RESUMING, and finally to RUNNING.

```
File Edit View Search Terminal Help
mysql> INSERT INTO wlslog(time_stamp,category,type,servername,code,msg) VALUES('Apr-8-2014-7:06:16-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000365','Server state changed to STANDBY');
Query OK, 1 row affected (0.00 sec)

mysql> INSERT INTO wlslog(time_stamp,category,type,servername,code,msg) VALUES('Apr-8-2014-7:06:17-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000365','Server state changed to STARTING');
Query OK, 1 row affected (0.00 sec)

mysql> INSERT INTO wlslog(time_stamp,category,type,servername,code,msg) VALUES('Apr-8-2014-7:06:18-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000365','Server state changed to ADMIN');
Query OK, 1 row affected (0.00 sec)

mysql> INSERT INTO wlslog(time_stamp,category,type,servername,code,msg) VALUES('Apr-8-2014-7:06:19-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000365','Server state changed to RESUMING');
Query OK, 1 row affected (0.00 sec)

mysql> INSERT INTO wlslog(time_stamp,category,type,servername,code,msg) VALUES('Apr-8-2014-7:06:20-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000361','Started WebLogic AdminServer');
Query OK, 1 row affected (0.00 sec)

mysql> INSERT INTO wlslog(time_stamp,category,type,servername,code,msg) VALUES('Apr-8-2014-7:06:21-PM-PDT','Notice','WebLogicServer','AdminServer','BEA-000365','Server state changed to RUNNING');
Query OK, 1 row affected (0.00 sec)

mysql>
```

Figure 11-8. Running INSERT SQL Statements

Run the following SQL query to list the data added.

```
SELECT * FROM wlslog;
```

The 6 rows of data get listed as shown in Figure 11-9.

```
mysql> select * from wlslog;
+-----+-----+-----+-----+
| time_stamp | category | type | servername | code |
| msg | | | | |
+-----+-----+-----+-----+
| Apr-8-2014-7:06:16-PM-PDT | Notice | WebLogicServer | AdminServer | BEA-0003
65 | Server state changed to STANDBY |
| Apr-8-2014-7:06:17-PM-PDT | Notice | WebLogicServer | AdminServer | BEA-0003
65 | Server state changed to STARTING |
| Apr-8-2014-7:06:18-PM-PDT | Notice | WebLogicServer | AdminServer | BEA-0003
65 | Server state changed to ADMIN |
| Apr-8-2014-7:06:19-PM-PDT | Notice | WebLogicServer | AdminServer | BEA-0003
65 | Server state changed to RESUMING |
| Apr-8-2014-7:06:20-PM-PDT | Notice | WebLogicServer | AdminServer | BEA-0003
61 | Started WebLogic AdminServer |
| Apr-8-2014-7:06:21-PM-PDT | Notice | WebLogicServer | AdminServer | BEA-0003
65 | Server state changed to RUNNING |
+-----+-----+-----+-----+
-----+
6 rows in set (0.00 sec)

mysql> █
```

Figure 11-9. Running a SQL Query

We need to create another database table for the sqoop export tool to export data from HDFS into MySQL database. Because the wlslog table already has data create another table called WLSLOG_COPY, which has the same table definition as the wlslog table. Run the following SQL script in MySQL CLI.

```
CREATE TABLE WLSLOG_COPY(time_stamp VARCHAR(255) PRIMARY KEY,category VARCHAR(255),type VARCHAR(255),servername VARCHAR(255), code VARCHAR(255),msg VARCHAR(255));
```

The WLSLOG_COPY table gets created as shown in Figure 11-10.

```
mysql> CREATE TABLE WLSLOG_COPY(time_stamp VARCHAR(255) PRIMARY KEY,category VAR  
CHAR(255),type VARCHAR(255),servername VARCHAR(255), code VARCHAR(255),msg VARCH  
AR(255));
Query OK, 0 rows affected (0.13 sec)

mysql> select * from WLSLOG_COPY;
Empty set (0.00 sec)

mysql> █
```

Figure 11-10. Creating MySQL Table WLSLOG_COPY

Adding MySQL JDBC Jar to Sqoop Classpath

We need to add the MySQL JDBC jar to the Apache Sqoop classpath. Start the interactive terminal for the cdh container if not already started.

```
sudo docker exec -it cdh bash
```

In the interactive shell, download the mysql-connector-java-5.1.37.jar and copy the jar to the /usr/lib/sqoop/lib directory.

```
wget http://central.maven.org/maven2/mysql/mysql-connector-java/5.1.37/
mysql-connector-java-5.1.37.jar
cp mysql-connector-java-5.1.37.jar /usr/lib/sqoop/lib
```

The output from the preceding commands is shown in Figure 11-11.

```
[ec2-user@ip-172-30-1-16 /]$ sudo docker exec -it cdh bash
root@6fba20d93011:/# wget http://central.maven.org/maven2/mysql/mysql-connector-
java/5.1.37/mysql-connector-java-5.1.37.jar
--2015-10-21 19:42:43-- http://central.maven.org/maven2/mysql/mysql-connector-j
ava/5.1.37/mysql-connector-java-5.1.37.jar
Resolving central.maven.org (central.maven.org)... 23.235.39.209
Connecting to central.maven.org (central.maven.org)|23.235.39.209|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 985600 (962K) [application/java-archive]
Saving to: 'mysql-connector-java-5.1.37.jar'

100%[=====] 985,600      --.-K/s   in 0.1s

2015-10-21 19:42:43 (7.50 MB/s) - 'mysql-connector-java-5.1.37.jar' saved [985600/985600]

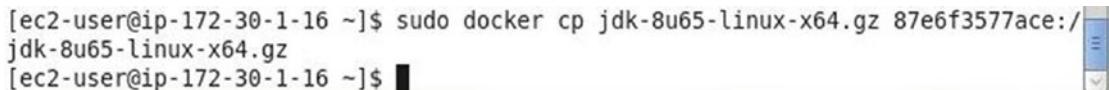
root@6fba20d93011:/# cp mysql-connector-java-5.1.37.jar /usr/lib/sqoop/lib
```

Figure 11-11. Adding MySQL JDBC Jar to Sqoop Classpath

Setting the JAVA_HOME Environment Variable

For the Apache Sqoop to run we need to set the JAVA_HOME environment variable. But, first we need to copy the jdk-8u65-linux-x64.gz file to the Docker container running the CDH frameworks including Apache Sqoop. We downloaded the jdk-8u65-linux-x64.gz earlier. Copy the jdk-8u65-linux-x64.gz file to the Docker container using the following command in which the container id is obtained from the output of the docker ps command in Figure 11-12.

```
sudo docker cp jdk-8u65-linux-x64.gz 49d774f8f1fe:/jdk-8u65-linux-x64.gz
```



```
[ec2-user@ip-172-30-1-16 ~]$ sudo docker cp jdk-8u65-linux-x64.gz 87e6f3577ace:/
jdk-8u65-linux-x64.gz
[ec2-user@ip-172-30-1-16 ~]$ █
```

Figure 11-12. Copying the JDK gz File to Docker Container

The jdk-8u65-linux-x64.gz file gets copied to the Docker container “cdh” as shown in Figure 11-12.

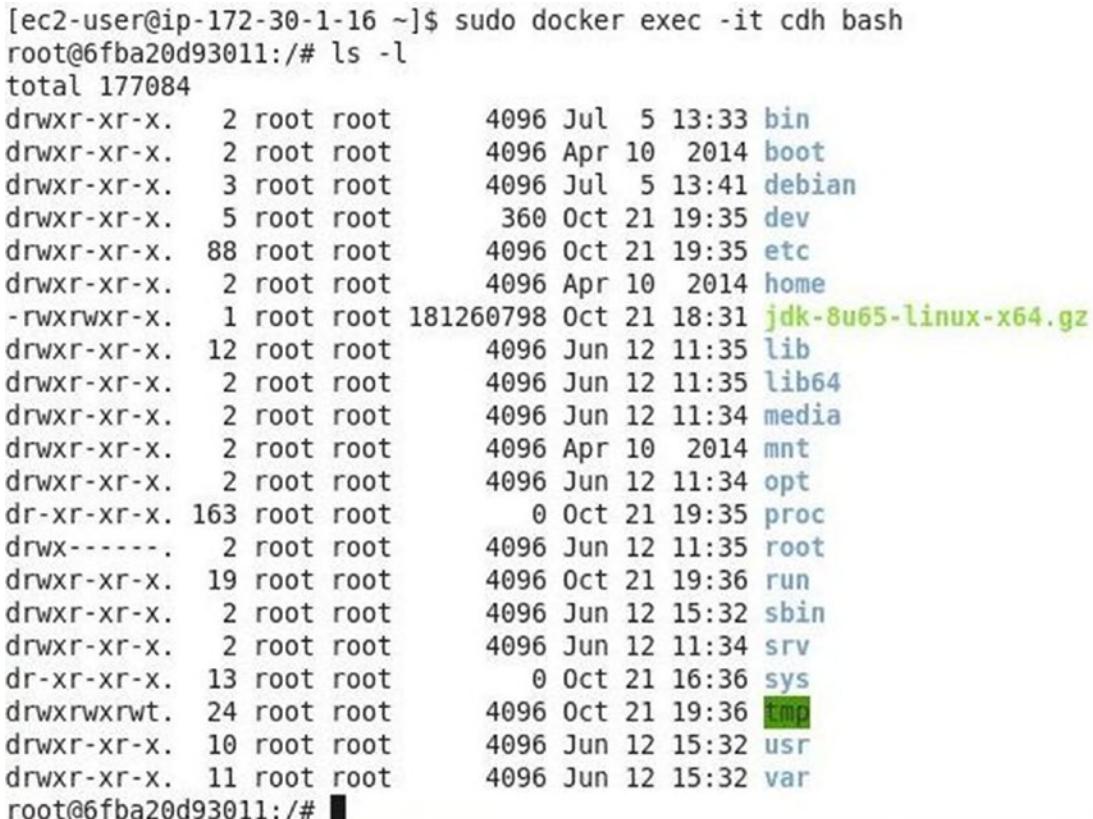
The preceding command is to be run from the Amazon EC2 instance. Start the interactive shell for the cdh container.

```
sudo docker exec -it cdh bash
```

List the files in the Docker container’s root directory with the following command.

```
ls -l
```

The jdk-8u65-linux-x64.gz file gets listed as shown in Figure 11-13.



```
[ec2-user@ip-172-30-1-16 ~]$ sudo docker exec -it cdh bash
root@6fba20d93011:/# ls -l
total 177084
drwxr-xr-x.  2 root root      4096 Jul  5 13:33 bin
drwxr-xr-x.  2 root root      4096 Apr 10 2014 boot
drwxr-xr-x.  3 root root      4096 Jul  5 13:41 debian
drwxr-xr-x.  5 root root      360 Oct 21 19:35 dev
drwxr-xr-x.  88 root root     4096 Oct 21 19:35 etc
drwxr-xr-x.  2 root root      4096 Apr 10 2014 home
-rwxrwxr-x.  1 root root 181260798 Oct 21 18:31 jdk-8u65-linux-x64.gz
drwxr-xr-x. 12 root root     4096 Jun 12 11:35 lib
drwxr-xr-x.  2 root root     4096 Jun 12 11:35 lib64
drwxr-xr-x.  2 root root     4096 Jun 12 11:34 media
drwxr-xr-x.  2 root root     4096 Apr 10 2014 mnt
drwxr-xr-x.  2 root root     4096 Jun 12 11:34 opt
dr-xr-xr-x. 163 root root      0 Oct 21 19:35 proc
drwx-----  2 root root     4096 Jun 12 11:35 root
drwxr-xr-x.  19 root root     4096 Oct 21 19:36 run
drwxr-xr-x.  2 root root     4096 Jun 12 15:32 sbin
drwxr-xr-x.  2 root root     4096 Jun 12 11:34 srv
dr-xr-xr-x.  13 root root      0 Oct 21 16:36 sys
drwxrwxrwt.  24 root root     4096 Oct 21 19:36 tmp
drwxr-xr-x.  10 root root     4096 Jun 12 15:32 usr
drwxr-xr-x.  11 root root     4096 Jun 12 15:32 var
root@6fba20d93011:/# █
```

Figure 11-13. Listing the files in Docker Container’s root Directory

Extract the jdk-8u65-linux-x64.gz file.

```
tar -xv jdk-8u65-linux-x64.gz
```

The .gz file gets extracted as shown in Figure 11-14.

```
File Edit View Search Terminal Help
jdk1.8.0_65/bin/jstat
jdk1.8.0_65/bin/ControlPanel
jdk1.8.0_65/bin/rmiregistry
jdk1.8.0_65/bin/appletviewer
jdk1.8.0_65/bin/javadoc
jdk1.8.0_65/bin/jdb
jdk1.8.0_65/bin/jjs
jdk1.8.0_65/bin/servertool
jdk1.8.0_65/bin/idlj
jdk1.8.0_65/bin/rmid
jdk1.8.0_65/bin/javapackager
jdk1.8.0_65/bin/policytool
jdk1.8.0_65/bin/javaws
jdk1.8.0_65/bin/unpack200
jdk1.8.0_65/bin/tnameserv
jdk1.8.0_65/bin/jmc.ini
jdk1.8.0_65/bin/jmap
jdk1.8.0_65/bin/serialver
jdk1.8.0_65/bin/wsgen
jdk1.8.0_65/bin/jrunscript
jdk1.8.0_65/bin/javah
jdk1.8.0_65/bin/javac
jdk1.8.0_65/bin/jvisualvm
jdk1.8.0_65/bin/jcontrol
jdk1.8.0_65/release
```

Figure 11-14. Extracting the JDK .gz File

We need to set the JAVA_HOME environment variable in the hadoop-env.sh file. To find the directory for the hadoop-env.sh file run the following command.

```
find -name hadoop-env.sh
```

The different directories containing the hadoop-env.sh file get listed as shown in Figure 11-15.


```
root@49d774f8f1fe:/# find -name hadoop-env.sh
./usr/lib/hadoop-0.20-mapreduce/example-confs/conf.secure/hadoop-env.sh
./etc/hadoop/conf.pseudo/hadoop-env.sh
./etc/profile.d/hadoop-env.sh
```

Figure 11-15. Finding the hadoop-env.sh File

Open the `./etc/hadoop/conf.psuedo/hadoop-env.sh` file in a vi editor and add the following export statement.

```
export JAVA_HOME=~/jdk1.8.0_65
```

The preceding statement in the `hadoop-env.sh` file is shown in Figure 11-16. Save the file with the `:wq` command.



```
# Licensed to the Apache Software Foundation (ASF) under one or more
# contributor license agreements. See the NOTICE file distributed with
# this work for additional information regarding copyright ownership.
# The ASF licenses this file to You under the Apache License, Version 2.0
# (the "License"); you may not use this file except in compliance with
# the License. You may obtain a copy of the License at
#
#     http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.
export JAVA_HOME=~/jdk1.8.0_65
# Set Hadoop-specific environment variables here.
# Forcing YARN-based mapreduce implementation.
# Make sure to comment out if you want to go back to the default or
# if you want this to be tweakable on a per-user basis
# export HADOOP_MAPRED_HOME=/usr/lib/hadoop-mapreduce

# The maximum amount of heap to use, in MB. Default is 1000.
#export HADOOP_HEAPSIZE=
:wq
```

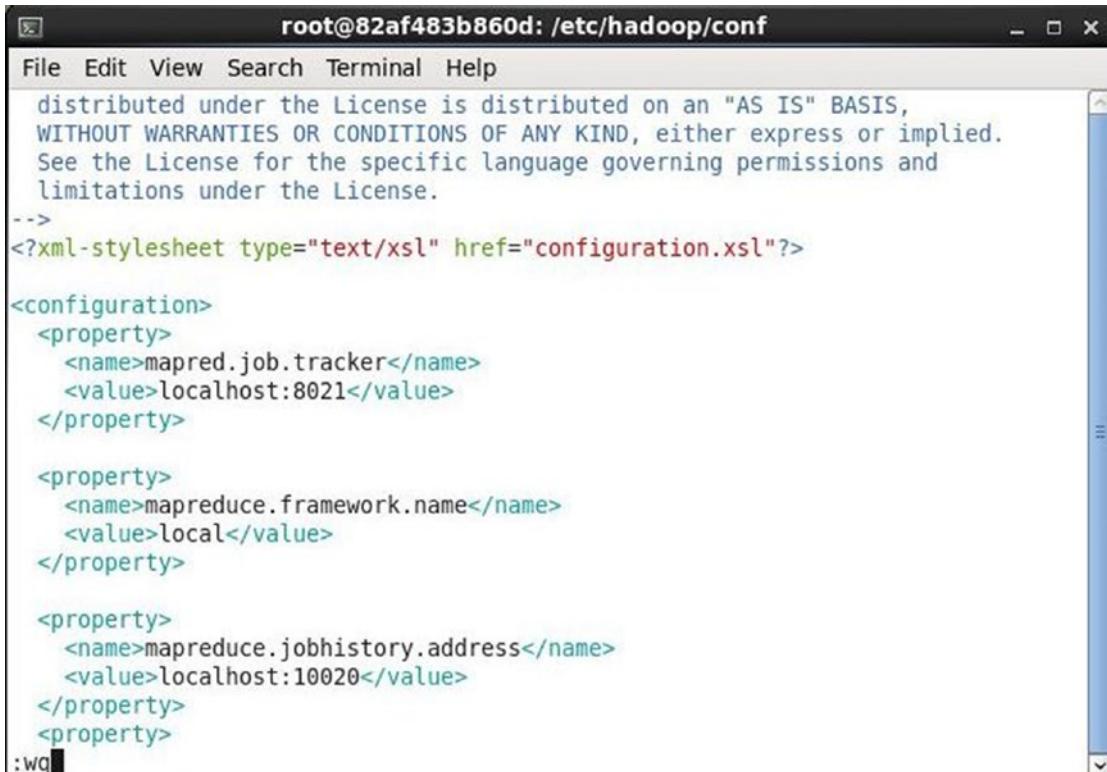
Figure 11-16. Setting the `JAVA_HOME` Environment Variable

Configuring Apache Hadoop

Apache Hadoop MapReduce framework may be started in one of the three modes: `local`, `classic` and `yarn`. In the “`local`” mode, MapReduce runs in a Java process. In the `classic` mode, MapReduce runs using the `MapReduce1` framework. With the `yarn` mode, MapReduce runs using the `MapReduce2` framework (also called `YARN`). The MapReduce framework to use is set in the `mapreduce.framework.name` setting in the `mapred-site.xml` configuration file, which is in the same directory as the `hadoop-env.sh`, the `./etc/hadoop/conf.psuedo` directory. As `yarn` and `classic` frameworks require more RAM than the `local`, set the `mapreduce.framework.name` to `local`.

```
<property>
    <name>mapreduce.framework.name</name>
    <value>local</value>
</property>
```

The `mapreduce.framework.name` setting is shown in Figure 11.17.



A terminal window titled "root@82af483b860d: /etc/hadoop/conf". The window contains the XML configuration for the Hadoop MapReduce framework. It includes a license notice, a reference to an XSL stylesheet, and several `<property>` blocks. One block sets the `mapreduce.framework.name` to `local`. The command `:wq` is visible at the bottom of the terminal window.

```
distributed under the License is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the License for the specific language governing permissions and
limitations under the License.

-->
<xm...-stylesheet type="text/xsl" href="configuration.xsl"?>

<configuration>
    <property>
        <name>mapred.job.tracker</name>
        <value>localhost:8021</value>
    </property>

    <property>
        <name>mapreduce.framework.name</name>
        <value>local</value>
    </property>

    <property>
        <name>mapreduce.jobhistory.address</name>
        <value>localhost:10020</value>
    </property>
    <property>
:wq
```

Figure 11-17. Setting the MapReduce Framework to local

Also set the following (Table 11-2) configuration properties in the `hdfs-site.xml` configuration file.

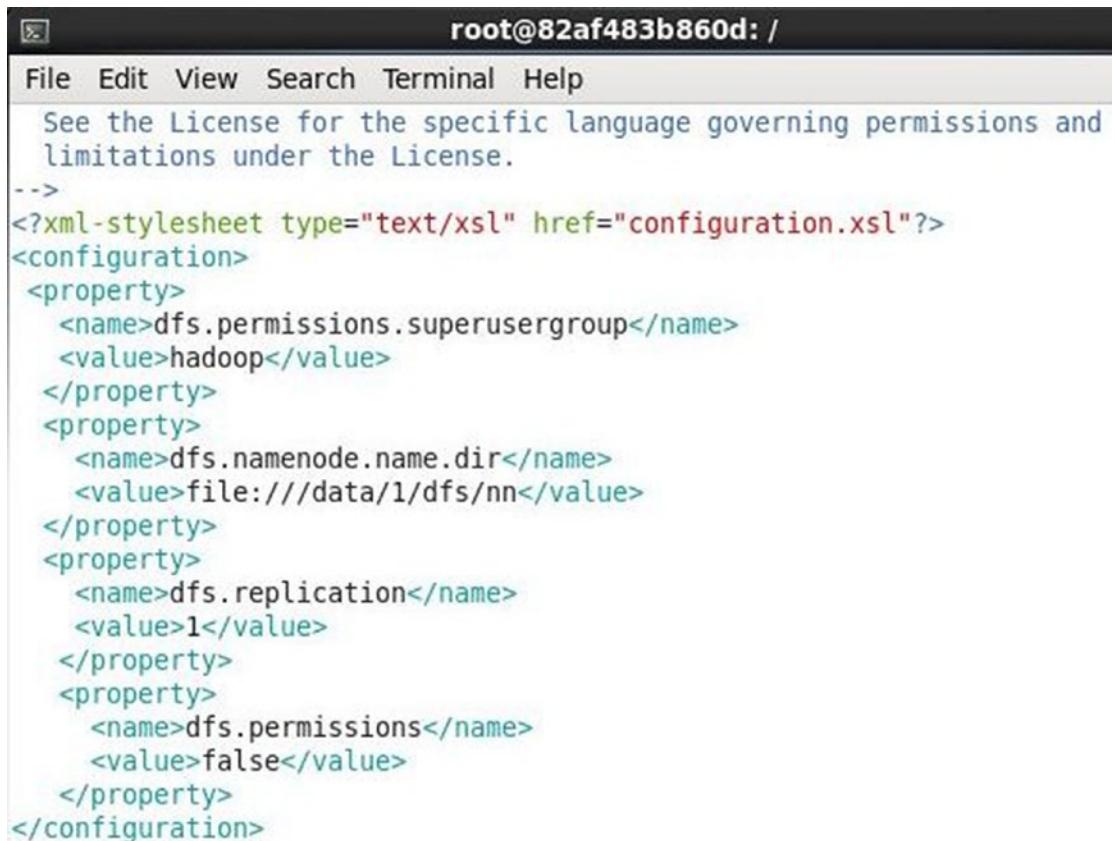
Table 11-2. Configuration Properties for `hdfs-site.xml`

Configuration Property	Description	Value
<code>dfs.permissions.superusergroup</code>	Sets the super user group	<code>hadoop</code>
<code>dfs.namenode.name.dir</code>	Sets the NameNode storage directory	<code>file:///data/1/dfs/nn</code>
<code>dfs.replication</code>	Sets the replication level	<code>1</code>
<code>dfs.permissions</code>	Whether permissions are to be checked	<code>false</code>

The `hdfs-site.xml` configuration settings are listed below.

```
<configuration>
<property>
  <name>dfs.permissions.superusergroup</name>
  <value>hadoop</value>
</property>
<property>
  <name>dfs.namenode.name.dir</name>
  <value>file:///data/1/dfs/nn</value>
</property>
<property>
  <name>dfs.replication</name>
  <value>1</value>
</property>
<property>
  <name>dfs.permissions</name>
  <value>false</value>
</property>
</configuration>
```

The `hdfs-site.xml` configuration file is shown in Figure 11-18.



The screenshot shows a terminal window with a black header bar containing the text "root@82af483b860d:/". Below the header is a light gray menu bar with options: File, Edit, View, Search, Terminal, Help. The main area of the terminal displays the XML configuration file. It starts with a note about license permissions, followed by the XML code itself. The XML code is identical to the one shown above, defining properties for DFS permissions, replication, and namenode directory.

```
See the License for the specific language governing permissions and
limitations under the License.

-->
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<configuration>
<property>
  <name>dfs.permissions.superusergroup</name>
  <value>hadoop</value>
</property>
<property>
  <name>dfs.namenode.name.dir</name>
  <value>file:///data/1/dfs/nn</value>
</property>
<property>
  <name>dfs.replication</name>
  <value>1</value>
</property>
<property>
  <name>dfs.permissions</name>
  <value>false</value>
</property>
</configuration>
```

Figure 11-18. The `hdfs-site.xml` Configuration File

We need to create the NameNode storage directory set in the `dfs.namenode.name.dir` property. Create the `/data/1/dfs/nn` directory and set its permissions to global (777).

```
sudo mkdir -p /data/1/dfs/nn
sudo chmod -R 777 /data/1/dfs/nn
```

Create the user group `hadoop` and a user `hadoop`.

```
groupadd hadoop
useradd hadoop
```

We need to set the following (Table 11-3) configuration properties in the `core-site.xml` file.

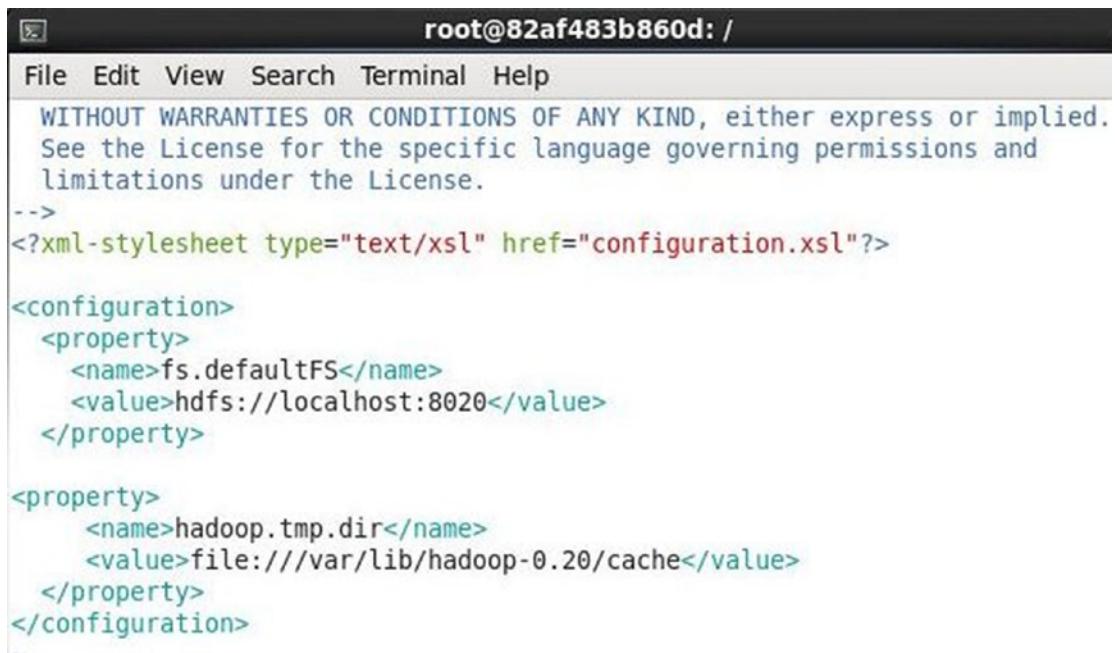
Table 11-3. Configuration Properties for `core-site.xml`

Configuration Property	Description	Value
<code>fs.defaultFS</code>	The NameNode URI	<code>hdfs://localhost:8020</code>
<code>hadoop.tmp.dir</code>	The Hadoop temporary directory	<code>file:///var/lib/hadoop-0.20/cache</code>

The `core-site.xml` configuration settings are listed:

```
<configuration>
<property>
  <name>fs.defaultFS</name>
  <value>hdfs://10.0.2.15:8020</value>
</property>
<property>
  <name>hadoop.tmp.dir</name>
  <value>file:///var/lib/hadoop-0.20/cache</value>
</property>
</configuration>
```

The `core-site.xml` file is shown in Figure 11-19. Save the file with :wq.



The screenshot shows a terminal window titled "root@82af483b860d: /". The window displays the contents of the core-site.xml configuration file. The file includes a license notice, XML declarations, and configuration properties for HDFS and Hadoop temporary directory settings.

```

WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the License for the specific language governing permissions and
limitations under the License.

-->
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>

<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://localhost:8020</value>
  </property>

  <property>
    <name>hadoop.tmp.dir</name>
    <value>file:///var/lib/hadoop-0.20/cache</value>
  </property>
</configuration>

```

Figure 11-19. The core-site.xml Configuration File

Create the directory set in the hadoop.tmp.dir directory and set its permissions to global (777).

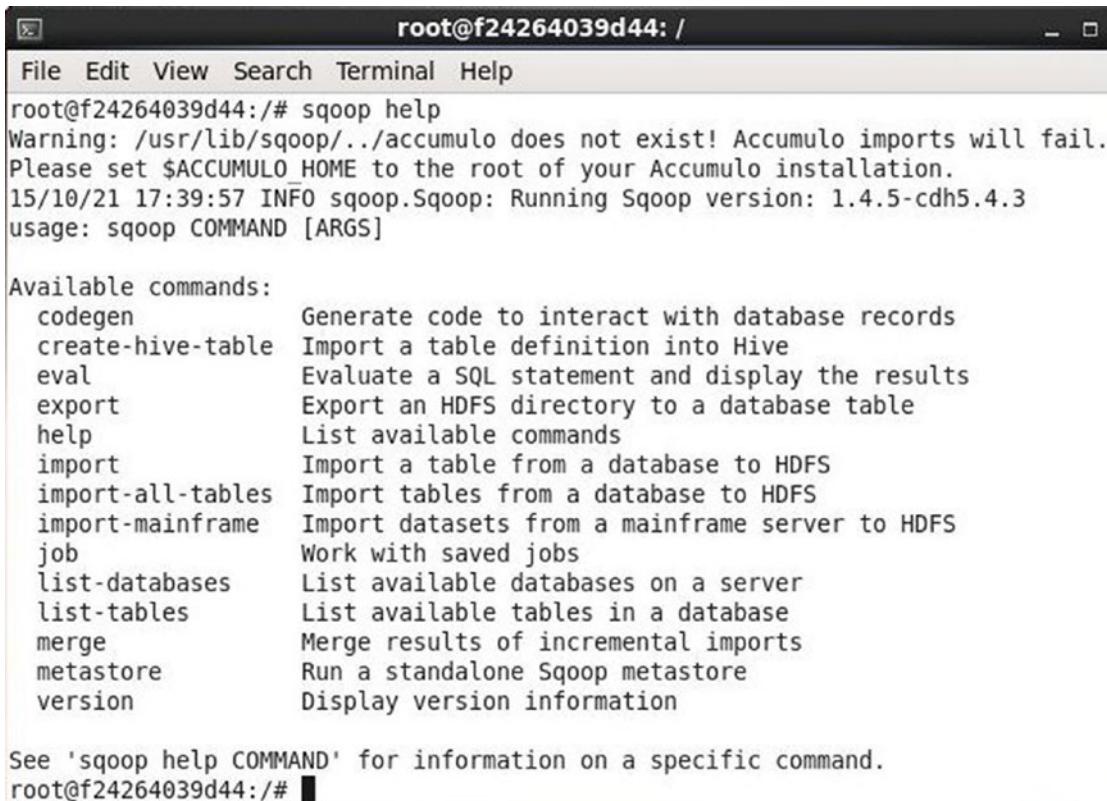
```
mkdir -p /var/lib/hadoop-0.20/cache
chmod -R 777 /var/lib/hadoop-0.20/cache
```

We also need to set the permissions of the / directory in HDFS to global (777) with the following command.

```
sudo -u hdfs hdfs dfs -chmod 777 /
```

Importing MySQL Table Data into HDFS with Sqoop

In this section we shall use the `sqoop import` command to import MySQL database table data to the HDFS. The different commands supported by the `sqoop` tool may be listed by running the `sqoop help` command from the interactive shell for the cdh container as shown in Figure 11-20. The `import` command is used to import a table from a relational database to HDFS.



```
root@f24264039d44: /root
File Edit View Search Terminal Help
root@f24264039d44:/# sqoop help
Warning: /usr/lib/sqoop/.../accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
15/10/21 17:39:57 INFO sqoop.Sqoop: Running Sqoop version: 1.4.5-cdh5.4.3
usage: sqoop COMMAND [ARGS]

Available commands:
codegen          Generate code to interact with database records
create-hive-table Import a table definition into Hive
eval             Evaluate a SQL statement and display the results
export           Export an HDFS directory to a database table
help             List available commands
import            Import a table from a database to HDFS
import-all-tables Import tables from a database to HDFS
import-mainframe  Import datasets from a mainframe server to HDFS
job              Work with saved jobs
list-databases   List available databases on a server
list-tables      List available tables in a database
merge            Merge results of incremental imports
metastore        Run a standalone Sqoop metastore
version          Display version information

See 'sqoop help COMMAND' for information on a specific command.
root@f24264039d44:/#
```

Figure 11-20. Running the `sqoop help` Command

Running the `sqoop import` command requires the code to be generated for accessing the relational database. The code may be generated directly while the `sqoop import` command is run or before the `sqoop import` command is run using the `sqoop codegen` command. Run the following `sqoop codegen` command to generate the code to interact with the database records.

```
sudo -u hdfs sqoop codegen --connect "jdbc:mysql://e414f8c41d0b:3306/mysqldb" --password "mysql" --username "mysql" --table "wlslog"
```

The `-u hdfs` specifies the user as `hdfs`. The command parameters are discussed in Table 11-4.

Table 11-4. Command Parameters for the hdfs Command

Parameter	Description	Value
--connect	The connection url to connect to MySQL database. The hostname is the container id in which MySQL is run.	"jdbc:mysql://e414f8c41d0b:3306/mysqlDb"
--password	Password to connect to MySQL. It is recommended to use a non-root user.	"mysql"
--username	Username to connect to MySQL.	"mysql"
--table	MySQL table from which to import from	"wlslog"

The code required to interact with the database gets generated in the wlslog.jar file as shown in Figure 11-21.

```
root@08b338cb2a90:/# sudo -u hdfs sqoop codegen --connect "jdbc:mysql://e414f8c41d0b:3306/mysqlDb" --password "mysql" --username "mysql" --table "wlslog"
Warning: /usr/lib/sqoop/../../accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
15/10/22 00:05:54 INFO sqoop.Sqoop: Running Sqoop version: 1.4.5-cdh5.4.3
15/10/22 00:05:54 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
15/10/22 00:05:55 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
15/10/22 00:05:55 INFO tool.CodeGenTool: Beginning code generation
15/10/22 00:05:56 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `wlslog` AS t LIMIT 1
15/10/22 00:05:56 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `wlslog` AS t LIMIT 1
15/10/22 00:05:57 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/lib/hadoop-mapreduce
Note: /tmp/sqoop-hdfs/compile/6348ef9539c8ad2bee9ba1875a62c923/wlslog.java uses
or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
15/10/22 00:06:03 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-hdfs/compile/6348ef9539c8ad2bee9ba1875a62c923/wlslog.jar
root@08b338cb2a90:/#
```

Figure 11-21. Output from the codegen Command

Next, run the sqoop import command as user hdfs. Add the wlslog.jar file in the classpath with the -libjars option.

```
sudo -u hdfs sqoop import -libjars /tmp/sqoop-hdfs/compile/6348ef9539c8ad2bee9ba1875a62c923/wlslog.jar --connect "jdbc:mysql://e414f8c41d0b:3306/mysqlDb" --password "mysql" --username "mysql" --table "wlslog" --columns "time_stamp,category,type,servername,code,msg" --target-dir "/mysql/import" -verbose
```

The other command parameters are discussed in Table 11-5.

Table 11-5. Command Parameters for *sqoop import*

Parameter	Description	Value
--connect	The connection url to connect to MySQL database. The hostname is the container id in which MySQL is run.	"jdbc:mysql://e414f8c41d0b:3306/mysldb"
--password	Password to connect to MySQL. It is recommended to use a non-root user.	"mysql"
--username	Username to connect to MySQL.	"mysql"
--columns	Columns to be imported	"time_stamp,category,type,servername,code,msg"
--table	MySQL table from which to import from	"wlslog"
--target-dir	The HDFS directory in which to import	"/mysql/import"

The output from the `sqoop import` command is shown in Figure 11-22.

```

root@08b338cb2a90: / 
File Edit View Search Terminal Help
15/10/22 00:07:28 INFO mapreduce.Job: Counters: 23
    File System Counters
        FILE: Number of bytes read=17796154
        FILE: Number of bytes written=18238016
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=0
        HDFS: Number of bytes written=615
        HDFS: Number of read operations=4
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=3
    Map -Reduce Framework
        Map input records=6
        Map output records=6
        Input split bytes=87
        Spilled Records=0
        Failed Shuffles=0
        Merged Map outputs=0
        GC time elapsed (ms)=306
        CPU time spent (ms)=0
        Physical memory (bytes) snapshot=0
        Virtual memory (bytes) snapshot=0
        Total committed heap usage (bytes)=138571776
    File Input Format Counters
        Bytes Read=0
    File Output Format Counters
        Bytes Written=615
15/10/22 00:07:28 INFO mapreduce.ImportJobBase: Transferred 615 bytes in 9.6688 seconds (63.6064 bytes/sec)
15/10/22 00:07:28 INFO mapreduce.ImportJobBase: Retrieved 6 records.
root@08b338cb2a90: # 

```

Figure 11-22. Output from *sqoop import*

The detailed output from the `sqoop import` command is listed:

```
root@08b338cb2a90:/# sudo -u hdfs sqoop import -libjars /tmp/sqoop-hdfs/compile/6348ef95
39c8ad2bee9ba1875a62c923/wlslog.jar --connect "jdbc:mysql://e414f8c41dob:3306/mysql"
--password "mysql" --username "mysql" --table "wlslog" --columns "time_stamp,category,type,s
ervername,code,msg" --target-dir "/mysql/import" -verbose
15/10/22 00:07:07 INFO sqoop.Sqoop: Running Sqoop version: 1.4.5-cdh5.4.3
ConnManager
15/10/22 00:07:10 INFO tool.CodeGenTool: Beginning code generation
15/10/22 00:07:10 DEBUG manager.SqlManager: Execute getColumnInfoRawQuery : SELECT t.* FROM
`wlslog` AS t LIMIT 1
15/10/22 00:07:10 DEBUG manager.SqlManager: No connection parameters specified. Using
regular API for making connection.
15/10/22 00:07:11 DEBUG manager.SqlManager: Using fetchSize for next query: -2147483648
15/10/22 00:07:11 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `wlslog`
AS t LIMIT 1
15/10/22 00:07:11 DEBUG manager.SqlManager: Found column time_stamp of type [12, 255, 0]
15/10/22 00:07:11 DEBUG manager.SqlManager: Found column category of type [12, 255, 0]
15/10/22 00:07:11 DEBUG manager.SqlManager: Found column type of type [12, 255, 0]
15/10/22 00:07:11 DEBUG manager.SqlManager: Found column servername of type [12, 255, 0]
15/10/22 00:07:11 DEBUG manager.SqlManager: Found column code of type [12, 255, 0]
15/10/22 00:07:11 DEBUG manager.SqlManager: Found column msg of type [12, 255, 0]
15/10/22 00:07:11 DEBUG orm.ClassWriter: selected columns:
15/10/22 00:07:11 DEBUG orm.ClassWriter:    time_stamp
15/10/22 00:07:11 DEBUG orm.ClassWriter:    category
15/10/22 00:07:11 DEBUG orm.ClassWriter:    type
15/10/22 00:07:11 DEBUG orm.ClassWriter:    servername
15/10/22 00:07:11 DEBUG orm.ClassWriter:    code
15/10/22 00:07:11 DEBUG orm.ClassWriter:    msg
15/10/22 00:07:11 DEBUG manager.SqlManager: Using fetchSize for next query: -2147483648
15/10/22 00:07:11 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `wlslog`
AS t LIMIT 1
15/10/22 00:07:11 DEBUG manager.SqlManager: Found column time_stamp of type VARCHAR
15/10/22 00:07:11 DEBUG manager.SqlManager: Found column category of type VARCHAR
15/10/22 00:07:11 DEBUG manager.SqlManager: Found column type of type VARCHAR
15/10/22 00:07:11 DEBUG manager.SqlManager: Found column servername of type VARCHAR
15/10/22 00:07:11 DEBUG manager.SqlManager: Found column code of type VARCHAR
15/10/22 00:07:11 DEBUG manager.SqlManager: Found column msg of type VARCHAR
15/10/22 00:07:11 DEBUG orm.ClassWriter: Writing source file: /tmp/sqoop-hdfs/compile/3c3425
a2eecf819af8fe8f4eabd40468/wlslog.java
15/10/22 00:07:11 DEBUG orm.ClassWriter: Table name: wlslog
15/10/22 00:07:11 DEBUG orm.ClassWriter: Columns: time_stamp:12, category:12, type:12,
servername:12, code:12, msg:12,
15/10/22 00:07:11 DEBUG orm.ClassWriter: sourceFilename is wlslog.java
15/10/22 00:07:11 DEBUG orm.CompilationManager: Found existing /tmp/sqoop-hdfs/compile/3c3425
a2eecf819af8fe8f4eabd40468/
15/10/22 00:07:11 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/lib/hadoop-mapreduce
15/10/22 00:07:11 DEBUG orm.CompilationManager: Returning jar file path /usr/lib/hadoop-
mapreduce/hadoop-mapreduce-client-core.jar:/usr/lib/hadoop-mapreduce/hadoop-mapreduce-
client-core-2.6.0-cdh5.4.3.jar
```

```

15/10/22 00:07:17 DEBUG orm.CompilationManager: Could not rename /tmp/sqoop-hdfs/compile/3c3
425a2eecf819af8fe8f4eabd40468/wlslog.java to ./wlslog.java
15/10/22 00:07:17 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-hdfs/compile/3c3
425a2eecf819af8fe8f4eabd40468/wlslog.jar
15/10/22 00:07:17 DEBUG orm.CompilationManager: Scanning for .class files in directory:
/tmp/sqoop-hdfs/compile/3c3425a2eecf819af8fe8f4eabd40468
15/10/22 00:07:17 DEBUG orm.CompilationManager: Got classfile: /tmp/sqoop-hdfs/compile/3c342
5a2eecf819af8fe8f4eabd40468/wlslog.class -> wlslog.class
15/10/22 00:07:17 DEBUG orm.CompilationManager: Finished writing jar file /tmp/sqoop-hdfs/co
mpile/3c3425a2eecf819af8fe8f4eabd40468/wlslog.jar
15/10/22 00:07:17 WARN manager.MySQLManager: It looks like you are importing from mysql.
15/10/22 00:07:17 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
15/10/22 00:07:17 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.
15/10/22 00:07:17 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull
(mysql)
15/10/22 00:07:17 DEBUG manager.MySQLManager: Rewriting connect string to jdbc:mysql://e414f8c41d0b:3306/mysql?zeroDateTimeBehavior=convertToNull
15/10/22 00:07:17 DEBUG manager.CatalogQueryManager: Retrieving primary key for table
'wlslog' with query SELECT column_name FROM INFORMATION_SCHEMA.COLUMNS WHERE TABLE_SCHEMA =
(SELECT SCHEMA()) AND TABLE_NAME = 'wlslog' AND COLUMN_KEY = 'PRI'
15/10/22 00:07:17 DEBUG manager.CatalogQueryManager: Retrieving primary key for table
'wlslog' with query SELECT column_name FROM INFORMATION_SCHEMA.COLUMNS WHERE TABLE_SCHEMA =
(SELECT SCHEMA()) AND TABLE_NAME = 'wlslog' AND COLUMN_KEY = 'PRI'
15/10/22 00:07:17 INFO mapreduce.ImportJobBase: Beginning import of wlslog
15/10/22 00:07:17 INFO Configuration.deprecation: mapred.job.tracker is deprecated. Instead,
use mapreduce.jobtracker.address
15/10/22 00:07:17 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use
mapreduce.job.jar
15/10/22 00:07:17 DEBUG db.DBConfiguration: Securing password into job credentials store
15/10/22 00:07:17 DEBUG mapreduce.DataDrivenImportJob: Using table class: wlslog
15/10/22 00:07:17 DEBUG mapreduce.DataDrivenImportJob: Using InputFormat: class com.
cloudera.sqoop.mapreduce.db.DataDrivenDBInputFormat
15/10/22 00:07:18 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead,
use mapreduce.job.maps
15/10/22 00:07:19 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker,
sessionId=
15/10/22 00:07:20 DEBUG db.DBConfiguration: Fetching password from job credentials store
15/10/22 00:07:20 INFO db.DBInputFormat: Using read committed transaction isolation
15/10/22 00:07:20 DEBUG db.DataDrivenDBInputFormat: Creating input split with lower bound
'1=1' and upper bound '1=1'
15/10/22 00:07:20 INFO mapreduce.JobSubmitter: number of splits:1
15/10/22 00:07:21 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_local2065078437_0001
15/10/22 00:07:25 INFO mapreduce.Job: The url to track the job: http://localhost:8080
15/10/22 00:07:25 INFO mapreduce.Job: Running job: job_local2065078437_0001
15/10/22 00:07:25 INFO mapred.LocalJobRunner: OutputCommitter set in config null
15/10/22 00:07:25 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 1
15/10/22 00:07:25 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.
mapreduce.lib.output.FileOutputCommitter
15/10/22 00:07:26 INFO mapred.LocalJobRunner: Waiting for map tasks

```

```

15/10/22 00:07:26 INFO mapred.LocalJobRunner: Starting task: attempt_local2065078437_0001
_m_000000_0
15/10/22 00:07:26 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 1
15/10/22 00:07:26 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
15/10/22 00:07:26 DEBUG db.DBConfiguration: Fetching password from job credentials store
15/10/22 00:07:26 INFO db.DBInputFormat: Using read committed transaction isolation
15/10/22 00:07:26 INFO mapred.MapTask: Processing split: 1=1 AND 1=1
15/10/22 00:07:26 DEBUG db.DataDrivenDBInputFormat: Creating db record reader for db product: MYSQL
15/10/22 00:07:26 INFO mapreduce.Job: Job job_local2065078437_0001 running in uber mode : false
15/10/22 00:07:26 INFO mapreduce.Job: map 0% reduce 0%
15/10/22 00:07:27 INFO db.DBRecordReader: Working on split: 1=1 AND 1=1
15/10/22 00:07:27 DEBUG db.DataDrivenDBRecordReader: Using query: SELECT `time_stamp`,
`category`, `type`, `servername`, `code`, `msg` FROM `wlslog` AS `wlslog` WHERE ( 1=1 ) AND ( 1=1 )
15/10/22 00:07:27 DEBUG db.DBRecordReader: Using fetchSize for next query: -2147483648
15/10/22 00:07:27 INFO db.DBRecordReader: Executing query: SELECT `time_stamp`, `category`,
`type`, `servername`, `code`, `msg` FROM `wlslog` AS `wlslog` WHERE ( 1=1 ) AND ( 1=1 )
15/10/22 00:07:27 DEBUG mapreduce.AutoProgressMapper: Instructing auto-progress thread to quit.
15/10/22 00:07:27 DEBUG mapreduce.AutoProgressMapper: Waiting for progress thread shutdown...
15/10/22 00:07:27 INFO mapreduce.AutoProgressMapper: Auto-progress thread is finished.
keepGoing=false
15/10/22 00:07:27 DEBUG mapreduce.AutoProgressMapper: Progress thread shutdown detected.
15/10/22 00:07:27 INFO mapred.LocalJobRunner:
15/10/22 00:07:27 INFO mapred.Task: Task:attempt_local2065078437_0001_m_000000_0 is done.
And is in the process of committing
15/10/22 00:07:27 INFO mapred.LocalJobRunner:
15/10/22 00:07:27 INFO mapred.Task: Task attempt_local2065078437_0001_m_000000_0 is allowed
to commit now
15/10/22 00:07:27 INFO output.FileOutputCommitter: Saved output of task 'attempt_loca
l2065078437_0001_m_000000_0' to hdfs://localhost:8020/mysql/import/_temporary/0/task_
local2065078437_0001_m_000000
15/10/22 00:07:27 INFO mapred.LocalJobRunner: map
15/10/22 00:07:27 INFO mapred.Task: Task 'attempt_local2065078437_0001_m_000000_0' done.
15/10/22 00:07:27 INFO mapred.LocalJobRunner: Finishing task: attempt_local2065078437_0001
_m_000000_0
15/10/22 00:07:27 INFO mapred.LocalJobRunner: map task executor complete.
15/10/22 00:07:28 INFO mapreduce.Job: map 100% reduce 0%
15/10/22 00:07:28 INFO mapreduce.Job: Job job_local2065078437_0001 completed successfully
15/10/22 00:07:28 INFO mapreduce.Job: Counters: 23

```

File System Counters

```

FILE: Number of bytes read=17796154
FILE: Number of bytes written=18238016
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=0
HDFS: Number of bytes written=615
HDFS: Number of read operations=4
HDFS: Number of large read operations=0
HDFS: Number of write operations=3

```

Map-Reduce Framework

```
Map input records=6
```

```

Map output records=6
Input split bytes=87
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=306
CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=138571776
File Input Format Counters
    Bytes Read=0
File Output Format Counters
    Bytes Written=615
15/10/22 00:07:28 INFO mapreduce.ImportJobBase: Transferred 615 bytes in 9.6688 seconds
(63.6064 bytes/sec)
15/10/22 00:07:28 INFO mapreduce.ImportJobBase: Retrieved 6 records.
root@08b338cb2a90:/#

```

Listing Data Imported into HDFS

To list the files generated with the `sqoop import` tool in the `/mysql/import` directory, run the following command.

```
sudo -u hdfs hdfs dfs -ls /mysql/import
```

Two files get listed: `_SUCCESS`, which indicates that the `sqoop import` command completed successfully, and `part-m-00000`, which has the data imported as shown in Figure 11-23.

```

root@08b338cb2a90:/# sudo -u hdfs hdfs dfs -ls /mysql/import
Found 2 items
-rw-r--r-- 1 hdfs supergroup      0 2015-10-22 00:07 /mysql/import/_SUCCESS
S
-rw-r--r-- 1 hdfs supergroup 615 2015-10-22 00:07 /mysql/import/part-m-00000
root@08b338cb2a90:/# █

```

Figure 11-23. Listing Files Generated by `sqoop import`

List the data in the data file `part-m-00000` with the following command.

```
sudo -u hdfs hdfs dfs -cat /mysql/import/part-m-00000
```

The data imported with the `sqoop import` tool gets listed as shown in Figure 11-24.

```
root@08b338cb2a90:/# sudo -u hdfs hdfs dfs -cat /mysql/import/part-m-00000
Apr-8-2014-7:06:16-PM-PDT,Notice,WebLogicServer,AdminServer,BEA-000365,Server state changed to STANDBY
Apr-8-2014-7:06:17-PM-PDT,Notice,WebLogicServer,AdminServer,BEA-000365,Server state changed to STARTING
Apr-8-2014-7:06:18-PM-PDT,Notice,WebLogicServer,AdminServer,BEA-000365,Server state changed to ADMIN
Apr-8-2014-7:06:19-PM-PDT,Notice,WebLogicServer,AdminServer,BEA-000365,Server state changed to RESUMING
Apr-8-2014-7:06:20-PM-PDT,Notice,WebLogicServer,AdminServer,BEA-000361,Started WebLogic AdminServer
Apr-8-2014-7:06:21-PM-PDT,Notice,WebLogicServer,AdminServer,BEA-000365,Server state changed to RUNNING
root@08b338cb2a90:/#
```

Figure 11-24. Listing Data imported by Sqoop

Exporting from HDFS to MySQL with Sqoop

Next, we shall export the data imported into HDFS back to MySQL database. In general the `sqoop export` tool exports a set of files from HDFS back to an RDBMS where the target table exists already in the database and the input files will be read and parsed into a set of records according to the delimiters specified in the "user-specified" values.

The code required to interact with the database may be generated during the `sqoop export` command or before the `sqoop export` command. We shall generate the code before running the `sqoop export` command using the `sqoop codegen` command as follows.

```
sudo -u hdfs sqoop codegen --connect "jdbc:mysql://e414f8c41d0b:3306/mysqlldb" --password "mysql" --username "mysql" --table "WLSLOG_COPY"
```

The command parameters are the same as for the `sqoop codegen` command run before the `sqoop import` command except the table name is `WLSLOG_COPY` instead of `wlslog`. The code required by the `sqoop export` command gets generated in the `WLSLOG_COPY.jar` file as shown in Figure 11-25.

```

root@08b338cb2a90:/# sudo -u hdfs sqoop codegen --connect "jdbc:mysql://e414f8c41d0b:3306/mysql" --password "mysql" --username "mysql" --table "WLSLOG_COPY"
Warning: /usr/lib/sqoop/../accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
15/10/22 00:12:29 INFO sqoop.Sqoop: Running Sqoop version: 1.4.5-cdh5.4.3
15/10/22 00:12:29 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
15/10/22 00:12:29 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
15/10/22 00:12:29 INFO tool.CodeGenTool: Beginning code generation
15/10/22 00:12:30 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `WLSLOG_COPY` AS t LIMIT 1
15/10/22 00:12:31 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `WLSLOG_COPY` AS t LIMIT 1
15/10/22 00:12:31 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/lib/hadoop-mapreduce
Note: /tmp/sqoop-hdfs/compile/047d0687acbb2298370a7b461cdfdd2e/WLSLOG_COPY.java
uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
15/10/22 00:12:37 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-hdfs/compile/047d0687acbb2298370a7b461cdfdd2e/WLSLOG_COPY.jar
root@08b338cb2a90:/#

```

Figure 11-25. Running the `sqoop codegen` Command

Next, run the `sqoop export` command adding the `WLSLOG_COPY.jar` in the classpath with the `-libjars` option. The other command parameters are the same as the `sqoop import` command except the `-table` being “`WLSLOG_COPY`” and the `--export-dir` option replacing the `--target-dir`. The directory in the `--export-dir` option should be the same as the directory in the `--data-dir` option for the `sqoop import` command.

```

sudo -u hdfs sqoop export -libjars /tmp/sqoop-hdfs/compile/047d0687acbb2298370a7b461cdfdd2e/WLSLOG_COPY.jar --connect "jdbc:mysql://e414f8c41d0b:3306/mysql" --password "mysql" --username "mysql" --export-dir "/mysql/import" --table "WLSLOG_COPY" --verbose

```

The output from the `sqoop export` command is shown in Figure 11-26.

```

root@08b338cb2a90: / 
File Edit View Search Terminal Help
15/10/22 00:14:13 INFO mapreduce.Job: Counters: 23
  File System Counters
    FILE: Number of bytes read=71190614
    FILE: Number of bytes written=72948608
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=4068
    HDFS: Number of bytes written=0
    HDFS: Number of read operations=86
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=0
  Map-Reduce Framework
    Map input records=6
    Map output records=6
    Input split bytes=576
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=0
    CPU time spent (ms)=0
    Physical memory (bytes) snapshot=0
    Virtual memory (bytes) snapshot=0
    Total committed heap usage (bytes)=576782336
  File Input Format Counters
    Bytes Read=0
  File Output Format Counters
    Bytes Written=0
15/10/22 00:14:13 INFO mapreduce.ExportJobBase: Transferred 3.9727 KB in 8.722 seconds (466.4067 bytes/sec)
15/10/22 00:14:13 INFO mapreduce.ExportJobBase: Exported 6 records.
root@08b338cb2a90:/# 

```

Figure 11-26. Output from the `sqoop export` command

The detailed output from the `sqoop export` command is listed:

```

root@08b338cb2a90:/# sudo -u hdfs sqoop export -libjars /tmp/sqoop-hdfs/compile/047d0687a
ccb2298370a7b461cdfdd2e/WLSLOG_COPY.jar --connect "jdbc:mysql://e414f8c41d0b:3306/mysql"
--password "mysql" --username "mysql" --export-dir "/mysql/import" --table "WLSLOG_COPY"
--verbose
15/10/22 00:13:52 INFO sqoop.Sqoop: Running Sqoop version: 1.4.5-cdh5.4.3
15/10/22 00:13:54 INFO tool.CodeGenTool: Beginning code generation
15/10/22 00:13:54 DEBUG manager.SqlManager: Execute getColumnInfoRawQuery : SELECT t.* FROM
`WLSLOG_COPY` AS t LIMIT 1
15/10/22 00:13:54 DEBUG manager.SqlManager: No connection parameters specified. Using
regular API for making connection.

```

```

15/10/22 00:13:55 DEBUG manager.SqlManager: Using fetchSize for next query: -2147483648
15/10/22 00:13:55 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM
`WLSLOG_COPY` AS t LIMIT 1
15/10/22 00:13:55 DEBUG manager.SqlManager: Found column time_stamp of type [12, 255, 0]
15/10/22 00:13:55 DEBUG manager.SqlManager: Found column category of type [12, 255, 0]
15/10/22 00:13:55 DEBUG manager.SqlManager: Found column type of type [12, 255, 0]
15/10/22 00:13:55 DEBUG manager.SqlManager: Found column servername of type [12, 255, 0]
15/10/22 00:13:55 DEBUG manager.SqlManager: Found column code of type [12, 255, 0]
15/10/22 00:13:55 DEBUG manager.SqlManager: Found column msg of type [12, 255, 0]
15/10/22 00:13:55 DEBUG orm.ClassWriter: selected columns:
15/10/22 00:13:55 DEBUG orm.ClassWriter:    time_stamp
15/10/22 00:13:55 DEBUG orm.ClassWriter:    category
15/10/22 00:13:55 DEBUG orm.ClassWriter:    type
15/10/22 00:13:55 DEBUG orm.ClassWriter:    servername
15/10/22 00:13:55 DEBUG orm.ClassWriter:    code
15/10/22 00:13:55 DEBUG orm.ClassWriter:    msg
15/10/22 00:13:55 DEBUG manager.SqlManager: Using fetchSize for next query: -2147483648
15/10/22 00:13:55 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM
`WLSLOG_COPY` AS t LIMIT 1
15/10/22 00:13:55 DEBUG manager.SqlManager: Found column time_stamp of type VARCHAR
15/10/22 00:13:55 DEBUG manager.SqlManager: Found column category of type VARCHAR
15/10/22 00:13:55 DEBUG manager.SqlManager: Found column type of type VARCHAR
15/10/22 00:13:55 DEBUG manager.SqlManager: Found column servername of type VARCHAR
15/10/22 00:13:55 DEBUG manager.SqlManager: Found column code of type VARCHAR
15/10/22 00:13:55 DEBUG manager.SqlManager: Found column msg of type VARCHAR
15/10/22 00:13:55 DEBUG orm.ClassWriter: Writing source file: /tmp/sqoop-hdfs/compile/715ce1
218221b63dffffd800222f863f0/WLSLOG_COPY.java
15/10/22 00:13:55 DEBUG orm.ClassWriter: Table name: WLSLOG_COPY
15/10/22 00:13:55 DEBUG orm.ClassWriter: Columns: time_stamp:12, category:12, type:12,
servername:12, code:12, msg:12,
15/10/22 00:13:55 DEBUG orm.ClassWriter: sourceFilename is WLSLOG_COPY.java
15/10/22 00:13:55 DEBUG orm.CompilationManager: Found existing /tmp/sqoop-hdfs/compile/715ce
1218221b63dffffd800222f863f0/
15/10/22 00:13:55 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/lib/hadoop-mapreduce
15/10/22 00:13:55 DEBUG orm.CompilationManager: Returning jar file path /usr/lib/hadoop-
mapreduce/hadoop-mapreduce-client-core.jar:/usr/lib/hadoop-mapreduce/hadoop-mapreduce-
client-core-2.6.0-cdh5.4.3.jar
15/10/22 00:14:02 INFO mapreduce.ExportJobBase: Beginning export of WLSLOG_COPY
15/10/22 00:14:02 INFO Configuration.deprecation: mapred.job.tracker is deprecated. Instead,
use mapreduce.jobtracker.address
15/10/22 00:14:02 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use
mapreduce.job.jar
15/10/22 00:14:04 DEBUG mapreduce.JobBase: Using InputFormat: class org.apache.sqoop.
mapreduce.ExportInputFormat
15/10/22 00:14:04 DEBUG db.DBConfiguration: Securing password into job credentials store
15/10/22 00:14:04 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker,
sessionId=
15/10/22 00:14:06 INFO input.FileInputFormat: Total input paths to process : 1
15/10/22 00:14:06 DEBUG mapreduce.ExportInputFormat: Target numMapTasks=4
15/10/22 00:14:06 DEBUG mapreduce.ExportInputFormat: Total input bytes=615
15/10/22 00:14:06 DEBUG mapreduce.ExportInputFormat: maxSplitSize=153

```

```

15/10/22 00:14:06 INFO input.FileInputFormat: Total input paths to process : 1
15/10/22 00:14:06 DEBUG mapreduce.ExportInputFormat: Generated splits:
15/10/22 00:14:06 DEBUG mapreduce.ExportInputFormat:   Paths:/mysql/import/
part-m-00000:0+153 Locations:08b338cb2a90:;
15/10/22 00:14:06 DEBUG mapreduce.ExportInputFormat:   Paths:/mysql/import/
part-m-00000:153+153 Locations:08b338cb2a90:;
15/10/22 00:14:06 DEBUG mapreduce.ExportInputFormat:   Paths:/mysql/import/
part-m-00000:306+153 Locations:08b338cb2a90:;
15/10/22 00:14:06 DEBUG mapreduce.ExportInputFormat:   Paths:/mysql/import/
part-m-00000:459+78,/mysql/import/part-m-00000:537+78 Locations:08b338cb2a90:;
15/10/22 00:14:06 INFO mapreduce.JobSubmitter: number of splits:4
15/10/22 00:14:06 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is
deprecated. Instead, use mapreduce.map.speculative
15/10/22 00:14:06 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_local1198888838_0001
15/10/22 00:14:11 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
15/10/22 00:14:11 INFO mapreduce.Job: Running job: job_local1198888838_0001
15/10/22 00:14:11 INFO mapred.LocalJobRunner: OutputCommitter set in config null
15/10/22 00:14:11 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.sqoop.mapreduce.
NullOutputCommitter
15/10/22 00:14:11 INFO mapred.LocalJobRunner: Waiting for map tasks
15/10/22 00:14:11 INFO mapred.LocalJobRunner: Starting task: attempt_local1198888838_0001
_m_000000_0
15/10/22 00:14:11 DEBUG mapreduce.CombineShimRecordReader: ChildSplit operates on:
hdfs://localhost:8020/mysql/import/part-m-00000
15/10/22 00:14:11 DEBUG db.DBConfiguration: Fetching password from job credentials store
15/10/22 00:14:12 DEBUG mapreduce.CombineShimRecordReader: ChildSplit operates on:
hdfs://localhost:8020/mysql/import/part-m-00000
15/10/22 00:14:12 DEBUG mapreduce.AutoProgressMapper: Instructing auto-progress thread to quit.
15/10/22 00:14:12 DEBUG mapreduce.AutoProgressMapper: Waiting for progress thread shutdown...
15/10/22 00:14:12 INFO mapreduce.AutoProgressMapper: Auto-progress thread is finished.
keepGoing=false
15/10/22 00:14:12 DEBUG mapreduce.AutoProgressMapper: Progress thread shutdown detected.
15/10/22 00:14:12 INFO mapred.LocalJobRunner:
15/10/22 00:14:12 DEBUG mapreduce.AsyncSqlOutputFormat: Committing transaction of 1 statements
15/10/22 00:14:12 INFO mapred.Task: Task:attempt_local1198888838_0001_m_000000_0 is done.
And is in the process of committing
15/10/22 00:14:12 INFO mapred.LocalJobRunner: map
15/10/22 00:14:12 INFO mapred.Task: Task 'attempt_local1198888838_0001_m_000000_0' done.
15/10/22 00:14:12 INFO mapred.LocalJobRunner: Finishing task: attempt_local1198888838_0001
_m_000000_0
15/10/22 00:14:12 INFO mapred.LocalJobRunner: Starting task: attempt_local1198888838_0001
_m_000001_0
15/10/22 00:14:12 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
15/10/22 00:14:12 INFO mapred.MapTask: Processing split: Paths:/mysql/import/
part-m-00000:0+153
15/10/22 00:14:12 DEBUG mapreduce.CombineShimRecordReader: ChildSplit operates on:
hdfs://localhost:8020/mysql/import/part-m-00000
15/10/22 00:14:12 DEBUG db.DBConfiguration: Fetching password from job credentials store
15/10/22 00:14:12 DEBUG mapreduce.AutoProgressMapper: Instructing auto-progress thread to quit.
15/10/22 00:14:12 DEBUG mapreduce.AutoProgressMapper: Waiting for progress thread shutdown...

```

```
15/10/22 00:14:12 INFO mapreduce.AutoProgressMapper: Auto-progress thread is finished.  
keepGoing=false  
15/10/22 00:14:12 DEBUG mapreduce.AutoProgressMapper: Progress thread shutdown detected.  
15/10/22 00:14:12 INFO mapred.LocalJobRunner:  
15/10/22 00:14:12 DEBUG mapreduce.AsyncSqlOutputFormat: Committing transaction of 1  
statements  
15/10/22 00:14:12 INFO mapred.Task: Task:attempt_local1198888838_0001_m_000001_0 is done.  
And is in the process of committing  
15/10/22 00:14:12 INFO mapred.LocalJobRunner: map  
15/10/22 00:14:12 INFO mapred.Task: Task 'attempt_local1198888838_0001_m_000001_0' done.  
15/10/22 00:14:12 INFO mapred.LocalJobRunner: Finishing task: attempt_local1198888838_0001  
_m_000001_0  
15/10/22 00:14:12 INFO mapred.LocalJobRunner: Starting task: attempt_local1198888838_0001  
_m_000002_0  
15/10/22 00:14:12 INFO mapreduce.Job: Job job_local1198888838_0001 running in uber mode : false  
15/10/22 00:14:12 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]  
15/10/22 00:14:12 INFO mapred.MapTask: Processing split: Paths:/mysql/import/  
part-m-00000:153+153  
15/10/22 00:14:12 DEBUG mapreduce.CombineShimRecordReader: ChildSplit operates on:  
hdfs://localhost:8020/mysql/import/part-m-00000  
15/10/22 00:14:12 INFO mapreduce.Job: map 100% reduce 0%  
15/10/22 00:14:12 DEBUG db.DBConfiguration: Fetching password from job credentials store  
15/10/22 00:14:12 DEBUG mapreduce.AutoProgressMapper: Instructing auto-progress thread to quit.  
15/10/22 00:14:12 DEBUG mapreduce.AutoProgressMapper: Waiting for progress thread shutdown...  
15/10/22 00:14:12 INFO mapreduce.AutoProgressMapper: Auto-progress thread is finished.  
keepGoing=false  
15/10/22 00:14:12 DEBUG mapreduce.AutoProgressMapper: Progress thread shutdown detected.  
15/10/22 00:14:12 INFO mapred.LocalJobRunner:  
15/10/22 00:14:12 DEBUG mapreduce.AsyncSqlOutputFormat: Committing transaction of 1  
statements  
15/10/22 00:14:12 INFO mapred.Task: Task:attempt_local1198888838_0001_m_000002_0 is done.  
And is in the process of committing  
15/10/22 00:14:12 INFO mapred.LocalJobRunner: map  
15/10/22 00:14:12 INFO mapred.Task: Task 'attempt_local1198888838_0001_m_000002_0' done.  
15/10/22 00:14:12 INFO mapred.LocalJobRunner: Finishing task: attempt_local1198888838_0001  
_m_000002_0  
15/10/22 00:14:12 INFO mapred.LocalJobRunner: Starting task: attempt_local1198888838_0001  
_m_000003_0  
15/10/22 00:14:12 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]  
15/10/22 00:14:12 INFO mapred.MapTask: Processing split: Paths:/mysql/import/  
part-m-00000:306+153  
15/10/22 00:14:12 DEBUG mapreduce.CombineShimRecordReader: ChildSplit operates on:  
hdfs://localhost:8020/mysql/import/part-m-00000  
15/10/22 00:14:12 DEBUG db.DBConfiguration: Fetching password from job credentials store  
15/10/22 00:14:12 DEBUG mapreduce.AutoProgressMapper: Instructing auto-progress thread to quit.  
15/10/22 00:14:12 DEBUG mapreduce.AutoProgressMapper: Waiting for progress thread shutdown...  
15/10/22 00:14:12 INFO mapreduce.AutoProgressMapper: Auto-progress thread is finished.  
keepGoing=false  
15/10/22 00:14:12 DEBUG mapreduce.AutoProgressMapper: Progress thread shutdown detected.  
15/10/22 00:14:12 INFO mapred.LocalJobRunner:  
15/10/22 00:14:12 DEBUG mapreduce.AsyncSqlOutputFormat: Committing transaction of 1 statements
```

```

15/10/22 00:14:12 INFO mapred.Task: Task:attempt_local1198888838_0001_m_000003_0 is done.
And is in the process of committing
15/10/22 00:14:12 INFO mapred.LocalJobRunner: map
15/10/22 00:14:12 INFO mapred.Task: Task 'attempt_local1198888838_0001_m_000003_0' done.
15/10/22 00:14:12 INFO mapred.LocalJobRunner: Finishing task: attempt_local1198888838_0001
_m_000003_0
15/10/22 00:14:12 INFO mapred.LocalJobRunner: map task executor complete.
15/10/22 00:14:13 INFO mapreduce.Job: Job job_local1198888838_0001 completed successfully
15/10/22 00:14:13 INFO mapreduce.Job: Counters: 23
    File System Counters
        FILE: Number of bytes read=71190614
        FILE: Number of bytes written=72948608
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=4068
        HDFS: Number of bytes written=0
        HDFS: Number of read operations=86
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=0
    Map-Reduce Framework
        Map input records=6
        Map output records=6
        Input split bytes=576
        Spilled Records=0
        Failed Shuffles=0
        Merged Map outputs=0
        GC time elapsed (ms)=0
        CPU time spent (ms)=0
        Physical memory (bytes) snapshot=0
        Virtual memory (bytes) snapshot=0
        Total committed heap usage (bytes)=576782336
    File Input Format Counters
        Bytes Read=0
    File Output Format Counters
        Bytes Written=0
15/10/22 00:14:13 INFO mapreduce.ExportJobBase: Transferred 3.9727 KB in 8.722 seconds
(466.4067 bytes/sec)
15/10/22 00:14:13 INFO mapreduce.ExportJobBase: Exported 6 records.
root@08b338cb2a90:/#


```

Querying Exported Data

Having exported from HDFS to MySQL, use the following SELECT statement in MySQL CLI to query the data exported.

```
select * from WLSLOG_COPY;
```

The six rows of data exported get listed as shown in Figure 11-27.



The screenshot shows a terminal window titled "ec2-user@ip-172-30-1-16:~". The user is running a MySQL command to select data from the "WLSLOG_COPY" table:

```
mysql> select * from WLSLOG_COPY;
```

The output displays 6 rows of log entries, each containing a timestamp, category, type, servername, and code. The log entries show the server transitioning through various states: STANDBY, STARTING, ADMIN, RESUMING, and RUNNING.

time_stamp	category	type	servername	code
Apr-8-2014-7:06:16-PM-PDT	Notice		WebLogicServer	AdminServer BEA-0003
65	Server state changed to STANDBY			
Apr-8-2014-7:06:17-PM-PDT	Notice		WebLogicServer	AdminServer BEA-0003
65	Server state changed to STARTING			
Apr-8-2014-7:06:18-PM-PDT	Notice		WebLogicServer	AdminServer BEA-0003
65	Server state changed to ADMIN			
Apr-8-2014-7:06:19-PM-PDT	Notice		WebLogicServer	AdminServer BEA-0003
65	Server state changed to RESUMING			
Apr-8-2014-7:06:20-PM-PDT	Notice		WebLogicServer	AdminServer BEA-0003
61	Started WebLogic AdminServer			
Apr-8-2014-7:06:21-PM-PDT	Notice		WebLogicServer	AdminServer BEA-0003
65	Server state changed to RUNNING			

6 rows in set (0.00 sec)

mysql> █

Figure 11-27. Querying Exported Data in WLSLOG_COPY

Stopping and Removing Docker Containers

To remove the mysqldb and cdh containers the containers have to be first stopped. Stop the mysqldb container with the docker stop command.

```
sudo docker stop mysqldb
```

Remove the mysqldb container with the docker rm command.

```
sudo docker rm mysqldb
```

The mysqldb container gets stopped and removed as shown in Figure 11-28.

```
[ec2-user@ip-172-30-1-16 ~]$ sudo docker stop mysql ldb  
mysql ldb  
[ec2-user@ip-172-30-1-16 ~]$ sudo docker rm mysql ldb  
mysql ldb
```

Figure 11-28. Stopping and Removing Docker Container for MySQL Database

Similarly stop and remove the cdh container.

```
sudo docker stop cdh  
sudo docker rm cdh
```

The cdh container gets stopped and removed as shown in Figure 11-29.

```
[ec2-user@ip-172-30-1-16 ~]$ sudo docker stop cdh  
cdh  
[ec2-user@ip-172-30-1-16 ~]$ sudo docker rm cdh  
cdh
```

Figure 11-29. Stopping and Removing Docker Container for CDH

Summary

In this chapter we used Docker images for CDH and MySQL database to run two separate, but linked, Docker containers. We created a MySQL database in the Docker container and ran the `sqoop import` tool in the CDH container to import data from MySQL to HDFS. Subsequently we ran the `sqoop export` tool to export from HDFS to MySQL database. In the next chapter we shall discuss Apache Kafka.

CHAPTER 12



Using Apache Kafka

Apache Kafka is a messaging system based on the publish-subscribe model. A Kafka cluster consists of one or more servers called brokers. Kafka keeps messages categorized by “topics”. Producers produce messages and publish the messages to topics. Consumers subscribe to specific topic/s and consume feeds of messages published to the topic/s. The messages published to a topic do not have to be consumed as produced and are stored in the topic for a configurable duration. A consumer may choose to consume the messages in a topic from the beginning. Apache ZooKeeper server is used to coordinate a Kafka cluster. The Kafka architecture is illustrated in Figure 12-1.

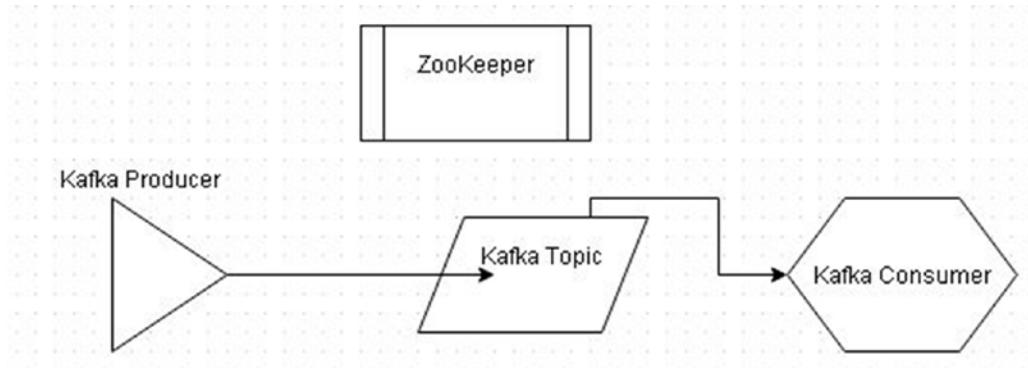


Figure 12-1. Apache Kafka Architecture

Apache Kafka is not directly based on Apache Hadoop nor does it make use of Apache Hadoop. But Kafka could be used as an Apache Flume source, channel, or sink. In this chapter we shall make use of a Docker image to run Apache Kafka in a Docker container. This chapter has the following sections.

Setting the Environment

Starting Docker Containers for Apache Kafka

Finding IP Addresses

Listing the Kafka Logs

Creating a Kafka Topic

Starting the Kafka Producer

- Starting the Kafka Consumer
- Producing and Consuming Messages
- Stopping and Removing the Docker Containers

Setting the Environment

The following software is required for this chapter.

- Docker (version 1.8)
- Docker image for Apache ZooKeeper (version latest)
- Docker image for Apache Kafka (version latest)

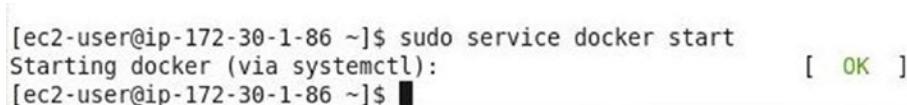
Connect to an Amazon EC2 instance on which the software is to be installed; the Public IP Address would be different for different users.

```
ssh -i "docker.pem" ec2-user@52.91.168.33
```

Install Docker and start the Docker service.

```
sudo service docker start
```

An OK message indicates that Docker has been started as shown in Figure 12-2.




```
[ec2-user@ip-172-30-1-86 ~]$ sudo service docker start
Starting docker (via systemctl):
[ OK ]
```

Figure 12-2. Starting Docker Service

Download the Docker image docker kafka/zookeeper for Apache ZooKeeper.

```
sudo docker pull docker kafka/zookeeper
```

The Docker image gets downloaded as shown in Figure 12-3.



The screenshot shows a terminal window titled "ec2-user@ip-172-30-1-86:~". The user has run the command "sudo docker pull dockerkafka/zookeeper". The output shows the image being pulled from the Docker registry. Most images are pulled complete, while some specific ones like "0b12df66f608" and "b9c1c5f483f7" are already present on the system. The final status message indicates that a newer image for "dockerkafka/zookeeper:latest" has been downloaded.

```
[ec2-user@ip-172-30-1-86 ~]$ sudo docker pull dockerkafka/zookeeper
Using default tag: latest
latest: Pulling from dockerkafka/zookeeper
fc7892a1e48d: Pull complete
84a782cf9c20: Pull complete
48b91d706c83: Pull complete
5d1adceefefa: Pull complete
cf2c3525a970: Pull complete
f9e3aa9a1c46: Pull complete
0bd0c18dd41b: Pull complete
6a8198f092e8: Pull complete
6d1c4364e27d: Pull complete
915a62a48764: Pull complete
dce0cb3abb42: Pull complete
d1df3324d542: Pull complete
0b12df66f608: Already exists
b9c1c5f483f7: Already exists
afc814718fc0: Already exists
3832c01b8067: Already exists
ba9b095d3962: Already exists
3f8a97ec9e67: Already exists
e05e4800a34e: Already exists
dc269adc7d73: Already exists
203ad8afa4f7: Already exists
cdcc5e2eb027: Already exists
4f75aa83c1e2: Already exists
ded5f5a542e6: Already exists
Digest: sha256:e6ed63a9c65a8a56113242ec14f423f61247bbcdc33b6e3f0fdb918f9180d552
Status: Downloaded newer image for dockerkafka/zookeeper:latest
[ec2-user@ip-172-30-1-86 ~]$
```

Figure 12-3. Downloading dockerkafka/zookeeper Docker Image

The dockerkafka/zookeeper image has been selected for download because a corresponding dockerkafka/kafka image is also available. Download the Docker image dockerkafka/kafka also.

```
sudo docker pull dockerkafka/kafka
```

Docker image dockerkafka/kafka gets downloaded as shown in Figure 12-4.

```
[ec2-user@ip-172-30-1-86 ~]$ sudo docker pull docker kafka/kafka
Using default tag: latest
latest: Pulling from docker kafka/kafka
0b12df66f608: Pull complete
b9c1c5f483f7: Pull complete
afc814718fc0: Pull complete
3832c01b8067: Pull complete
ba9b095d3962: Pull complete
3f8a97ec9e67: Pull complete
e05e4800a34e: Pull complete
dc269adc7d73: Pull complete
203ad8afa4f7: Pull complete
cdcc5e2eb027: Pull complete
4f75aa83c1e2: Pull complete
ded5f5a542e6: Pull complete
e5e30918403e: Pull complete
a8bce79e2431: Pull complete
cb16f24e3554: Pull complete
7696a5ea6c29: Pull complete
776a0ad98c89: Pull complete
04dfaf166381: Pull complete
44bbba2bdad8: Pull complete
a98340b715b2: Pull complete
dba5e17df2c6: Pull complete
Digest: sha256:57be79cd7617b96db9ba6856b0a872fd951a57728825b958475c300dbccc6a7e
Status: Downloaded newer image for docker kafka/kafka:latest
[ec2-user@ip-172-30-1-86 ~]$
```

Figure 12-4. Downloading the docker kafka/kafka Docker Image

Starting Docker Containers for Apache Kafka

We need to start both Apache ZooKeeper and Apache Kafka containers as both are required for a Kafka cluster. First, start a Docker container for Apache ZooKeeper using the following `docker run` command in which the port for ZooKeeper is set to 2181. The Docker container is started in detached mode with the `-d` option.

```
sudo docker run -d --name zookeeper -p 2181:2181 docker kafka/zookeeper
```

Next, start the Docker container for the Kafka server using the `docker kafka/kafka` image. Specify the port for the Kafka server as 9092 and link the Kafka container with the container running the ZooKeeper using `-link` parameter.

```
sudo docker run --name kafka -p 9092:9092 --link zookeeper:zookeeper docker kafka/kafka
```

List the running containers with the `docker ps` command.

```
sudo docker ps
```

The two containers, one for Apache ZooKeeper and the other for Apache Kafka get listed as shown in Figure 12-5.

```
[root@localhost ~]# ssh -i "docker.pem" ec2-user@54.152.8.111
Last login: Sat Oct 17 15:39:59 2015 from d75-157-54-139.bchsys.telus.net
[ec2-user@ip-172-30-1-86 ~]$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            NAMES
f0c0e53ecffb      docker kafka/kafka   "/docker-entrypoint.s"   About a minute ago   kafka
lab1be535f44       docker kafka/zookeeper   "/docker-entrypoint.s"   About a minute ago   zookeeper
[ec2-user@ip-172-30-1-86 ~]$
```

Figure 12-5. Listing Running Docker Containers

Finding IP Addresses

To run the Kafka Producer and Consumer, we need to find the IP address of the Docker container running the ZooKeeper and IP address of the Docker container running the Kafka server. Run the following two commands to export the ZK_IP and KAFKA_IP environment variables.

```
export ZK_IP=$(sudo docker inspect --format '{{ .NetworkSettings.IPAddress }}' zookeeper)
export KAFKA_IP=$(sudo docker inspect --format '{{ .NetworkSettings.IPAddress }}' kafka)
```

Subsequently, echo the ZK_IP and KAFKA_IP variables. The ZK_IP is output as 172.17.0.1 and the KAFKA_IP is output as 172.17.0.2 as shown in Figure 12-6. We shall use these IP addresses in subsequent sections.

```
[ec2-user@ip-172-30-1-86 ~]$ export ZK_IP=$(sudo docker inspect --format '{{ .NetworkSettings.IPAddress }}' zookeeper)
[ec2-user@ip-172-30-1-86 ~]$ export KAFKA_IP=$(sudo docker inspect --format '{{ .NetworkSettings.IPAddress }}' kafka)
[ec2-user@ip-172-30-1-86 ~]$ echo $ZK_IP
172.17.0.1
[ec2-user@ip-172-30-1-86 ~]$ echo $KAFKA_IP
172.17.0.2
[ec2-user@ip-172-30-1-86 ~]$
```

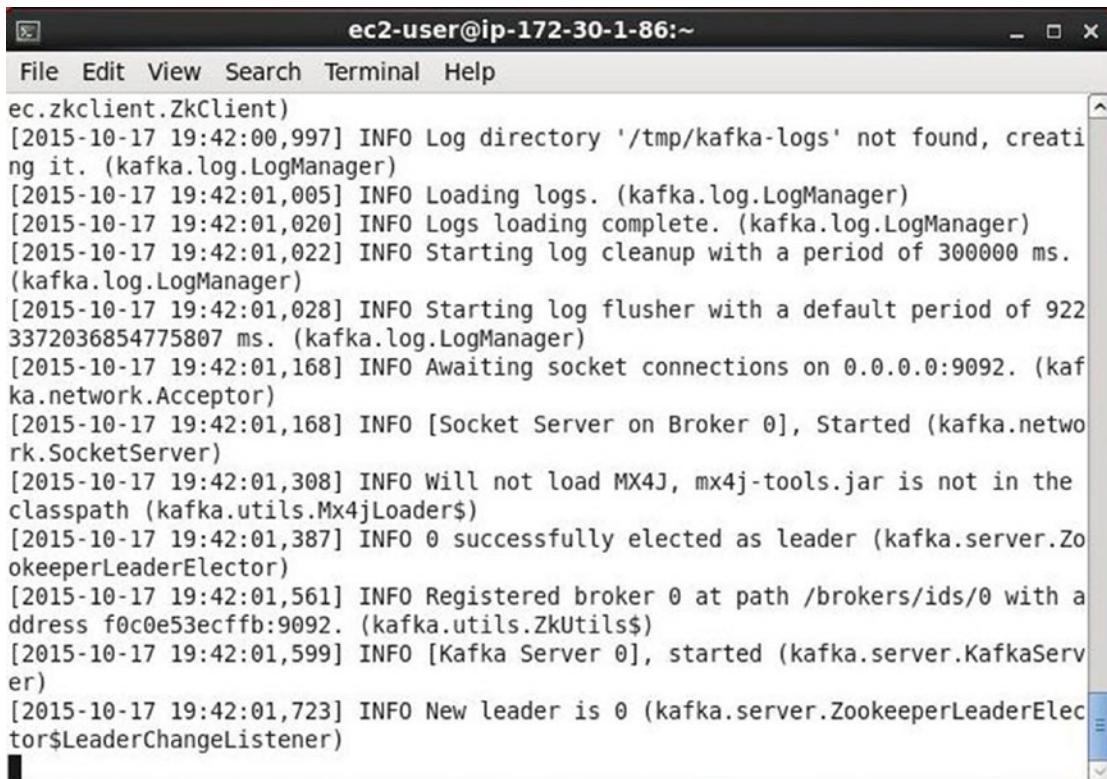
Figure 12-6. Finding IP Addresses for Zookeeper and Kafka Servers

Listing the Kafka Logs

Output the logs for the Docker container “kafka” with the docker logs command.

```
sudo docker logs -f kafka
```

The output indicates that the Kafka server got started as shown in Figure 12-7.



A screenshot of a terminal window titled "ec2-user@ip-172-30-1-86:~". The window contains the following log output:

```
File Edit View Search Terminal Help
ec.zkclient.ZkClient)
[2015-10-17 19:42:00,997] INFO Log directory '/tmp/kafka-logs' not found, creating it. (kafka.log.LogManager)
[2015-10-17 19:42:01,005] INFO Loading logs. (kafka.log.LogManager)
[2015-10-17 19:42:01,020] INFO Logs loading complete. (kafka.log.LogManager)
[2015-10-17 19:42:01,022] INFO Starting log cleanup with a period of 300000 ms. (kafka.log.LogManager)
[2015-10-17 19:42:01,028] INFO Starting log flusher with a default period of 922 3372036854775807 ms. (kafka.log.LogManager)
[2015-10-17 19:42:01,168] INFO Awaiting socket connections on 0.0.0.0:9092. (kafka.network.Acceptor)
[2015-10-17 19:42:01,168] INFO [Socket Server on Broker 0], Started (kafka.network.SocketServer)
[2015-10-17 19:42:01,308] INFO Will not load MX4J, mx4j-tools.jar is not in the classpath (kafka.utils.Mx4jLoader$)
[2015-10-17 19:42:01,387] INFO 0 successfully elected as leader (kafka.server.ZookeeperLeaderElector)
[2015-10-17 19:42:01,561] INFO Registered broker 0 at path /brokers/ids/0 with address f0c0e53ecffb:9092. (kafka.utils.ZkUtils$)
[2015-10-17 19:42:01,599] INFO [Kafka Server 0], started (kafka.server.KafkaServer)
[2015-10-17 19:42:01,723] INFO New leader is 0 (kafka.server.ZookeeperLeaderElector$LeaderChangeListener)
```

Figure 12-7. Listing Kafka Logs

In subsequent sections we shall create a Kafka topic, start a Kafka producer, start a Kafka consumer and produce messages at the Kafka Producer to be published at a Kafka topic, and consume the messages at the Kafka Consumer.

Creating a Kafka Topic

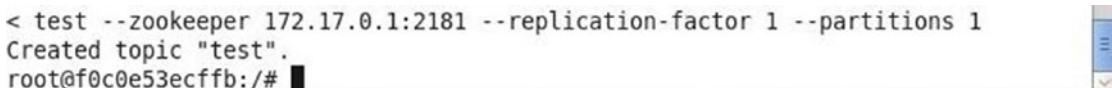
First, we need to create a Kafka topic to publish messages to. Start the interactive terminal with the following command.

```
sudo docker exec -it kafka bash
```

Create a Kafka topic in the interactive terminal with the `kafka-topics.sh --create` command. Specify the topic to create with the `-topic` option as “test”. Specify the ZooKeeper address as the IP address for the ZooKeeper obtained earlier and set in the environment variable `ZK_IP`. Specify the ZooKeeper port as 2181. The number of partitions is set to 1 with the `--partitions` option, and the replication factor is set to 1 with the `--replication-factor` option.

```
kafka-topics.sh --create --topic test --zookeeper 172.17.0.1:2181 --replication-factor 1
--partitions 1
```

The output from the command is Created topic “test” as shown in Figure 12-8.



```
< test --zookeeper 172.17.0.1:2181 --replication-factor 1 --partitions 1
Created topic "test".
root@f0c0e53ecffb:/# █
```

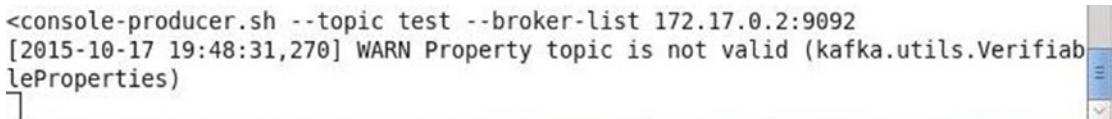
Figure 12-8. Creating a Kafka topic

Starting the Kafka Producer

Next, start the Kafka producer with the following command from an interactive terminal for the “kafka” container running the Kafka server. The broker list is specified as 172.17.0.2:9092 in which the IP address is the environment variable `KAFKA_IP` exported earlier. The port Kafka server listens on is 9092. The topic to which the messages are to be published is set with the `-topic` option as “test”.

```
kafka-console-producer.sh --topic test --broker-list 172.17.0.2:9092
```

Kafka producer console gets started as shown in Figure 12-9.



```
<console-producer.sh --topic test --broker-list 172.17.0.2:9092
[2015-10-17 19:48:31,270] WARN Property topic is not valid (kafka.utils.VerifiableProperties)
[] █
```

Figure 12-9. Starting the Kafka Producer

Starting the Kafka Consumer

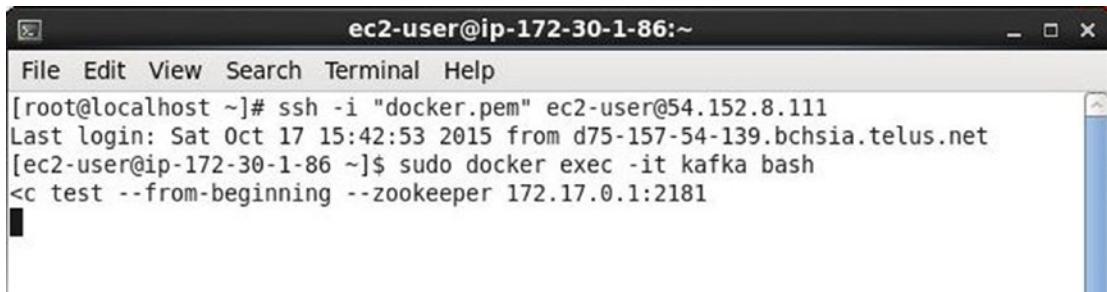
For the Kafka Consumer console we need to start another interactive terminal for the “kafka” container.

```
sudo docker exec -it kafka bash
```

Run the following command to start the Kafka consumer console to consume messages published to the “test” topic as specified with the `-topic` option. The ZooKeeper host:port is set with the `-zookeeper` option to 172.17.0.1:2181 in which the IP Address is the environment variable `ZK_IP` and the port is 2181. The `--from-beginning` option implies that messages are to be consumed from the beginning.

```
kafka-console-consumer.sh --topic test --from-beginning --zookeeper 172.17.0.1:2181
```

The Kafka consumer console gets started as shown in Figure 12-10.



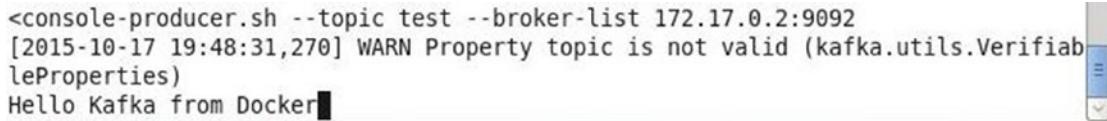
```
[root@localhost ~]# ssh -i "docker.pem" ec2-user@54.152.8.111
Last login: Sat Oct 17 15:42:53 2015 from d75-157-54-139.bchsysa.telus.net
[ec2-user@ip-172-30-1-86 ~]$ sudo docker exec -it kafka bash
<test --from-beginning --zookeeper 172.17.0.1:2181
```

Figure 12-10. Starting the Kafka Consumer

Producing and Consuming Messages

In this section we shall publish messages from the Kafka Producer to the Kafka topic “test” configured when we started the Producer, and consume the messages at the Kafka consumer also subscribed to the “test” topic.

Publish a message “Hello Kafka from Docker” at the Producer console as shown in Figure 12-11. Click on Enter to navigate to the next line in the console.



```
<console-producer.sh --topic test --broker-list 172.17.0.2:9092
[2015-10-17 19:48:31,270] WARN Property topic is not valid (kafka.utils.VerifiableProperties)
Hello Kafka from Docker
```

Figure 12-11. Producing a Message at the Kafka Producer

The message published to the “test” topic gets consumed at the Kafka Consumer and gets output in the Consumer console as shown in Figure 12-12.



```
[root@localhost ~]# ssh -i "docker.pem" ec2-user@54.152.8.111
Last login: Sat Oct 17 15:42:53 2015 from d75-157-54-139.bchsysa.telus.net
[ec2-user@ip-172-30-1-86 ~]$ sudo docker exec -it kafka bash
<test --from-beginning --zookeeper 172.17.0.1:2181
Hello Kafka from Docker
```

Figure 12-12. Consuming Messages at the Consumer

Similarly, publish more messages to the “test” topic from the Kafka Producer as shown in Figure 12-13.

```
<console-producer.sh --topic test --broker-list 172.17.0.2:9092
[2015-10-17 19:48:31,270] WARN Property topic is not valid (kafka.utils.VerifiableProperties)
Hello Kafka from Docker
HELLO KAFKA from DOCKER
HELLO kaFKA FR0m DocKER
HeLlO KAFka fR0m dOCKER
heLLO kaFKA FROM dockER
```

Figure 12-13. Producing More Messages at the Producer

The messages get output at the Kafka Consumer console as shown in Figure 12-14.



```
ec2-user@ip-172-30-1-86:~
File Edit View Search Terminal Help
[root@localhost ~]# ssh -i "docker.pem" ec2-user@54.152.8.111
Last login: Sat Oct 17 15:42:53 2015 from d75-157-54-139.bchsia.telus.net
[ec2-user@ip-172-30-1-86 ~]$ sudo docker exec -it kafka bash
<c test --from-beginning --zookeeper 172.17.0.1:2181
Hello Kafka from Docker
HELLO KAFKA from DOCKER
HELLO kaFKA FR0m DocKER
HeLlO KAFka fR0m dOCKER
heLLO kaFKA FROM dockER
```

Figure 12-14. Consumming Messages

Stopping and Removing the Docker Containers

To stop the Docker containers, run the `docker stop` command. Stop the “kafka” container as follows.

```
sudo docker stop kafka
```

The “kafka” container may be removed with the `docker rm` command.

```
sudo docker rm kafka
```

Similarly, stop and remove the Docker container “zookeeper”.

```
sudo docker stop zookeeper
sudo docker rm zookeeper
```

Summary

In this chapter we used Docker containers for Apache ZooKeeper and Apache Kafka to run a Kafka server process linked to an Apache ZooKeeper process. We created a Kafka Topic, started a Kafka producer, started a Kafka Consumer, published messages to the topic from the Kafka producer and consumed the messages at the Consumer. In the next chapter we shall discuss using Apache Solr with Docker.

CHAPTER 13



Using Apache Solr

Apache Solr is an open source search platform built on Apache Lucene, a text search engine library. Apache Solr is scalable and reliable and provides indexing and querying service. Cloudera Search is based on Apache Solr. In this chapter we shall use the official Docker image for Apache Solr to run Apache Solr in a Docker container. This chapter has the following sections.

- Setting the Environment
- Starting Docker Container for Apache Solr Server
- Starting Interactive Shell
- Logging in to the Solr Admin Console
- Creating a Core Admin Index
- Loading Sample Data
- Querying Apache Solr in Solr Admin Console
- Querying Apache Solr using REST API Client
- Deleting Data
- Listing Logs
- Stopping Apache Solr Server

Setting the Environment

The following software is required for this chapter.

- Docker Engine (version 1.8)
- Docker image for Apache Solr

We will use an Amazon EC2 instance based on the Ubuntu Server 14.04 LTS (HVM), SSD Volume Type - ami-d05e75b8. Login to the Amazon EC2 instance with the user name “ubuntu” and the public IP address of the Amazon EC2 instance.

```
ssh -i "docker.pem" ubuntu@54.208.53.110
```

Ubuntu instance on Amazon EC2 gets logged in to as shown in Figure 13-1.

```

ubuntu@ip-172-30-1-190: ~
File Edit View Search Terminal Help
[root@localhost ~]# ssh -i "docker.pem" ubuntu@54.208.53.110
Welcome to Ubuntu 14.04.3 LTS (GNU/Linux 3.13.0-66-generic x86_64)

 * Documentation:  https://help.ubuntu.com/
 
 System information as of Fri Oct 23 13:44:42 UTC 2015

 System load:  0.02          Processes:           154
 Usage of /:   76.1% of 7.74GB  Users logged in:      1
 Memory usage: 46%           IP address for eth0:  172.30.1.190
 Swap usage:   0%             IP address for docker0: 172.17.42.1

 Graph this data and manage this system at:
 https://landscape.canonical.com/
 
 Get cloud support with Ubuntu Advantage Cloud Guest:
 http://www.ubuntu.com/business/services/cloud

12 packages can be updated.
6 updates are security updates.

Last login: Fri Oct 23 13:43:45 2015 from d75-157-54-139.bchsia.telus.net
ubuntu@ip-172-30-1-190:~$ █

```

Figure 13-1. Logging in to Ubuntu on AmazonEC2

Install Docker on Ubuntu as discussed in Chapter 1. Start the Docker service. If Docker is already started, a message “start: Job is already running: docker” gets output.

```
sudo service docker start
```

Docker service status may be output with the following command.

```
sudo service docker status
```

A message indicating that a docker process is running gets output as shown in Figure 13-2.

```

ubuntu@ip-172-30-1-190:~$ sudo service docker start
start: Job is already running: docker
ubuntu@ip-172-30-1-190:~$ sudo service docker status
docker start/running, process 898
ubuntu@ip-172-30-1-190:~$ █

```

Figure 13-2. Starting and Finding Docker Service Status

Next, download the official Docker image for Apache Solr with the docker pull command.

```
sudo docker pull solr
```

The Docker image gets downloaded as shown in Figure 13-3.

```
ubuntu@ip-172-30-1-190: ~
File Edit View Search Terminal Help
41d087dfc152: Pull complete
91e60c645519: Pull complete
999bfec8c1c0: Pull complete
fe71f7e2e724: Pull complete
3aff24e7e0cf: Pull complete
a2c988a24043: Pull complete
89f566edf57e: Pull complete
43db431fe52c: Pull complete
33f61a6e7e2d: Pull complete
2907972f275b: Pull complete
634bab7a910f: Pull complete
0b12ca9eae90: Pull complete
714f1e1c798b: Pull complete
869cc5a38aa9: Pull complete
6a8591d747a6: Pull complete
61ffde14c319: Pull complete
2bc40cd89ee9: Pull complete
e7db410e45aa: Pull complete
8b1a14bf57bc: Pull complete
38730e98ec07: Pull complete
6af2834071ca: Pull complete
f30488539021: Pull complete
5683815705ca: Pull complete
da0cb77ff0cb: Pull complete
a866495d197c: Pull complete
Digest: sha256:13af873d09a10c486860cb16fadac4d4f693f525ab0e64716b161b8bd3042764
Status: Downloaded newer image for solr:latest
ubuntu@ip-172-30-1-190:~$
```

Figure 13-3. Downloading Docker Image solr

Starting Docker Container for Apache Solr Server

To start Apache Solr server run the docker run command with -p as 8983. Specify the container name with --name option as “solr_on_docker,” which is arbitrary. The -d command parameter makes the Docker container run in a detached mode.

```
sudo docker run -p 8983:8983 -d --name solr_on_docker solr
```

List the running Docker containers with the docker ps command.

```
sudo docker ps
```

The Docker container running Apache Solr get listed including the container id assigned to the container as shown in Figure 13-4.


```
ubuntu@ip-172-30-1-190:~$ sudo docker run -p 8983:8983 -d --name solr_on_docker solr
8061f79d1f1631ad824f659e54c2d3ee08e11f5da296f1377c3a0a78263bebd4
ubuntu@ip-172-30-1-190:~$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS
8061f79d1f16        solr                "/opt/solr/bin/solr -"   4 seconds ago
Up 3 seconds         0.0.0.0:8983->8983/tcp    solr_on_docker
ubuntu@ip-172-30-1-190:~$
```

Figure 13-4. Starting Docker Container for Apache Solr

Run the docker logs command to output the logs for the Docker container. Either the container name or container id may be used in docker commands.

```
sudo docker logs -f 8061f79d1f16
```

The container logs indicate that the Apache Solr server has started as shown in Figure 13-5.



The screenshot shows a terminal window titled "ubuntu@ip-172-30-1-190: ~". The window displays the log output of an Apache Solr Docker container. The log includes various INFO and WARN messages from different Solr components like SolrResourceLoader, HttpShardHandlerFactory, and LogWatcher, indicating the startup process and configuration loading.

```

1234 WARN (main) [ ] o.a.s.c.SolrResourceLoader No files added to classloader from lib: lib (resolved as: /opt/solr/server/solr/lib).
1260 INFO (main) [ ] o.a.s.h.c.HttpShardHandlerFactory created with socketTim eout : 600000,connTimeout : 60000,maxConnectionsPerHost : 20,maxConnections : 10 000,corePoolSize : 0,maximumPoolSize : 2147483647,maxThreadIdleTime : 5,sizeOfQu eue : -1,fairnessPolicy : false,useRetries : false,
1583 INFO (main) [ ] o.a.s.u.UpdateShardHandler Creating UpdateShardHandler H TTP client with params: socketTimeout=600000&connTimeout=60000&retry=true
1585 INFO (main) [ ] o.a.s.l.LogWatcher SLF4J impl is org.slf4j.impl.Log4jLog gerFactory
1586 INFO (main) [ ] o.a.s.l.LogWatcher Registering Log Listener [Log4j (org. slf4j.impl.Log4jLoggerFactory)]
1588 INFO (main) [ ] o.a.s.c.CoreContainer Security conf doesn't exist. Skipp ing setup for authorization module.
1588 INFO (main) [ ] o.a.s.c.CoreContainer No authentication plugin used.
1626 INFO (main) [ ] o.a.s.c.CoresLocator Looking for core definitions undern eat /opt/solr/server/solr
1632 INFO (main) [ ] o.a.s.c.CoresLocator Found 0 core definitions
1635 INFO (main) [ ] o.a.s.s.SolrDispatchFilter user.dir=/opt/solr/server
1635 INFO (main) [ ] o.a.s.s.SolrDispatchFilter SolrDispatchFilter.init() don e
1648 INFO (main) [ ] o.e.j.s.h.ContextHandler Started o.e.j.w.WebAppContext@5 7ffffcd7{/solr,file:/opt/solr/server/solr-webapp/webapp/,AVAILABLE}{/opt/solr/ser ver/solr-webapp/webapp}
1659 INFO (main) [ ] o.e.j.s.ServerConnector Started ServerConnector@ad48334{ HTTP/1.1}{0.0.0.0:8983}
1659 INFO (main) [ ] o.e.j.s.Server Started @2075ms

```

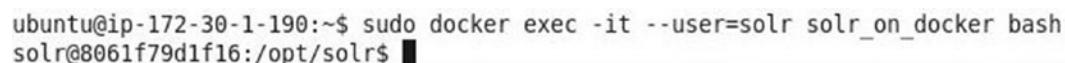
Figure 13-5. Listing Docker Container Log

Starting the Interactive Shell

Start the interactive shell for the Docker container as user “solr”:

```
sudo docker exec -it --user=solr solr_on_docker bash
```

The interactive shell (or tty) gets started as shown in Figure 13-6.



The screenshot shows a terminal window titled "ubuntu@ip-172-30-1-190:~\$". The command "sudo docker exec -it --user=solr solr_on_docker bash" is entered, followed by a prompt "solr@8061f79d1f16:/opt/solr\$". This indicates that the interactive shell has been successfully started within the Docker container as the "solr" user.

Figure 13-6. Starting TTY

Apache Solr commands may be run in the interactive terminal.

Logging in to the Solr Admin Console

If the Docker container running the Apache Solr server is running on a different host than the Admin Console, use the public DNS name of the Amazon EC2 instance running the Docker engine and the Docker container. Obtain the public DNS from the Amazon EC2 Management Console. The public DNS is `ec2-54-208-53-110.compute-1.amazonaws.com` as shown in Figure 13-7.

The screenshot shows the AWS EC2 Management Console interface. On the left, there's a sidebar with navigation links for EC2 Dashboard, Events, Tags, Reports, Limits, Instances (selected), Spot Requests, Reserved Instances, Images, AMIs, Bundle Tasks, Elastic Block Store, Volumes, Snapshots, Network & Security, Security Groups, Elastic IPs, Placement Groups, and Key Pairs. The main area displays a table of instances. One instance, named "Docker" with Instance ID i-024c364d, is selected. The table columns include Name, Instance ID, Instance Type, Availability Zone, Instance State, Status Checks, Alarm Status, Public DNS, and Public IP. The Public DNS for the selected instance is listed as "ec2-54-208-53-110.compute-1.amazonaws.com". Below the table, a detailed view for the selected instance shows its description, status checks (2/2 checks passed), monitoring, and tags. It also lists its Instance ID (i-024c364d), Instance state (running), and Instance type (t2.micro). The Public DNS and Public IP are also displayed here.

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS	Public IP
i-036b01bc	m3.medium	us-east-1b	stopped	None				
Docker	i-0259f40	t2.micro	us-east-1b	stopped	None			
Docker	i-016ee63a	t2.micro	us-east-1b	stopped	None			
Docker	i-036e693c	m3.medium	us-east-1b	stopped	None			
Docker	i-005265d	m3.medium	us-east-1b	stopped	None			
Docker	i-0ed571bc	m3.medium	us-east-1b	stopped	None			
Docker	i-024c364d	t2.micro	us-east-1b	running	2/2 checks...	None	ec2-54-208-53-110.com...	54.208.53.110
Ubuntu test	i-0247029d	t2.micro	us-east-1b	stopped	None			

Instance: i-024c364d (Docker) Public DNS: ec2-54-208-53-110.compute-1.amazonaws.com

Description	Status Checks	Monitoring	Tags
Instance ID: i-024c364d	Public DNS: ec2-54-208-53-110.compute-1.amazonaws.com	Public IP: 54.208.53.110	Elastic IP: -
Instance state: running	2/2 checks passed		
Instance type: t2.micro			

Figure 13-7. Finding the Public DNS

Use the URL <http://ec2-54-208-53-110.compute-1.amazonaws.com:8983/> to access the Apache Solr Admin Console. The Dashboard is shown in Figure 13-8.

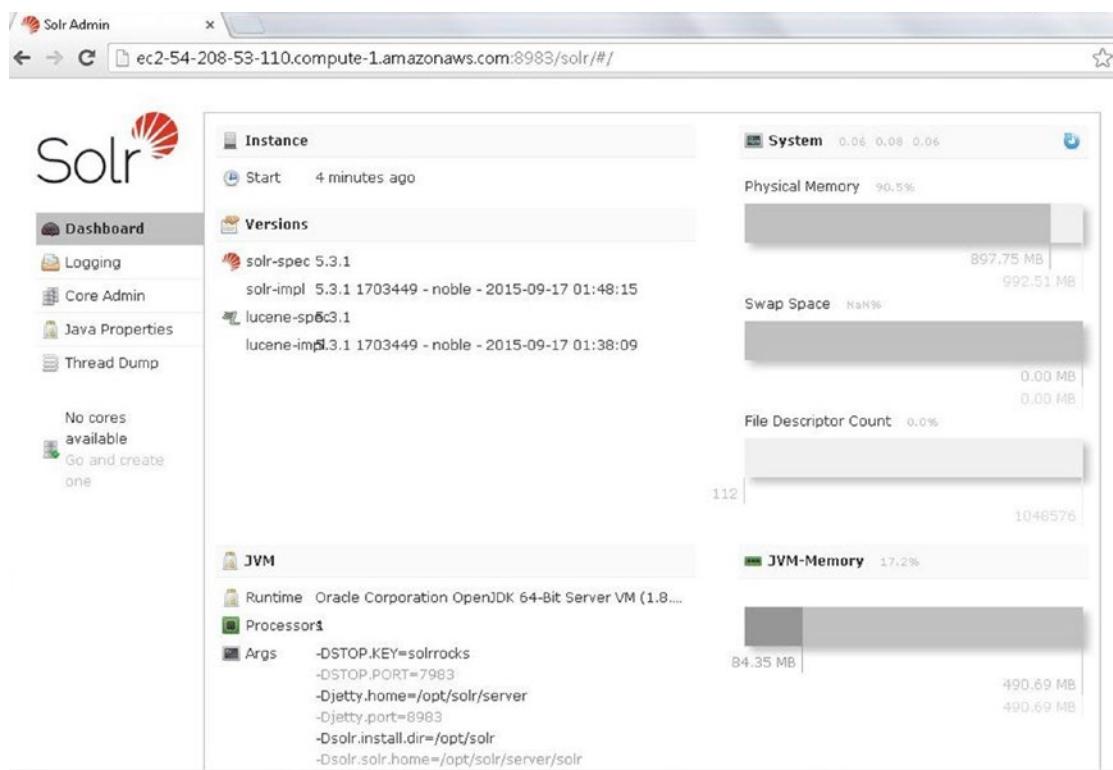


Figure 13-8. Logging in to Solr Admin Console

Creating a Core Index

Next, create a core, an index for the data to be stored in Apache Solr. From the tty run the `bin/solr create_core` command to create a core called `gettingstarted`.

```
bin/solr create_core -c gettingstarted
```

A new core called “`gettingstarted`” gets created as shown in Figure 13-9.

```
ubuntu@ip-172-30-1-190:~$ sudo docker exec -it --user=solr solr_on_docker bash
solr@8061f79d1f16:/opt/solr$ bin/solr create_core -c gettingstarted

Setup new core instance directory:
/opt/solr/server/solr/gettingstarted

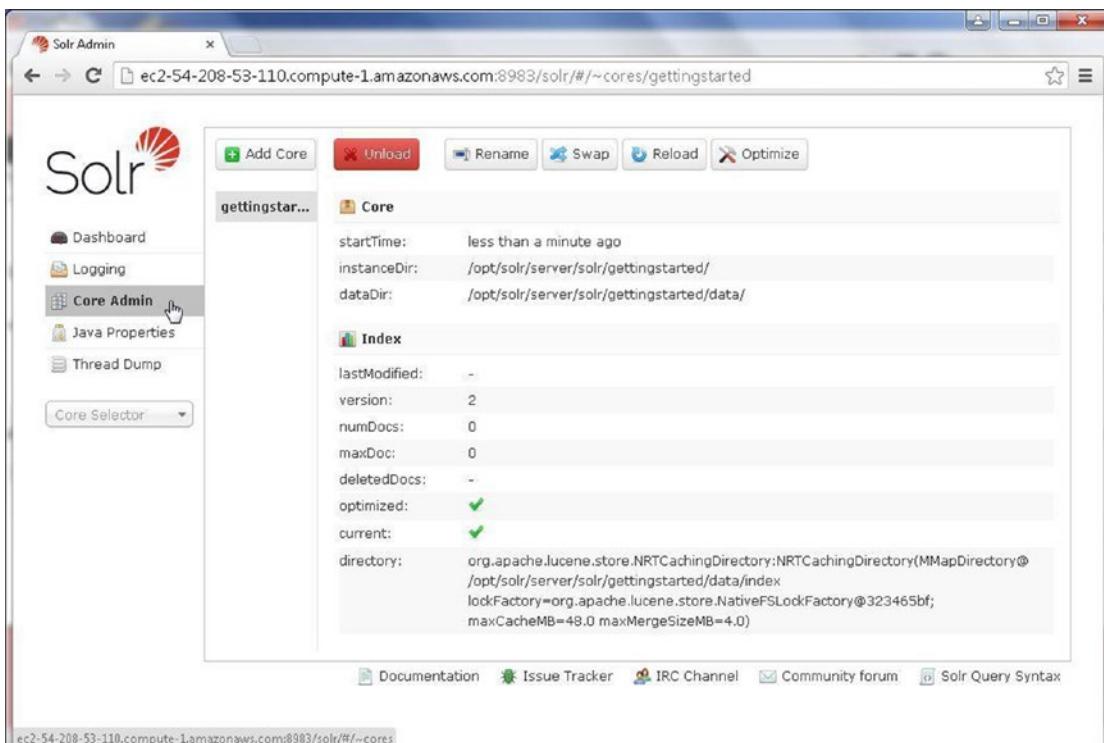
Creating new core 'gettingstarted' using command:
http://localhost:8983/solr/admin/cores?action=CREATE&name=gettingstarted&instanceDir=gettingstarted

{
  "responseHeader": {
    "status": 0,
    "QTime": 2102,
    "core": "gettingstarted"
  }
}

solr@8061f79d1f16:/opt/solr$
```

Figure 13-9. Creating a Core called “gettingstarted”

In the Solr Admin Console, select Core Admin as shown in Figure 13-10.

**Figure 13-10.** Selecting Core Admin

In the Core Selector, select the `gettingstarted` core as shown in Figure 13-11.

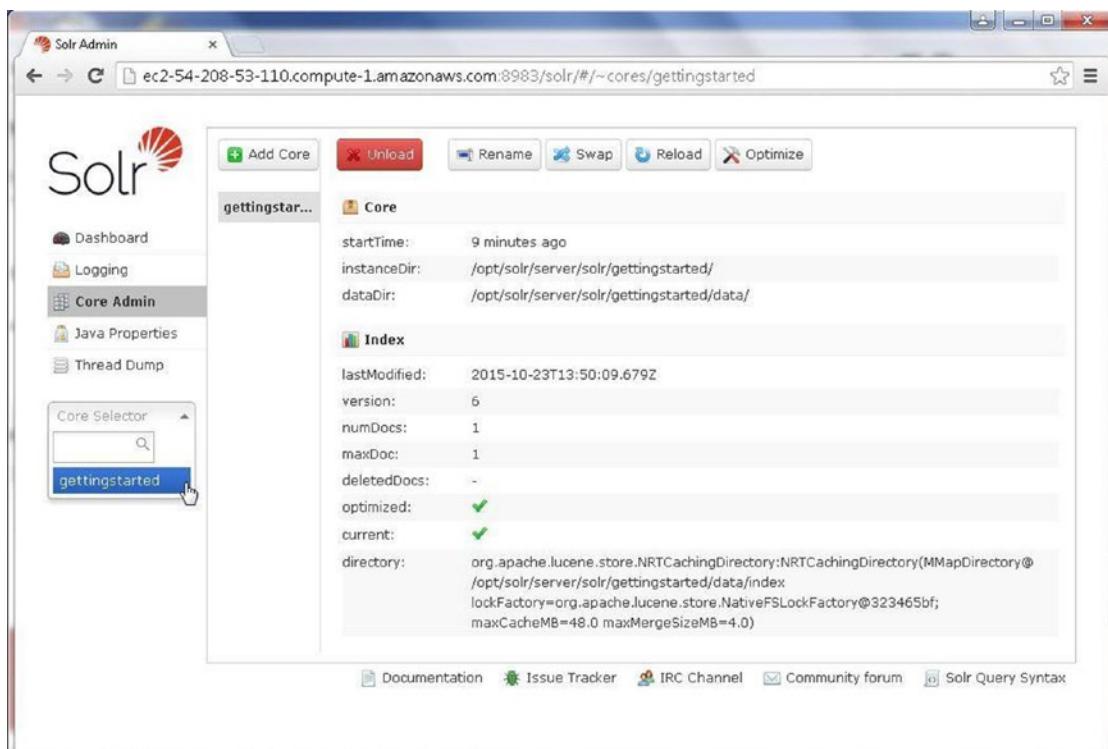


Figure 13-11. Selecting the `gettingstarted` Core

Select Overview tab in the margin as shown in Figure 13-12. The index stats get listed such as the version, Num Docs, Max Doc, and Deleted.

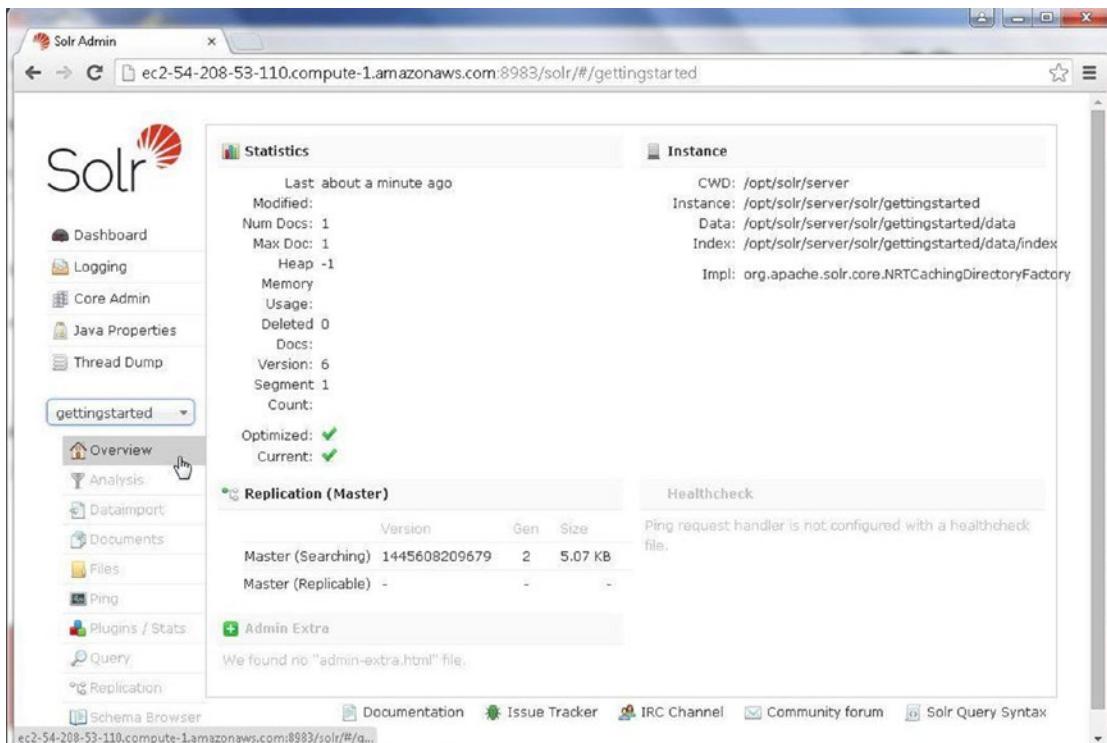


Figure 13-12. Displaying the Overview of the gettingstarted Core

Loading Sample Data

Apache Solr supports indexing of documents in XML, JSON and CSV formats. We shall index using the XML format. The root element is required to be `<add>` and each document must be enclosed in the `<doc>` element. The `id` field is required. We shall index the following XML format document. Store the document as `solr.xml`.

```

<add>
<doc>
  <field name="id">SOLR1000</field>
  <field name="name">Solr, the Enterprise Search Server</field>
  <field name="manu">Apache Software Foundation</field>
  <field name="cat">software</field>
  <field name="cat">search</field>
  <field name="features">Advanced Full-Text Search Capabilities using Lucene</field>
  <field name="features">Optimized for High Volume Web Traffic</field>
  <field name="features">Standards Based Open Interfaces - XML and HTTP</field>
  <field name="features">Comprehensive HTML Administration Interfaces</field>
  <field name="features">Scalability - Efficient Replication to other Solr Search Servers</field>
  <field name="features">Flexible and Adaptable with XML configuration and Schema</field>
  <field name="features">Good unicode support: h&#xE9;llo (hello with an accent over the e)</field>
  <field name="price">0</field>

```

```


<field name="popularity">10</field>
<field name="inStock">true</field>
<field name="incubationdate_dt">2006-01-17T00:00:00.000Z</field>
</doc>
</add>

```

Copy the solr.xml to the /opt/solr directory in the Docker container. Run the following docker cp command from the Ubuntu host, not the Docker container, to copy the solr.xml document to the Docker container with id 8061f79d1f16, which is running the Apache Solr server. The container id may be obtained from the output of the docker ps command.

```
sudo docker cp solr.xml 8061f79d1f16:/opt/solr/solr.xml
```

The solr.xml document gets copied to the /opt/solr directory in the Docker container as shown in Figure 13-13.




```
ubuntu@ip-172-30-1-190:~$ sudo docker cp solr.xml 8061f79d1f16:/opt/solr/solr.xml
l
ubuntu@ip-172-30-1-190:~$ █
```

Figure 13-13. Copying solr.xml to DockerContainer

Start the interactive terminal (tty) with the following command.

```
sudo docker exec -it -user=solr solr_on_docker bash
```

From the /opt/solr directory run the following command to list the files and directories in the directory. The solr.xml should get listed as shown in Figure 13-14.



```
solr@8061f79d1f16:/opt/solr$ ls -l
total 1164
-rw-r--r-- 1 solr solr 503614 Sep 16 19:07 CHANGES.txt
-rw-r--r-- 1 solr solr 12646 Aug 12 09:16 LICENSE.txt
-rw-r--r-- 1 solr solr 566457 Sep 9 11:31 LUCENE_CHANGES.txt
-rw-r--r-- 1 solr solr 26529 Aug 12 09:16 NOTICE.txt
-rw-r--r-- 1 solr solr 7167 Aug 12 09:16 README.txt
drwxr-xr-x 3 solr solr 4096 Oct 16 23:08 bin
drwxr-xr-x 13 solr solr 4096 Sep 16 20:20 contrib
drwxr-xr-x 4 solr solr 4096 Oct 16 23:08 dist
drwxr-xr-x 19 solr solr 4096 Oct 16 23:08 docs
drwxr-xr-x 7 solr solr 4096 Oct 16 23:08 example
drwxr-xr-x 2 solr solr 36864 Oct 16 23:08 licenses
drwxr-xr-x 13 solr solr 4096 Oct 23 13:42 server
-rwxrwxr-x 1 root root 1048 Oct 23 13:44 solr.xml
solr@8061f79d1f16:/opt/solr$ █
```

Figure 13-14. Listing the solr.xml File in Docker Container

Run the following command to post the solr.xml to the gettingstarted index.

```
bin/post -c gettingstarted ./solr.xml
```

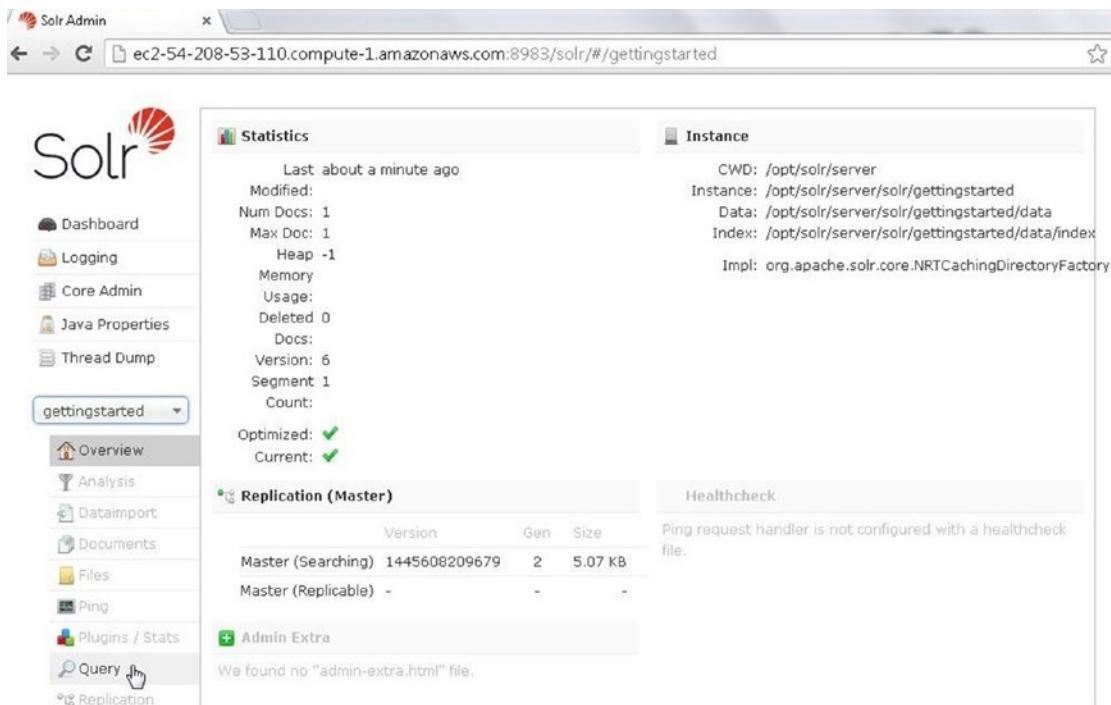
The solr.xml file gets indexed as shown in Figure 13-15.

```
solr@8061f79d1f16:/opt/solr$ bin/post -c gettingstarted ./solr.xml
java -classpath /opt/solr/dist/solr-core-5.3.1.jar -Dauto=yes -Dc=gettingstarted
-Ddata=files org.apache.solr.util.SimplePostTool ./solr.xml
SimplePostTool version 5.0.0
Posting files to [base] url http://localhost:8983/solr/gettingstarted/update...
Entering auto mode. File endings considered are xml,json,csv,pdf,doc,docx,ppt,pp
tx,xls,xlsx,odt,odp,ods,ott,otp,ots,rtf,htm,html,txt,log
POSTing file solr.xml (application/xml) to [base]
1 files indexed.
COMMITting Solr index changes to http://localhost:8983/solr/gettingstarted/updat
e...
Time spent: 0:00:00.402
solr@8061f79d1f16:/opt/solr$
```

Figure 13-15. indexing solr.xml

Querying Apache Solr in Solr Admin Console

The indexed document may be queried from the Solr Admin console. Select the Query tab as shown in Figure 13-16.



The screenshot shows the Apache Solr Admin interface. The left sidebar has a navigation menu with items like Dashboard, Logging, Core Admin, Java Properties, Thread Dump, and several tabs for the 'gettingstarted' core: Overview (selected), Analysis, Dataimport, Documents, Files, Ping, Plugins / Stats, Query (selected), and Replication.

The main content area is divided into several sections:

- Statistics:** Last about a minute ago. Modified: Num Docs: 1, Max Doc: 1, Heap: -1, Memory, Usage, Deleted: 0, Docs: Version: 6, Segment 1, Count: Optimized: ✓, Current: ✓
- Instance:** CWD: /opt/solr/server, Instance: /opt/solr/server/solr/gettingstarted, Data: /opt/solr/server/solr/gettingstarted/data, Index: /opt/solr/server/solr/gettingstarted/data/index, Impl: org.apache.solr.core.NRTCachingDirectoryFactory
- Replication (Master):** Shows a table with columns Version, Gen, Size. One row is visible: Master (Searching) 1445608209679 2 5.07 KB. Another row is listed as Master (Replicable) - - -
- Healthcheck:** Ping request handler is not configured with a healthcheck file.
- Admin Extra:** We found no "admin-extra.html" file.

Figure 13-16. Selecting the Query tab

The Request-Handler (qt) should be set to /select and the query should be set to * * to select all documents in the index as shown in Figure 13-17. The start index is set to 0 and the number of rows to select is set to 10. The wt (response writer) is set to json to return the queried documents in JSON format. Other supported formats are XML and CSV.

The screenshot shows the Solr Admin interface for a Solr instance at `ec2-54-208-53-110.compute-1.amazonaws.com:8983/solr/#/gettingstarted/query`. The left sidebar shows navigation links like Dashboard, Logging, Core Admin, Java Properties, Thread Dump, and sections for gettingstarted (Overview, Analysis, Dataimport, Documents, Files, Ping, Plugins / Stats, Query, Replication, Schema Browser). The 'Query' link is currently selected. The main panel displays the configuration for the '/select' Request Handler. The 'Request-Handler (qt)' dropdown is set to '/select'. The 'q' field contains the query `*:*`. The 'fq' field is empty. The 'sort' field is empty. The 'start, rows' section has '0' in the start field and '10' in the rows field. The 'fl' field is empty. The 'df' field is empty. The 'Raw Query Parameters' section contains the key-value pair `key1=val1&key2=val2`. The 'wt' (writer type) dropdown is set to 'json'. Two checkboxes are present: 'indent' (checked) and 'debugQuery' (unchecked).

Figure 13-17. The /select Request Handler

Click on Execute Query as shown in Figure 13-18.




Figure 13-18. Clicking on Execute Query

The query result gets returned as JSON as shown in Figure 13-19.


The screenshot shows the Solr Admin interface for a 'gettingstarted' core. On the left, there's a sidebar with various navigation links like Dashboard, Logging, Core Admin, Java Properties, Thread Dump, and others under the 'gettingstarted' section. The main area is titled 'Request-Handler (qt)' and shows a form for a '/select' query. The 'q' field contains '*:*'. The 'response' tab is selected, displaying the JSON response. The response header includes status 0, QTime 24, and parameters like q=* and wt=json. The response body shows one document with id=SOLR1000, name="Solr, the Enterprise Search Server", manu="Apache Software Foundation", cat="software, search", and features="Advanced Full-Text Search Capabilities using Lucene".

```
http://ec2-54-208-53-110.compute-1.amazonaws.com:8983/solr/gettingstarted/select?q=*&wt=json&indent=true

{
  "responseHeader": {
    "status": 0,
    "QTime": 24,
    "params": {
      "q": "*:*",
      "indent": "true",
      "wt": "json",
      "_": "1445608349663"
    }
  },
  "response": {
    "numFound": 1,
    "start": 0,
    "docs": [
      {
        "id": "SOLR1000",
        "name": [
          "Solr, the Enterprise Search Server"
        ],
        "manu": [
          "Apache Software Foundation"
        ],
        "cat": [
          "software",
          "search"
        ],
        "features": [
          "Advanced Full-Text Search Capabilities using Lucene"
        ]
      }
    ]
  }
}
```

Figure 13-19. JSON Response from Query

The *version* field gets added to the JSON document returned as shown in Figure 13-20.



The screenshot shows the Solr Admin interface with the URL `ec2-54-208-53-110.compute-1.amazonaws.com:8983/solr/#/gettingstarted/query`. The left sidebar has links for Plugins / Stats, Query (selected), Replication, Schema Browser, and Segments info. The main area has a dropdown menu for 'wt' set to 'json'. Below it are several checkboxes: 'indent' (checked), 'debugQuery', 'dismax', 'edismax', 'hl', 'facet', 'spatial', and 'spellcheck'. A blue 'Execute Query' button is at the bottom. The right pane displays a JSON response with the following content:

```
{
  "cat": [
    "software",
    "search"
  ],
  "features": [
    "Advanced Full-Text Search Capabilities using Lucene",
    "Optimized for High Volume Web Traffic",
    "Standards Based Open Interfaces - XML and HTTP",
    "Comprehensive HTML Administration Interfaces",
    "Scalability - Efficient Replication to other Solr Search Servers",
    "Flexible and Adaptable with XML configuration and Schema",
    "Good unicode support: héllø (hello with an accent over the e)"
  ],
  "price": [
    0
  ],
  "popularity": [
    10
  ],
  "inStock": [
    true
  ],
  "incubationdate_dt": "2006-01-17T00:00:00Z",
  "_version_": 1515830073872089000
}
```

Figure 13-20. The `_version_` field added automatically

Querying Apache Solr using REST API Client

The Apache Solr indexed documents may also be accessed using the REST client such as `curl`. For example, query all the documents in the `gettingstarted` index using the following `curl` command run from the interactive terminal for the “solr” container.

```
curl http://ec2-54-208-53-110.compute-1.amazonaws.com:8983/solr/gettingstarted/select?q=*%3A*&wt=json&indent=true
```

All the documents indexed in the `gettingstarted` index get output as shown in Figure 13-21.

```
<.com:8983/solr/gettingstarted/select?q=*&wt=json&indent=true"
{
  "responseHeader": {
    "status": 0,
    "QTime": 1,
    "params": {
      "q": "*:*",
      "indent": "true",
      "wt": "json" {}
    }
  },
  "response": {
    "numFound": 1,
    "start": 0,
    "docs": [
      {
        "id": "SOLR1000",
        "name": ["Solr, the Enterprise Search Server"],
        "manu": ["Apache Software Foundation"],
        "cat": ["software", "search"],
        "features": ["Advanced Full-Text Search Capabilities using Lucene", "Optimized for High Volume Web Traffic", "Standards Based Open Interfaces - XML and HTTP", "Comprehensive HTML Administration Interfaces", "Scalability - Efficient Replication to other Solr Search Servers", "Flexible and Adaptable with XML configuration and Schema", "Good unicode support: héllø (hello with an accent over the e)"],
        "price": [0],
        "popularity": [10],
        "inStock": [true],
        "incubationdate_dt": "2006-01-17T00:00:00Z",
        "_version_": 1515830073872089088
      }
    ]
  }
}
solr@8061f79d1f16:/opt/solr$
```

Figure 13-21. Running a REST Client Query

As another example, query all documents with “Lucene” in the document.

```
curl "http://ec2-54-208-53-110.compute-1.amazonaws.com:8983/solr/gettingstarted/select?wt=json&indent=true&q=Lucene"
```

As the single document indexed has “Lucene” in it the document gets returned as shown in Figure 13-22.

```
<.com:8983/solr/gettingstarted/select?wt=json&indent=true&q=Lucene"
{
  "responseHeader": {
    "status": 0,
    "QTime": 9,
    "params": {
      "q": "Lucene",
      "indent": "true",
      "wt": "json"}},
  "response": {"numFound": 1, "start": 0, "docs": [
    {
      "id": "SOLR1000",
      "name": ["Solr, the Enterprise Search Server"],
      "manu": ["Apache Software Foundation"],
      "cat": ["software",
        "search"],
      "features": ["Advanced Full-Text Search Capabilities using Lucene",
        "Optimized for High Volume Web Traffic",
        "Standards Based Open Interfaces - XML and HTTP",
        "Comprehensive HTML Administration Interfaces",
        "Scalability - Efficient Replication to other Solr Search Servers",
        "Flexible and Adaptable with XML configuration and Schema",
        "Good unicode support: héllø (hello with an accent over the e)"],
      "price": [0],
      "popularity": [10],
      "inStock": [true],
      "incubationdate_dt": "2006-01-17T00:00:00Z",
      "_version_": 1515830073872089088}]}
}
solr@8061f79d1f16:/opt/solr$
```

Figure 13-22. Running a REST Client Query using term ‘Lucene’

To query for a document with text in a specific field use the `field=text` format in the `q` parameter. For example, search for all documents with “Lucene” in the “name” field.

```
curl "http://ec2-54-208-53-110.compute-1.amazonaws.com:8983/solr/gettingstarted/select?wt=json&indent=true&q=name:Lucene"
```

As the name field of the single document in the index does not include “Lucene” no document gets returned as shown in Figure 13-23.

```
<.com:8983/solr/gettingstarted/select?wt=json&indent=true&q=name:Lucene"
{
  "responseHeader":{
    "status":0,
    "QTime":0,
    "params":{
      "q":"name:Lucene",
      "indent":"true",
      "wt":"json"}},
  "response":{"numFound":0,"start":0,"docs":[]}
}
solr@8061f79d1f16:/opt/solr$ █
```

Figure 13-23. Running a REST Client Query with “Lucene” in “name” Field

A phrase search may also be performed using the REST client. For example search for the phrase “Enterprise Search”.

```
curl "http://ec2-54-208-53-110.compute-1.amazonaws.com:8983/solr/gettingstarted/select?wt=json&indent=true&q=\"Enterprise+Search\""
```

As the single document has ‘Enterprise Search’ in it, the document gets returned as shown in Figure 13-24.

```
<ted/select?wt=json&indent=true&q=\"Enterprise+Search\""
{
  "responseHeader": {
    "status":0,
    "QTime":8,
    "params": {
      "q": "\"Enterprise Search\"",
      "indent": "true",
      "wt": "json"}},
  "response": {"numFound":1,"start":0,"docs": [
    {
      "id": "SOLR1000",
      "name": ["Solr, the Enterprise Search Server"],
      "manu": ["Apache Software Foundation"],
      "cat": ["software",
              "search"],
      "features": ["Advanced Full-Text Search Capabilities using Lucene",
                  "Optimized for High Volume Web Traffic",
                  "Standards Based Open Interfaces - XML and HTTP",
                  "Comprehensive HTML Administration Interfaces",
                  "Scalability - Efficient Replication to other Solr Search Servers",
                  "Flexible and Adaptable with XML configuration and Schema",
                  "Good unicode support: héllø (hello with an accent over the e)"],
      "price": [0],
      "popularity": [10],
      "inStock": [true],
      "incubationdate_dt": "2006-01-17T00:00:00Z",
      "_version_": 1515830073872089088}]}
}
solr@8061f79d1f16:/opt/solr$ █
```

Figure 13-24. Running a REST Query using a Phrase

Deleting Data

To delete a document run the same tool, the post tool, as used to post a document. Specify the document id to delete using the XML <delete><id>SOLR1000</id></delete>. The index to delete from is specified with the -c option.

```
bin/post -c gettingstarted -d "<delete><id>SOLR1000</id></delete>"
```

The single document indexed, which has the id SOLR1000, gets deleted as shown in Figure 13-25.

```
<$ bin/post -c gettingstarted -d "<delete><id>SOLR1000</id></delete>"  
java -classpath /opt/solr/dist/solr-core-5.3.1.jar -Dauto=yes -Dc=gettingstarted  
-Ddata=args org.apache.solr.util.SimplePostTool <delete><id>SOLR1000</id></dele  
te>  
SimplePostTool version 5.0.0  
POSTing args to http://localhost:8983/solr/gettingstarted/update...  
COMMITTING Solr index changes to http://localhost:8983/solr/gettingstarted/updat  
e...  
Time spent: 0:00:00.043  
solr@8061f79d1f16:/opt/solr$ █
```

Figure 13-25. Deleting a Single Document

Subsequently, run the same curl command as run before to search for all documents.

```
curl http://ec2-54-208-53-110.compute-1.amazonaws.com:8983/solr/gettingstarted/select?q=%3A*&wt=json&indent=true
```

No document gets found as shown in Figure 13-26 as the only document indexed has been deleted.

```
<$ bin/post -c gettingstarted -d "<delete><id>SOLR1000</id></delete>"  
java -classpath /opt/solr/dist/solr-core-5.3.1.jar -Dauto=yes -Dc=gettingstarted  
-Ddata=args org.apache.solr.util.SimplePostTool <delete><id>SOLR1000</id></dele  
te>  
SimplePostTool version 5.0.0  
POSTing args to http://localhost:8983/solr/gettingstarted/update...  
COMMITTING Solr index changes to http://localhost:8983/solr/gettingstarted/updat  
e...  
Time spent: 0:00:00.043  
<.com:8983/solr/gettingstarted/select?q=%3A*&wt=json&indent=true"  
{  
  "responseHeader":{  
    "status":0,  
    "QTime":0,  
    "params":{  
      "q": "*:*",  
      "indent":"true",  
      "wt":"json"}},  
  "response":{ "numFound":0,"start":0,"docs":[]}  
}  
solr@8061f79d1f16:/opt/solr$ █
```

Figure 13-26. REST Query does not list any Document after deleting the only document

Run a query in the Solr Admin Console after deleting the only indexed document and no document gets returned as indicated by numFound field value of 0 in the JSON document returned as shown in Figure 13-27.

The screenshot shows the Solr Admin interface. On the left is a sidebar with various navigation links: Dashboard, Logging, Core Admin, Java Properties, Thread Dump, gettingstarted (selected), Overview, Analysis, Dataimport, Documents, Files, Ping, Plugins / Stats, Query (selected), Replication, Schema Browser, and Segments info. The main area has a title bar with the URL "http://ec2-54-208-53-110.compute-1.amazonaws.com:8983/solr/gettingstarted/select?q=*&wt=json". Below the title bar is a "Request-Handler (qt)" dropdown set to "/select". The search parameters are as follows:

- q**: `*:*`
- fq**: (empty)
- sort**: (empty)
- start, rows**: 0, 10
- fl**: (empty)
- df**: (empty)
- Raw Query Parameters**: key1=val1&key2=val2
- wt**: json
- indent**: checked
- debugQuery**: unchecked
- dismax**: unchecked

The JSON response is displayed on the right:

```

{
  "responseHeader": {
    "status": 0,
    "QTime": 0,
    "params": {
      "q": "*:*",
      "indent": "true",
      "wt": "json",
      "_": "1445609167976"
    }
  },
  "response": {
    "numFound": 0,
    "start": 0,
    "docs": []
  }
}

```

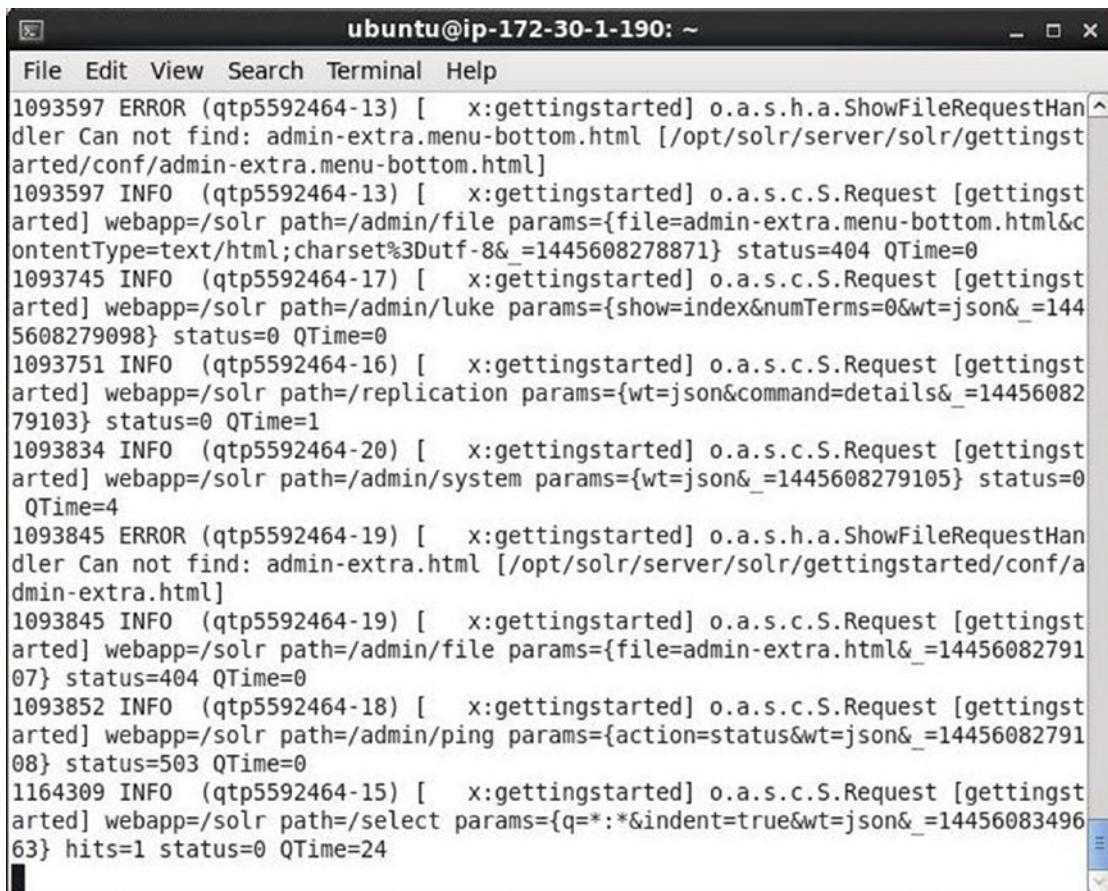
Figure 13-27. Query in Sole Admin Console does not list any document after Deleting the only Document

Listing Logs

The Docker container logs for all commands run on the Apache Solr server may be output using the docker logs command.

```
sudo docker logs -f solr_on_docker
```

The Docker container logs get output as shown in Figure 13-28.



The screenshot shows a terminal window titled "ubuntu@ip-172-30-1-190: ~". The window displays a series of log entries from an Apache Solr container. The logs include various ERROR and INFO messages related to file requests, replication, and system status checks. The log entries are timestamped and show details like file paths, parameters, and response times.

```

1093597 ERROR (qtp5592464-13) [ x:gettingstarted] o.a.s.h.a.ShowFileRequestHandler Can not find: admin-extra.menu-bottom.html [/opt/solr/server/solr/gettingstarted/conf/admin-extra.menu-bottom.html]
1093597 INFO (qtp5592464-13) [ x:gettingstarted] o.a.s.c.S.Request [gettingstarted] webapp=/solr path=/admin/file params={file=admin-extra.menu-bottom.html&contentType=text/html;charset%3Dutf-8&_=1445608278871} status=404 QTime=0
1093745 INFO (qtp5592464-17) [ x:gettingstarted] o.a.s.c.S.Request [gettingstarted] webapp=/solr path=/admin/luke params={show=index&numTerms=0&wt=json&_=1445608279098} status=0 QTime=0
1093751 INFO (qtp5592464-16) [ x:gettingstarted] o.a.s.c.S.Request [gettingstarted] webapp=/solr path=/replication params={wt=json&command=details&_=1445608279103} status=0 QTime=1
1093834 INFO (qtp5592464-20) [ x:gettingstarted] o.a.s.c.S.Request [gettingstarted] webapp=/solr path=/admin/system params={wt=json&_=1445608279105} status=0 QTime=4
1093845 ERROR (qtp5592464-19) [ x:gettingstarted] o.a.s.h.a.ShowFileRequestHandler Can not find: admin-extra.html [/opt/solr/server/solr/gettingstarted/conf/admin-extra.html]
1093845 INFO (qtp5592464-19) [ x:gettingstarted] o.a.s.c.S.Request [gettingstarted] webapp=/solr path=/admin/file params={file=admin-extra.html&_=1445608279107} status=404 QTime=0
1093852 INFO (qtp5592464-18) [ x:gettingstarted] o.a.s.c.S.Request [gettingstarted] webapp=/solr path=/admin/ping params={action=status&wt=json&_=1445608279108} status=503 QTime=0
1164309 INFO (qtp5592464-15) [ x:gettingstarted] o.a.s.c.S.Request [gettingstarted] webapp=/solr path=/select params={q=*&indent=true&wt=json&_=1445608349663} hits=1 status=0 QTime=24

```

Figure 13-28. Listing Docker Container Logs

Stopping Apache Solr Server

The running Docker containers may be listed with the `docker ps` command. The `solr_on_docker` container is listed as running as shown in Figure 13-29.

ubuntu@ip-172-30-1-190:~\$ sudo docker ps				
CONTAINER ID	IMAGE	COMMAND		CREATED
STATUS	PORTS	NAMES		
8061f79d1f16	solr	"/opt/solr/bin/solr -"	solr on docker	36 minutes ago
Up 36 minutes	0.0.0.0:8983->8983/tcp			

Figure 13-29. Listing Running Docker Containers

To stop the `solr_on_docker` container run the `docker stop` command as shown in Figure 13-30.

```
sudo docker stop solr_on_docker
```

```
ubuntu@ip-172-30-1-190:~$ sudo docker stop solr_on_docker
solr_on_docker
ubuntu@ip-172-30-1-190:~$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND       CREATED
STATUS              PORTS              NAMES
```

Figure 13-30. Stopping Docker Container for Apache Solr

Run the `docker ps` command to list the running Docker containers again. The `solr_on_docker` container does not get listed.

The Docker image still gets listed with the `docker images` command as shown in Figure 13-31.

```
ubuntu@ip-172-30-1-190:~$ sudo docker images
REPOSITORY          TAG      IMAGE ID      CREATED
VIRTUAL SIZE
couchbase           latest   ff61ecf3bacb    6 days ago
371.3 MB
solr                latest   a866495d197c    6 days ago
740.2 MB
hello-world         latest   0a6ba66e537a    9 days ago
960 B
ubuntu@ip-172-30-1-190:~$ █
```

Figure 13-31. Listing Docker Image for a stopped Docker Container

If the Docker image is to be removed, first the Docker container `solr_on_docker` has to be removed after being stopped.

```
sudo docker rm solr_on_docker
sudo docker rm solr
```

Summary

In this chapter we used the official Docker image for Apache Solr to run the Apache Solr server in a Docker container. We created a core index and posted a document to the index. Subsequently, we queried the document from the Solr Admin Console and also the REST client tool curl. In the next chapter we shall discuss Apache Spark with Docker.

CHAPTER 14



Using Apache Spark

Apache Spark is a data processing engine for large data sets. Apache Spark is much faster (up to 100 times faster in memory) than Apache Hadoop MapReduce. In cluster mode, Spark applications run as independent processes coordinated by the `SparkContext` object in the driver program, which is the main program. The `SparkContext` may connect to several types of cluster managers to allocate resources to Spark applications. The supported cluster managers include the Standalone cluster manager, Mesos and YARN. Apache Spark is designed to access data from varied data sources including the HDFS, Apache HBase and NoSQL databases such as Apache Cassandra and MongoDB. In this chapter we shall use the same CDH Docker image that we used for several of the Apache Hadoop frameworks including Apache Hive and Apache HBase. We shall run an Apache Spark Master in cluster mode using the YARN cluster manager in a Docker container.

Setting the Environment

Running the Docker Container for CDH

Running Apache Spark Job in yarn-cluster Mode

Running Apache Spark Job in yarn-client Mode

Running the Apache Spark Shell

Setting the Environment

The following software is required for this chapter.

-Docker Engine (version 1.8)

-Docker image for Apache Spark

Connect to an Amazon EC2 instance using the public IP address for the instance. The public IP address may be found from the Amazon EC2 Console as explained in Appendix A.

```
ssh -i "docker.pem" ec2-user@54.208.146.254
```

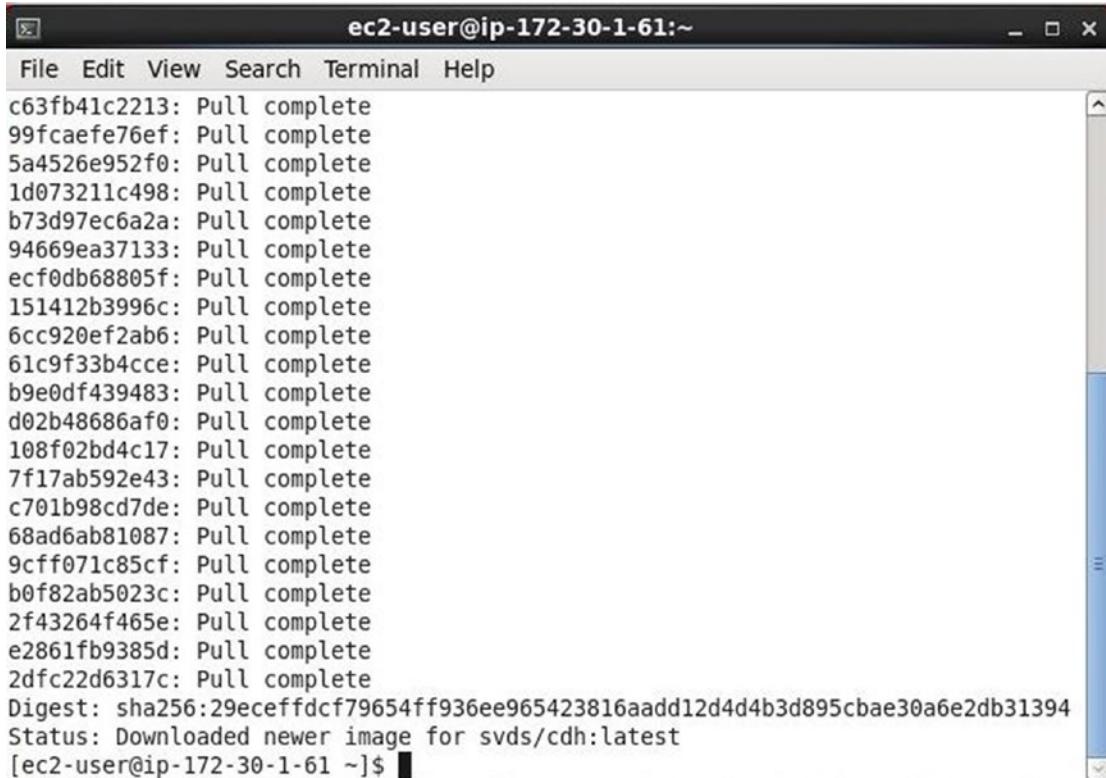
Start the Docker service and verify status as started.

```
sudo service docker start  
sudo service docker status
```

Download the Docker image for CDH, the svds/cdh image if not already downloaded for an earlier chapter.

```
sudo docker pull svds/cdh
```

Docker image svds/cdh gets downloaded as shown in Figure 14-1.



The screenshot shows a terminal window titled "ec2-user@ip-172-30-1-61:~". The window contains a list of Docker image IDs followed by the message "Pull complete". At the bottom, it shows the command used ("Digest: sha256:29ecefdfc79654ff936ee965423816aadd12d4d4b3d895cbae30a6e2db31394 Status: Downloaded newer image for svds/cdh:latest [ec2-user@ip-172-30-1-61 ~]\$").

```
c63fb41c2213: Pull complete
99fcaefe76ef: Pull complete
5a4526e952f0: Pull complete
1d073211c498: Pull complete
b73d97ec6a2a: Pull complete
94669ea37133: Pull complete
ecf0db68805f: Pull complete
151412b3996c: Pull complete
6cc920ef2ab6: Pull complete
61c9f33b4cce: Pull complete
b9e0df439483: Pull complete
d02b48686af0: Pull complete
108f02bd4c17: Pull complete
7f17ab592e43: Pull complete
c701b98cd7de: Pull complete
68ad6ab81087: Pull complete
9cff071c85cf: Pull complete
b0f82ab5023c: Pull complete
2f43264f465e: Pull complete
e2861fb9385d: Pull complete
2dfc22d6317c: Pull complete
Digest: sha256:29ecefdfc79654ff936ee965423816aadd12d4d4b3d895cbae30a6e2db31394
Status: Downloaded newer image for svds/cdh:latest
[ec2-user@ip-172-30-1-61 ~]$
```

Figure 14-1. Downloading svds/cdh Docker Image

Running the Docker Container for CDH

Start a Docker container for the CDH frameworks using the Apache Spark Master port as 8088.

```
sudo docker run -p 8088 -d --name cdh svds/cdh
```

List the running Docker containers.

```
sudo docker ps
```

CDH processes including Apache Spark get started and the container cdh gets listed as running as shown in Figure 14-2.

```
[ec2-user@ip-172-30-1-61 ~]$ sudo docker run -p 8088 -d --name cdh svds/cdh
4b4780802318f23004530e41322d7f01f9719c59c7310a58dc03a15612cb9755
[ec2-user@ip-172-30-1-61 ~]$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED          STATUS              PORTS
NAMES
4b4780802318        svds/cdh           "cdh_startup_script.s"   11 seconds ago
Up 10 seconds       8020/tcp, 8888/tcp, 9090/tcp, 11000/tcp, 11443/tcp, 198
88/tcp, 0.0.0.0:32768->8088/tcp   cdh
```

Figure 14-2. Starting Docker Container for CDH including Apache Spark

Start an interactive terminal for the cdh container.

```
sudo docker exec -it cdh bash
```

The interactive terminal gets started as shown in Figure 14-3.

```
[ec2-user@ip-172-30-1-61 ~]$ sudo docker exec -it cdh bash
root@4fbf325c6ecd:/# █
```

Figure 14-3. Starting the TTY

In YARN mode, a Spark application may be submitted to a cluster in *yarn-cluster* mode or *yarn-client* mode. In the *yarn-cluster* mode, the Apache Spark driver runs inside an Application Master, which is managed by the YARN. In *yarn-client* mode, the Spark driver runs in the client process outside of YARN and the Application Master is used only for requesting resources from YARN. The *--master* parameter is *yarn-cluster* or *yarn-client* based on the mode of application submission. In *yarn-client* mode the Spark driver logs to the console.

We shall run a Spark application using each of the application submission modes. We shall use the example application `org.apache.spark.examples.SparkPi`.

Running Apache Spark Job in *yarn-cluster* Mode

To submit the Spark application `SparkPi` in *yarn-cluster* mode using 1000 iterations, run the following `spark-submit` command with the *--master* parameter as *yarn-cluster*.

```
spark-submit --master yarn-cluster --class org.apache.spark.examples.SparkPi /usr/lib/spark/
examples/lib/spark-examples-1.3.0-cdh5.4.7-hadoop2.6.0-cdh5.4.7.jar 1000
```

The preceding command is run from the interactive terminal as shown in Figure 14-4.

```
[ec2-user@ip-172-30-1-61 ~]$ sudo docker exec -it cdh bash
root@4b4780802318:/# spark-submit --master yarn-cluster --class org.apache.spark.examples.SparkPi /usr/lib/spark/examples/lib/spark-examples-1.3.0-cdh5.4.7-hadoop2.6.0-cdh5.4.7.jar 1000
```

Figure 14-4. Submitting the Spark Application in *yarn-cluster* Mode

The output from the Spark application is shown in Figure 14-5.

```
root@4b4780802318: /
File Edit View Search Terminal Help
15/10/23 19:13:46 INFO yarn.Client: Application report for application_144562752
1793_0001 (state: RUNNING)
15/10/23 19:13:47 INFO yarn.Client: Application report for application_144562752
1793_0001 (state: RUNNING)
15/10/23 19:13:48 INFO yarn.Client: Application report for application_144562752
1793_0001 (state: RUNNING)
15/10/23 19:13:49 INFO yarn.Client: Application report for application_144562752
1793_0001 (state: RUNNING)
15/10/23 19:13:50 INFO yarn.Client: Application report for application_144562752
1793_0001 (state: RUNNING)
15/10/23 19:13:51 INFO yarn.Client: Application report for application_144562752
1793_0001 (state: FINISHED)
15/10/23 19:13:51 INFO yarn.Client:
    client token: N/A
    diagnostics: N/A
    ApplicationMaster host: 4b4780802318
    ApplicationMaster RPC port: 0
    queue: root.root
    start time: 1445627587658
    final status: SUCCEEDED
    tracking URL: http://4b4780802318:8088/proxy/application_1445627521793_
0001/A
    user: root
root@4b4780802318:/#
```

Figure 14-5. Output from Spark Job in *yarn-cluster* Mode

A more detailed output from the `spark-submit` command is listed:

```
spark-submit --master yarn-cluster --class org.apache.spark.examples.SparkPi /usr/lib/spark/examples/lib/spark-examples-1.3.0-cdh5.4.7-hadoop2.6.0-cdh5.4.7.jar 1000
15/10/23 19:12:52 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
15/10/23 19:12:54 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
15/10/23 19:12:56 INFO yarn.Client: Requesting a new application from cluster with 1
NodeManagers
15/10/23 19:12:56 INFO yarn.Client: Verifying our application has not requested more than
the maximum memory capability of the cluster (8192 MB per container)
```

```
15/10/23 19:12:56 INFO yarn.Client: Will allocate AM container, with 896 MB memory including  
384 MB overhead  
15/10/23 19:12:56 INFO yarn.Client: Setting up container launch context for our AM  
15/10/23 19:12:56 INFO yarn.Client: Preparing resources for our AM container  
15/10/23 19:12:59 WARN shortcircuit.DomainSocketFactory: The short-circuit local reads  
feature cannot be used because libhadoop cannot be loaded.  
15/10/23 19:12:59 INFO yarn.Client: Uploading resource file:/usr/lib/spark/lib/spark-  
assembly-1.3.0-cdh5.4.7-hadoop2.6.0-cdh5.4.7.jar -> hdfs://localhost:8020/user/root/.  
sparkStaging/application_1445627521793_0001/spark-assembly-1.3.0-cdh5.4.7-hadoop2.6.0-  
cdh5.4.7.jar  
15/10/23 19:13:05 INFO yarn.Client: Uploading resource file:/usr/lib/spark/examples/lib/  
spark-examples-1.3.0-cdh5.4.7-hadoop2.6.0-cdh5.4.7.jar -> hdfs://localhost:8020/user/root/.  
sparkStaging/application_1445627521793_0001/spark-examples-1.3.0-cdh5.4.7-hadoop2.6.0-  
cdh5.4.7.jar  
15/10/23 19:13:06 INFO yarn.Client: Setting up the launch environment for our AM container  
15/10/23 19:13:07 INFO spark.SecurityManager: Changing view acls to: root  
15/10/23 19:13:07 INFO spark.SecurityManager: Changing modify acls to: root  
15/10/23 19:13:07 INFO spark.SecurityManager: SecurityManager: authentication disabled;  
ui acls disabled; users with view permissions: Set(root); users with modify permissions:  
Set(root)  
15/10/23 19:13:07 INFO yarn.Client: Submitting application 1 to ResourceManager  
15/10/23 19:13:08 INFO impl.YarnClientImpl: Submitted application  
application_1445627521793_0001  
15/10/23 19:13:09 INFO yarn.Client: Application report for application_1445627521793_0001  
(state: ACCEPTED)  
15/10/23 19:13:09 INFO yarn.Client:  
    client token: N/A  
    diagnostics: N/A  
    ApplicationMaster host: N/A  
    ApplicationMaster RPC port: -1  
    queue: root.root  
    start time: 1445627587658  
    final status: UNDEFINED  
    tracking URL: http://4b4780802318:8088/proxy/application\_1445627521793\_0001/  
    user: root  
15/10/23 19:13:10 INFO yarn.Client: Application report for application_1445627521793_0001  
(state: ACCEPTED)  
15/10/23 19:13:11 INFO yarn.Client: Application report for application_1445627521793_0001  
(state: ACCEPTED)  
15/10/23 19:13:24 INFO yarn.Client: Application report for application_1445627521793_0001  
(state: RUNNING)  
15/10/23 19:13:24 INFO yarn.Client:  
    client token: N/A  
    diagnostics: N/A  
    ApplicationMaster host: 4b4780802318  
    ApplicationMaster RPC port: 0  
    queue: root.root  
    start time: 1445627587658  
    final status: UNDEFINED  
    tracking URL: http://4b4780802318:8088/proxy/application\_1445627521793\_0001/  
    user: root
```

```
15/10/23 19:13:25 INFO yarn.Client: Application report for application_1445627521793_0001  
(state: RUNNING)  
15/10/23 19:13:26 INFO yarn.Client: Application report for  
15/10/23 19:13:51 INFO yarn.Client: Application report for application_1445627521793_0001  
(state: FINISHED)  
15/10/23 19:13:51 INFO yarn.Client:  
    client token: N/A  
    diagnostics: N/A  
    ApplicationMaster host: 4b4780802318  
    ApplicationMaster RPC port: 0  
    queue: root.root  
    start time: 1445627587658  
    final status: SUCCEEDED  
    tracking URL: http://4b4780802318:8088/proxy/application\_1445627521793\_0001/A  
    user: root
```

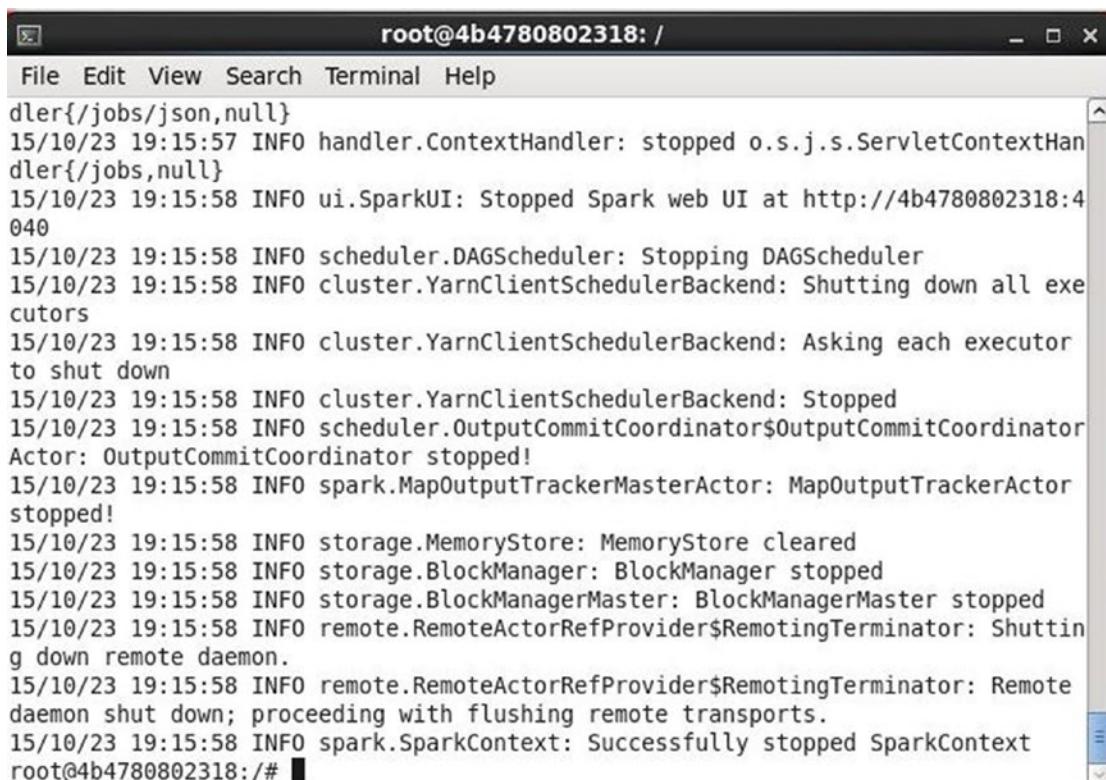
In yarn-cluster mode, the Spark application result is not output to the console and has to be accessed from the YARN container logs accessible from the ResourceManager using the tracking URL http://4b4780802318:8088/proxy/application_1445627521793_0001/A in a browser if the final status is SUCCEEDED.

Running Apache Spark Job in yarn-client Mode

To submit the Spark application SparkPi in yarn-client mode using 1000 iterations, run the following spark-submit command with the --master parameter as yarn-client.

```
spark-submit  
  --master yarn-client  
  --class org.apache.spark.examples.SparkPi  
  /usr/lib/spark/examples/lib/spark-examples-1.3.0-cdh5.4.7-hadoop2.6.0-cdh5.4.7.jar  
  1000
```

The output from the spark-submit command is shown in Figure 14-6.



The screenshot shows a terminal window titled "root@4b4780802318:/". The window contains a list of log messages from an Apache Spark application. The messages are timestamped and show the progression of the application's execution, including the stopping of various components like the DAGScheduler, YarnClientSchedulerBackend, and MapOutputTrackerMasterActor, and the successful shutdown of the SparkContext.

```

root@4b4780802318:/#
File Edit View Search Terminal Help
dler{/jobs/json,null}
15/10/23 19:15:57 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHan
dler{/jobs,null}
15/10/23 19:15:58 INFO ui.SparkUI: Stopped Spark web UI at http://4b4780802318:4
040
15/10/23 19:15:58 INFO scheduler.DAGScheduler: Stopping DAGScheduler
15/10/23 19:15:58 INFO cluster.YarnClientSchedulerBackend: Shutting down all exe
cutors
15/10/23 19:15:58 INFO cluster.YarnClientSchedulerBackend: Asking each executor
to shut down
15/10/23 19:15:58 INFO cluster.YarnClientSchedulerBackend: Stopped
15/10/23 19:15:58 INFO scheduler.OutputCommitCoordinator$OutputCommitCoordinator
Actor: OutputCommitCoordinator stopped!
15/10/23 19:15:58 INFO spark.MapOutputTrackerMasterActor: MapOutputTrackerActor
stopped!
15/10/23 19:15:58 INFO storage.MemoryStore: MemoryStore cleared
15/10/23 19:15:58 INFO storage.BlockManager: BlockManager stopped
15/10/23 19:15:58 INFO storage.BlockManagerMaster: BlockManagerMaster stopped
15/10/23 19:15:58 INFO remote.RemoteActorRefProvider$RemotingTerminator: Shuttin
g down remote daemon.
15/10/23 19:15:58 INFO remote.RemoteActorRefProvider$RemotingTerminator: Remote
daemon shut down; proceeding with flushing remote transports.
15/10/23 19:15:58 INFO spark.SparkContext: Successfully stopped SparkContext
root@4b4780802318:/# █

```

Figure 14-6. Submitting Spark Application in *yarn-client* Mode

A more detailed output from the Apache Spark application is as follows and includes the value of Pi calculated approximately.

```

spark-submit --master yarn-client --class org.apache.spark.examples.SparkPi
/usr/lib/spark/examples/lib/spark-examples-1.3.0-cdh5.4.7-hadoop2.6.0-cdh5.4.7.jar 1000
15/10/23 19:15:19 INFO spark.SparkContext: Running Spark version 1.3.0
15/10/23 19:15:43 INFO cluster.YarnScheduler: Adding task set 0.0 with 1000 tasks
15/10/23 19:15:43 INFO scheduler.TaskSetManager: Starting task 0.0 in stage 0.0
(TID 0, 4b4780802318, PROCESS_LOCAL, 1353 bytes)
15/10/23 19:15:43 INFO scheduler.TaskSetManager: Starting task 1.0 in stage 0.0
(TID 1, 4b4780802318, PROCESS_LOCAL, 1353 bytes)
15/10/23 19:15:57 INFO scheduler.TaskSetManager: Finished task 999.0 in stage 0.0 (TID 999)
in 22 ms on 4b4780802318 (999/1000)
15/10/23 19:15:57 INFO scheduler.TaskSetManager: Finished task 998.0 in stage 0.0 (TID 998)
in 28 ms on 4b4780802318 (1000/1000)
15/10/23 19:15:57 INFO cluster.YarnScheduler: Removed TaskSet 0.0, whose tasks have all
completed, from pool
15/10/23 19:15:57 INFO scheduler.DAGScheduler: Stage 0 (reduce at SparkPi.scala:35)
finished in 14.758 s
15/10/23 19:15:57 INFO scheduler.DAGScheduler: Job 0 finished: reduce at SparkPi.scala:35,
took 15.221643 s

```

Pi is roughly 3.14152984

Running the Apache Spark Shell

The Apache Spark shell is started in yarn-client mode as follows.

```
spark-shell --master yarn-client
```

The `scala>` command prompt gets displayed as shown in Figure 14-7. A Spark context gets created and becomes available as 'sc'. A SQL context also becomes available as 'sqlContext'.

```
root@4b4780802318: /
File Edit View Search Terminal Help
er
15/10/23 19:17:33 INFO storage.BlockManagerMasterActor: Registering block manager 4b4780802318:44758 with 265.4 MB RAM, BlockManagerId(<driver>, 4b4780802318, 44758)
15/10/23 19:17:33 INFO storage.BlockManagerMaster: Registered BlockManager
15/10/23 19:17:44 INFO cluster.YarnClientSchedulerBackend: Registered executor: Actor[akka.tcp://sparkExecutor@4b4780802318:39462/user/Executor#145546620] with ID 1
15/10/23 19:17:44 INFO storage.BlockManagerMasterActor: Registering block manager 4b4780802318:47082 with 530.3 MB RAM, BlockManagerId(1, 4b4780802318, 47082)
15/10/23 19:17:44 INFO cluster.YarnClientSchedulerBackend: Registered executor: Actor[akka.tcp://sparkExecutor@4b4780802318:35432/user/Executor#-1808868577] with ID 2
15/10/23 19:17:45 INFO cluster.YarnClientSchedulerBackend: SchedulerBackend is ready for scheduling beginning after reached minRegisteredResourcesRatio: 0.8
15/10/23 19:17:45 INFO repl.SparkILoop: Created spark context..
Spark context available as sc.
15/10/23 19:17:45 INFO repl.SparkILoop: Created sql context (with Hive support).
.
SQL context available as sqlContext.
15/10/23 19:17:45 INFO storage.BlockManagerMasterActor: Registering block manager 4b4780802318:48279 with 530.3 MB RAM, BlockManagerId(2, 4b4780802318, 48279)

scala> 
```


Figure 14-7. The `scala>` Command Prompt

A more detailed output from the `spark-shell` command is as follows.

```
root@4b4780802318:/# spark-shell --master yarn-client
15/10/23 19:17:16 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
15/10/23 19:17:16 INFO spark.SecurityManager: Changing view acls to: root
15/10/23 19:17:16 INFO spark.SecurityManager: Changing modify acls to: root
15/10/23 19:17:16 INFO spark.SecurityManager: SecurityManager: authentication disabled;
ui acls disabled; users with view permissions: Set(root); users with modify permissions:
Set(root)
15/10/23 19:17:16 INFO spark.HttpServer: Starting HTTP Server
15/10/23 19:17:16 INFO server.Server: jetty-8.y.z-SNAPSHOT
```

```
15/10/23 19:17:16 INFO server.AbstractConnector: Started SocketConnector@0.0.0.0:56899
15/10/23 19:17:16 INFO util.Utils: Successfully started service 'HTTP class server' on port
56899.
```

Welcome to



version 1.3.0

Using Scala version 2.10.4 (OpenJDK 64-Bit Server VM, Java 1.7.0_79)

Type in expressions to have them evaluated.

Type :help for more information.

```
15/10/23 19:17:22 INFO spark.SparkContext: Running Spark version 1.3.0
```

```
15/10/23 19:17:45 INFO repl.SparkILoop: Created spark context..
```

Spark context available as `sc`.

```
15/10/23 19:17:45 INFO repl.SparkILoop: Created sql context (with Hive support)..
```

SQL context available as `sqlContext`.

```
15/10/23 19:17:45 INFO storage.BlockManagerMasterActor: Registering block manager
```

```
4b4780802318:48279 with 530.3 MB RAM, BlockManagerId(2, 4b4780802318, 48279)
```

`scala>`

Run the following Scala script consisting of a `HelloWorld` module in the Spark shell for a Hello World program.

```
object HelloWorld {
    def main(args: Array[String]) {
        println("Hello, world!")
    }
}
HelloWorld.main(null)
```

The output from the Scala script is shown in Figure 14-8.

```
scala> object HelloWorld {
|   def main(args: Array[String]) {
|     println("Hello, world!")
|   }
| }
defined module HelloWorld

scala> HelloWorld.main(null)
Hello, world!

scala> ■
```

Figure 14-8. Output from Scala Script

Summary

In this chapter, we ran Apache Spark applications on a YARN cluster in a Docker container using the spark-submit command. We submitted the example application in yarn-cluster and yarn-client modes. We also ran a HelloWorld Scala script in a Spark shell.

This chapter concludes the book on Docker. In addition to running some of the commonly used software on Docker, we discussed the main Docker administrative tasks such as installing Docker, downloading a Docker image, creating and running a Docker container, starting an interactive shell, running commands in an interactive shell, listing Docker containers, listing Docker container logs, stopping a Docker container, and removing a Docker container and a Docker image. Only a few of the software applications could be discussed in the scope of this book. Several more Docker images are available on the Docker hub at <https://hub.docker.com/>.

APPENDIX A



Using the Amazon EC2

Amazon Web Services (AWS) provides various services and Amazon Elastic Compute Cloud (Amazon EC2) is one of the services. Amazon EC2 may be used to create a virtual host server. Amazon EC2 provides a wide selection of instance AMIs (Amazon Machine Images) to choose from when creating a virtual server. In this Appendix we shall discuss creating and configuring Amazon EC2 instance/s for installing Docker and Docker images. Amazon EC2 instance is not a requirement to run Docker software and an alternative platform, local or remote, may be used instead.

Creating an Amazon EC2 Instance

Create a Key Pair

Starting an Amazon EC2 Instance

Connecting to an Amazon EC2 Instance

Finding the Public IP Address

Finding the Public DNS

Adding the default Security Group

Stopping an Amazon EC2 Instance

Changing the Instance Type

Creating an Amazon EC2 Instance

We have used Amazon EC2 instances based on Linux for deploying Docker and Docker images. Amazon EC2 is not a requirement and an alternative such as a local Linux installation may be used instead. The Linux platform is required to support 64 bit software. We have made use of two different 64 bit (required) AMIs:

1. Ubuntu Server 14.04 LTS (HVM), SSD Volume Type - ami-d05e75b8 64 bit
2. Red Hat Enterprise Linux version 7.1 (HVM), EBS General Purpose (SSD) Volume Type (ami-12663b7a) 64 bit

To create an Amazon EC2 Instance, an Amazon Web Services Account is required, which may be created at https://aws.amazon.com/getting-started/?nc2=h_12_cc. To create an Amazon EC2 instance, navigate to <https://aws.amazon.com/ec2/> and click on Sign In to the Console. Select EC2 from the listed Amazon Web Services. Click on INSTANCES ► Instances to list the Amazon EC2 instances already created in the account. Click on Launch Instance to create a new Amazon EC2 instance as shown in Figure A-1.

APPENDIX A ■ USING THE AMAZON EC2




Figure A-1. Launching an Amazon EC2 Instance

Select an AMI to create a virtual server from. Some of the AMIs are eligible for the Free tier. For example, select the Ubuntu AMI as shown in Figure A-2.




Figure A-2. Selecting an AMI

In Choose an Instance Type different types are available differing by features such as supported capacity and virtual CPUs (vCPUs). Select one of the Instance Types, for example the General Purpose ➤ t2.micro and click on Review and Launch as shown in Figure A-3.

Family	Type	vCPUs	Memory (GiB)	Instance Storage (GB)	EBS Optimized Available	Network Performance
General purpose	t2.micro <small>Free tier eligible</small>	1	1	EBS only	-	Low to Moderate
General purpose	t2.small	1	2	EBS only	-	Low to Moderate
General purpose	t2.medium	2	4	EBS only	-	Low to Moderate
General purpose	t2.large	2	8	EBS only	-	Low to Moderate

Figure A-3. Review and Launch

Click on Launch in Review Instance Launch as shown in Figure A-4.

Step 7: Review Instance Launch

Please review your instance launch details. You can go back to edit changes for each section. Click **Launch** to assign a key pair to your instance and complete the launch process.

AMI Details

Ubuntu Server 14.04 LTS (HVM), SSD Volume Type - ami-d05e75b8
Free tier eligible
 Ubuntu Server 14.04 LTS (HVM), EBS General Purpose (SSD) Volume Type. Support available from Canonical (<http://www.ubuntu.com/cloud/services>).
 Root Device Type: ebs Virtualization type: hvm

Instance Type

Instance Type	ECUs	vCPUs	Memory (GiB)	Instance Storage (GB)	EBS-Optimized Available	Network Performance
t2.micro	Variable	1	1	EBS only	-	Low to Moderate

Security Groups

Security group name: launch-wizard-8

Figure A-4. Launch

APPENDIX A ■ USING THE AMAZON EC2

A dialog gets displayed to create or select an existing key pair. A key pair is required for authorization. To create a new key pair, select the “Create a new key pair” option as shown in Figure A-5.

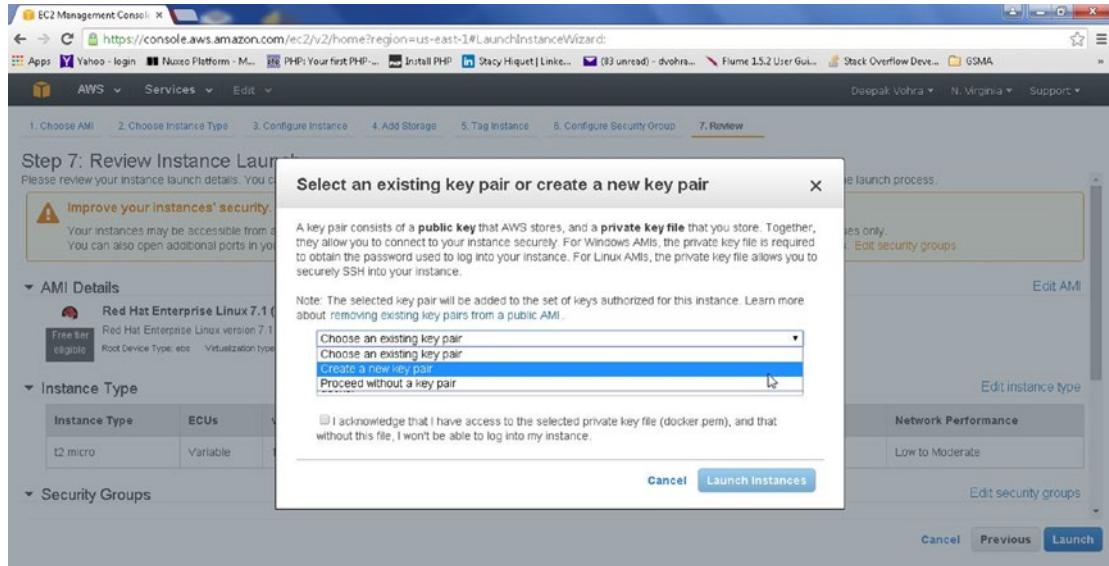


Figure A-5. Selecting “Create a new key pair”

Specify a Key pair name and click on Download Key Pair as shown in Figure A-6. The Key pair gets created and downloaded. The key pair selected for an Amazon EC2 instance when creating the instance is required when connecting to the instance, as is discussed later in this Appendix.

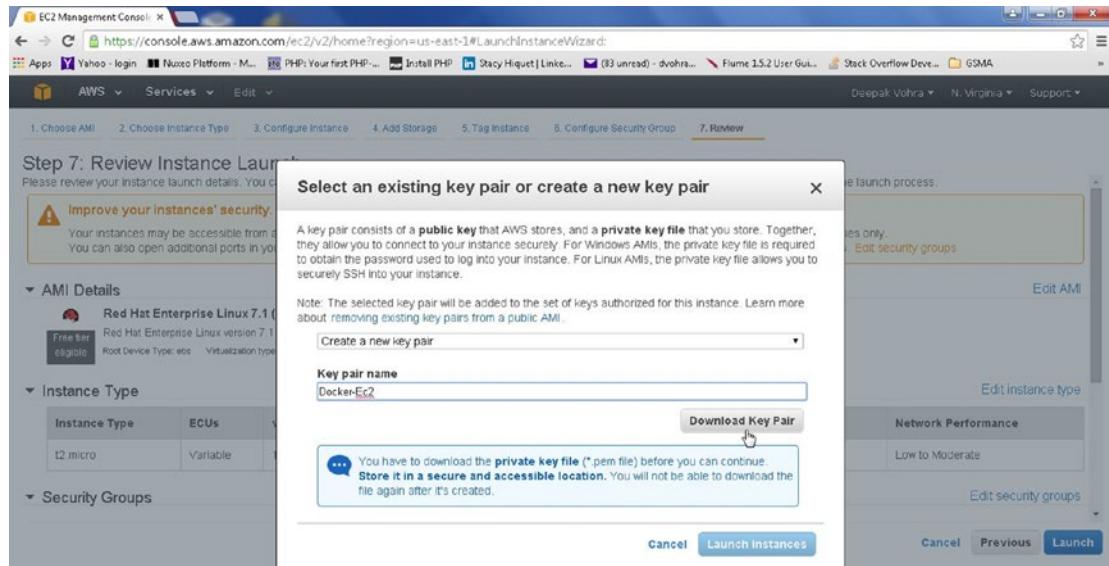


Figure A-6. Download Key Pair

Alternatively, select the option “Choose an existing key pair” and click on Launch Instances as shown in Figure A-7.

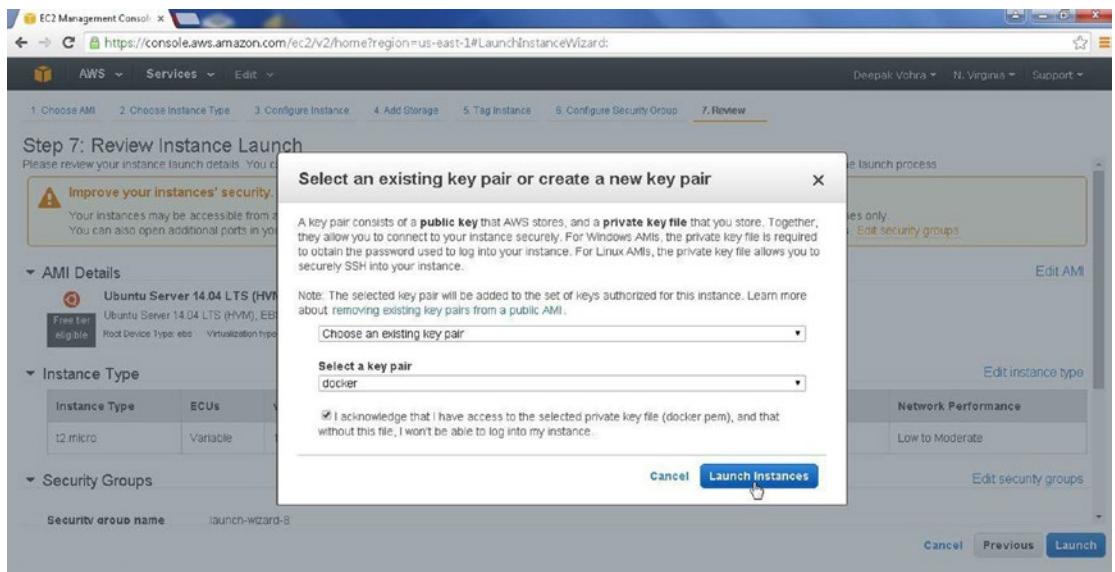


Figure A-7. Choose an existing Key Pair

The Launch Status gets displayed. Click on the instance id to display the instance as shown in Figure A-8.

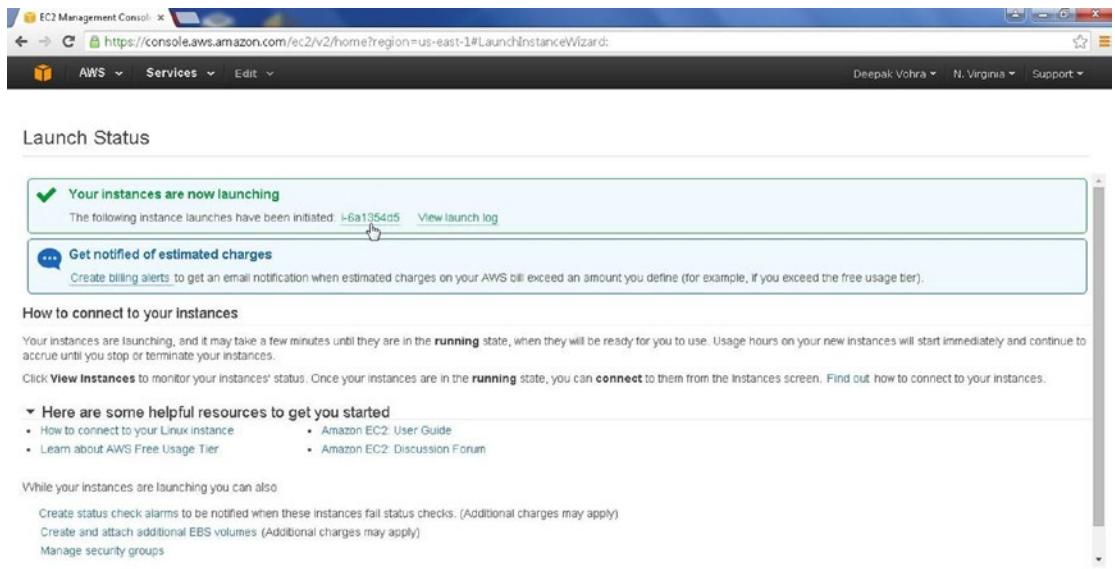


Figure A-8. Launch Status

APPENDIX A ■ USING THE AMAZON EC2

The instance gets listed and is initially in the “pending” state as shown in Figure A-9.

The screenshot shows the AWS Management Console with the EC2 service selected. The left sidebar shows navigation links for EC2 Dashboard, Events, Tags, Reports, Limits, Instances, Images, Elastic Block Store, Network & Security, and more. The main pane displays a table of instances. One instance, with the ID i-6a1354d5, is highlighted. Its details are shown in a modal dialog. The instance's state is listed as "pending".

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS	Public IP
	i-6a1354d5	t2.micro	us-east-1b	pending	Initializing	None	ec2-52-23-208-194.com...	52.23.208.194

Instance: i-6a1354d5 Public DNS: ec2-52-23-208-194.compute-1.amazonaws.com

Description	Status Checks	Monitoring	Tags
Instance ID: i-6a1354d5			
Instance state: pending			
Instance type: t2.micro			
Public DNS: ec2-52-23-208-194.compute-1.amazonaws.com			
Public IP: 52.23.208.194			
Elastic IP: -			

Figure A-9. Amazon EC2 Instance in Pending State

When an instance has launched completely, the Instance State becomes “running” as shown in Figure A-10.

The screenshot shows the AWS Management Console with the EC2 service selected. The left sidebar shows navigation links for EC2 Dashboard, Events, Tags, Reports, Limits, Instances, Images, Elastic Block Store, Network & Security, and more. The main pane displays a table of instances. One instance, with the ID i-6a1354d5, is highlighted. Its details are shown in a modal dialog. The instance's state is listed as "running".

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS	Public IP
	i-036b01bc	m3.medium	us-east-1b	stopped		None	ec2-52-23-208-194.com...	52.23.208.194
Docker	i-e25afa40	t2.micro	us-east-1b	stopped		None	ec2-52-23-208-194.com...	52.23.208.194
Docker	i-816ee83e	t2.micro	us-east-1b	stopped		None	ec2-52-23-208-194.com...	52.23.208.194
Docker	i-836fe93c	m3.medium	us-east-1b	stopped		None	ec2-52-23-208-194.com...	52.23.208.194
Docker	i-f05265d	m3.medium	us-east-1b	stopped		None	ec2-52-23-208-194.com...	52.23.208.194
Docker	i-2ed3718c	m3.medium	us-east-1b	stopped		None	ec2-52-23-208-194.com...	52.23.208.194
	i-24c364d	t2.micro	us-east-1b	running	2/2 checks...	None	ec2-54-210-177-249.co...	54.210.177.249
Docker	i-6a1354d5	t2.micro	us-east-1b	pending	Initializing	None	ec2-52-23-208-194.com...	52.23.208.194
Ubuntu test	i-2247029d	t2.micro	us-east-1b	stopped		None	ec2-52-23-208-194.com...	52.23.208.194

Instance: i-6a1354d5 Public DNS: ec2-52-23-208-194.compute-1.amazonaws.com

Figure A-10. Running Instance

Creating a Key Pair

As mentioned previously, a key pair is required to connect to a Amazon EC2 instance. A key pair may be created while creating an instance or separately. To create a key pair separately select Network & Security ► Key Pairs as shown in Figure A-11.

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS	Public IP
Docker	i-036501b	m3.medium	us-east-1b	stopped		None		
Docker	i-e25afa40	t2.micro	us-east-1b	stopped		None		
Docker	i-816e683e	t2.micro	us-east-1b	stopped		None		
Docker	i-836fe93c	m3.medium	us-east-1b	stopped		None		
Docker	i-055255d	m3.medium	us-east-1b	stopped		None		
Docker	i-2ed371bc	m3.medium	us-east-1b	stopped		None		
	i-024c3f4d	t2.micro	us-east-1b	running	2/2 checks...	None	ec2-54-210-177-248.co...	54.210.177.2
	i-6a1354b5	m3.medium	us-east-1b	running	2/2 checks...	None	ec2-54-152-57-191.com...	54.152.57.19
Ubuntu test	i-2247029d	t2.micro	us-east-1b	stopped		None		

Figure A-11. Network & Security ► Key Pairs

The key pairs already created get listed. A key pair may be deleted by selecting the key pair and clicking on Delete. Click on Yes in the dialog as shown in Figure A-12.

Figure A-12. Delete Key Pair

To create a new key pair, click on Create Key Pair as shown in Figure A-13.

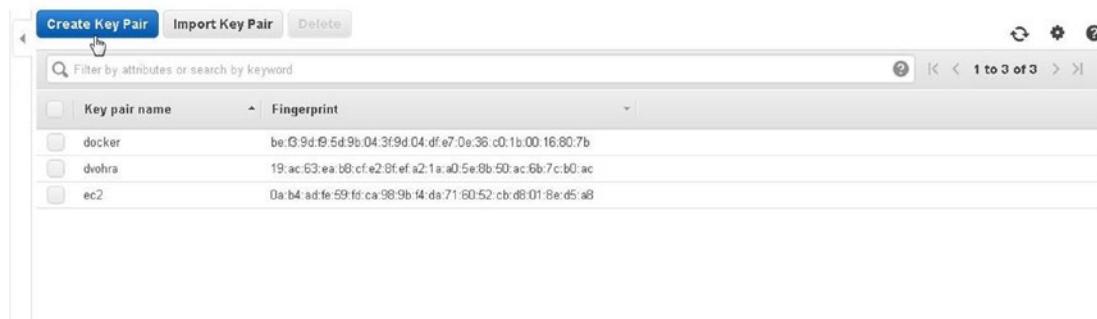


Figure A-13. Create Key Pair

Specify a Key pair name and click on Create button as shown in Figure A-14.



Figure A-14. Create Button

A new key pair gets created as shown in Figure A-15.



Figure A-15. New Key Pair

Starting an Amazon EC2 Instance

When a new Amazon EC2 instance is created and Launch is selected, the instance gets started. A stopped instance may be started by selecting the checkbox adjacent to the instance and selecting Actions ▶ Instance State ▶ Start as shown in Figure A-16.




Figure A-16. Actions ▶ Instance State ▶ Start

In Start Instances dialog click on Yes, Start as shown in Figure A-17.




Figure A-17. Starting an instance

Connecting to an Amazon EC2 Instance

An instance that has been started may be connected to from a local machine such as a local Linux instance without as much RAM and a different Linux distribution than the instance being connected to. The ssh command to use to connect to a running instance may be obtained by clicking on Connect as shown in Figure A-18.

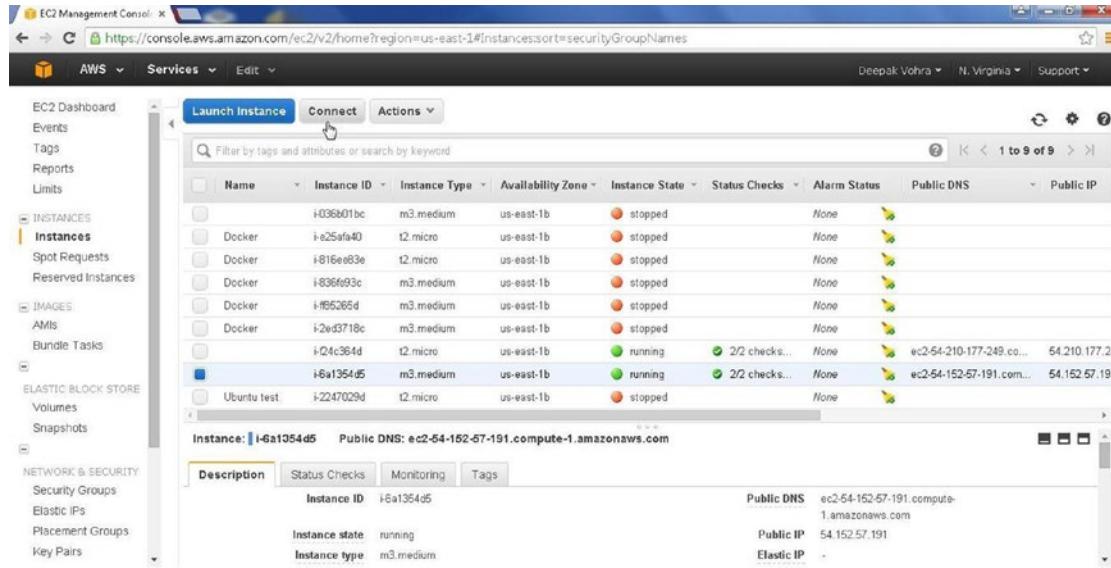


Figure A-18. Connect

In the Connect To Your Instance dialog, the ssh command is displayed. The “docker.pem” is the key pair used to create an instance and also downloaded to the local instance from which the Amazon EC2 instance is to be connected. The username for an Ubuntu instance is “ubuntu” as shown in Figure A-19 and for a Red Hat instance is “ec2-user”.

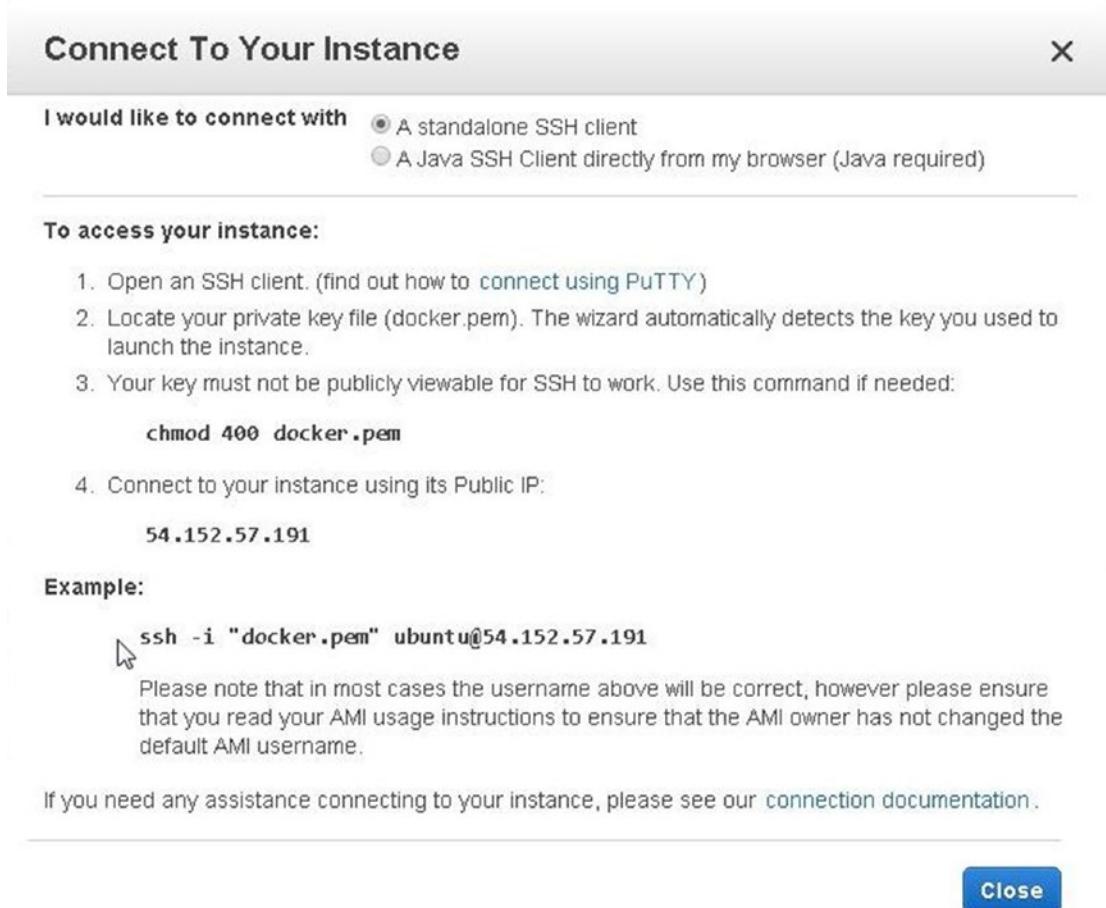


Figure A-19. Connect To Your Instance dialog

The IP Address shown in the ssh command is the Public IP Address of the Amazon EC2 instance.

Finding the Public IP Address

The Public IP Address may also be obtained from the EC2 Console as shown in Figure A-20.

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS	Public IP	Key Name
Decker	i-016ee83e	t2.micro	us-east-1b	stopped		None			docker
Decker	i-036fe83c	m3.medium	us-east-1b	stopped		None			docker
Decker	i-0f5265d	m3.medium	us-east-1b	stopped		None			docker
Decker	i-0ed3718c	m3.medium	us-east-1b	stopped		None			docker
Decker	i-024c364d	t2.micro	us-east-1b	running	2/2 checks...	None	ec2-52-91-72-225.compute-1.amazonaws.com	52.91.72.225	docker
Ubuntu test	i-0247029d	t2.micro	us-east-1b	stopped		None			docker

Instance: i-024c364d (Decker) Public DNS: ec2-52-91-72-225.compute-1.amazonaws.com

Description Status Checks Monitoring Tags

Instance ID: i-024c364d
Instance state: running
Instance type: t2.micro

Public DNS: ec2-52-91-72-225.compute-1.amazonaws.com
Public IP: 52.91.72.225
Elastic IP:

Figure A-20. Public IP Address

Finding the Public DNS

To connect to an Amazon EC2 instance process such as the HelloWorld application in Chapter 1 from a remote browser, the Public DNS is required. The Public DNS may also be obtained from the EC2 Management Console as shown in Figure A-21.

The screenshot shows the AWS EC2 Management Console. At the top, there are buttons for 'Launch Instance', 'Connect', and 'Actions'. Below this is a search bar with placeholder text 'Filter by tags and attributes or search by keyword'. The main area is a table listing instances. One instance, 'i-024c364d', is selected and highlighted with a blue border. The table columns include Name, Instance ID, Instance Type, Availability Zone, Instance State, Status Checks, Alarm Status, Public DNS, and Public IP. The instance details for 'i-024c364d' are shown in a modal window at the bottom. The modal has tabs for 'Description', 'Status Checks', 'Monitoring', and 'Tags'. The 'Description' tab is selected, showing the Instance ID (i-024c364d), Instance state (running), and Instance type (t2.micro). To the right, the Public DNS (ec2-54-210-177-249.compute-1.amazonaws.com) and Public IP (54.210.177.249) are listed.

Figure A-21. Public DNS

The Public DNS may not get displayed initially. To display the Public DNS, select Services ► VPC in the EC2 Management Console as shown in Figure A-22. VPC is a virtual private cloud assigned to a user.

The screenshot shows the AWS Management Console. The top navigation bar includes links for 'AWS', 'Services', and 'Edit'. A dropdown menu is open under 'Services', showing options like 'History', 'EC2', 'VPC', 'Console Home', 'Billing', and 'CloudWatch'. The 'VPC' option is highlighted with a blue background and white text. To the right of the menu, there are four service cards: 'VPC' (described as Amazon Virtual Private Cloud), 'Direct Connect' (described as AWS Direct Connect), and 'Route 53' (described as Amazon Route 53). The URL in the browser is https://console.aws.amazon.com/ec2/v2/home?region=us-east-1#instances:sort=securityGroupNames.

Figure A-22. Services ► VPC

In the VPC Dashboard, select Your VPCs as shown in Figure A-23.

The screenshot shows the AWS VPC Management Console interface. At the top, there's a navigation bar with links for AWS, Services, and Edit. Below it is the VPC Dashboard. On the left, a sidebar lists various VPC components: Virtual Private Cloud, Your VPCs (which is selected and highlighted in yellow), Subnets, Route Tables, Internet Gateways, DHCP Options Sets, Elastic IPs, Endpoints, Peering Connections, Security, Network ACLs, Security Groups, and VPN Connections. In the main content area, there's a "Resources" section with two buttons: "Start VPC Wizard" and "Launch EC2 Instances". It includes a note about launching instances in the US East (N. Virginia) region and a summary of resources in the US East (N. Virginia) region:

1 VPC	1 Internet Gateway
4 Subnets	1 Route Table
1 Network ACL	1 Elastic IP
0 VPC Peering Connections	0 Endpoints
11 Security Groups	1 Running Instance
0 VPN Connections	0 Virtual Private Gateways
0 Customer Gateways	

Below this is a "VPN Connections" section with a note about using Amazon VPC to connect isolated resources within the AWS cloud to a datacenter. It shows a table with four columns: VPN Connections, Customer Gateways, VPC ID, and Status. The table is currently empty, displaying the message "You do not have any VPNs." There is also a "Create VPN Connection" button.

Figure A-23. Your VPCs

Select the VPC listed as shown in Figure A-24.

The screenshot shows the AWS VPC Management Console. On the left, there's a sidebar with options like Virtual Private Cloud, Your VPCs, Subnets, Route Tables, Internet Gateways, DHCP Options Sets, Elastic IPs, Endpoints, Peering Connections, Security, Network ACLs, Security Groups, and VPN Connections. The main area shows a table of VPCs. One VPC, 'vpc-3b12ef5f', is selected and highlighted with a blue border. A context menu is open over this VPC entry, with the 'Edit DNS Hostnames' option highlighted in red. Below the table, a detailed view of the selected VPC is shown, including its Summary, Flow Logs, and Tags. The summary section lists the VPC ID, State, VPC CIDR, DHCP options set, Route table, Network ACL, Tenancy, and DNS hostnames.

Name	VPC ID	State	VPC CIDR	DHCP options set	Route table	Network ACL	Tenancy	Def
vpc-3b12ef5f	available	172.30.0.0/16	dopt-09b9476c	rtb-cb399af	acl-a92ad2cd	Default	No	

vpc-3b12ef5f (172.30.0.0/16)

Summary

VPC ID: vpc-3b12ef5f
State: available
VPC CIDR: 172.30.0.0/16
DHCP options set: dopt-09b9476c
Route table: rtb-cb399af

Network ACL: acl-a92ad2cd
Tenancy: Default
DNS resolution: yes
DNS hostnames: yes

Figure A-24. Selecting the VPC

From Actions, select Edit DNS Hostnames as shown in Figure A-25.

This screenshot is similar to Figure A-24, showing the VPC Management Console. The context menu over the selected VPC 'vpc-3b12ef5f' now has the 'Edit DNS Hostnames' option highlighted in red. The rest of the interface, including the table of VPCs and the detailed view below, remains the same as in Figure A-24.

Name	VPC ID	State	VPC CIDR	DHCP options set	Route table	Network ACL	Tenancy	Def
vpc-3b12ef5f	available	172.30.0.0/16	dopt-09b9476c	rtb-cb399af	acl-a92ad2cd	Default	No	

vpc-3b12ef5f (172.30.0.0/16)

Summary

VPC ID: vpc-3b12ef5f
State: available
VPC CIDR: 172.30.0.0/16
DHCP options set: dopt-09b9476c
Route table: rtb-cb399af

Network ACL: acl-a92ad2cd
Tenancy: Default
DNS resolution: yes
DNS hostnames: yes

Figure A-25. Edit DNS Hostnames

In the Edit DNS Hostnames dialog, select Yes for the DNS Hostnames, and click on Save as shown in Figure A-26.




Figure A-26. Edit DNS Hostnames Dialog

Adding the default Security Group

To be able to connect from a remote browser, the Inbound and Outbound rules are required to be set to allow all traffic using any protocol on all ports in the range 0-65535 from any source. The “default” security group is configured by default to allow all traffic. We need to assign the “default” security group to the Amazon EC2 instance running Docker. Select the instance and select Actions ➤ Networking ➤ Change Security Groups as shown in Figure A-27.

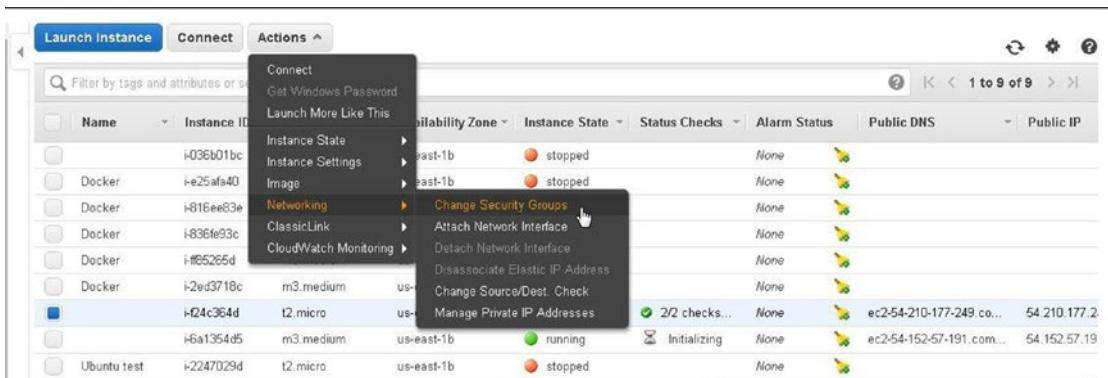


Figure A-27. Actions ➤ Networking ➤ Change Security Groups

In the Change Security Groups panel, the “default” group might not be selected as shown in Figure A-28.

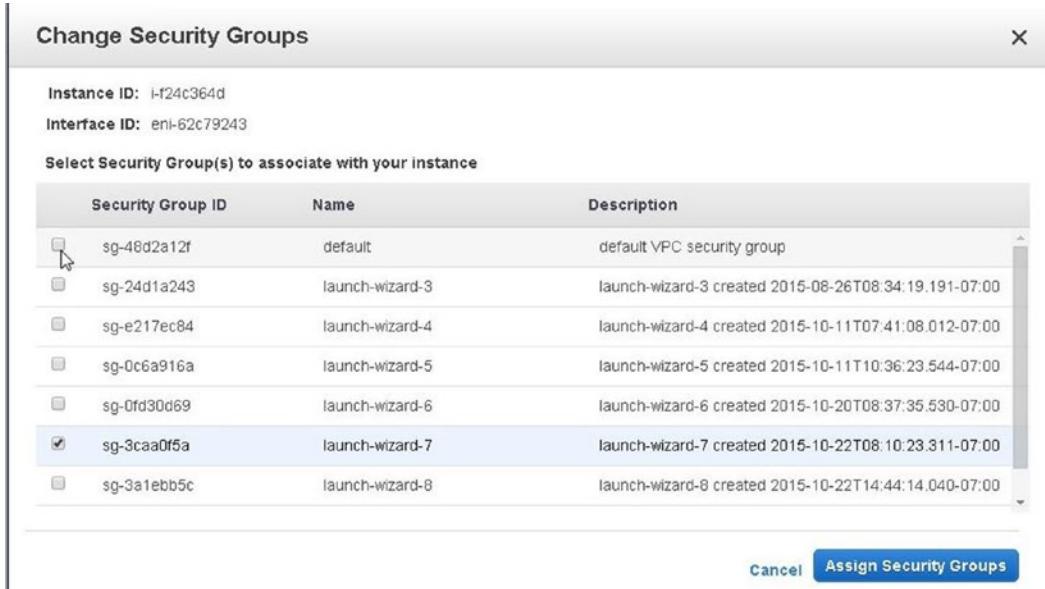


Figure A-28. The “default” group not selected

Select the checkbox for the “default” security group and click on Assign Security Groups as shown in Figure A-29.




Figure A-29. Assign Security Groups

APPENDIX A ■ USING THE AMAZON EC2

The default security group gets assigned to the Amazon EC2 instance. To find the available security groups and their inbound/outbound rules, click on Network & Security ► Security Groups as shown in Figure A-30.

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS	Public IP
Docker	i-816ee3e	t2.micro	us-east-1b	stopped		None		
Docker	i-839fe9c	m3.medium	us-east-1b	stopped		None		
Docker	i-f85265d	m3.medium	us-east-1b	stopped		None		
Docker	i-2ed3718c	m3.medium	us-east-1b	stopped		None		
Docker	i-24c364d	t2.micro	us-east-1b	running	2/2 checks	None	ec2-54-210-177-249.co...	54.210.177.249
Docker	i-6a1354d5	m3.medium	us-east-1b	running	2/2 checks	None	ec2-54-152-57-191.com...	54.152.57
Ubuntu test	i-2247029d	t2.micro	us-east-1b	stopped		None		

Instance: i-24c364d (Docker) Public DNS: ec2-54-210-177-249.compute-1.amazonaws.com

Description Status Checks Monitoring Tags

Instance ID: i-24c364d Public DNS: ec2-54-210-177-249.compute-1.amazonaws.com
Instance state: running Public IP: 54.210.177.249
Instance type: t2.micro Elastic IP:

Figure A-30. Network & Security ► Security Groups

The “default” security group should be listed. Select the “default” group. Select the Inbound tab. The Type should be listed as “All Traffic”, the Protocol as “All”, the Port Range as All and Source as 0.0.0.0. To edit the inbound rules, click on Inbound ► Edit as shown in Figure A-31.

Name	Group ID	Group Name	VPC ID	Description
sg-0cb9a916a	launch-wizard-5	vpc-3b12e5f	launch-wizard-5 created 2015-10-11T10:36:23.544-07:00	
sg-0fd30d69	launch-wizard-6	vpc-3b12e5f	launch-wizard-6 created 2015-10-20T08:37:35.530-07:00	
sg-24dfa243	launch-wizard-3	vpc-3b12e5f	launch-wizard-3 created 2015-08-26T08:34:19.191-07:00	
sg-3a1eb25c	launch-wizard-8	vpc-3b12e5f	launch-wizard-8 created 2015-10-22T14:44:14.040-07:00	
sg-3caad05a	launch-wizard-7	vpc-3b12e5f	launch-wizard-7 created 2015-10-22T08:10:23.311-07:00	
sg-48d2a12f	default	vpc-3b12e5f	default VPC security group	
sg-53dc1c0e	launch-wizard-2	vpc-3b12e5f	launch-wizard-2 created 2015-01-26T16:29:09.921-08:00	
sg-560c5e09	TrustySG	vpc-3b12e5f	Security Group for Ubuntu Trusty	
sg-a8d554c6	quicklaunch-1	vpc-3b12e5f	quicklaunch-1	
sg-ce8a53a6	default	vpc-3b12e5f	default group	

Description Inbound Outbound Tags

Edit

Type	Protocol	Port Range	Source
All traffic	All	All	0.0.0.0/0

Figure A-31. Inbound ► Edit

The inbound rules get displayed and should be kept as the default settings as shown in Figure A-32. Click on Save.

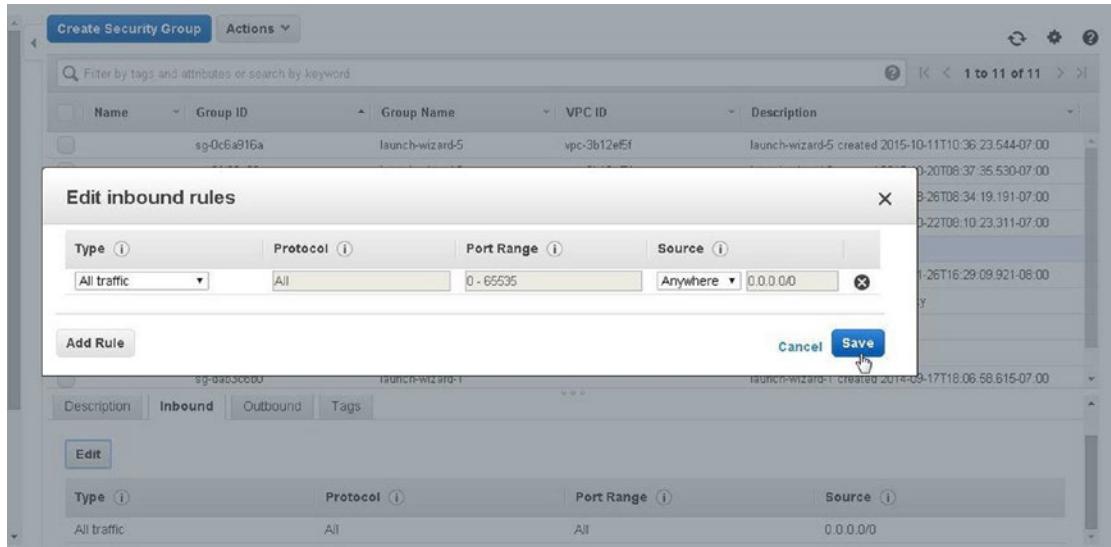


Figure A-32. Edit inbound rules dialog

Similarly, select the Outbound tab. The Type should be listed as “All Traffic”, the Protocol as “All”, the Port Range as All and Destination as 0.0.0.0. Click on Edit as shown in Figure A-33.




Figure A-33. Outbound ► Edit

APPENDIX A ■ USING THE AMAZON EC2

The default settings for the Outbound rules get displayed and should be kept as the default as shown in Figure A-34. Click on Save.

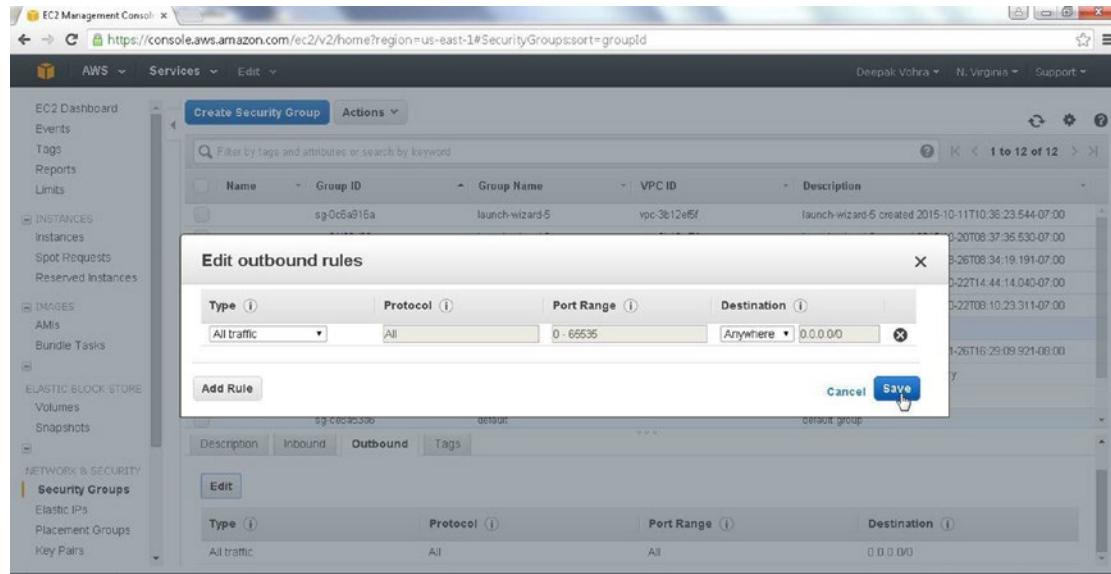


Figure A-34. Edit outbound rules dialog

The security groups assigned to an instance are listed in the Security Groups column as shown in Figure A-35.

Launch Instance Connect Actions							
Filter by tags and attributes or search by keyword							
Status Checks	Alarm Status	Public DNS	Public IP	Key Name	Monitoring	Launch Time	Security Groups
None	🟡			docker	disabled	October 22, 2015 at 6:22:42 ...	default, launch-wizard-5
None	🟡			docker	disabled	October 20, 2015 at 7:12:55 ...	launch-wizard-4
None	🟡			docker	disabled	October 20, 2015 at 6:20:33 ...	launch-wizard-4
None	🟡			docker	disabled	October 20, 2015 at 7:40:00 ...	launch-wizard-4
None	🟡			docker	disabled	October 20, 2015 at 4:47:46 ...	launch-wizard-5
None	🟡			docker	disabled	October 20, 2015 at 6:56:08 ...	launch-wizard-5
2/2 checks...	🟢	ec2-54-210-177-249.co...	54.210.177.249	docker	disabled	October 22, 2015 at 8:10:29 ...	launch-wizard-7, ...
Initializing	🟡	ec2-54-152-57-191.com...	54.152.57.191	docker	disabled	October 22, 2015 at 2:47:54 ...	launch-wizard-8
	🟡			docker	disabled	October 22, 2015 at 10:28:0...	TrustySG

Figure A-35. Security Groups column

Stopping an Amazon EC2 Instance

To stop an Amazon EC2 instance select the instance and select Actions ► Instance State ► Stop as shown in Figure A-36.

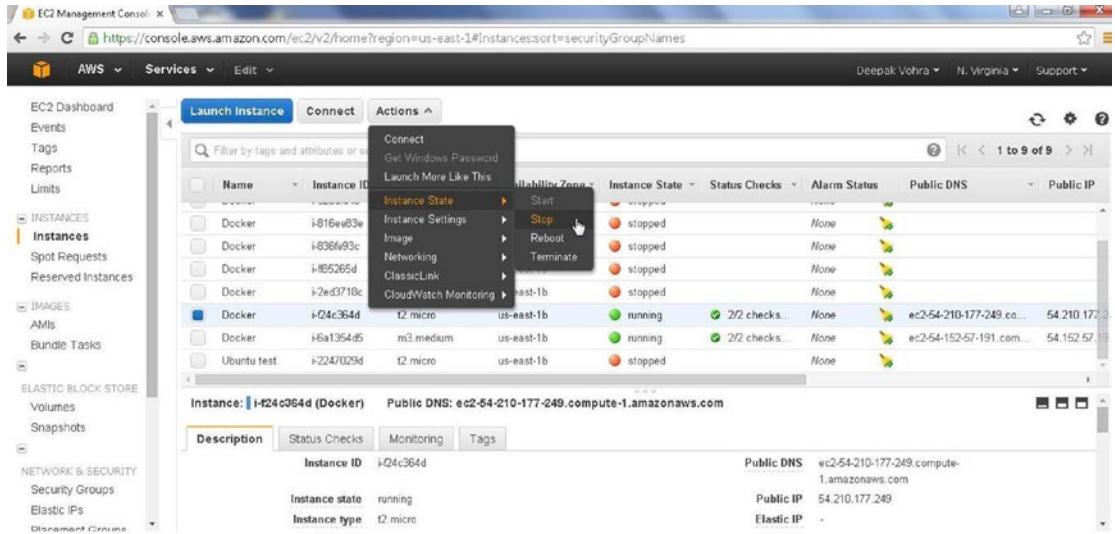


Figure A-36. Actions ► Instance State ► Stop

Multiple instances may be selected and stopped together as shown in Figure A-37.

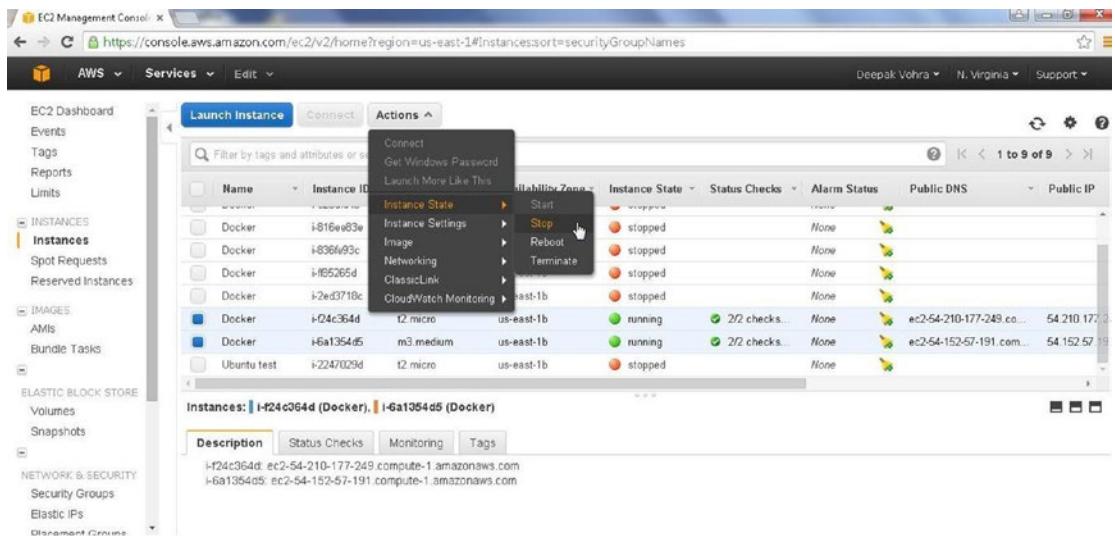


Figure A-37. Stopping Multiple Instances

In the Stop Instance dialog, click on Yes, Stop as shown in Figure A-38.

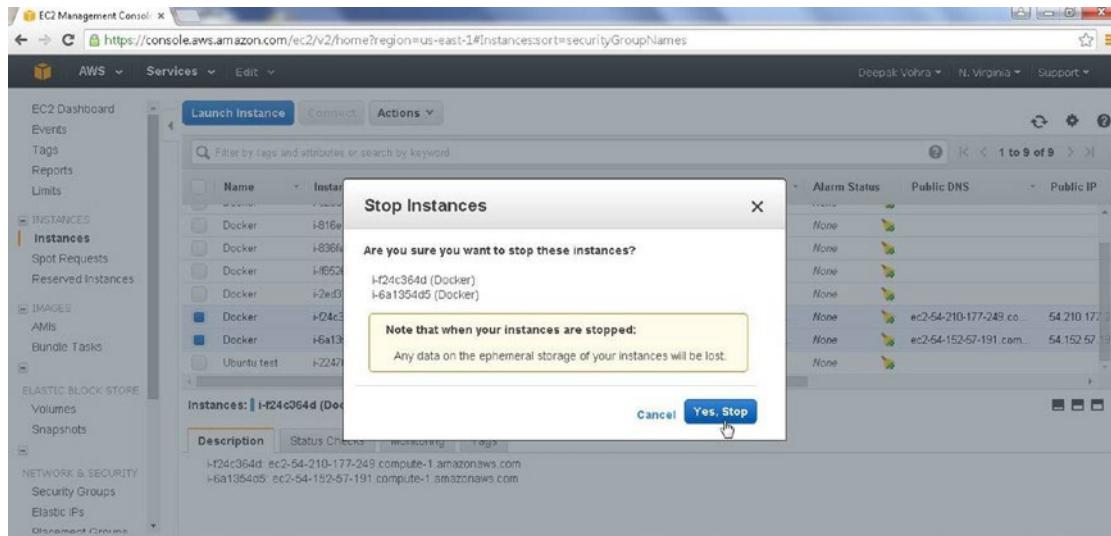


Figure 38. Stop Instance dialog

The instance/s get stopped.

Changing the Instance Type

To increase or decrease the capacity of an instance, it may be required to change the instance type, such as from a micro instance to a medium instance. An instance must first be stopped before changing its type and later restarted after modifying the type. To change the instance type, select the instance and select Actions ► Instance Settings ► Change Instance Type as shown in Figure A-39.

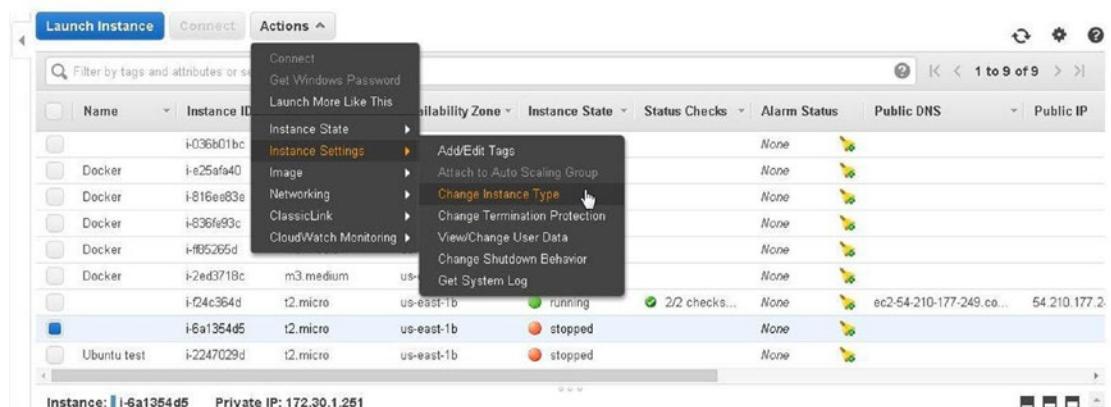


Figure 39. Actions ► Instance Settings ► Change Instance Type

In the Change Instance Type dialog, select the Instance Type to apply, for example, m3.medium as shown in Figure A-40.

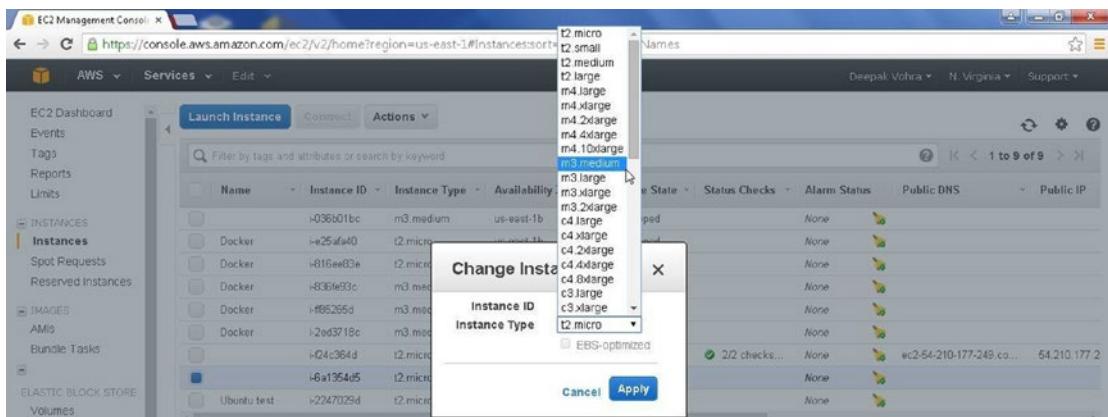


Figure A-40. Change Instance Type dialog

Click on Apply as shown in Figure A-41.

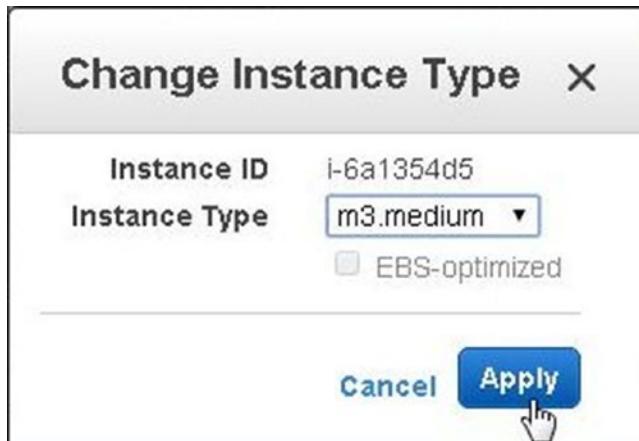


Figure A-41. Applying a new Instance Type

The instance type gets upgraded to m3.medium as shown in Figure A-42. Keep in consideration that upgrading an instance type could make the instance not eligible for the free tier.

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS	Public IP
	i-036bb01bc	m3.medium	us-east-1b	stopped	None			
Docker	i-e25afa40	t2.micro	us-east-1b	stopped	None			
Docker	i-816ee83e	t2.micro	us-east-1b	stopped	None			
Docker	i-836fe93c	m3.medium	us-east-1b	stopped	None			
Docker	i-f85265d	m3.medium	us-east-1b	stopped	None			
Docker	i-2ed3718c	m3.medium	us-east-1b	stopped	None			
	i-Q4c364d	t2.micro	us-east-1b	running	2/2 checks...	None		ec2-54-210-177-249.co... 54.210.177.2
<input checked="" type="checkbox"/>	i-6a1354d5	m3.medium	us-east-1b	stopped	None			
Ubuntu test	i-2247029d	t2.micro	us-east-1b	stopped	None			

Figure A-42. Upgraded Instance Type

Summary

In Appendix A we discussed creating an Amazon EC2 instance based on an AMI, starting an instance, connecting to an instance, finding the Public IP Address, finding the Public DNS, changing the instance type and stopping an instance.

Index

A, B

- Amazon EC2 instance, 229
 - apply instance type, 251
 - 64 bits, 229
 - change instance type, 250
 - creation
 - keypair, 232
 - launch status, 229, 233
 - pending state, 234
 - review and launch, 231
 - running instance, 234
 - Ubuntu AMI, 230
 - default security group, 244
 - DNS Hostnames dialog, 244
 - inbound rules, 247
 - key pair, 235
 - creation, 236
 - delete, 235
 - network and security, 235
 - network security, 246
 - outbound rules, 247
 - Public DNS, 240
 - Public IP Address, 240
 - security groups, 248
 - ssh command, 238
 - start instances, 237
 - stop command, 249
 - Ubuntu instance, 239
 - upgrade instance type, 252
 - VPC Dashboard, 242
- Amazon EC2 instances
 - Linux architecture, 4
 - public IP address, 3
 - RHEL instance, 4
 - Ubuntu AMI, 2
 - Ubuntu instance, 4
- Amazon Elastic Compute Cloud (Amazon EC2), 229
- Amazon Machine Images (AMIs), 2
- Apache Cassandra
 - CQL shell, 85
 - detached mode, 83
 - Docker containers, 84
 - environment settings, 82
 - exit command, 91
 - interactive terminal (tty), 84
 - keyspace
 - alter statement, 86
 - creation, 85
 - USE statement, 86
 - multiple instances, 92
 - overview, 81
 - stop command, 92
 - table
 - creation, 87
 - DELETE statement, 89
 - DROP COLUMN FAMILY statement, 90
 - DROP KEYSPACE statement, 91
 - INSERT DML statement, 87
 - SELECT statement, 88
 - TRUNCATE statement, 90
- Apache Hadoop
 - CDH framework, 128
 - configuration files, 129
 - HDFS commands, 128
 - environment settings, 117
 - image sequenceiq/hadoop-docker, 119
 - interactive shell (tty), 120
 - MapReduce Word Count application, 121
 - HADOOP_PREFIX
 - Directory, 121
 - HDFS directory, 124
 - output, 125, 127
 - wq command, 122
 - YARN framework, 124
 - overview, 117
 - pull command, 118
 - run command, 119
 - sequenceiq/hadoop-docker image, 119
 - stop command, 128

■ INDEX

- Apache HBase, 141
 - Docker container, 143
 - environment settings, 141
 - interactive terminal (tty), 143
 - shell command, 144
 - tables
 - creation, 144
 - get command, 147
 - lists, 146
 - scan command, 148
 - stop command, 149
- Apache Hive
 - Beeline CLI, 132
 - cdh process, 132
 - !connect command, 134
 - environment settings, 131
 - HiveServer2, 133
 - load data, 136
 - overview, 131
 - query, 138
 - table creation, 135
 - use default command, 135
- Apache Kafka
 - consumer console, 191
 - docker image, 186
 - environment settings, 186
 - IP address, 189
 - logs command, 190
 - messages, 192
 - overview, 185
 - producer console, 191
 - run command, 188
 - sh-create command, 191
 - stop command, 193
- Apache Solr, 195
 - admin console login, 200
 - Amazon EC2 instance, 195
 - core index creation, 201
 - delete command, 214
 - docker container, 198
 - interactive shell, 199
 - logs command, 198, 216
 - pull command, 197
 - query tab, 206
 - execute query, 208
 - JSON Response, 209
 - Request-Handler, 207
 - version field, 210
 - REST API client, 210
 - service status, 196
 - stop command, 217
 - XML format document, 204
- Apache Spark, 219
 - CDH frameworks, 220
 - environment settings, 219
 - scala> command, 226
 - shell command, 226
 - yarn-client mode, 224
 - yarn-cluster mode, 221
- Apache Sqoop, 151
 - cdh and mysql containers, 154
 - core-site.xml configuration, 166
 - data transfer paths, 151
 - environment settings, 152
- HDFS
 - command parameters, 168
 - listed data, 174
 - scoop codegen command, 168
 - scoop help command, 168
 - scoop import command, 168, 170
 - hdfs-site.xml configuration, 165
 - INSERT SQL statements, 157
 - interactive terminals (tty), 155
 - JAVA_HOME environment, 160
 - MapReduce framework, 163
 - MySQL CLI, 155
 - MySQL JDBC jar, 160
 - preceding commands, 153
 - privileges settings, 156
 - remove command, 183
 - run command, 154
 - SELECT statement, 181
 - SQL query, 158
 - scoop export command, 175
 - stop command, 182
 - table creation, 157
 - variables, 154
 - wlslog, 156
 - WLSLOG_COPY table, 159
 - wlslog.jar file, 169
- Apache ZooKeeper, 188

■ C

- Cloudera Hadoop distribution (CDH), 128
- command-line interface tools (CLI), 114
- Couchbase Server, 95
 - add documents
 - JSON Document, 112
 - CLI tools, 114
 - cluster configuration, 101
 - buckets section, 103
 - default buckets, 108

disk storage section, 101
 flush, enable, 104
 IP Address, 108
 notifications, 105
 overview tab, 107
 RAM configuration, 102
 servers summary, 108
 username and password, 106
Docker image, 97
 document creation, 109
 environment settings, 95
 exit command, 115
 interactive terminal (tty), 114
 logs command, 99
 Ubuntu Server AMI, 95
 unlimit settings, 98
Web Console, 99
 admin, 100
 login, 100

■ D, E, F, G, H, I, J, K

Docker
 Amazon EC2 instances, 2
 Linux architecture, 4
 public DNS, 15
 public IP address, 3
 RHEL instance, 4
 Ubuntu AMI, 2
 Ubuntu instance, 4
container, 13
 listing port, 13
 list running, 13
 stop command, 16
containerremove
 command, 17
curl tool, 14
find status, 9
 Hello word application, 9
 image download, 11
 installation, 7
 overview, 1
Red Hat, 5
 installation, 5
 updated service, 6
rmi command, 17
run command, 12
stop command, 18
systemctl, 8
 Ubuntu version, 7
 uninstallation, 7

■ L

Linux
container
 attached mode, 23
 detached mode, 22
 inspect command, 24
 lists, 23
 top command, 25
 create command, 28
 environment settings, 19
 image download, 21
 image list, 21
 interactive shell, 26
 exit command, 28
 files and directories, 27
 Oracle Server, 27
 -i and -t options, 26
 remove command, 29
 stop command, 29

■ M, N

MongoDB, 57
 backup data, 73
 restore database, 73
 test database, 73
 batch of documents, 69
 collection, 63–64
 database creation, 63
 Docker container, 59
 document creation, 65
 drop() method, 69
 environment settings, 58
 exit command, 80
 find() method, 66
 findOne() method, 68
 insert () method, 66
 interactive terminal, 60
 JSON format, 72
 remove documents, 75
 save() method, 71
 shell command, 60
 stop command, 78
MySQL database
 CLI shell, 46
 command parameters, 45
 data directory, 44
 default database, 53
 environment settings, 42
 exit command, 49

■ INDEX

INSERT statement, 47
listing commands, 48
login command, 52
logs command, 54
overview, 41
password, 50
run command, 45
SELECT statement, 48
stop command, 49
table creation, 47
use command, 46

■ O, P, Q, R, S, T, U, V, W, X, Y, Z

Oracle Database, 31
container logs, 34
environment settings, 31
orcldb container, 33
remove command, 39
SQL*Plus, 37
table creation, 38
user creation, 37